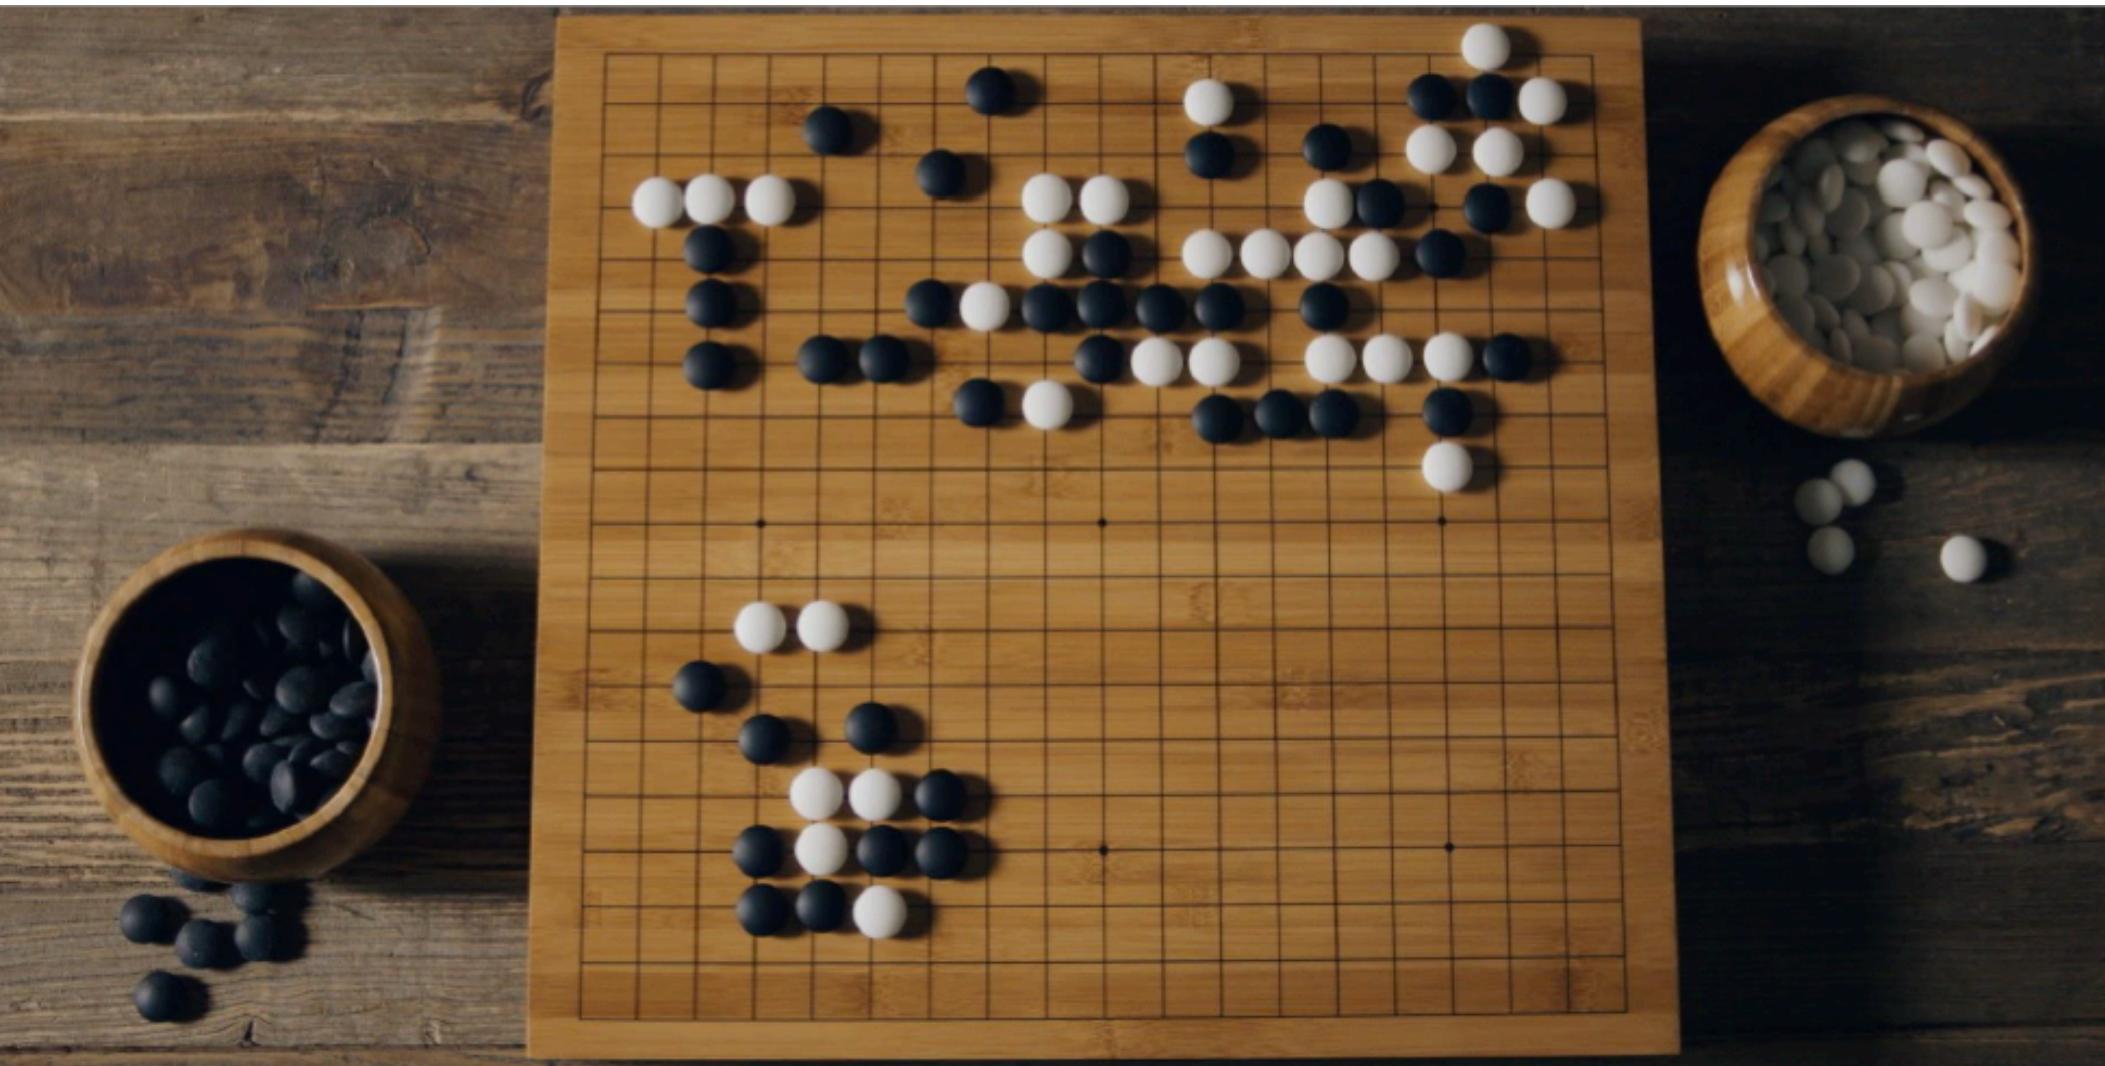


AlphaGo

Go (board game)



Invented more than 2,000 years ago, in China



Growth in East Asia

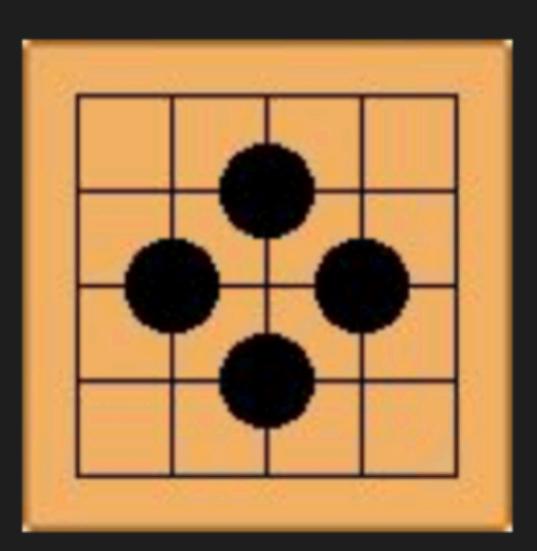
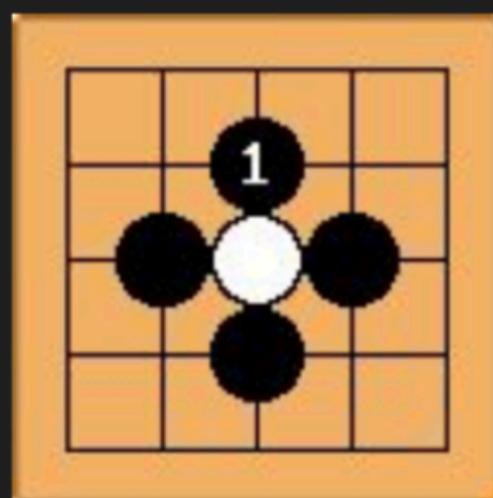
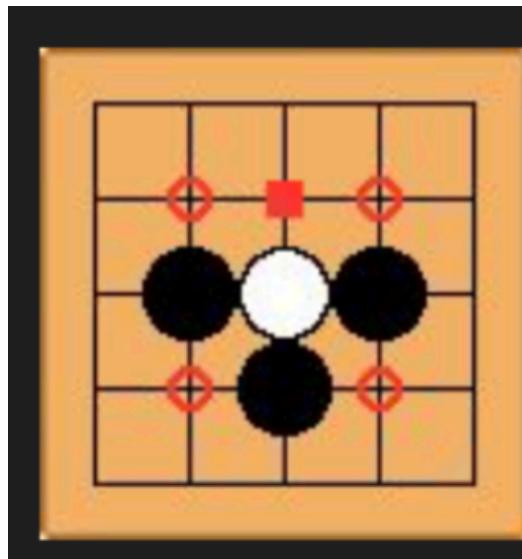


Currently over 40 million players



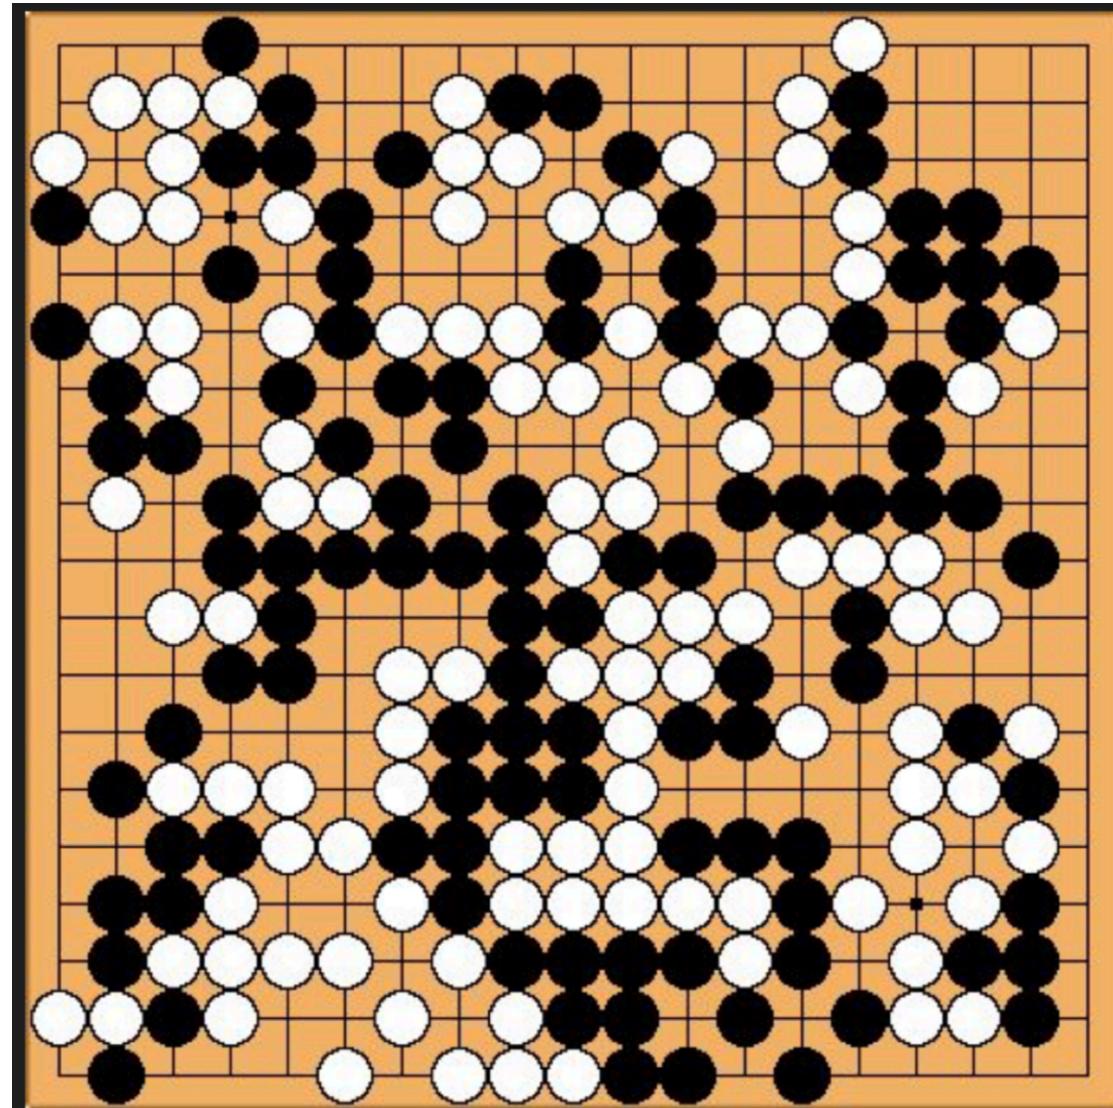
Rule of Go

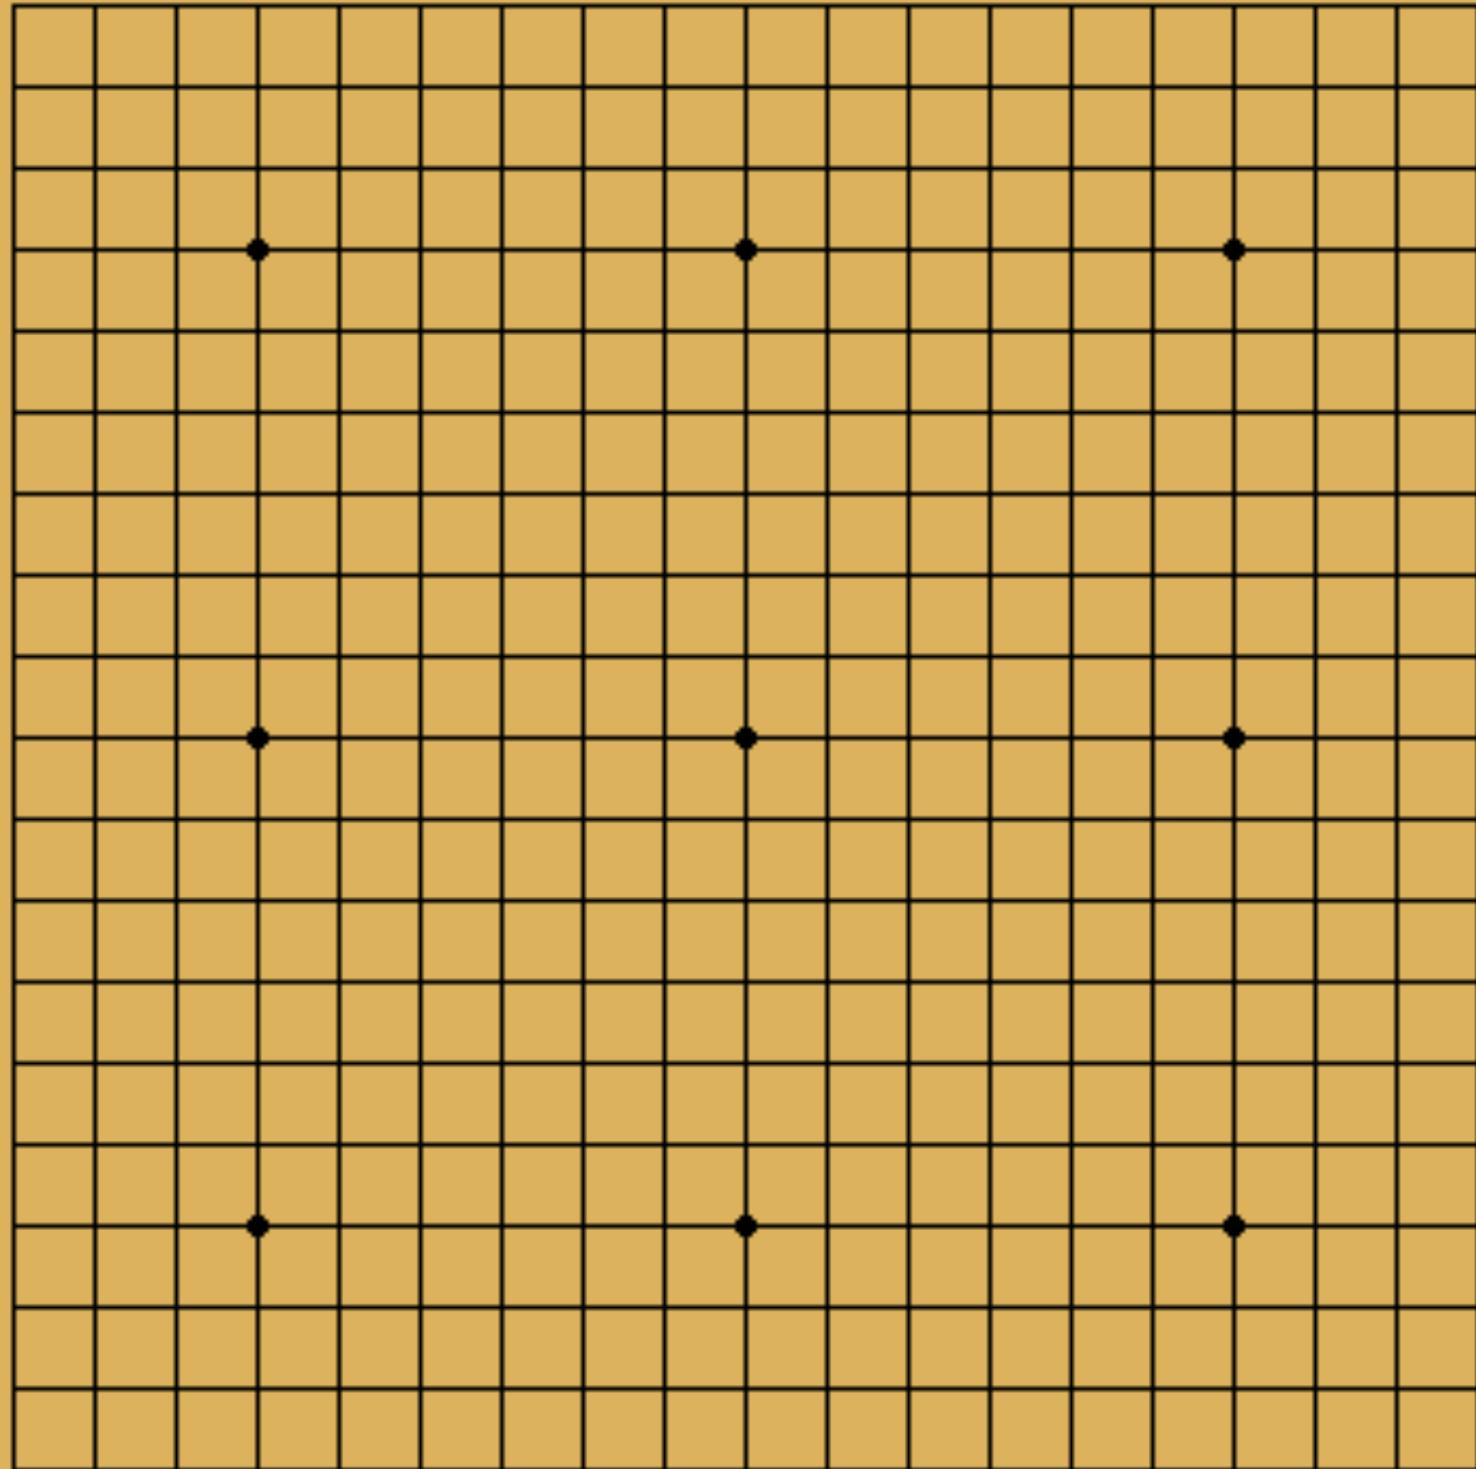
- Played on a 19x19 board
- Two players, black and white, each place one stone per turn
- Capture the opponent's stones by surrounding them

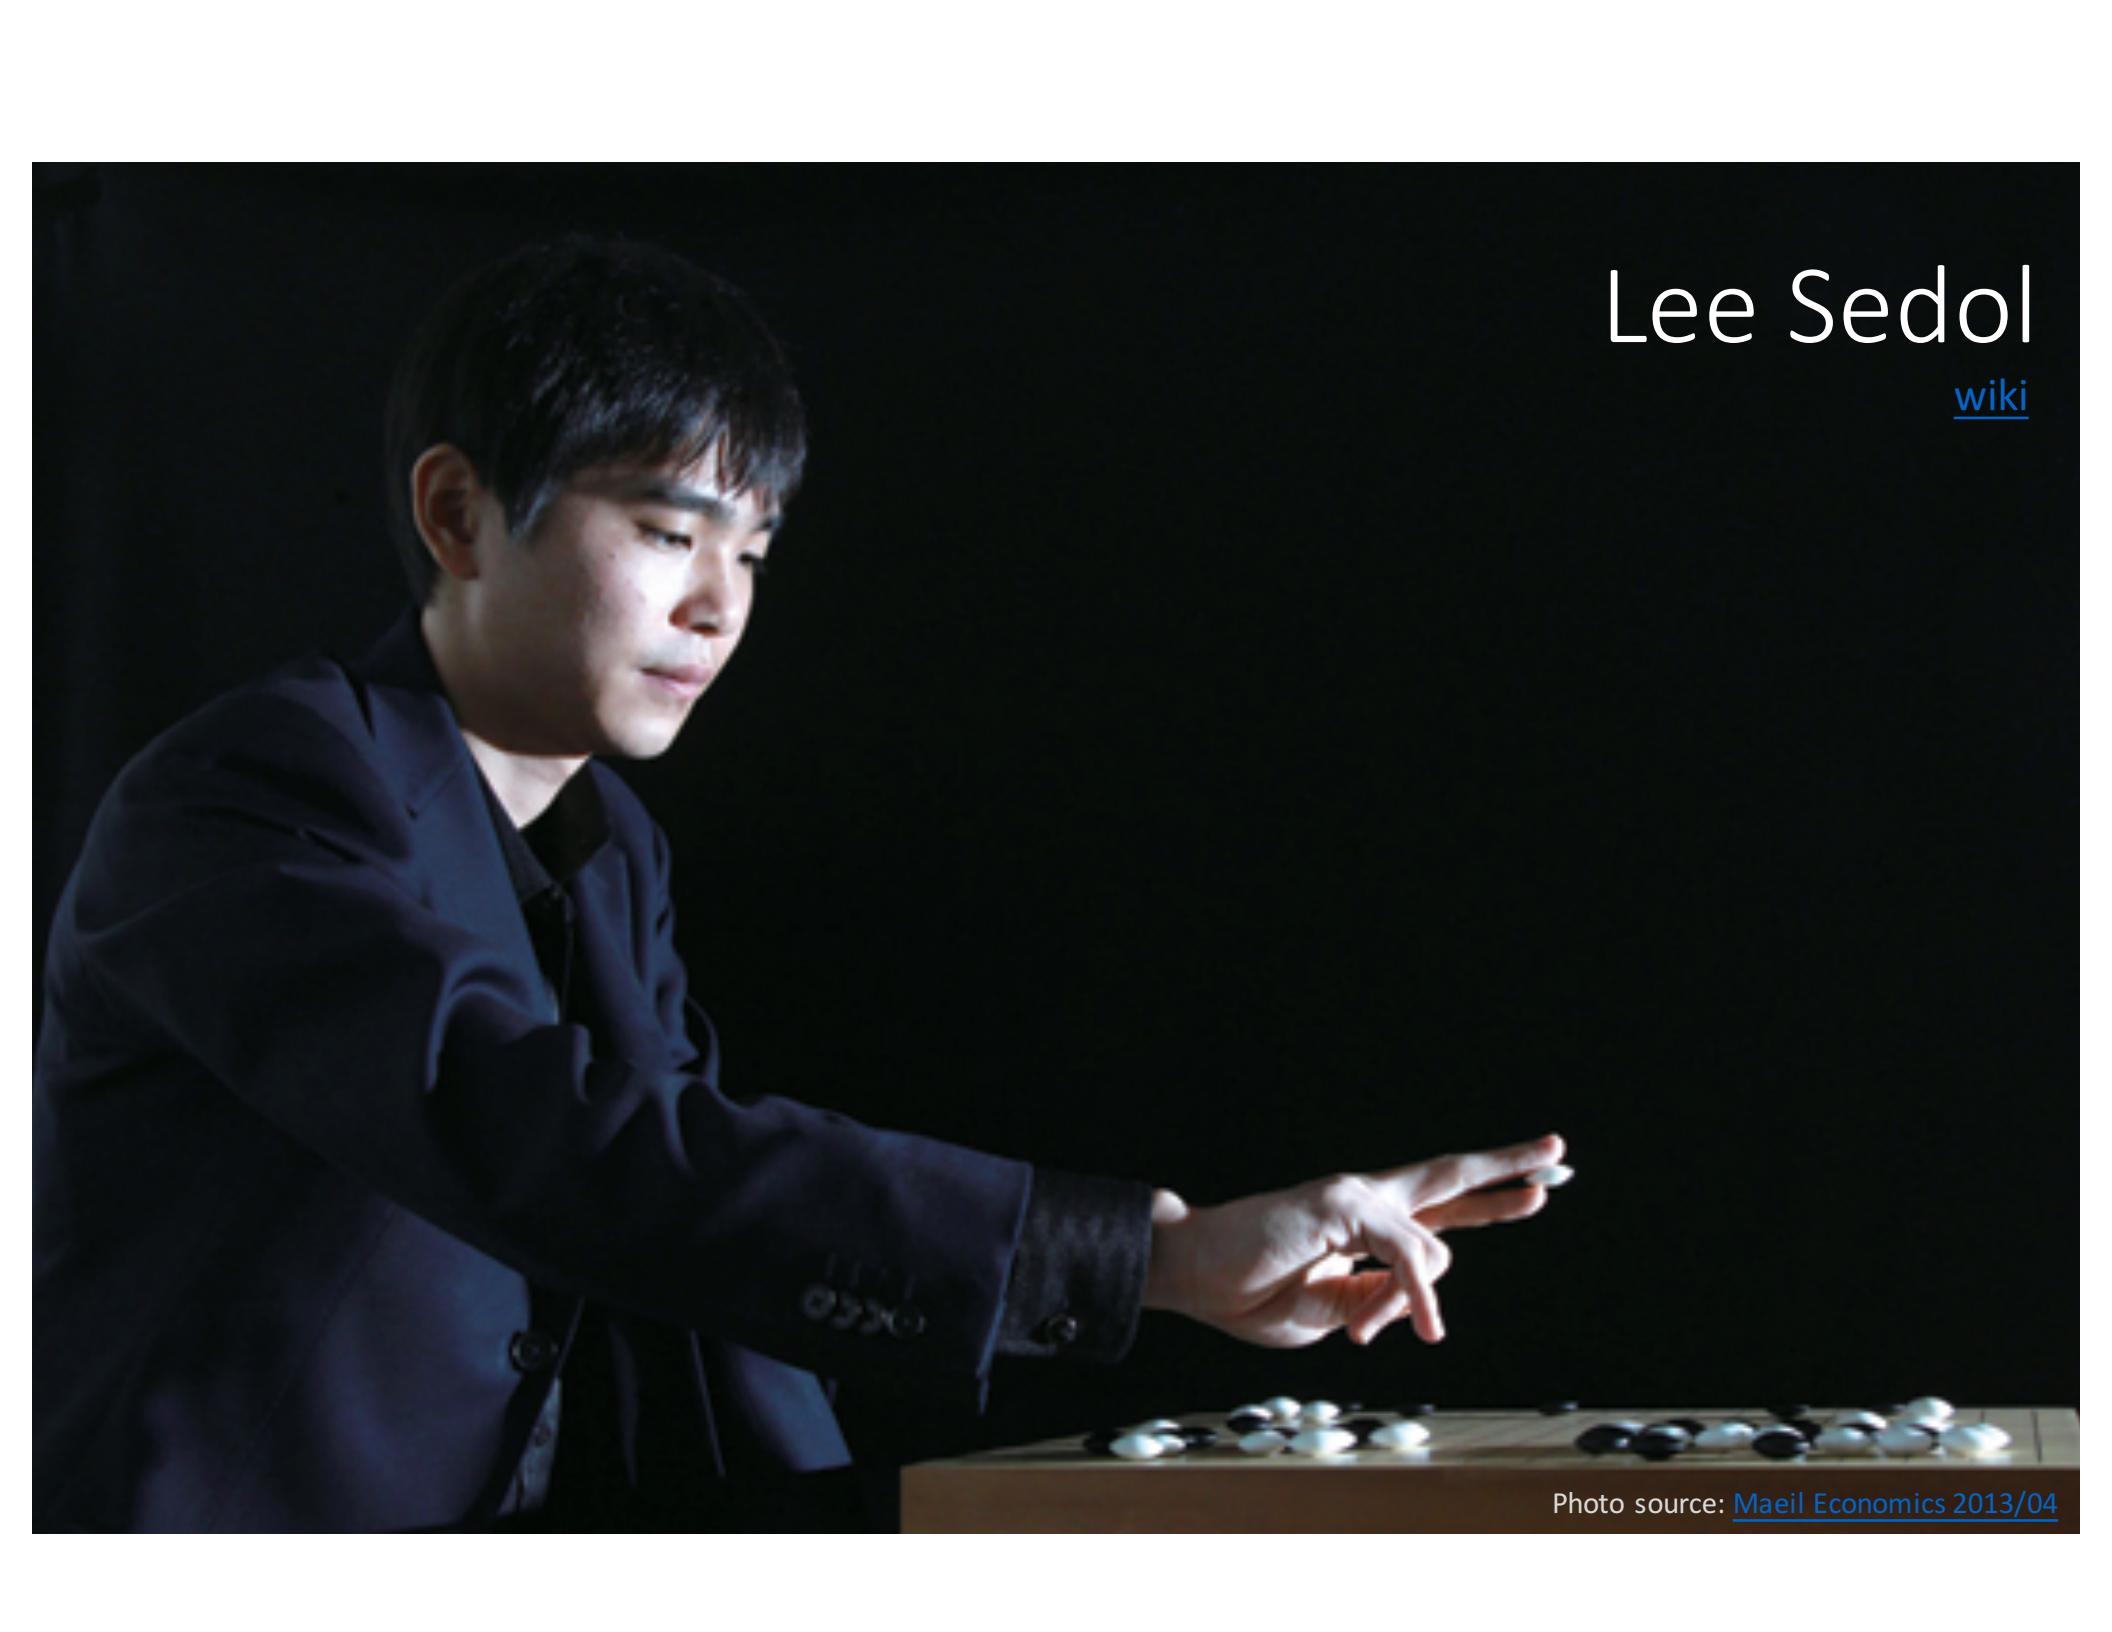


Rule of Go

- Goal is to control as much territory as possible.







Lee Sedol

[wiki](#)

Photo source: [Maeil Economics 2013/04](#)

Lee Sedol



Go History

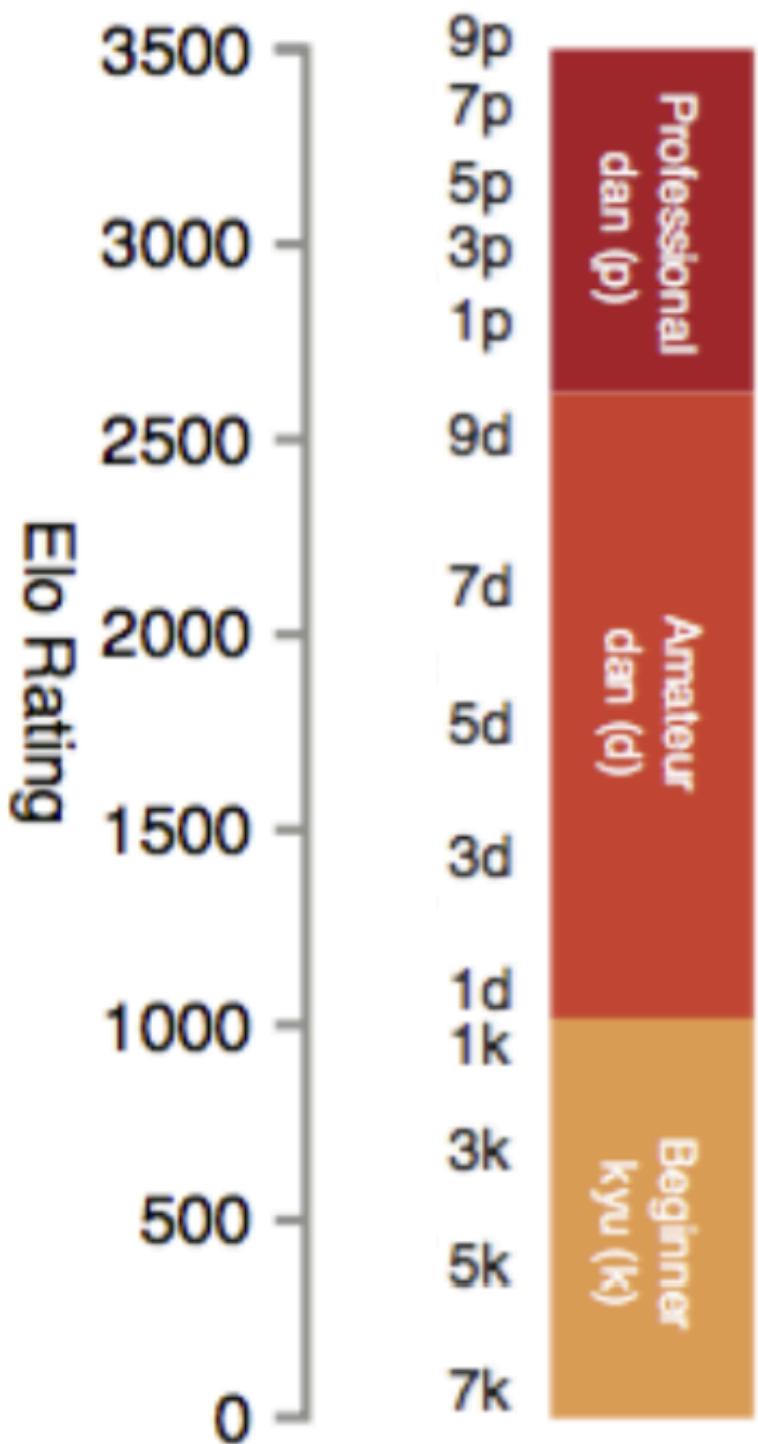


 South Korea
Lee Sedol (26)
#1 WHR holder
for 1352 days

Top 10 strongest Go players by WHR on
Feb 3, 2010



<https://www.youtube.com/watch?v=oRvlyEpOQ-8>

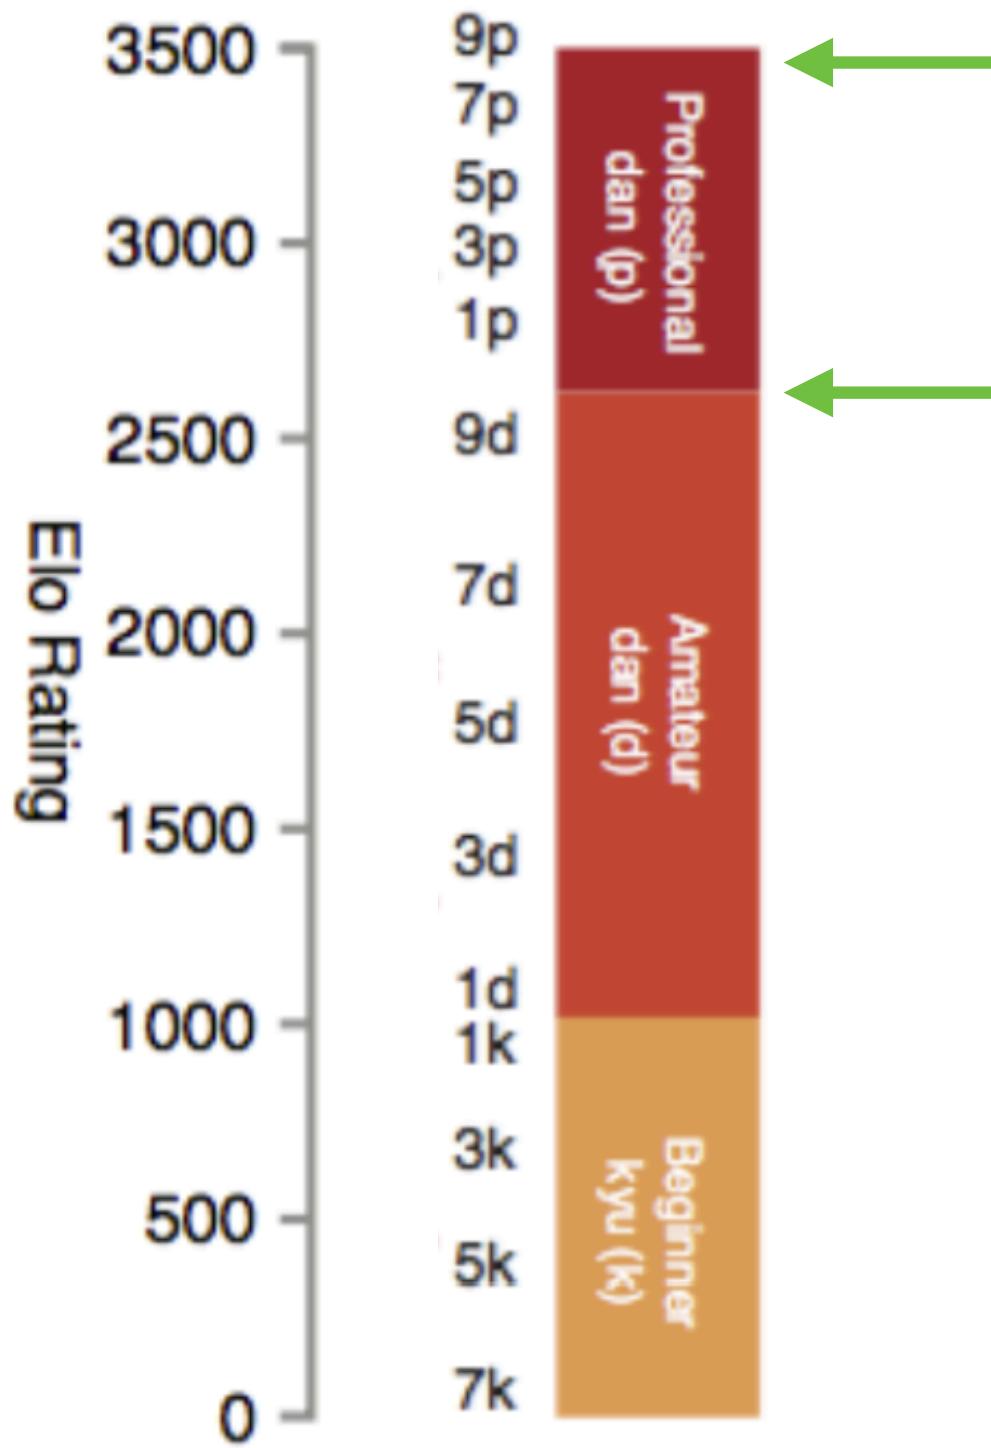




Lee Se-dol (1983)

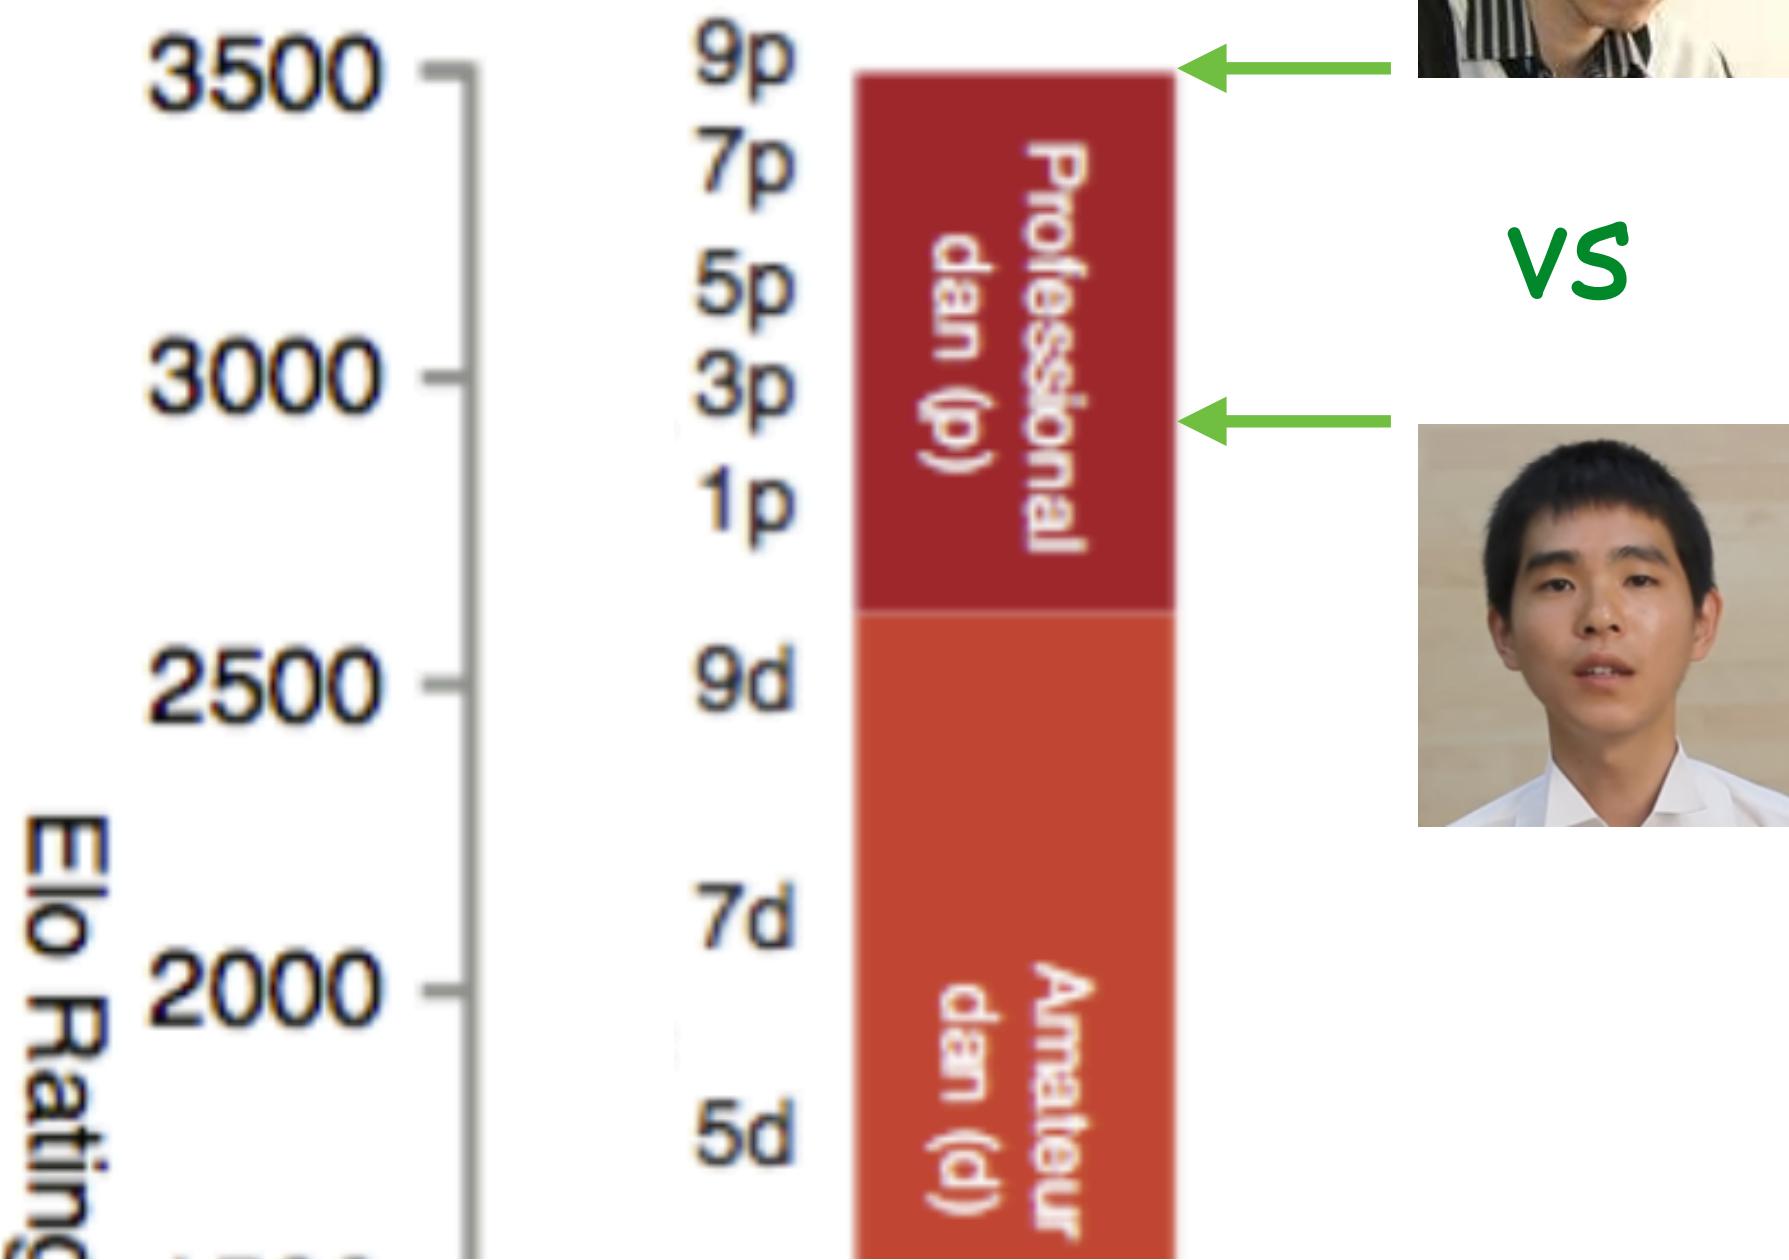
Rank	Date	
1 dan	2 July 1995	Promoted to professional dan rank after passing quality test.
2 dan	1 January 1998	
3 dan	1 January 2000	
6 dan	28 March 2003	Won the 7th LG Cup against Lee Changho (4 dan and 1 p).
7 dan	6 May 2003	Runner up in the KT Cup against Yoo Changhyuk .
9 dan	7 July 2003	Won 16th Fujitsu Cup against Song Taekon (8 dan and 1 p).

- In 2003, advanced to 9p, became the youngest 9p (21) and the fastest player to reach 9p from 1p (**8 years**).



use 8 years
fastest
youngest

2001



Lee Chang ho (9p dan)



LG CUP (2001)



9 dan vs 3 dan



LG CUP (2001)



LG CUP (2003)

Years	Nat.	Winner	Score	Nat.	Runner-up
1996-1997		Lee Changho	3-0		
1997-1998		O Rissei	3-2		Yoo Changhyuk
1998-1999		Lee Changho	3-0		Ma Xiaochun
1999-2000		Yu Bin	3-1		Yoo Changhyuk
2000-2001		Lee Changho	3-2		Lee Sedol
2001-2002		Yoo Changhyuk	3-2		Cho Hunhyun
2002-2003		Lee Sedol	3-1		Lee Changho

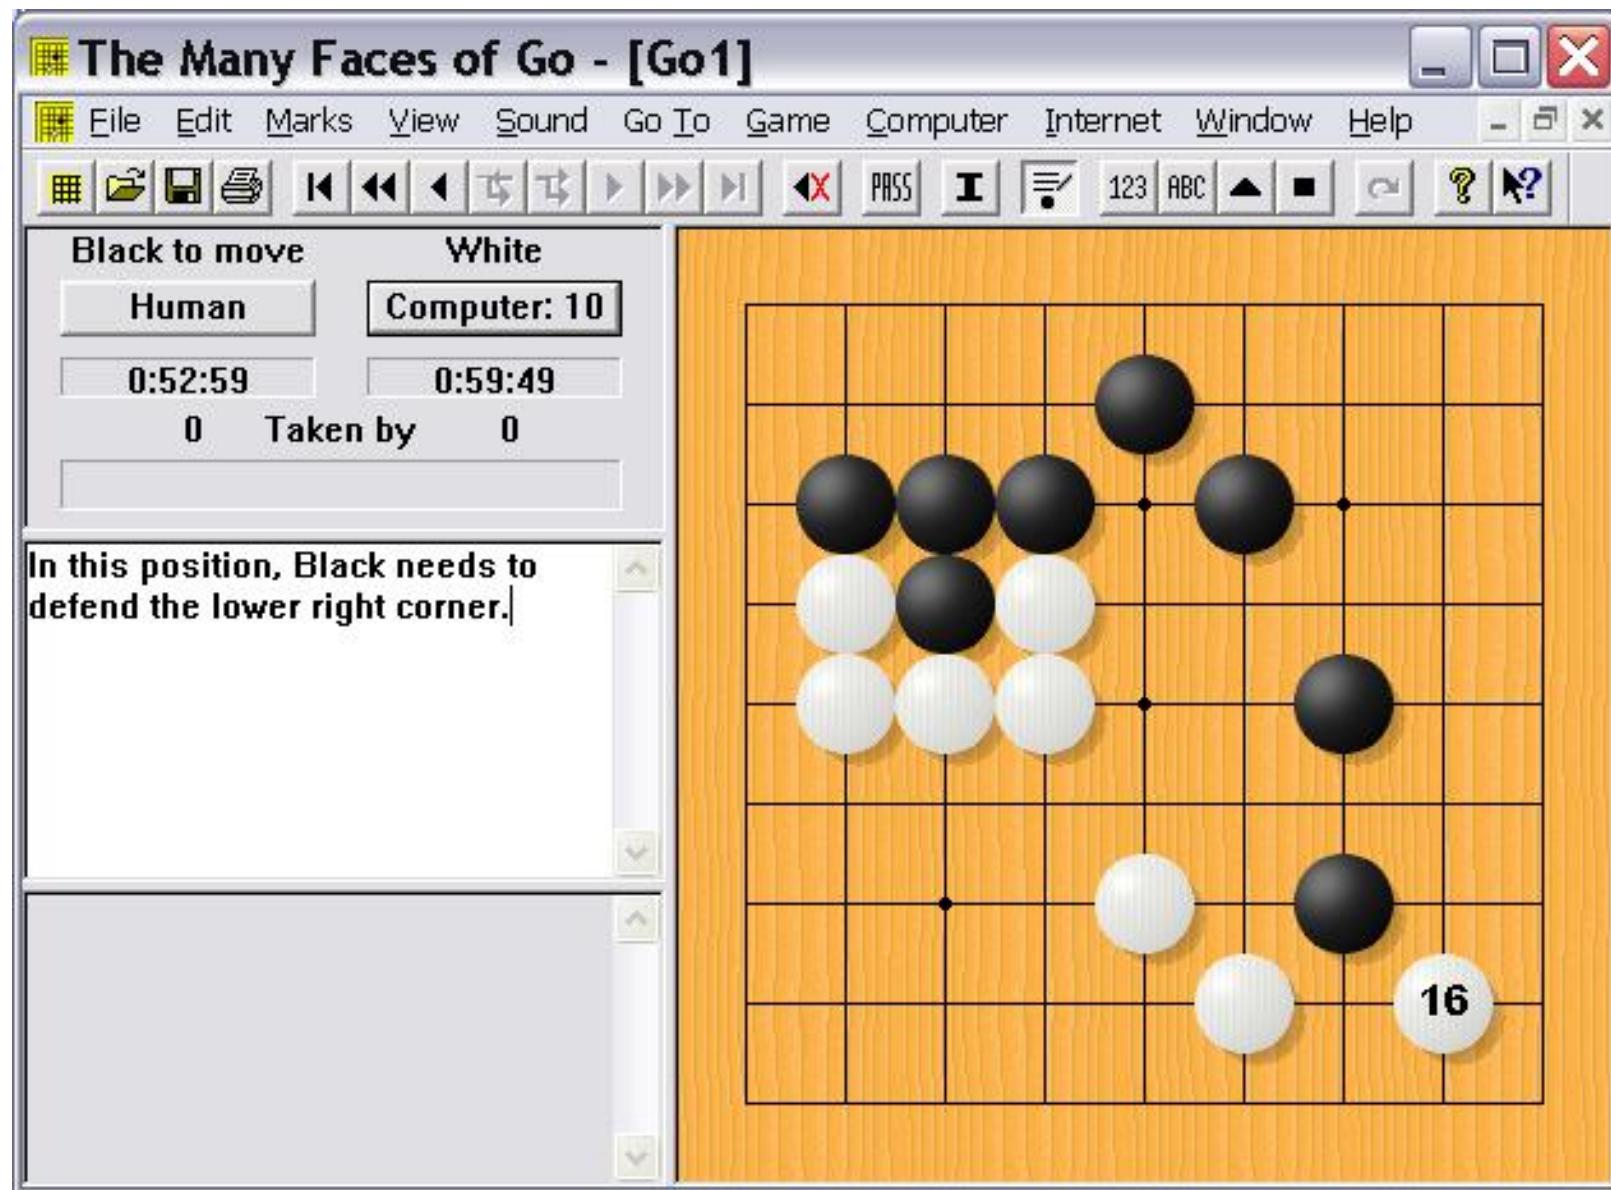
Lee Sedol

[wiki](#)



Photo source: [Maeil Economics 2013/04](#)

Computer (AI)



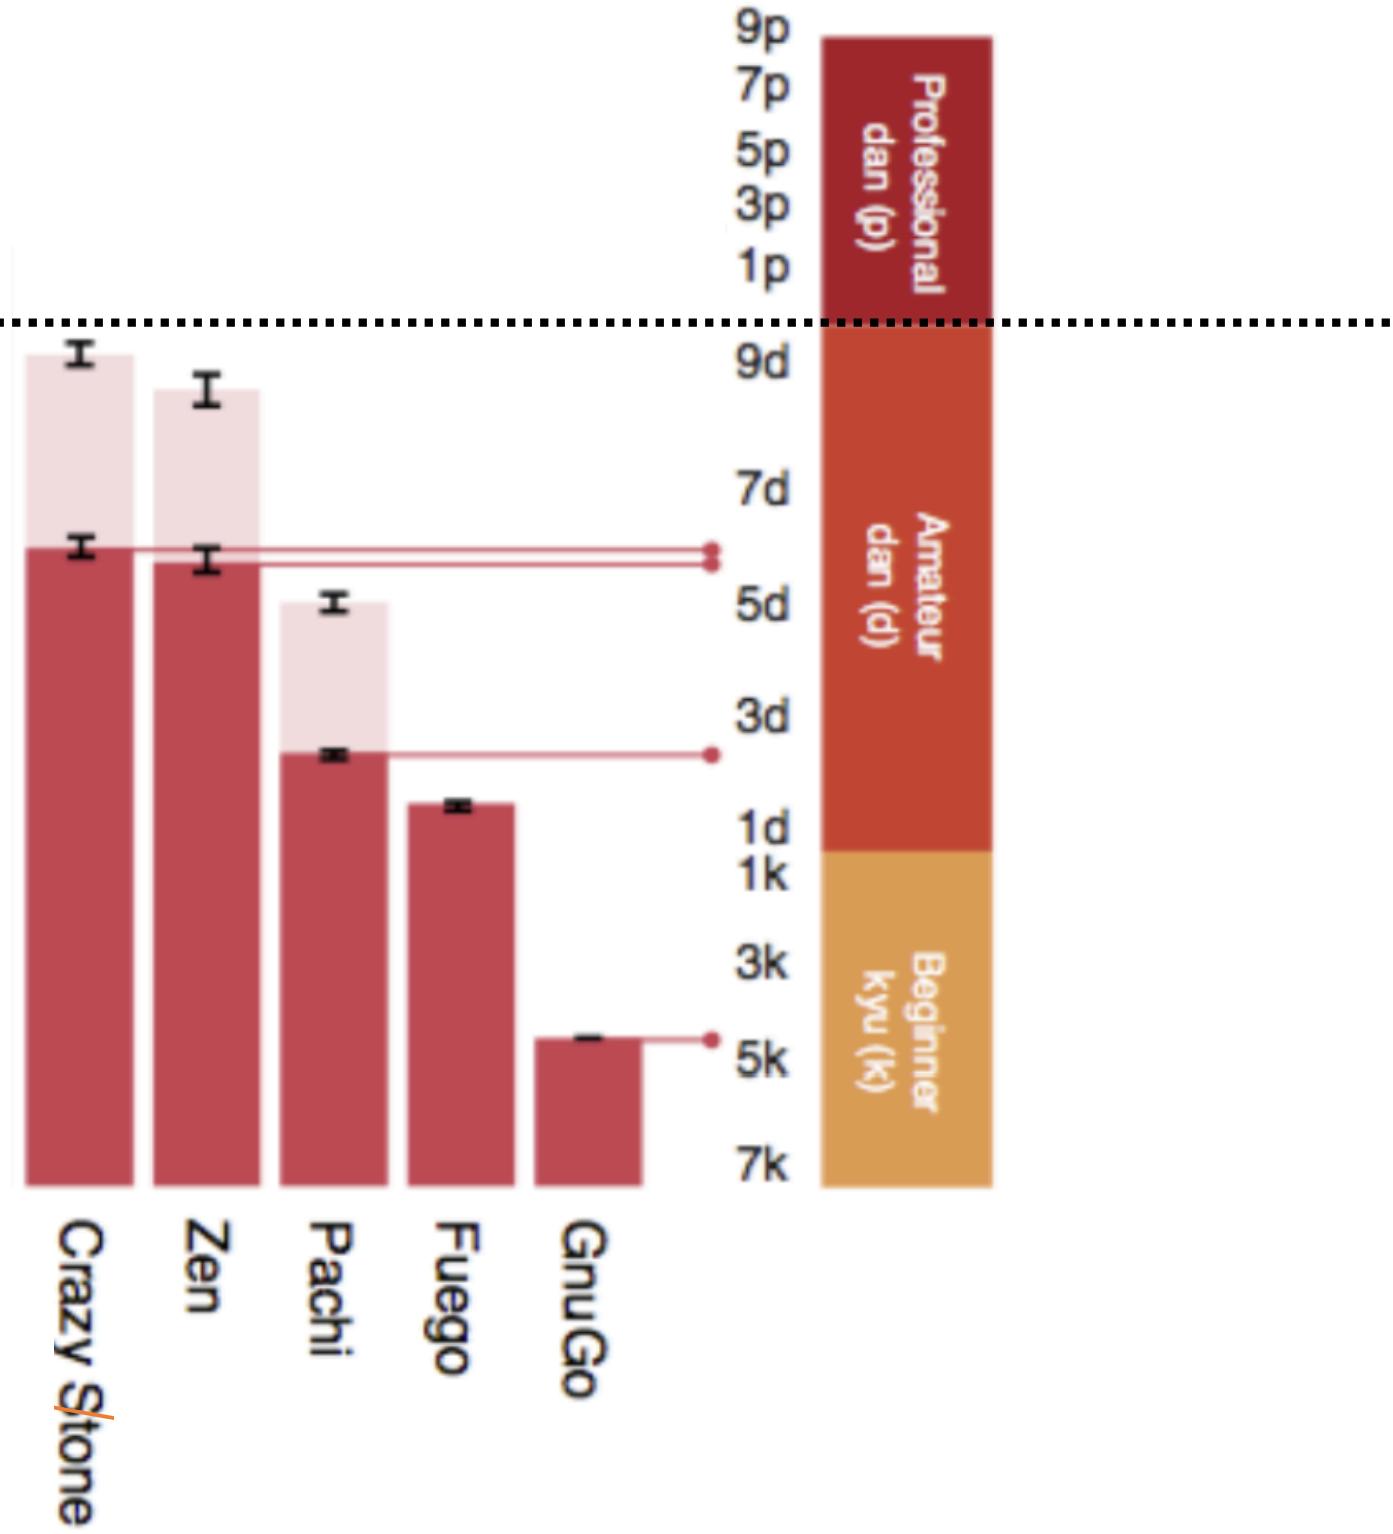
Human Expert vs Computer AI



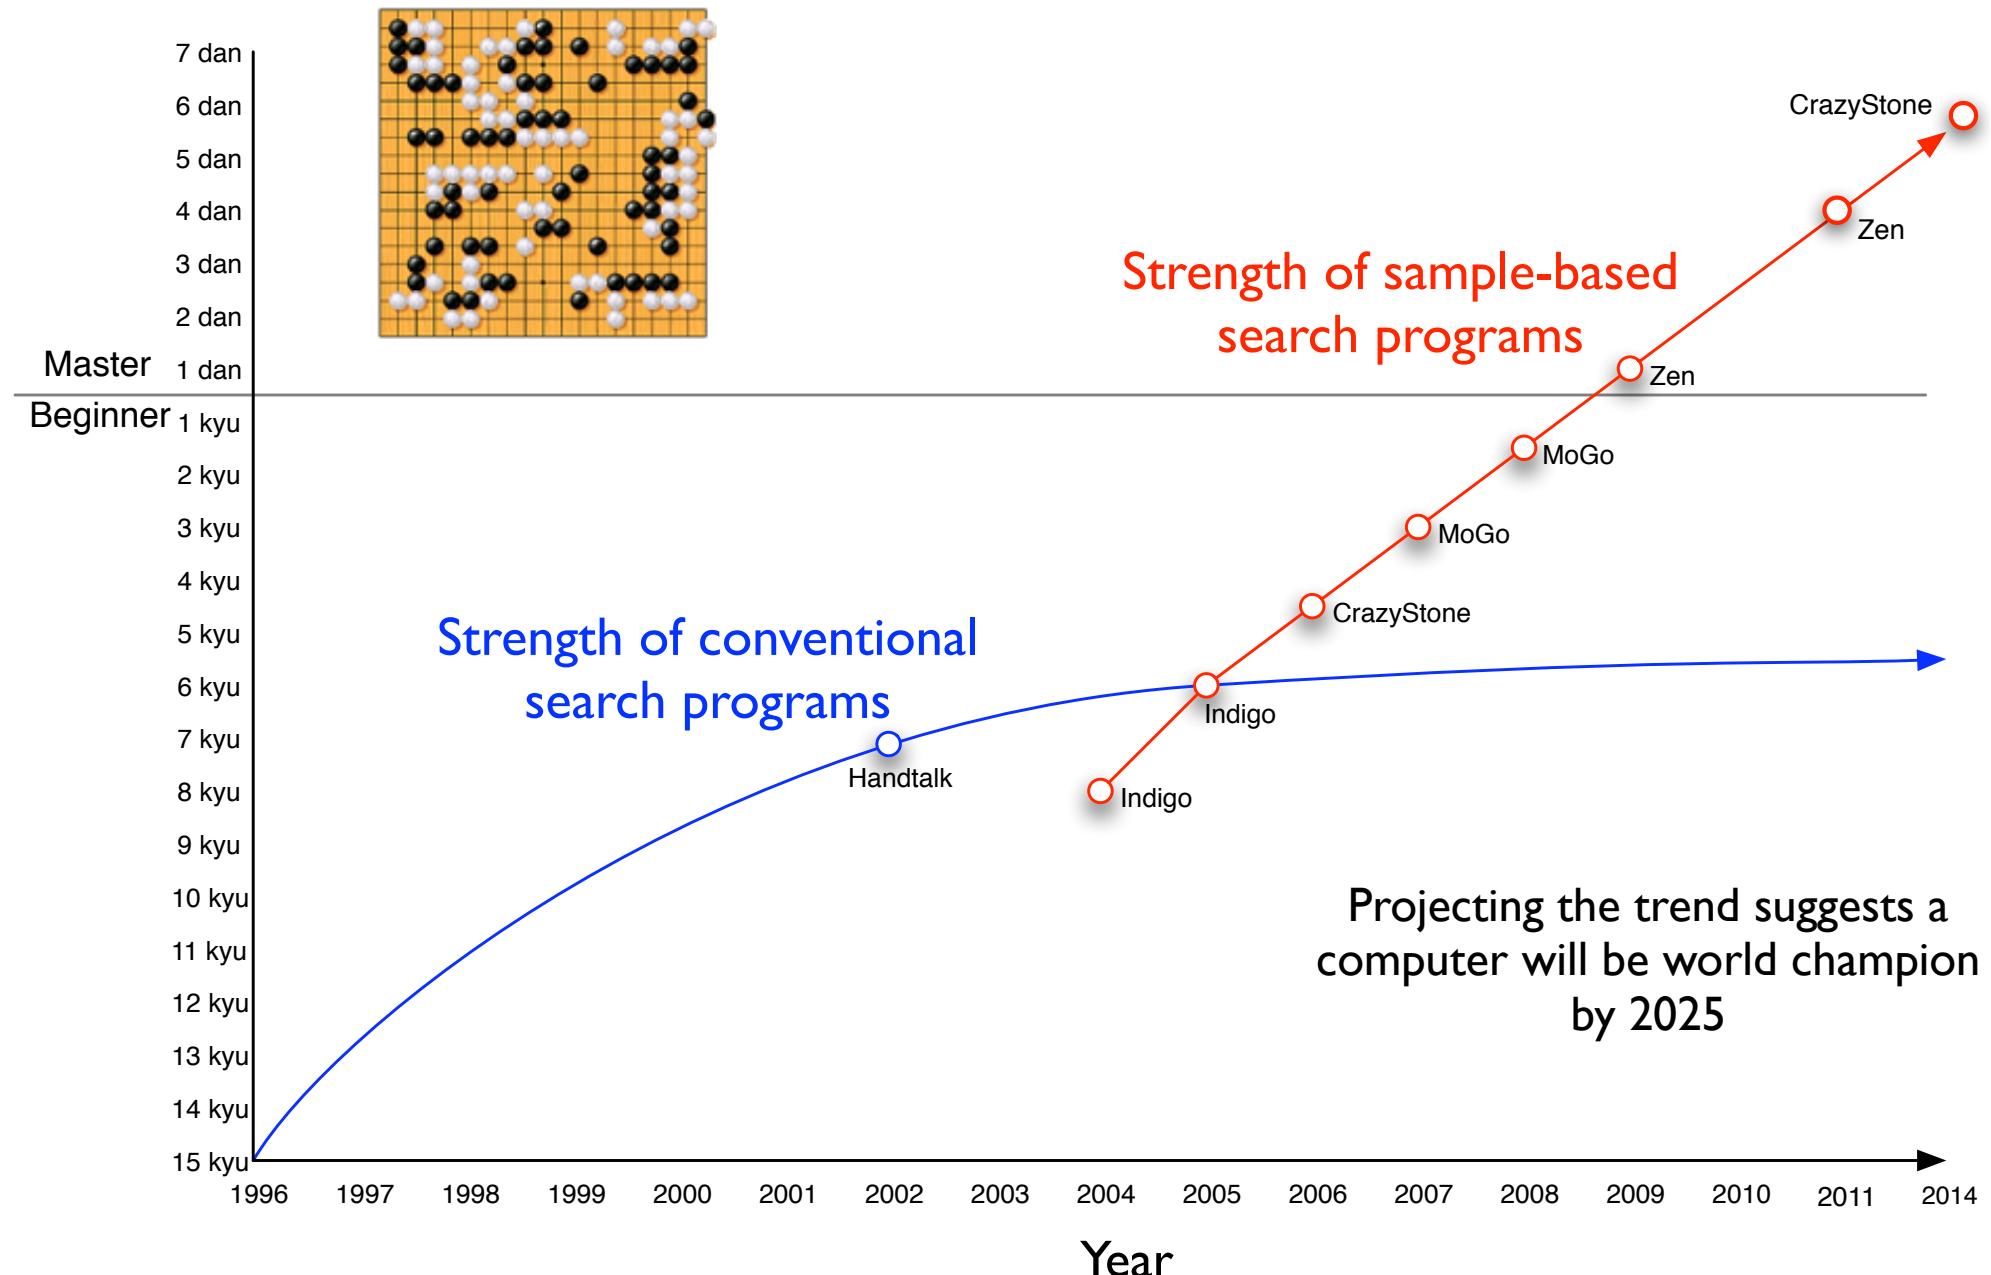
History of AI in Go

- 1997: Super human chess w/ Alpha-Beta + fast computer
- 2005: Computer Go is impossible!
- 2006: Monte-Carlo Tree Search applied to 9x9 Go (bit of learning)
- 2007: Human master level achieved at 9x9 Go (more learning)
- 2008: Human grandmaster level achieved at 9x9 Go (even more learning)
- 2012: Zen program beats former international champion with only 4 stone handicap in 19x19

if 4 handicap moves
are given



Steady, exponential improvement (since MCTS, 2005) in the strength of the best computer Go programs



2015



AlphaGo vs European Champion (Fan Hui 2-Dan)*



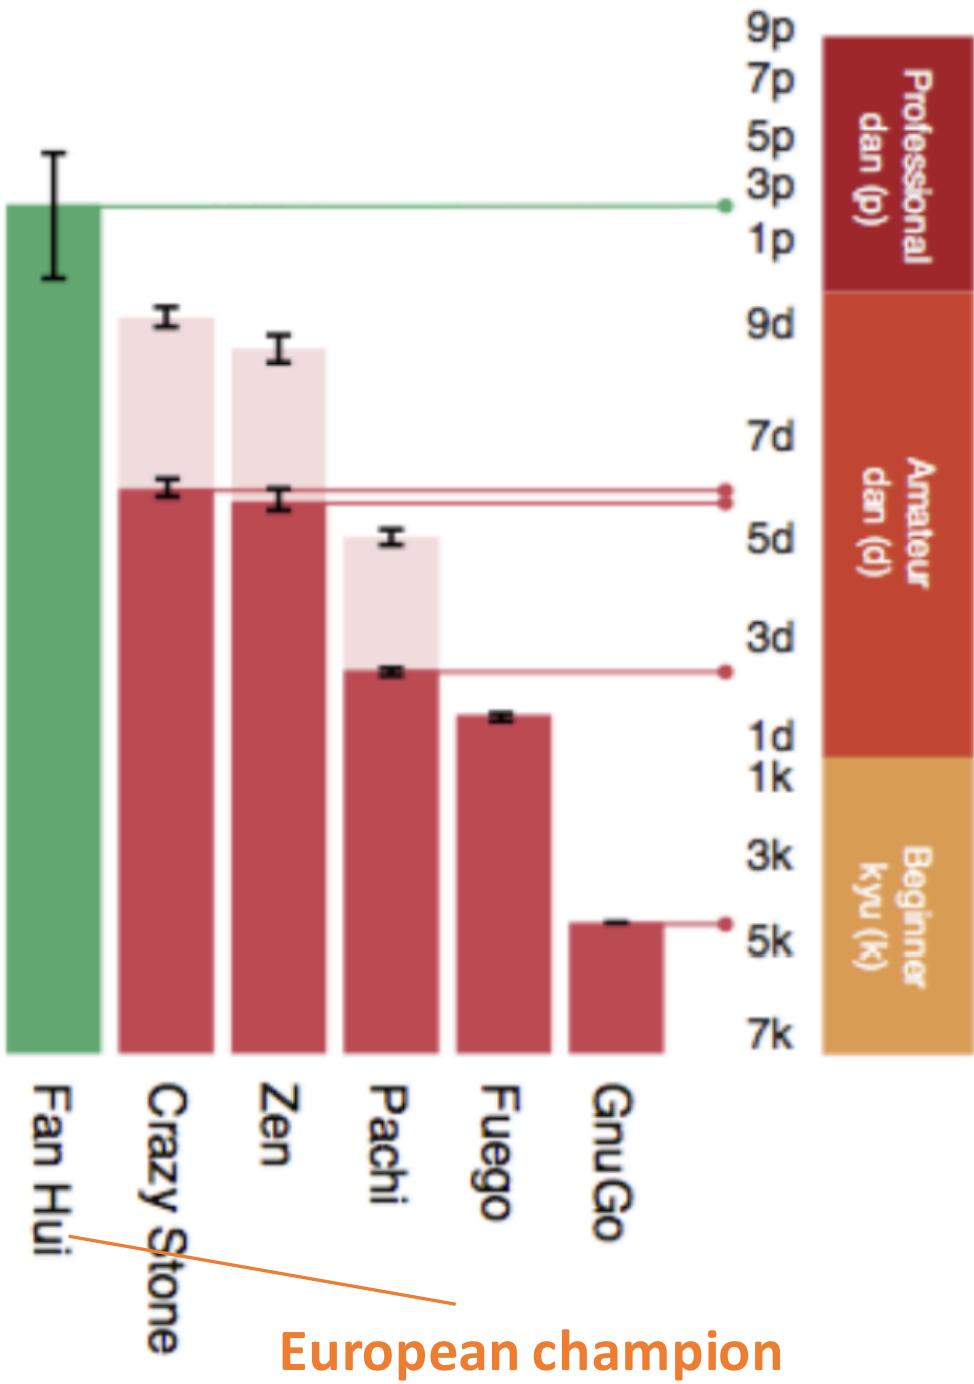
October 5 – 9, 2015

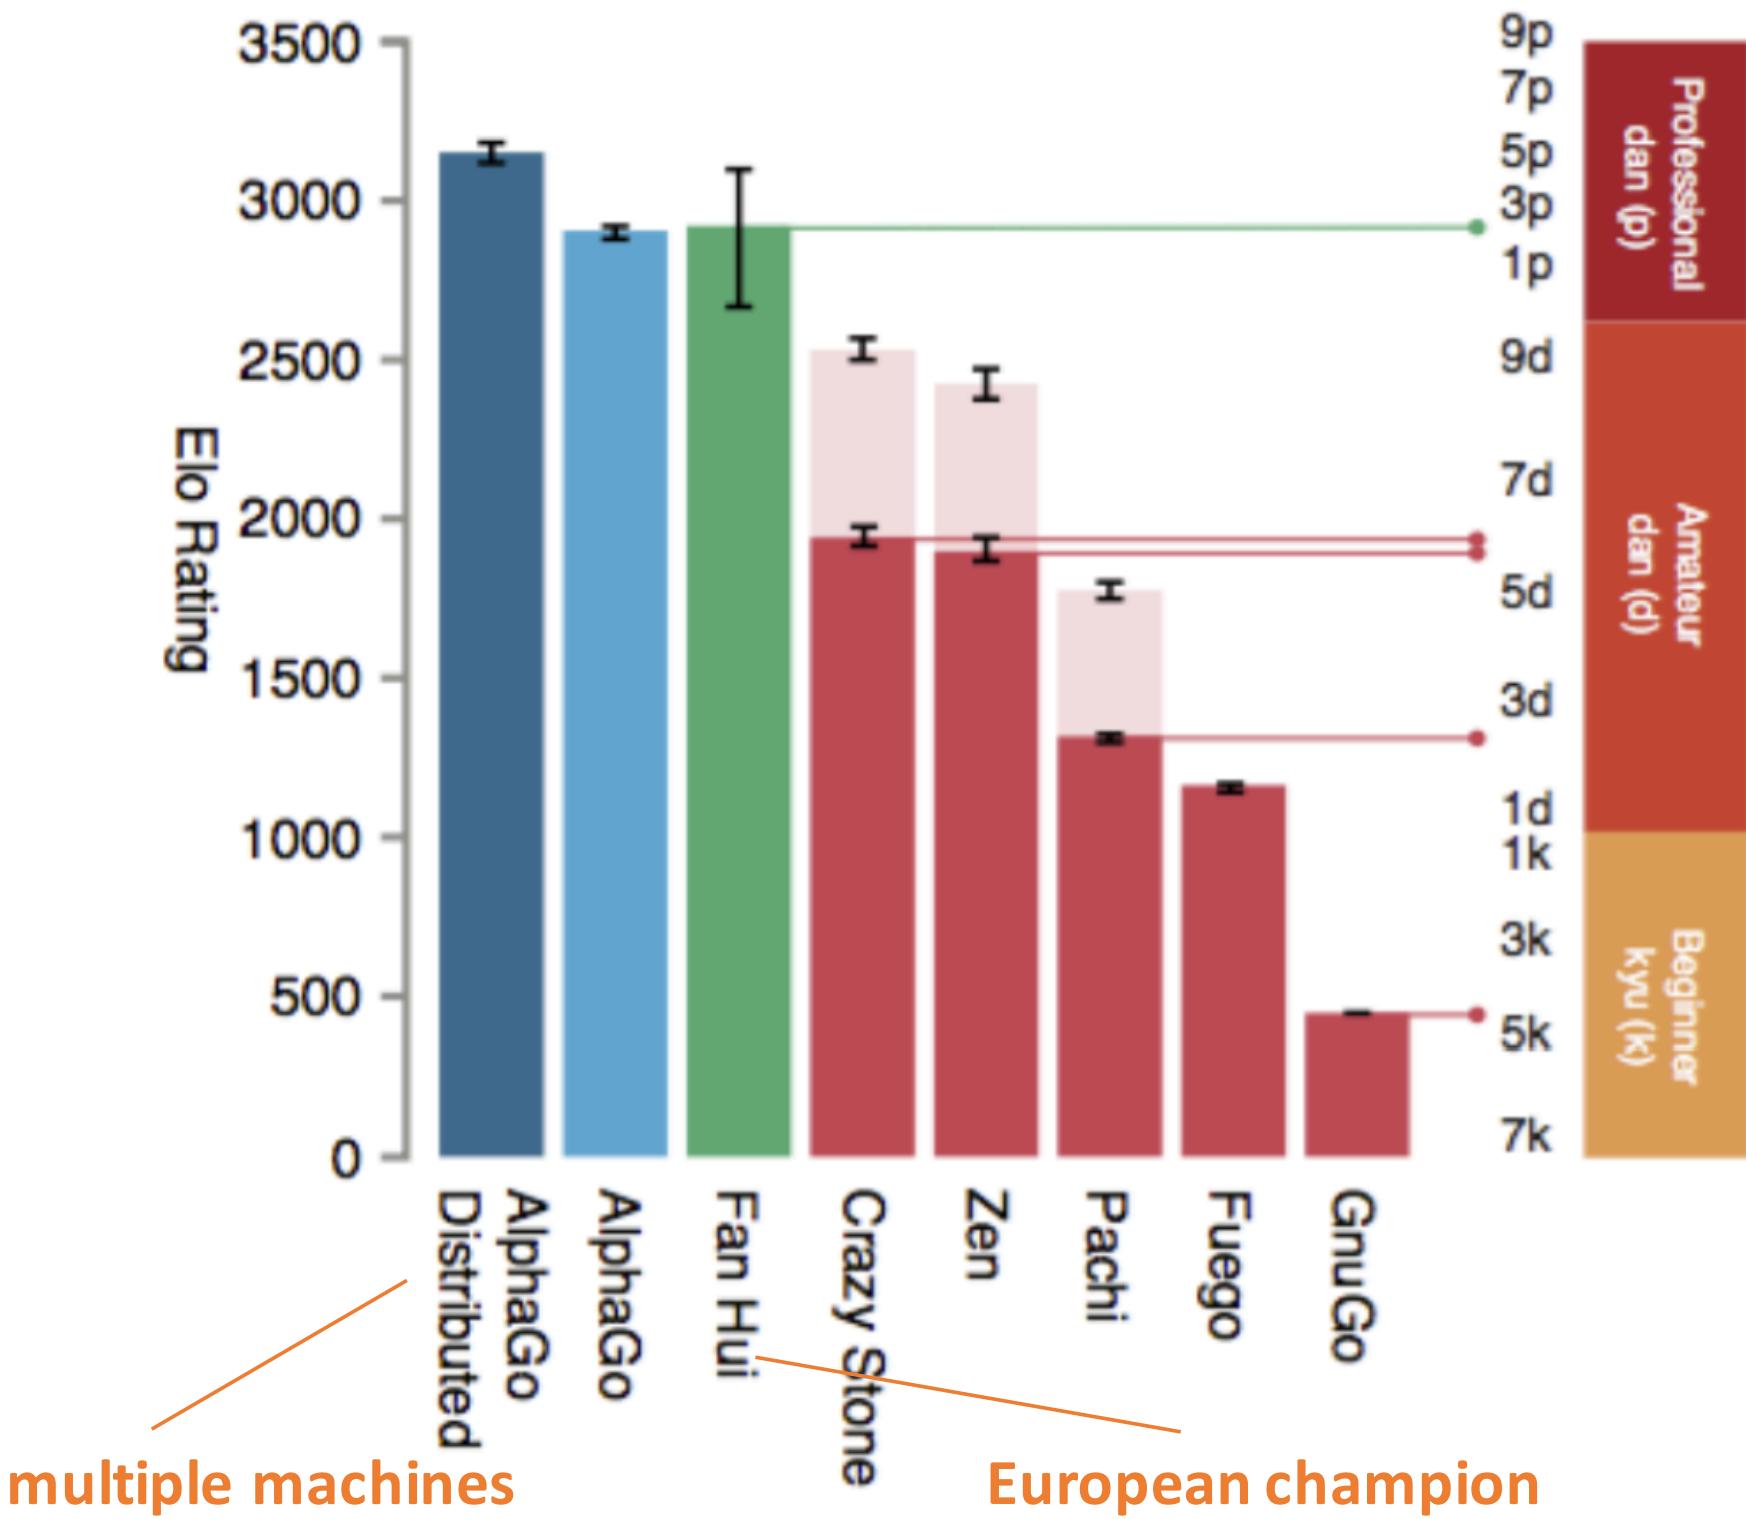
<Official match>

- Time limit: 1 hour
- AlphaGo Wins (5:0)

<https://www.youtube.com/watch?v=53YLZBSS0cc>

<https://www.youtube.com/watch?v=g-dKXOlsf98>





AlphaGo vs World Champion (Lee Sedol 9-Dan)



March 9 – 15, 2016

<Official match>

- Time limit: 2 hours

Venue: Seoul, Four Seasons Hotel

https://www.youtube.com/watch?v=8tqIC8spV_g

THE ULTIMATE GO CHALLENGE

AlphaGo

2015

BORN

1983

~5 dan pro

RANK

9 dan pro

0

TITLES

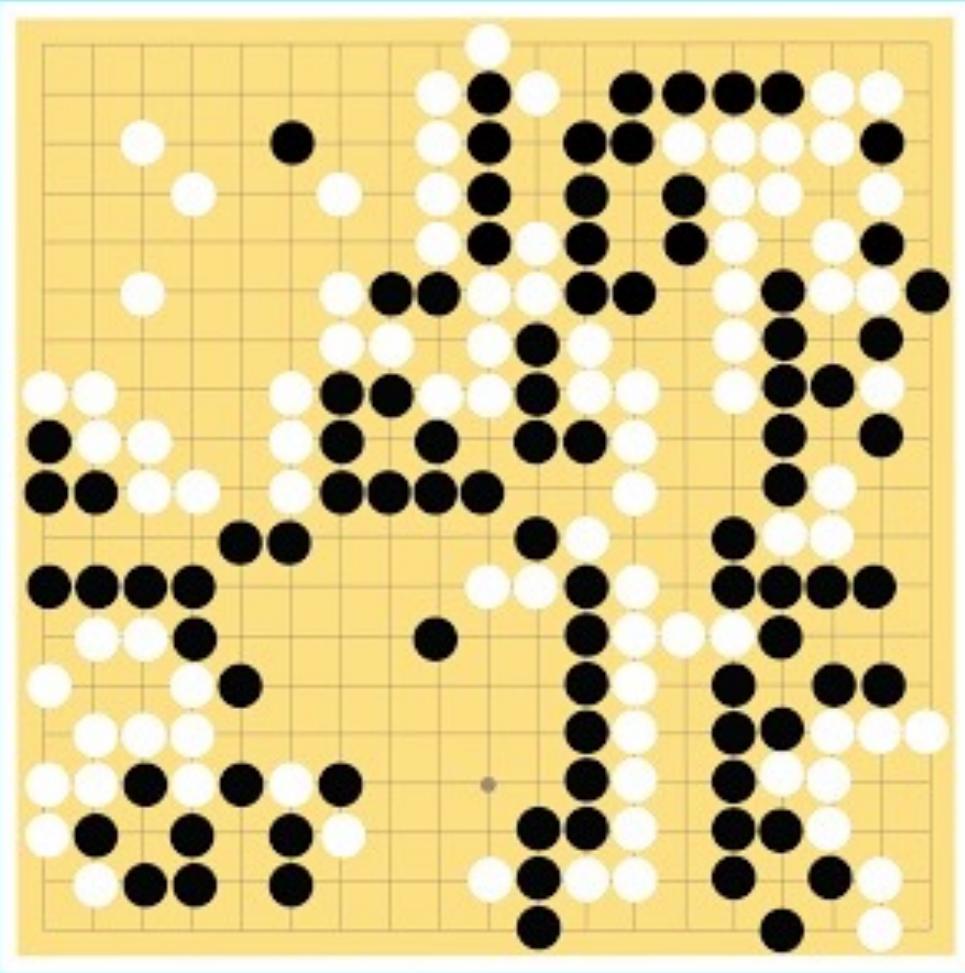
47

Lee Sedol



The Game







THE ULTIMATE GO CHALLENGE

GAME 2 OF 5

10 MARCH 2016



VS
AlphaGo
Won 2 of 5



Lee Sedol
Won 0 of 5

RESULT

B+
Res

NUMBER
OF MOVES

211

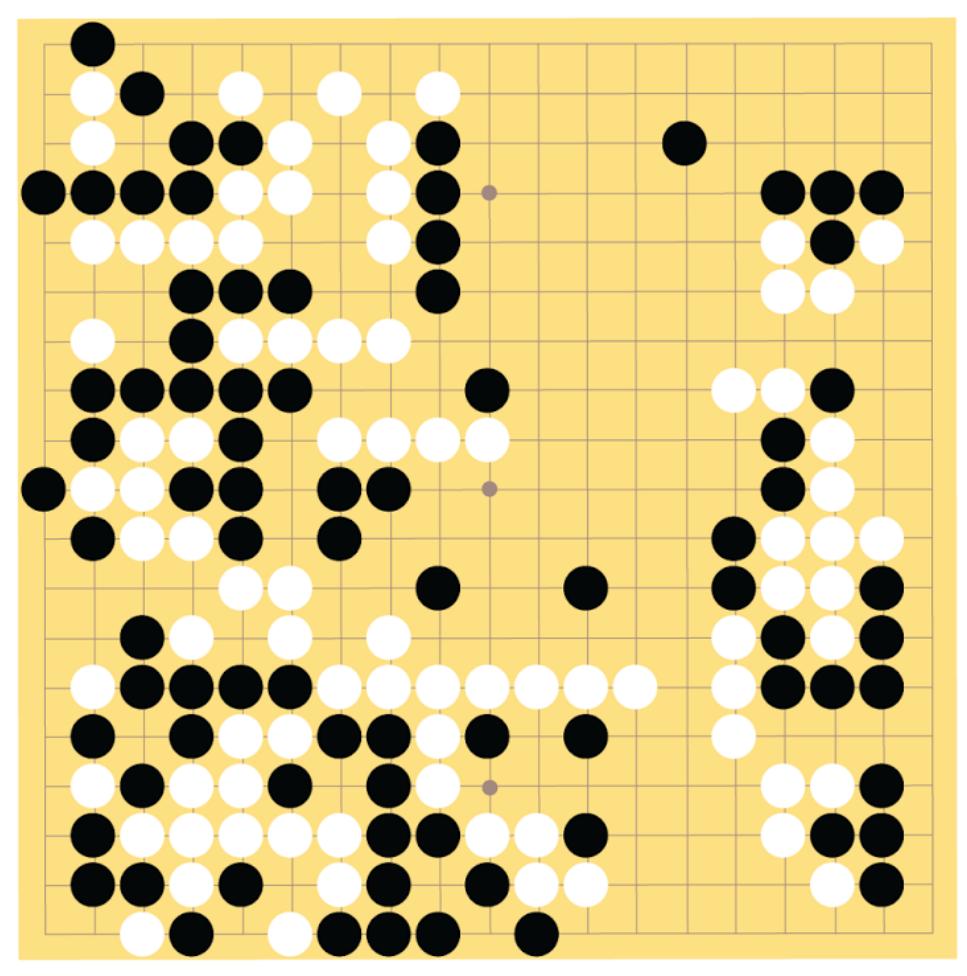
TIME
BLACK

2h+

TIME
WHITE

2h+





THE ULTIMATE GO CHALLENGE

GAME 3 OF 5

12 MARCH 2016



AlphaGo
Won 3 of 5



Lee Sedol
Won 0 of 5

RESULT

**W+
Res**

NUMBER
OF MOVES

176

TIME
WHITE

**1h
51m**

TIME
BLACK

2h+





© Google

THE ULTIMATE GO CHALLENGE

GAME 4 OF 5

13 MARCH 2016



AlphaGo
Won 3 of 5



Lee Sedol
Won 1 of 5



RESULT

**W+
Res**

NUMBER
OF MOVES

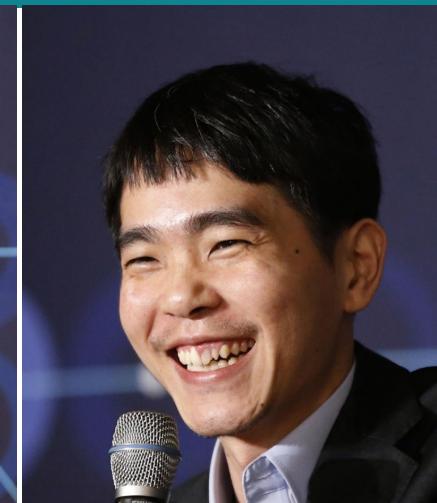
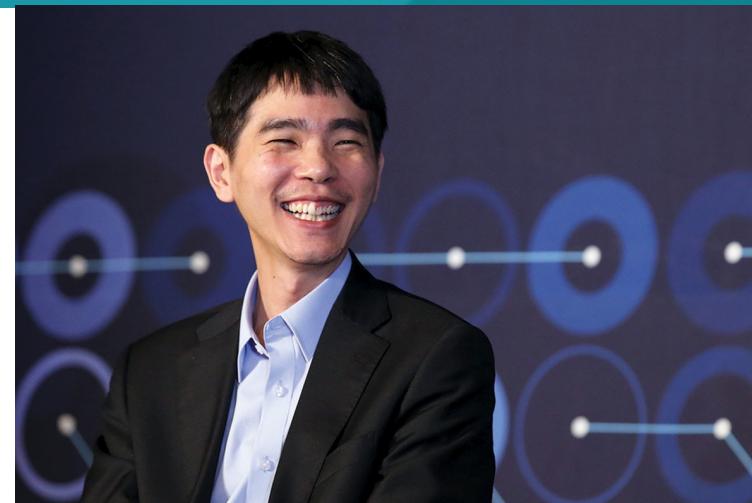
180

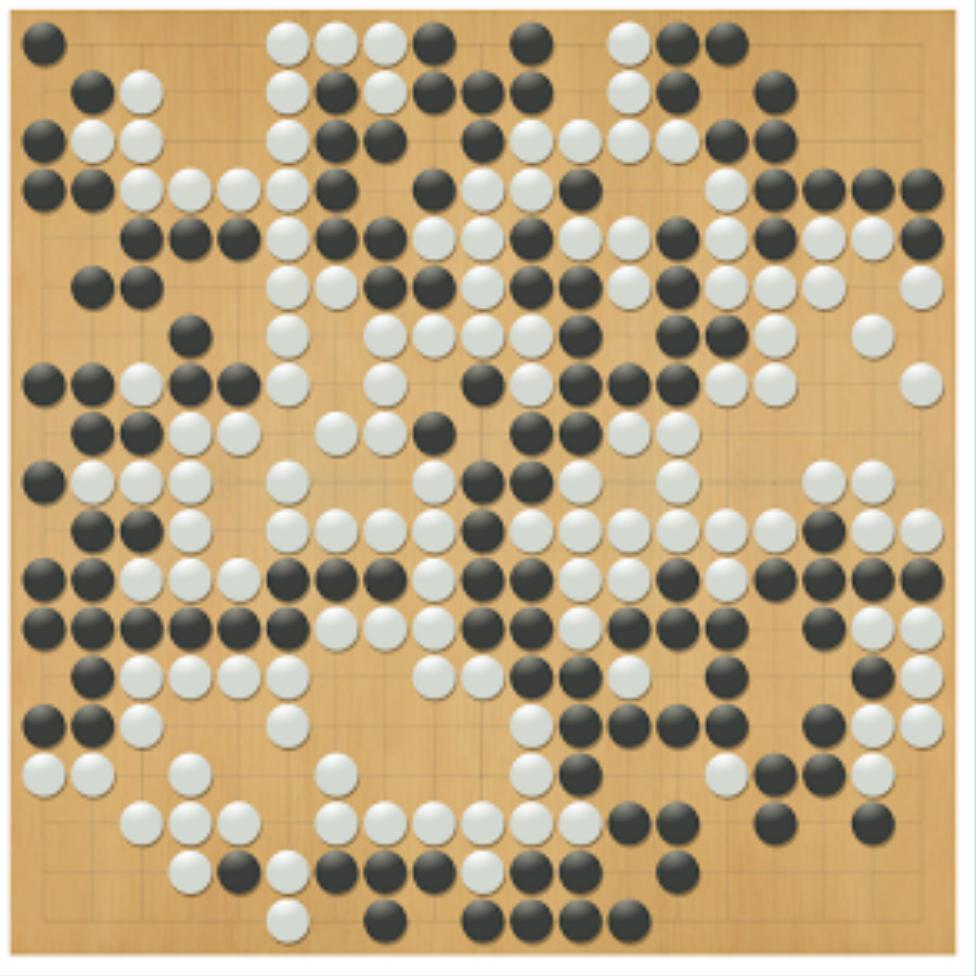
TIME
BLACK

**1h
59m 6s**

TIME
WHITE

2h+





THE ULTIMATE GO CHALLENGE

GAME 5 OF 5

15 MARCH 2016



vs



Lee Sedol
Won 1 of 5



AlphaGo
Won 4 of 5

RESULT

**W+
Res**

NUMBER
OF MOVES

280

TIME
WHITE

2h+

TIME
BLACK

2h+

Final Result

Game	Date	Black	White
1	9 March 2016	Lee Sedol	AlphaGo
2	10 March 2016	AlphaGo	Lee Sedol
3	12 March 2016	Lee Sedol	AlphaGo
4	13 March 2016	AlphaGo	Lee Sedol
5	15 March 2016	Lee Sedol ^[note 1]	AlphaGo

Result:

AlphaGo 4 – 1 Lee Sedol



Go Elo Ranking

For older ratings, check the [History](#) page. There is also a [History of top ladies](#).

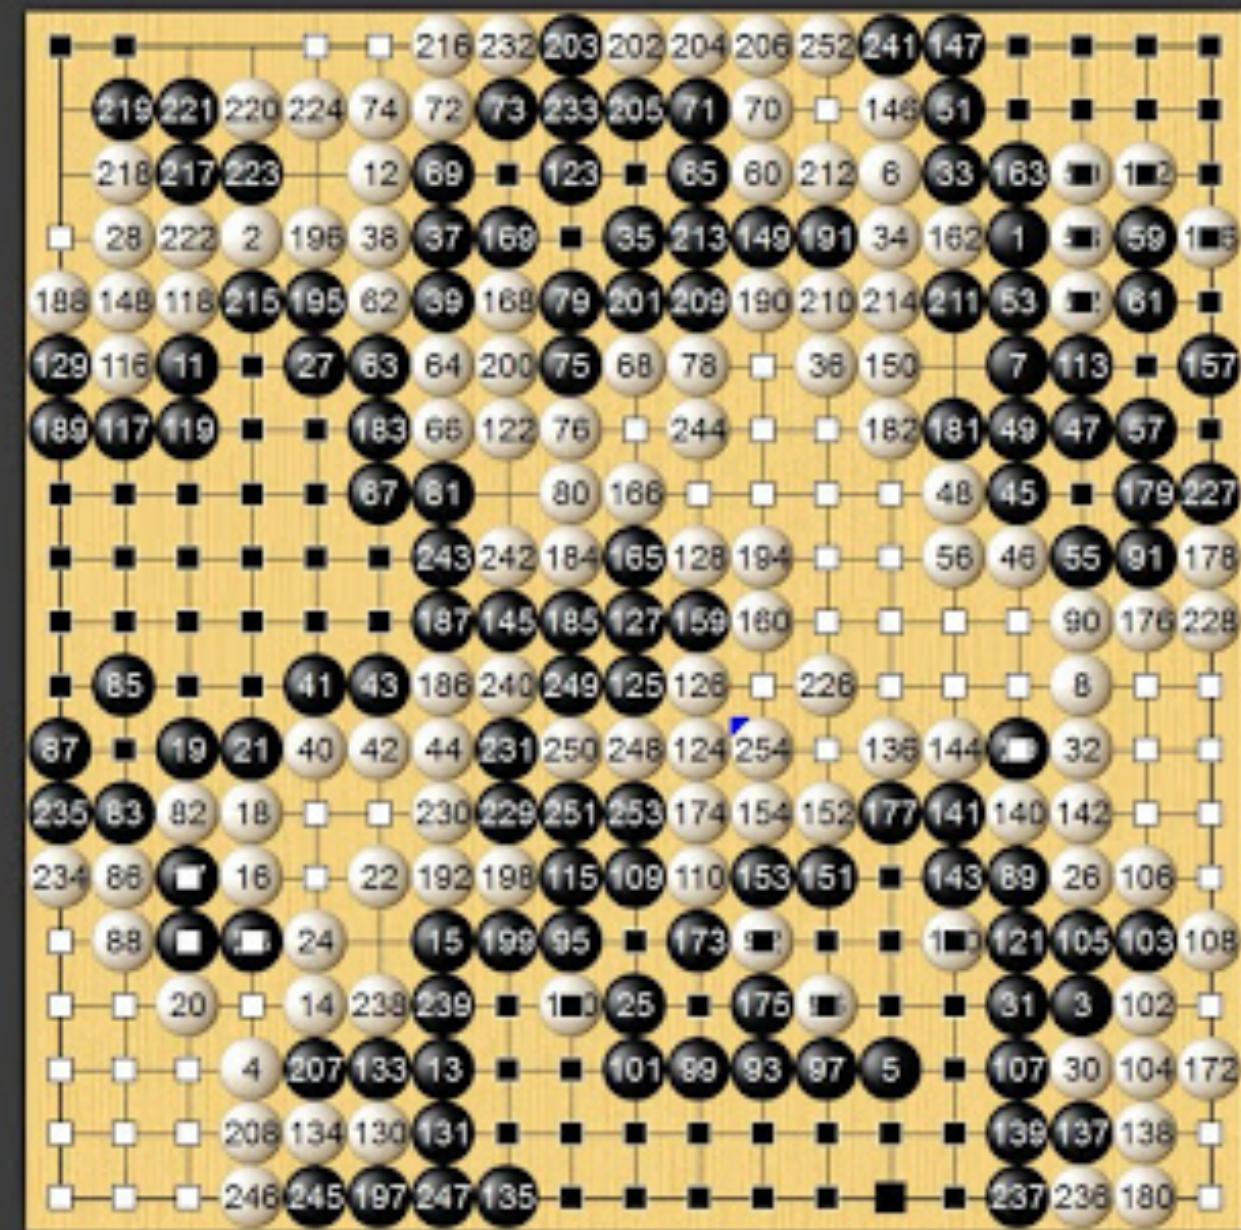


Rank	Name	♂ ♀	Flag	Elo
1	Ke Jie	♂		3622
2	Google DeepMind AlphaGo			3595
3	Park Junghwan	♂		3551
4	Lee Sedol	♂		3532
5	Shi Yue	♂		3529
6	Iyama Yuta	♂		3527
7	Kim Jiseok	♂		3513
8	Zhou Ruiyang	♂		3512
9	Park Yeonghun	♂		3510
10	Mi Yuting	♂		3507
11	Lian Xiao	♂		3504
12	Tang Weixing	♂		3487
13	Tuo Jiaxi	♂		3484



Lee Sedol VS Ke Jie





聂卫平 九段 VS 哥 Buster
6 0 5 1

用户名	08950	胜	负	对局
对	聂卫平	0	1	九段
对	Buster	24	0	高强赛

比赛成绩
*** 1月28日陈诗阳 满级联赛 1111 VS 香港000225
[1111]的竞赛过程开始 ***
通知 'Buster' 向 '聂卫平' 请求清除死子
开始清除死子
*** 陈诗阳棋圣金紫双生啦
*** <强者冠军> - 1111号陈诗阳 [1111]单身晋级 [1111] 95.34% mxfan [9段] 赢得陈诗阳
区域赛开始 ***
通知 'Buster' 取 7日手链 '聂卫平'
通知 华丽棋圣金紫魅影获得勇者无畏奖励100万加
速度赛2个1步人机挑战赛奖励一名 [土豪看我赛] 2
通知 赢者Buster获得活动分数3222分，失去聂卫
平棋圣等级分数1000分，在月人气排行榜中可查看目前的名次
通知 华丽棋圣陈诗阳获得中运得宝石1个 [土豪看
我赛] 1
通知 [土豪看我赛] mxfan [1111] 95.34% 陈诗阳晋级 [1111]
的竞赛过程开始 ***
*** <棋圣赛> - [2483号聂卫平] 聂卫平 [9段] 赛
始讲解 ***
通知Buster [1111]: 聂卫平老师

对局研究 对局研究 派出房间
形势判断 抽奖 例句 100棋



60 wins streak

0 loss

用户信息

Master
他很懒，什么都没有留下

所在地/棋友会 韩国

棋力/胜率指数/内功 82811 / 42223 / 40588

狐币 1,310,370,182 高天上

有效期至 终身

当前状态: 拒绝邀请

9 段 最近战绩 11胜 0负 棋谱

一胜▲ 一胜▲ 15负▼ 一负▼ 升降级规则

000000000000

战绩统计

总战绩	11胜 / 0负
当前战绩	11胜 / 0负
现段位	11胜 / 0负
升降级	11胜 / 0负
友谊	0胜 / 0负
比赛	0胜 / 0负

最高纪录

段位/胜率指数/内功	9段 / 42223 / 40588
音猜段位/狐币	9段 / 1,310,370,182

非职业竞猜 > 10,000+

职业竞猜榜 > 118

每日竞猜榜 > 1,000

竞猜额 0

获奖励 0

排行榜 月人气 ▶ 2

Rank	Name	♂	♀	Flag	Elo
1	Ke Jie	↑			3622
2	Google DeepMind AlphaGo				3595
3	Park Junghwan	↑			3551
4	Lee Sedol	↑			3532
5	Shi Yue	↑			3529
6	Iyama Yuta	↑			3527
7	Kim Jiseok	↑			3513
8	Zhou Ruiyang	↑			3512
9	Park Yeonghun	↑			3510
10	Mi Yuting	↑			3507
11	Lian Xiao	↑			3504
12	Tang Weixing	↑			3487
13	Tuo Jiaxi	↑			3484

2017



KE JIE VS. ALPHAGO
THE ULTIMATE WEIQI MATCH!

99:15:126x14
2147
65
3717
35

2017

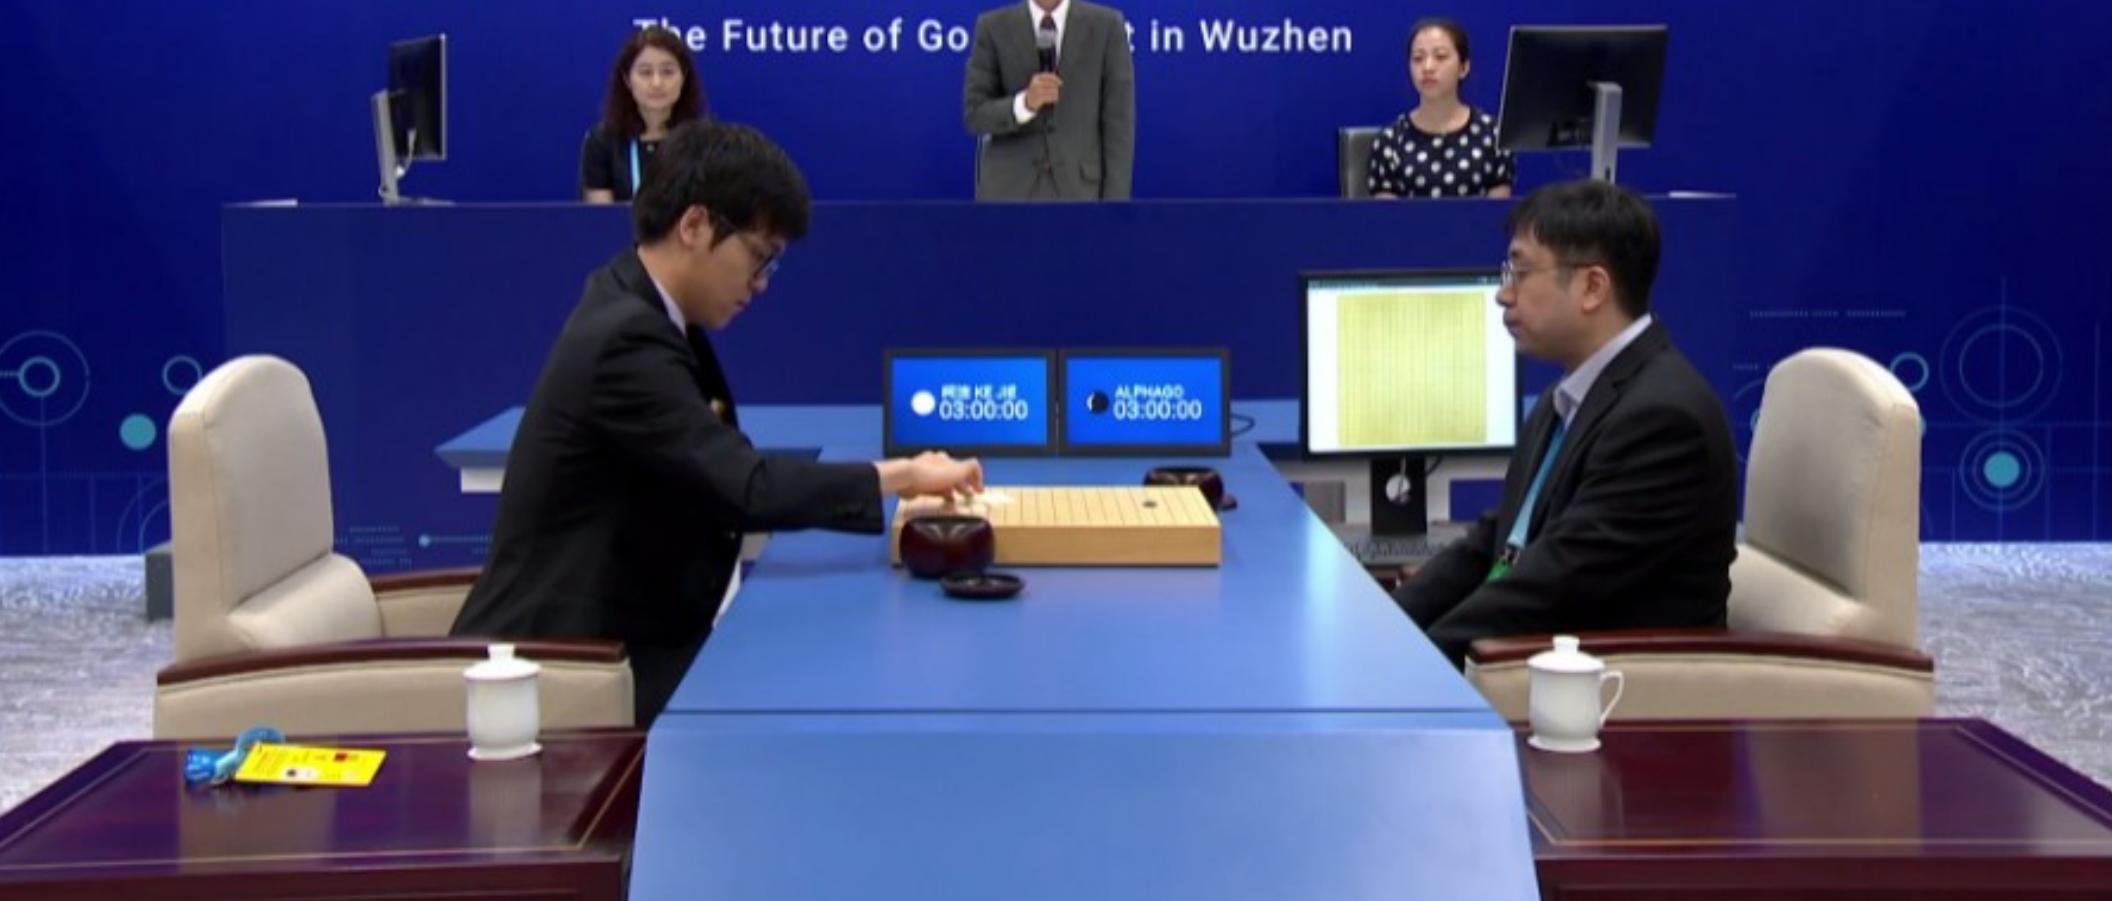
中国围棋协会

Google

浙江省体育局

中国乌镇围棋峰会

The Future of Go Meet in Wuzhen





柯洁 KE JIE
00:46:57

ALPHAGO
02:08:18

AlphaGo vs Ke Jie

3 : 0

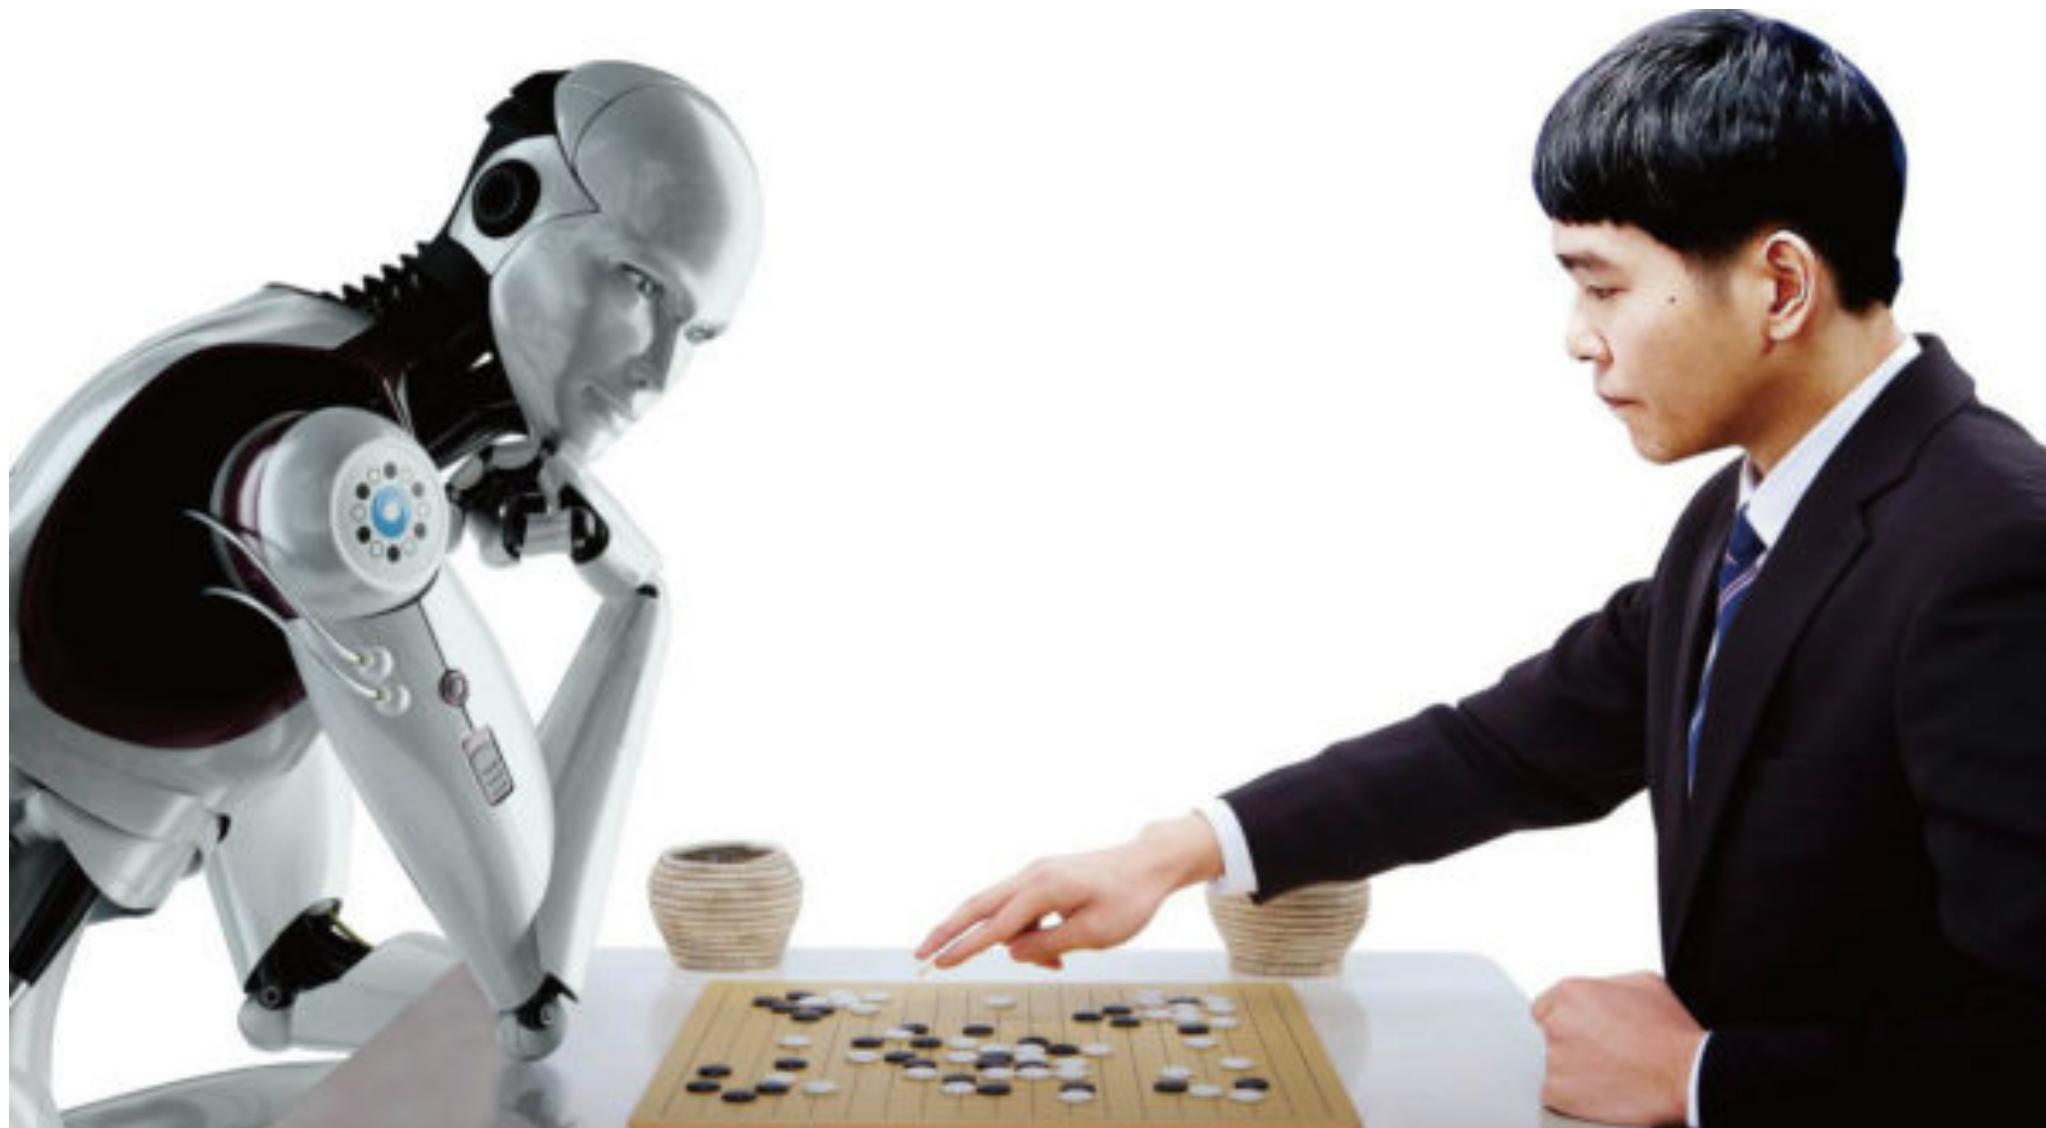


After the game (2017)

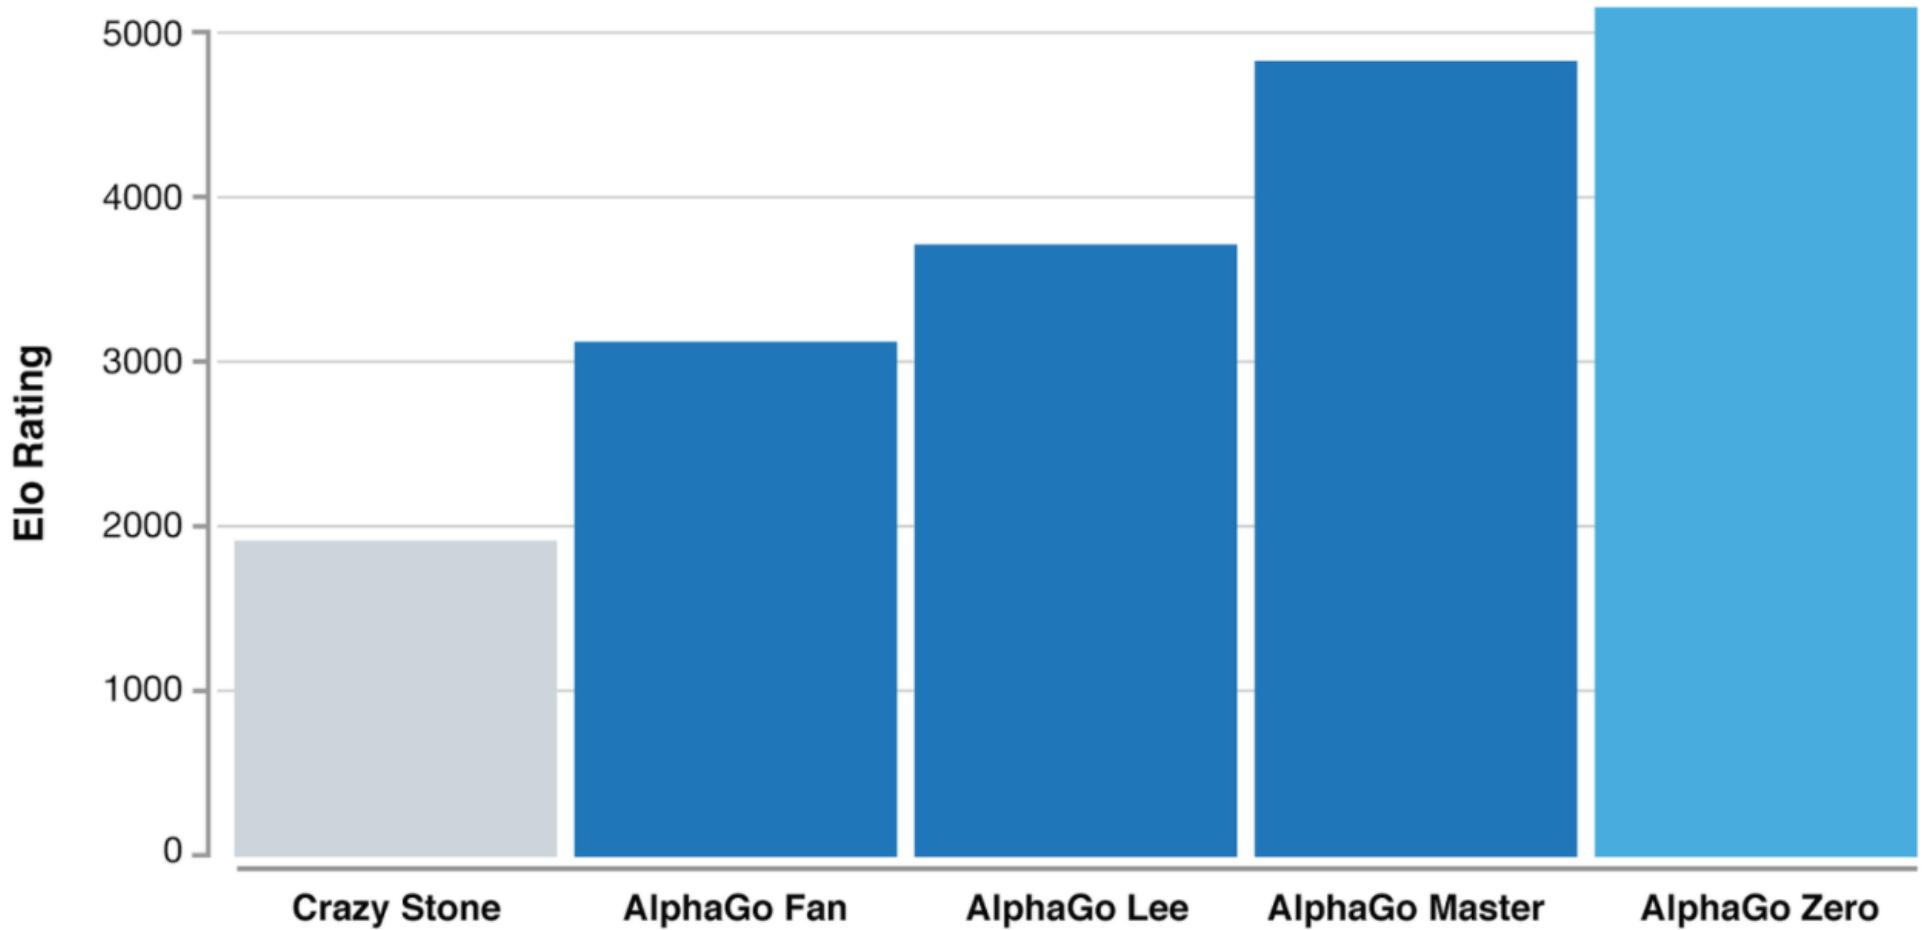
“AlphaGo is almost like the God of Go”



Lee Se-dol won 1 game



Different versions of AlphaGo

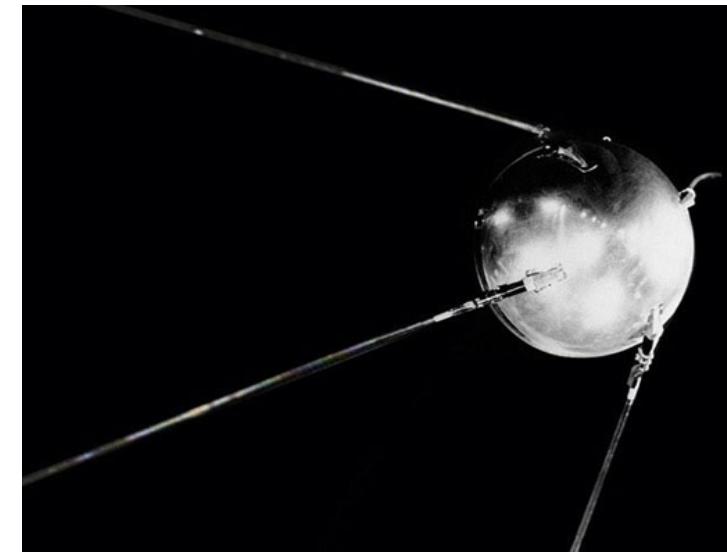


Elo ratings - a measure of the relative skill levels of players in competitive games such as Go - show how AlphaGo has become progressively stronger during its development



sputnik
moment

1957



Saturday Evening, October 5, 1957

(UP)—Means United Press

Price: Five Cents

Russians Win Race To Launch Earth Satellite

Man On Threshold Of Space Travel

By DANIEL F. GILMORE

United Press Staff Correspondent

LONDON (UP)—The pulsating radio "beep" of the first manmade earth satellite signalled today to the world that man had crossed the threshold into the age of travel through space.

The Soviet Union announced it had won the race into space by launching an earth satellite Friday, a 184-pound, 22-inch globe now orbiting the earth at 18,000 miles an hour, 560 miles up.

Millions of persons throughout the world heard the "beep...beep...beep..." rebroadcast today by local stations and realized that man had taken his first faltering steps into the new era.

Launching of the satellite was a tremendous victory for science. It was a more tremendous victory for Soviet propaganda to be able to trumpet to the world the Russians were the first to break through the frontiers of space.

Bolsters ICBM Claims
Russian claims to

How To Spot Satellite

By UNITED PRESS

Here's how to look for the Russian earth satellite which will be whizzing through the sky at 18,000 miles an hour.

The best time to spot it is at dawn or dusk when the sky is semi-dark. There is a chance that it could be seen if it travels across the face of the moon at night.

The best instruments to use are ordinary binoculars or telescopes. Powerful telescopes won't pick it up because of their narrow fields.

Through optical instruments, the satellite will look like the faintest star which can be seen with the naked eye.

Keep a sharp eye out. The satellite travels so fast it may appear on the horizon for only seconds and chances of spotting it have been estimated at one in a hundred.

U.S. May Speed Up Satellite Program

By JOSEPH L. MYLES

United Press Staff Correspondent

WASHINGTON (UP)—American scientists, caught flatfooted Russia's epic launching of the man-made moon, indicated the United States may speed its own earth satellite program. Leaders of the U.S. satellite program also said that it appears Russia rocketed its heavy pound satellite into a globe circling orbit with a rocket "to" an intercontinental ballistic missile.

That could mean Russia only has beaten this country frontiers of space, but also it has been called the "useful weapon" for modern day war. ICBM. This country has tested a successful ICBM. American diplomats claim Russia had scored a notable

— WEATHER —

WEST VIRGINIA—Partly cloudy with highest in the 60s today and Sunday. Lowest tonight 50 west and 40 east portions.

VIRGINIA—Fair with lowest 45 to 50 west and north and 50 to 55 southeast portions tonight, Sunday mostly sunny and a little warmer. Tides on the coast and lower bay will run a foot or two above normal.

Eric Maki

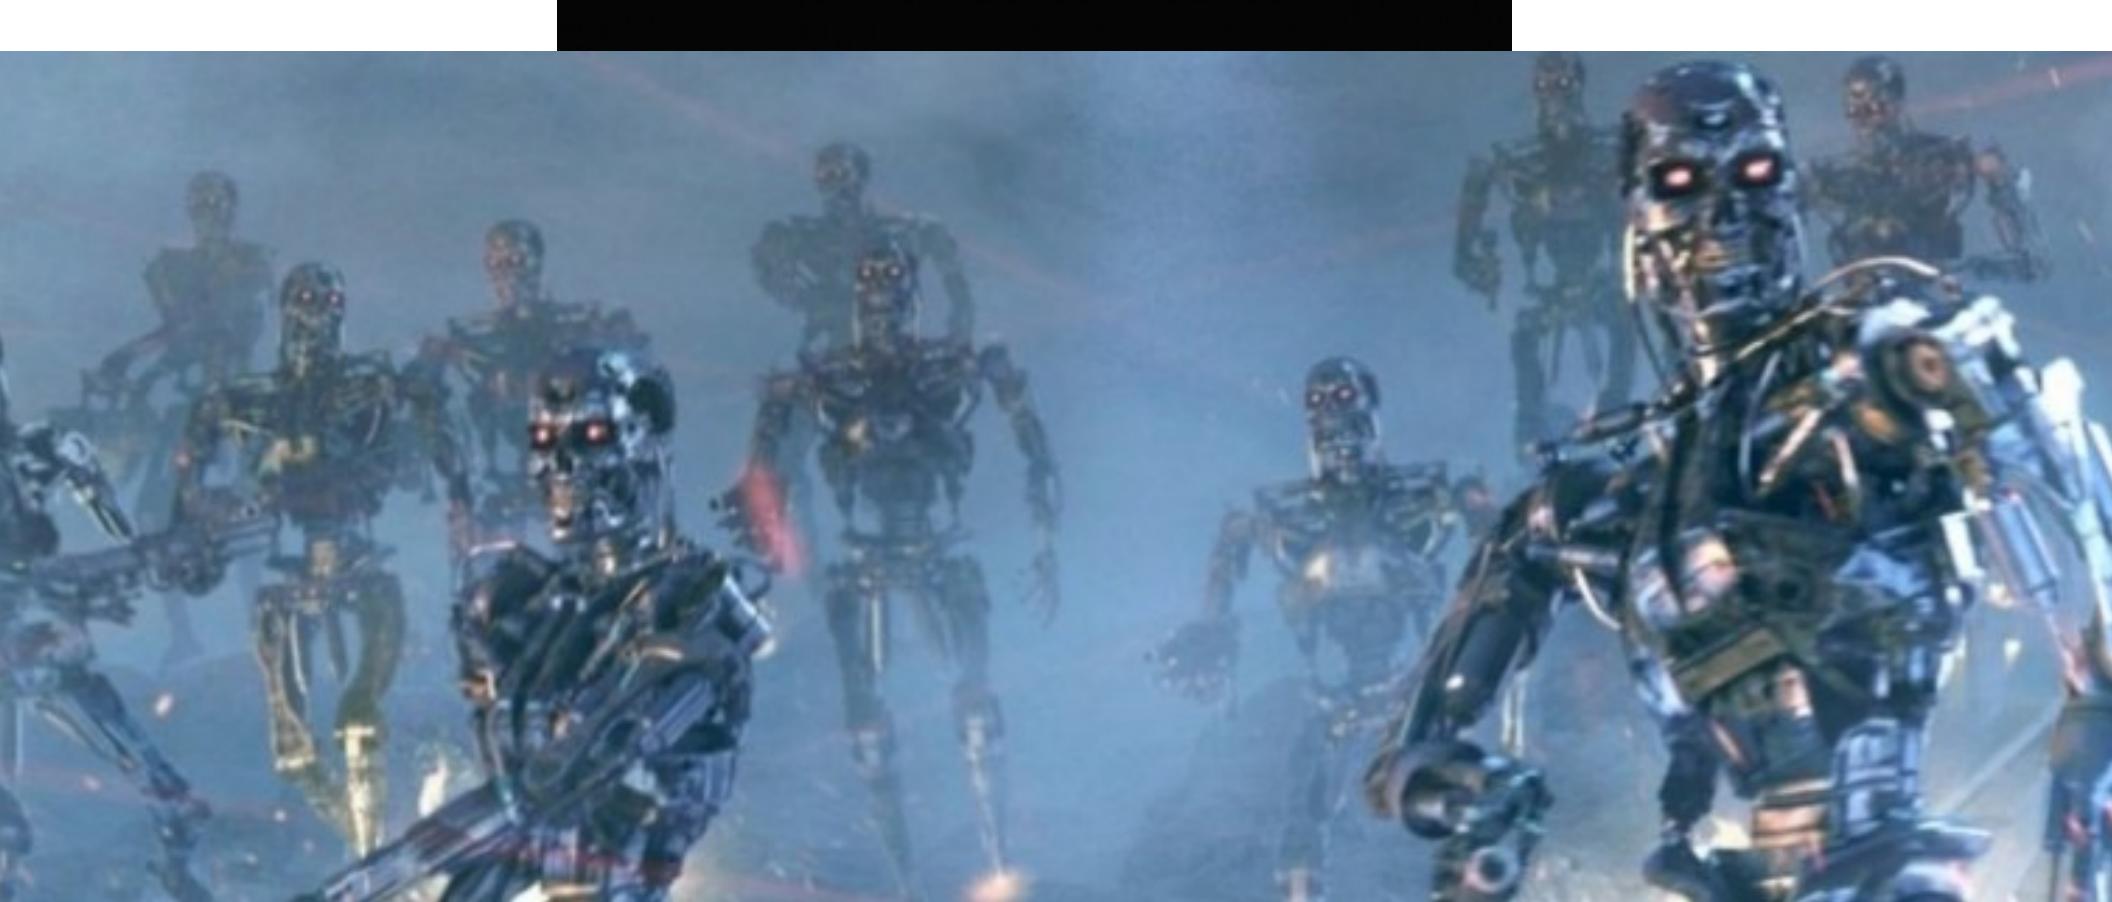
Go Community



AlphaGo is the Skynet?



AlphaGo is the Skynet?



1994

Data

Problem



Jerry Yang and David Filo

Webpages (URLs)

Jerry and David's Guide
to the World Wide Web

Searching Webpage

Yahoo

[Y Top](#) | [Up](#) | [Search](#) | [mail](#) | [Add](#) | [Help](#)

- [Art \(619\) NEW](#)
- [Business \(8546\) NEW](#)
- [Computers \(3266\) NEW](#)
- [Economy \(898\) NEW](#)
- [Education \(1839\) NEW](#)
- [Entertainment \(8814\) NEW](#)
- [Environment and Nature \(268\) NEW](#)
- [Events \(64\) NEW](#)
- [Government \(1226\) NEW](#)
- [Health \(548\) NEW](#)
- [Humanities \(226\) NEW](#)
- [Law \(221\) NEW](#)
- [News \(301\) NEW](#)
- [Politics \(184\) NEW](#)
- [Reference \(495\) NEW](#)
- [Regional Information \(4597\) NEW](#)
- [Science \(3289\) NEW](#)
- [Social Science \(115\) NEW](#)
- [Society and Culture \(933\) NEW](#)

There are currently 31897 entries in the Yahoo database

Some Other General Internet Directories:

[[WWW Virtual Library](#) * [EINet Galaxy](#) * [University of Michigan Clearinghouse](#)]
[[GNN - Whole Internet Catalog](#) * [Planet Earth](#) * [Yanoff's Connections](#)]

1997

Data

Webpages (URLs)

Problem

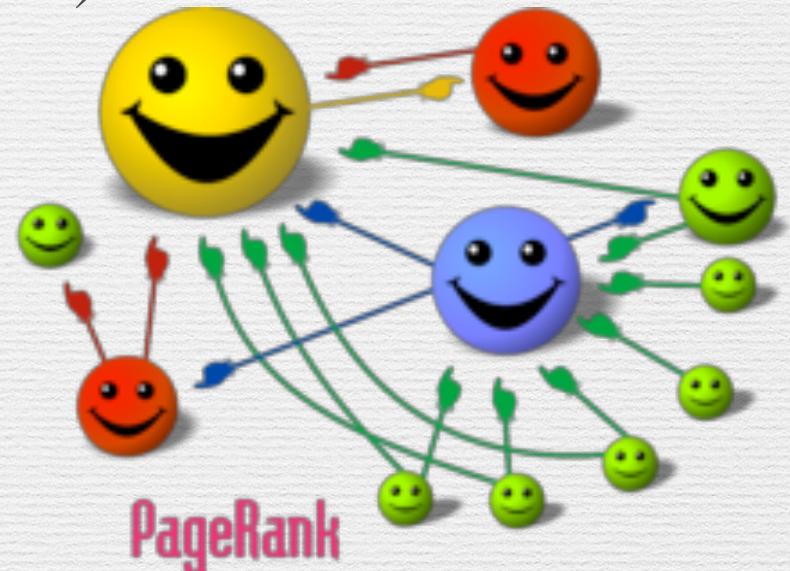
Find URLs

Idea: Algorithm

PageRank



Larry Page and Sergey Brin



Google!

BETA

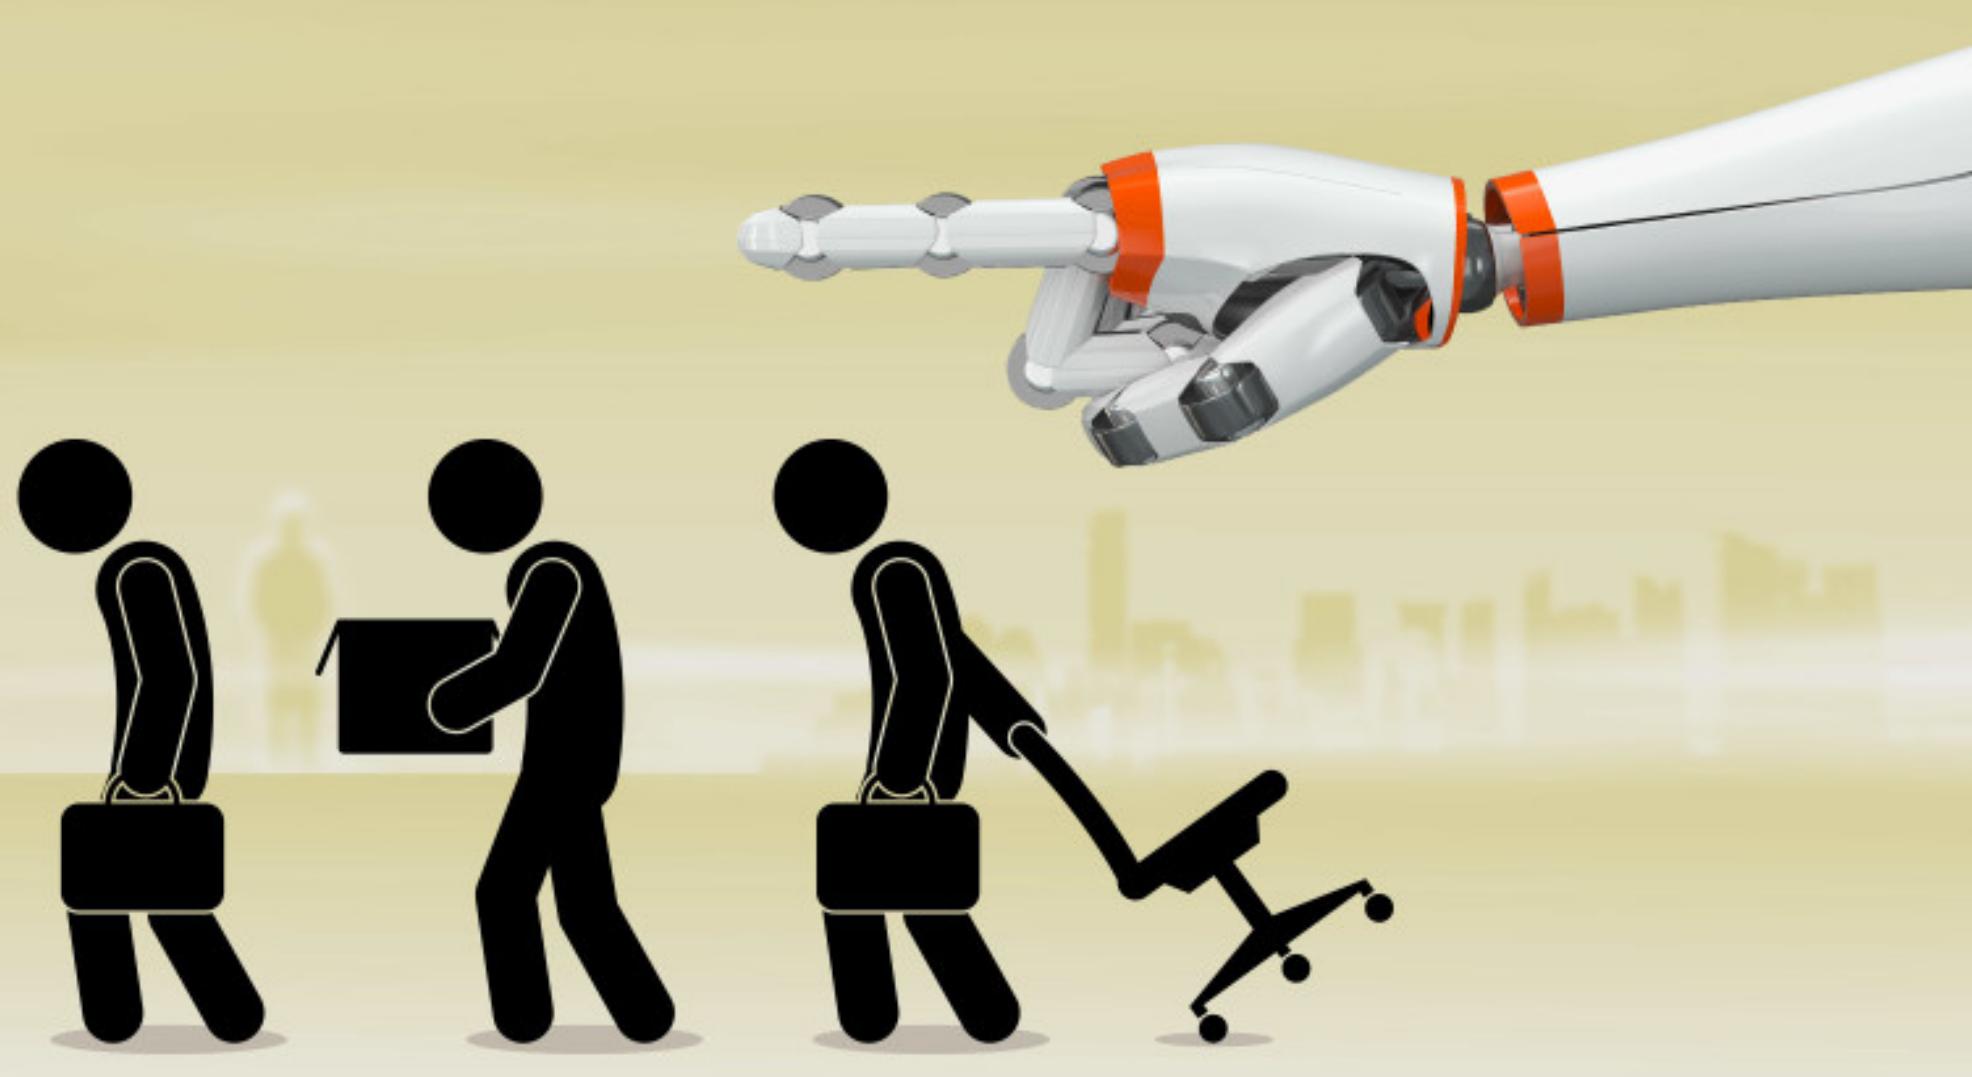
Search the web using Google!

Special Searches
[Stanford Search](#)
[Linux Search](#)

Help!
[About Google](#)
[Company Info](#)
[Google! Logos](#)

Get Google! updates monthly:

Copyright ©1998 Google Inc.





DEEPMIND TECHNOLOGIES LIMITED

Statement of Financial Position As at 31 December 2017

	Notes	2017
		£
Fixed assets		
Tangible assets	9	1,133,220
Intangible assets	10	1,157,140
Financial assets	11	765,095
		<hr/>
		3,055,455
Current assets		
Debtors: amounts falling due within one year	12	99,173,584
Cash and cash equivalents		<hr/>
		975,277
		<hr/>
		100,148,861
Current Liabilities		
Creditors: amounts falling due within one year	13	(123,939,156)
		<hr/>
		(23,790,295)
		<hr/>
Total assets less current liabilities		(20,734,840)
Creditors: amounts falling due after one year	14	(435,318,927)
		<hr/>
Net Liabilities		(456,053,767)
		<hr/>
Capital and reserves		
Called up share capital	15	1,641
Share premium		41,490,727
Retained loss		(497,546,135)
		<hr/>
Total shareholders' deficit		(456,053,767)

How about Other Games?



FINAL 16 PLAYERS

INVITED	INVITED	SOUTH KOREA QUALIFIERS	SOUTH KOREA QUALIFIERS	SOUTH KOREA QUALIFIERS	EUROPEAN QUALIFIERS	EUROPEAN QUALIFIERS	AUSTRALIAN QUALIFIERS
Flash South Korea	PartinG South Korea	Solar South Korea	Rogue South Korea	Journey South Korea	First South Korea	YoDa South Korea	iaguz Australia
TAIWAN QUALIFIERS	TAIWAN QUALIFIERS	AMERICAN QUALIFIERS	AMERICAN QUALIFIERS	#2 REPLACEMENT	APAC QUALIFIERS	SEA QUALIFIERS	SEA QUALIFIERS
Sen Taiwan	Has Taiwan	HuK Canada	viOLEt South Korea	Ian Taiwan	Cloudy China	EnDerr Philippines	Blysk Singapore

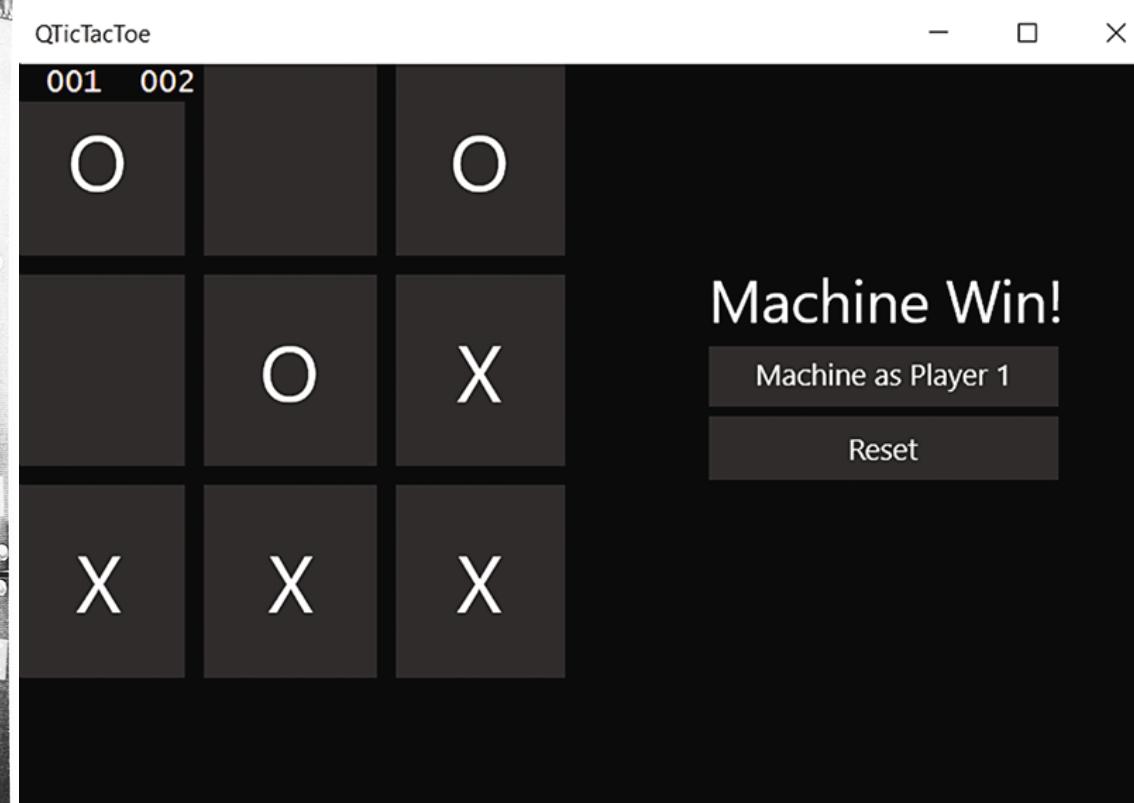
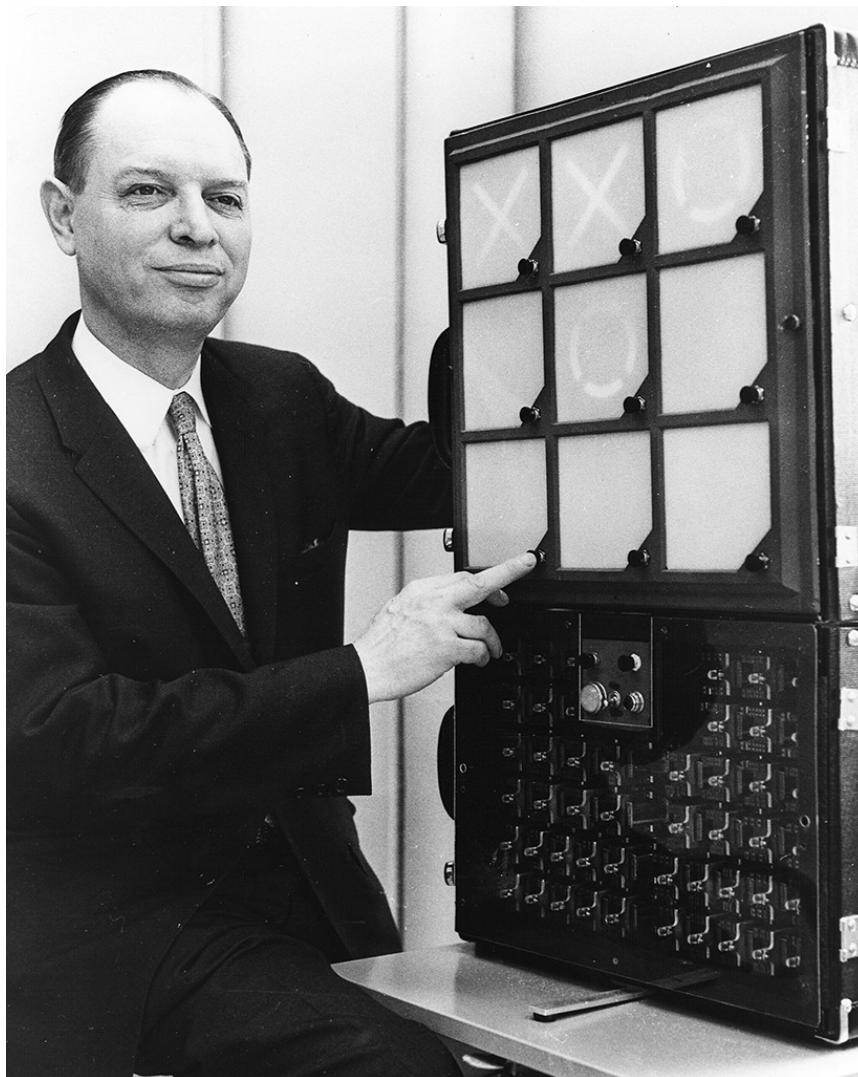
Aug, 2017



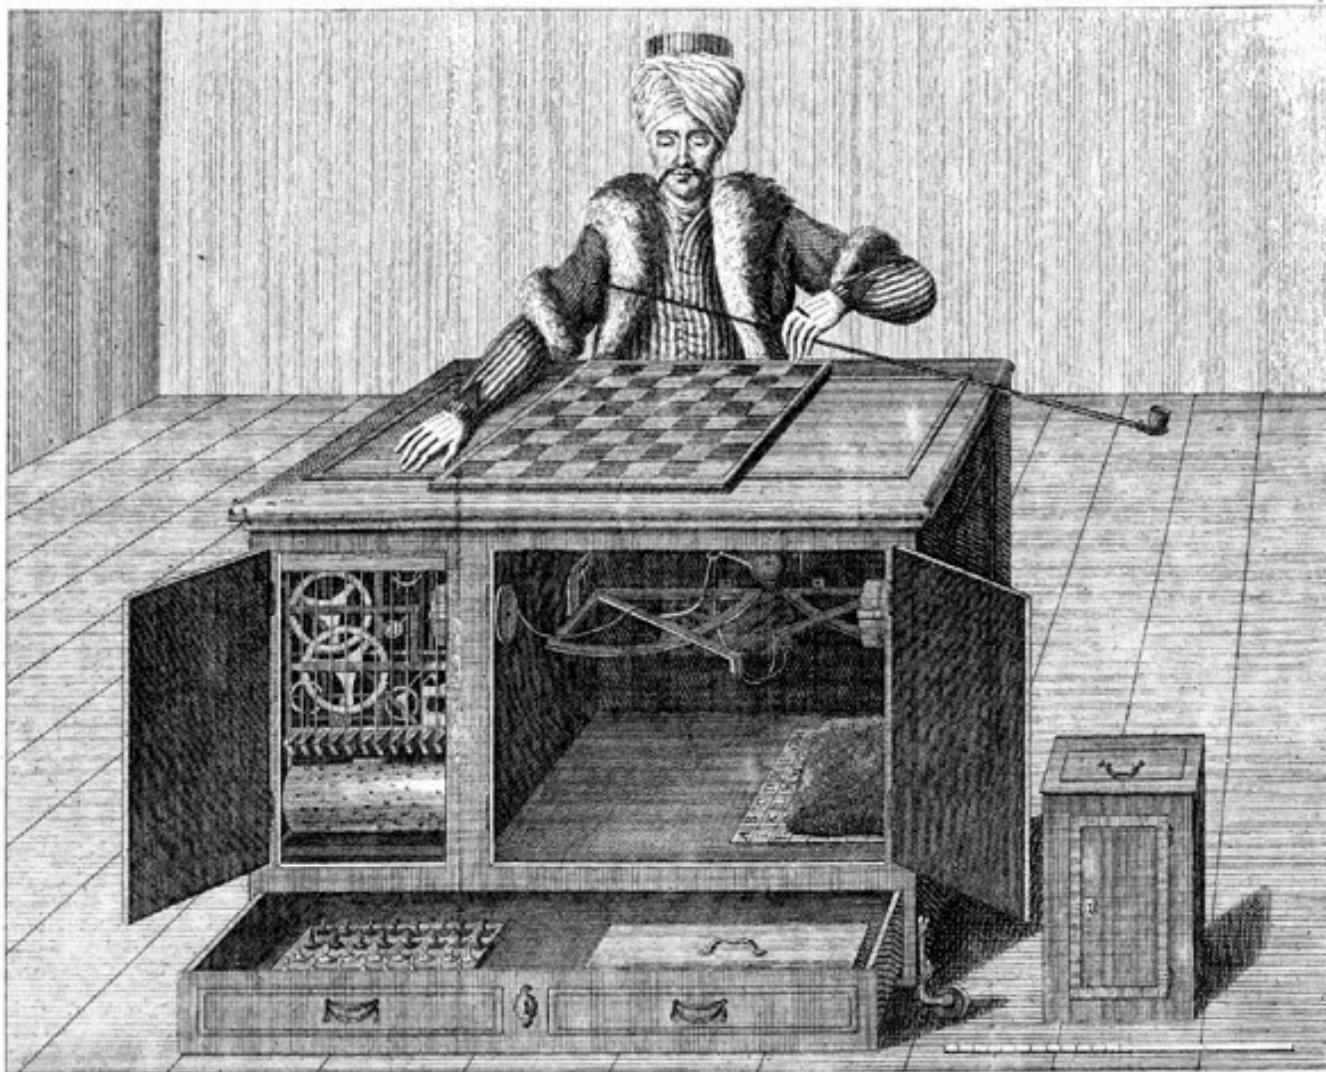
<https://youtu.be/7U4-wvhgx0w?t=2m10s>

Quiz

Tic Tac Toe



Chess



W. de Kempenaer del.

Chr. a Mechel excud. Basileæ.

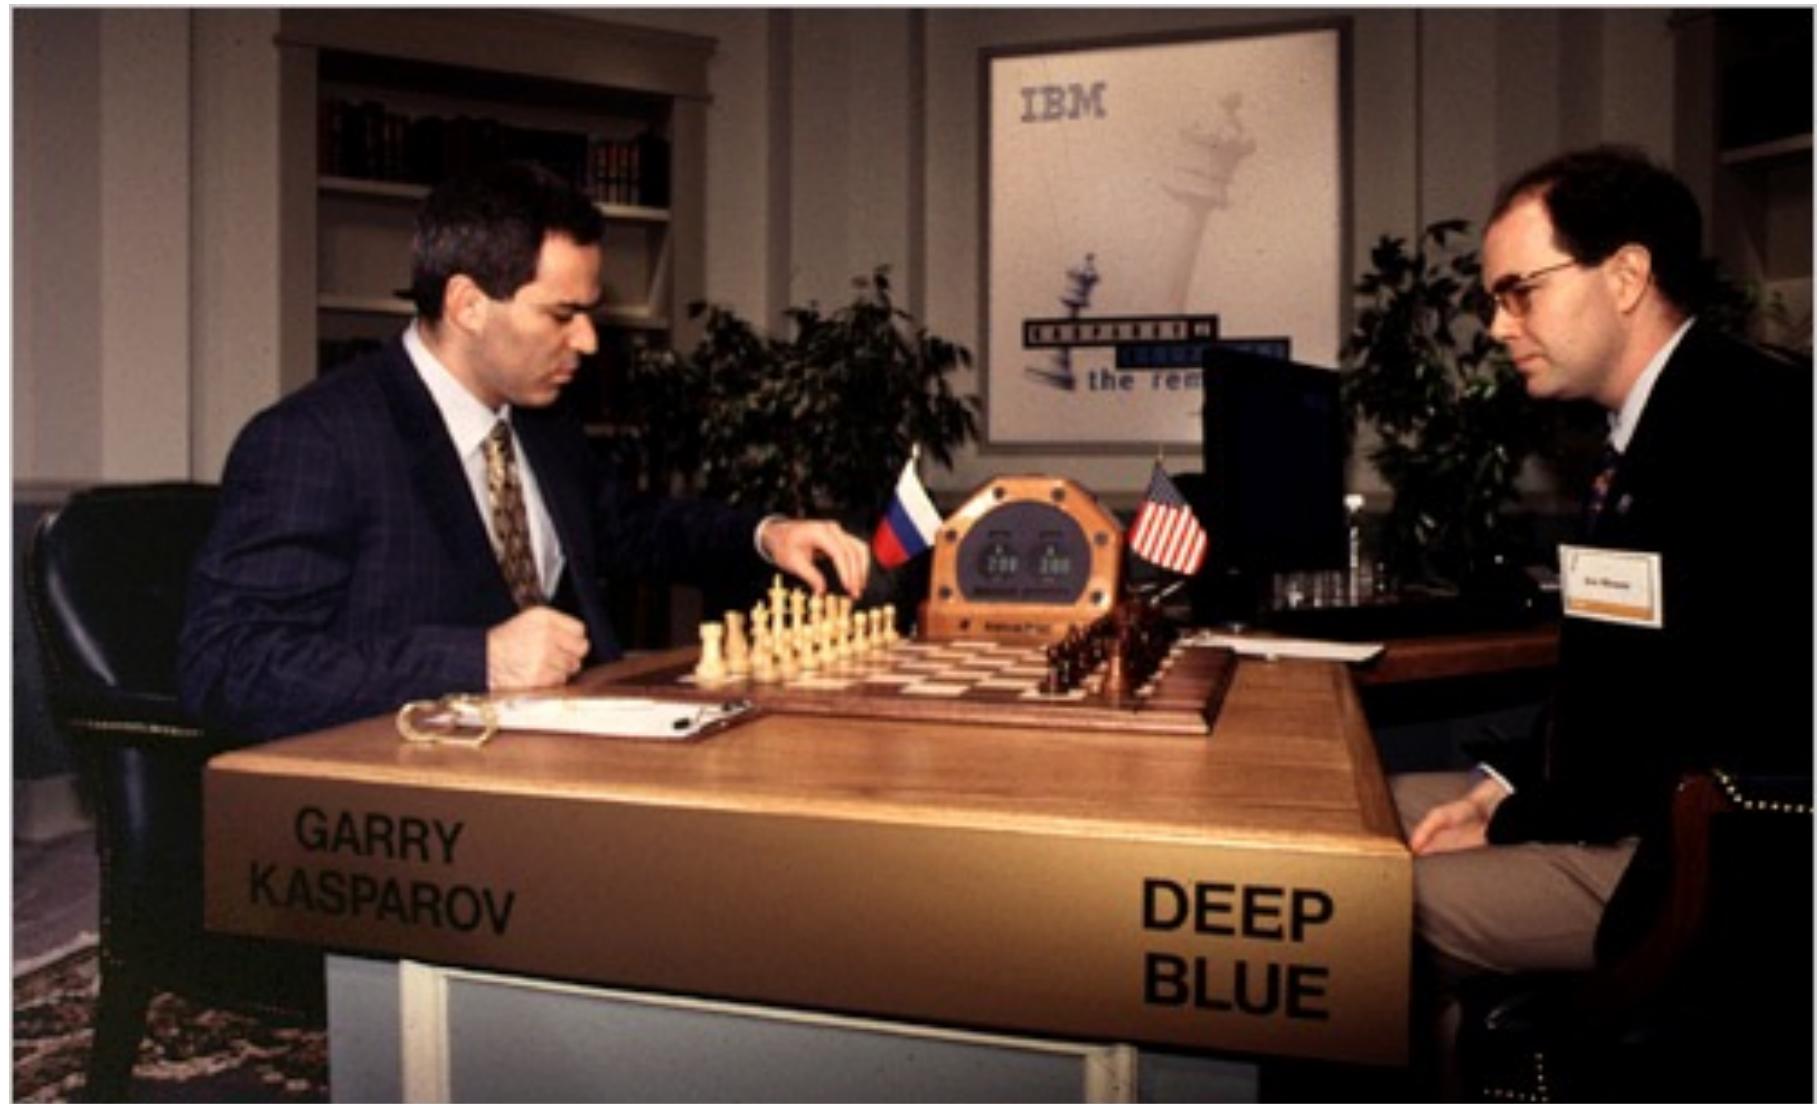
P. G. Piatz sc:

Der Schachspieler, wie er vor dem Spiele gesetzt wird von einem Le. Joueu de checs, tel qu'on le montre avant le jeu, par devant.

Chess (1996)



Deep Blue (1996)



The 1996 match

Game #	White	Black	Result	Comment
1	Deep Blue	Kasparov	1-0	
2	Kasparov	Deep Blue	1-0	
3	Deep Blue	Kasparov	½–½	Draw by mutual agreement
4	Kasparov	Deep Blue	½–½	Draw by mutual agreement
5	Deep Blue	Kasparov	0–1	Kasparov offered a draw after the 23rd move.
6	Kasparov	Deep Blue	1–0	
Result: Kasparov–Deep Blue: 4–2				

The 1997 rematch

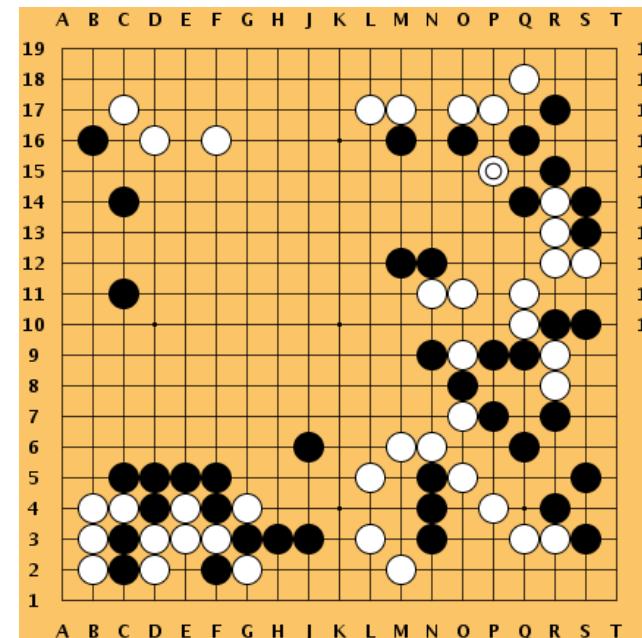
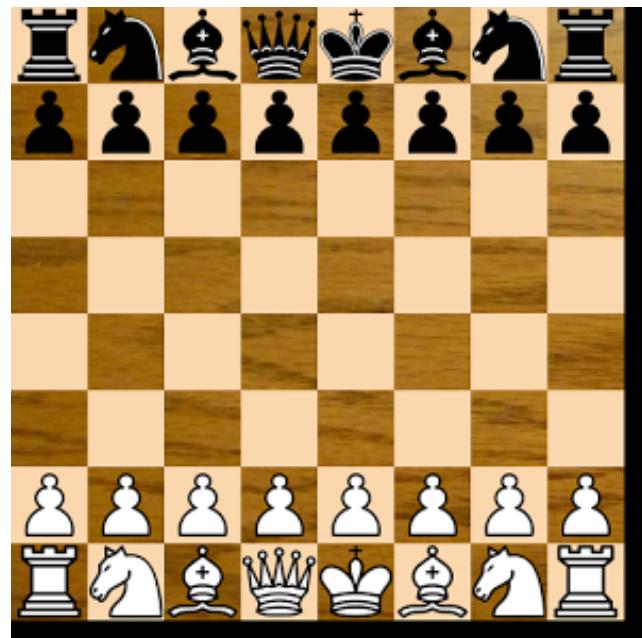
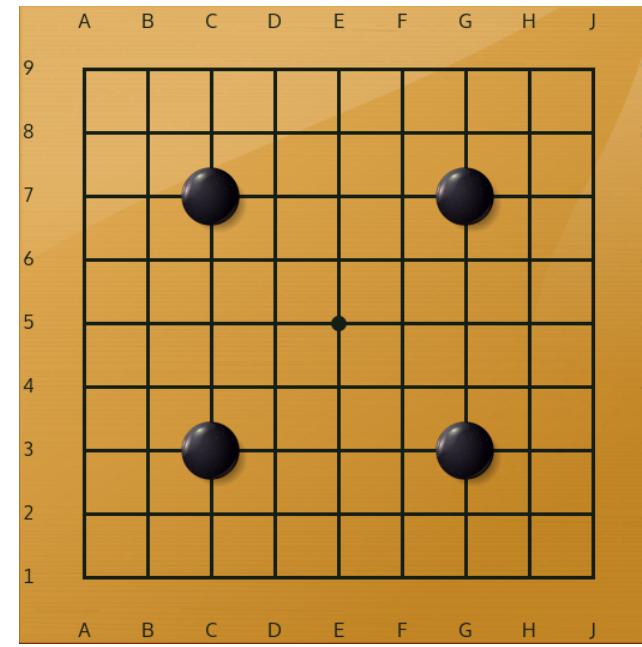
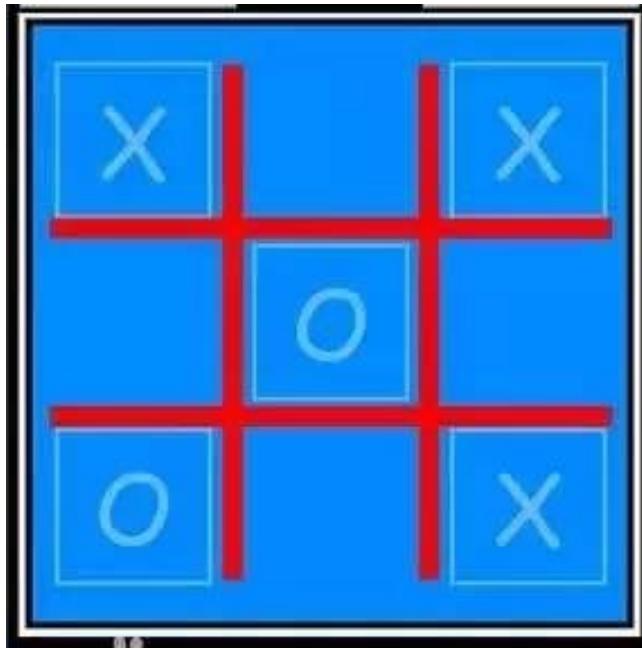
Game #	White	Black	Result	Comment
1	Kasparov	Deep Blue	1–0	
2	Deep Blue	Kasparov	1–0	
3	Kasparov	Deep Blue	½–½	Draw by mutual agreement
4	Deep Blue	Kasparov	½–½	Draw by mutual agreement
5	Kasparov	Deep Blue	½–½	Draw by mutual agreement
6	Deep Blue	Kasparov	1–0	
Result: Deep Blue–Kasparov: 3½–2½				

High Complexity

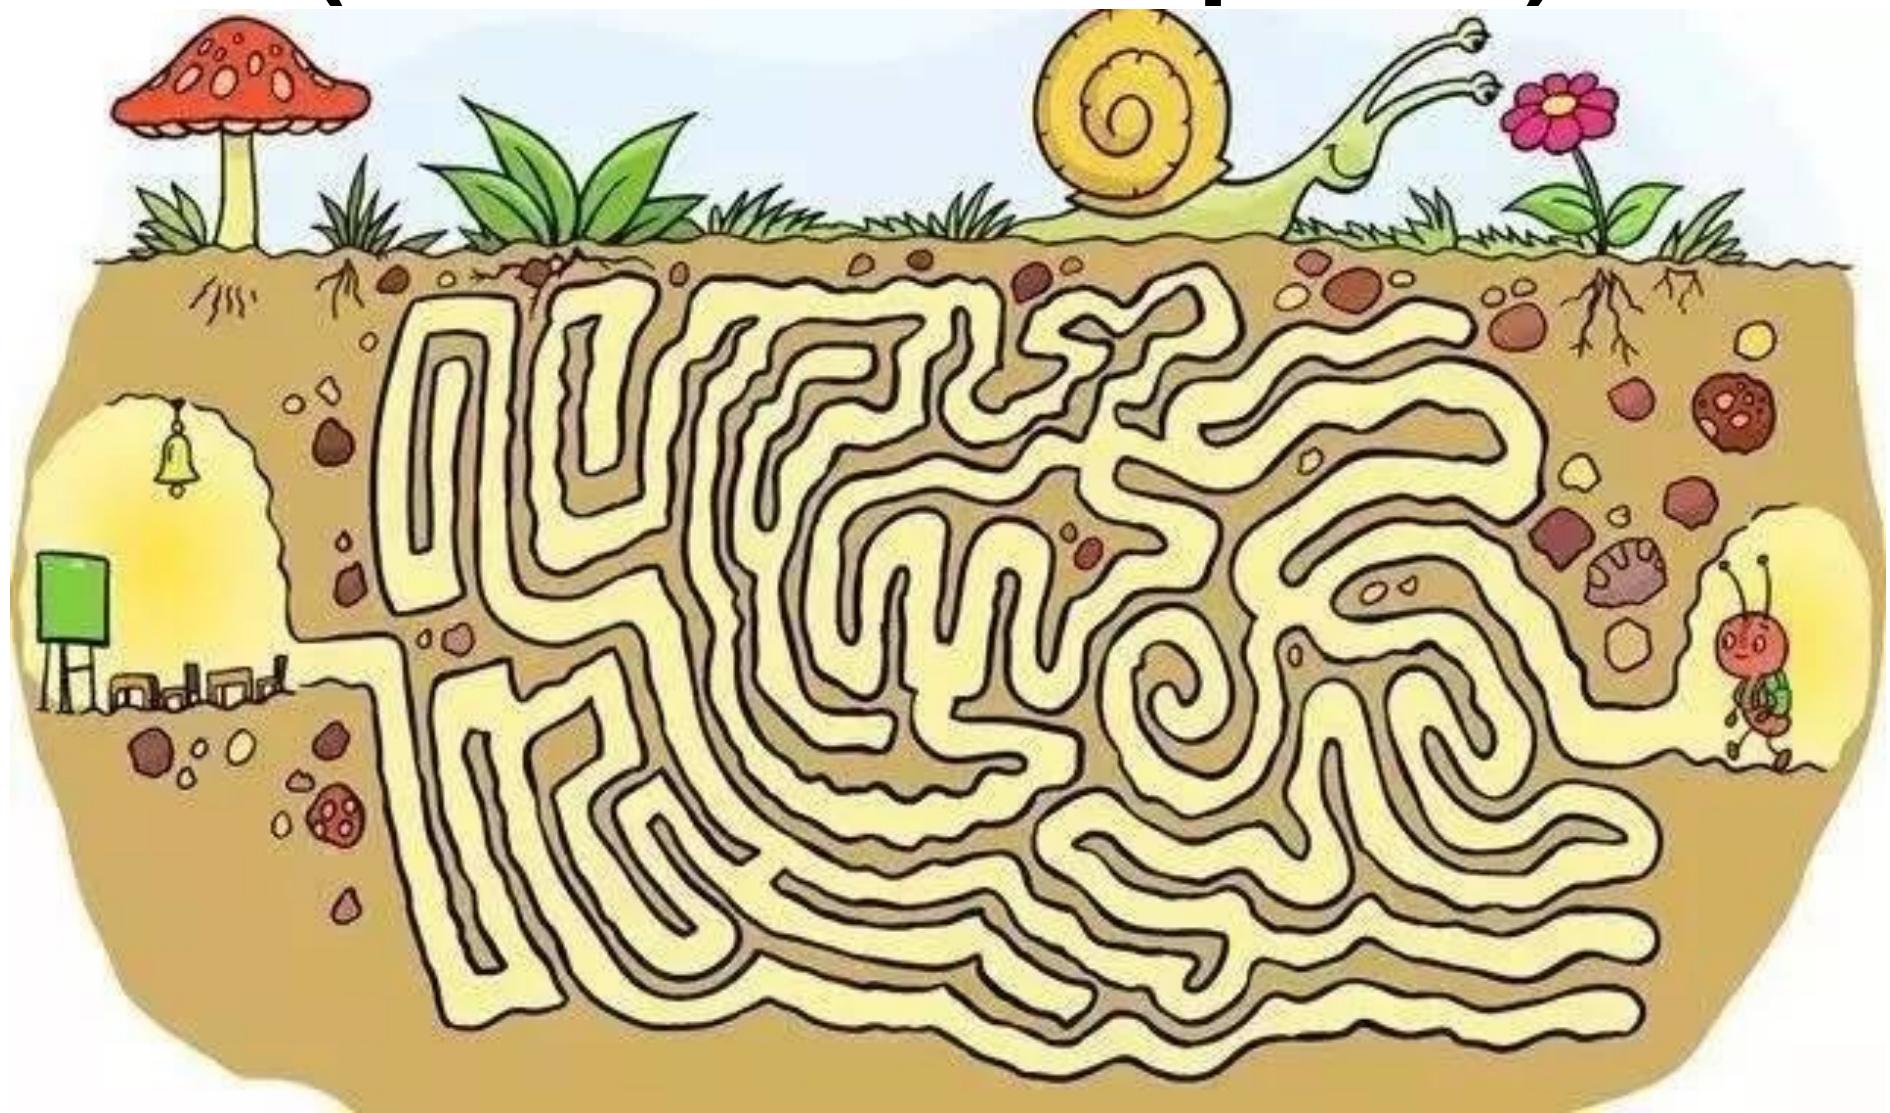
Complication



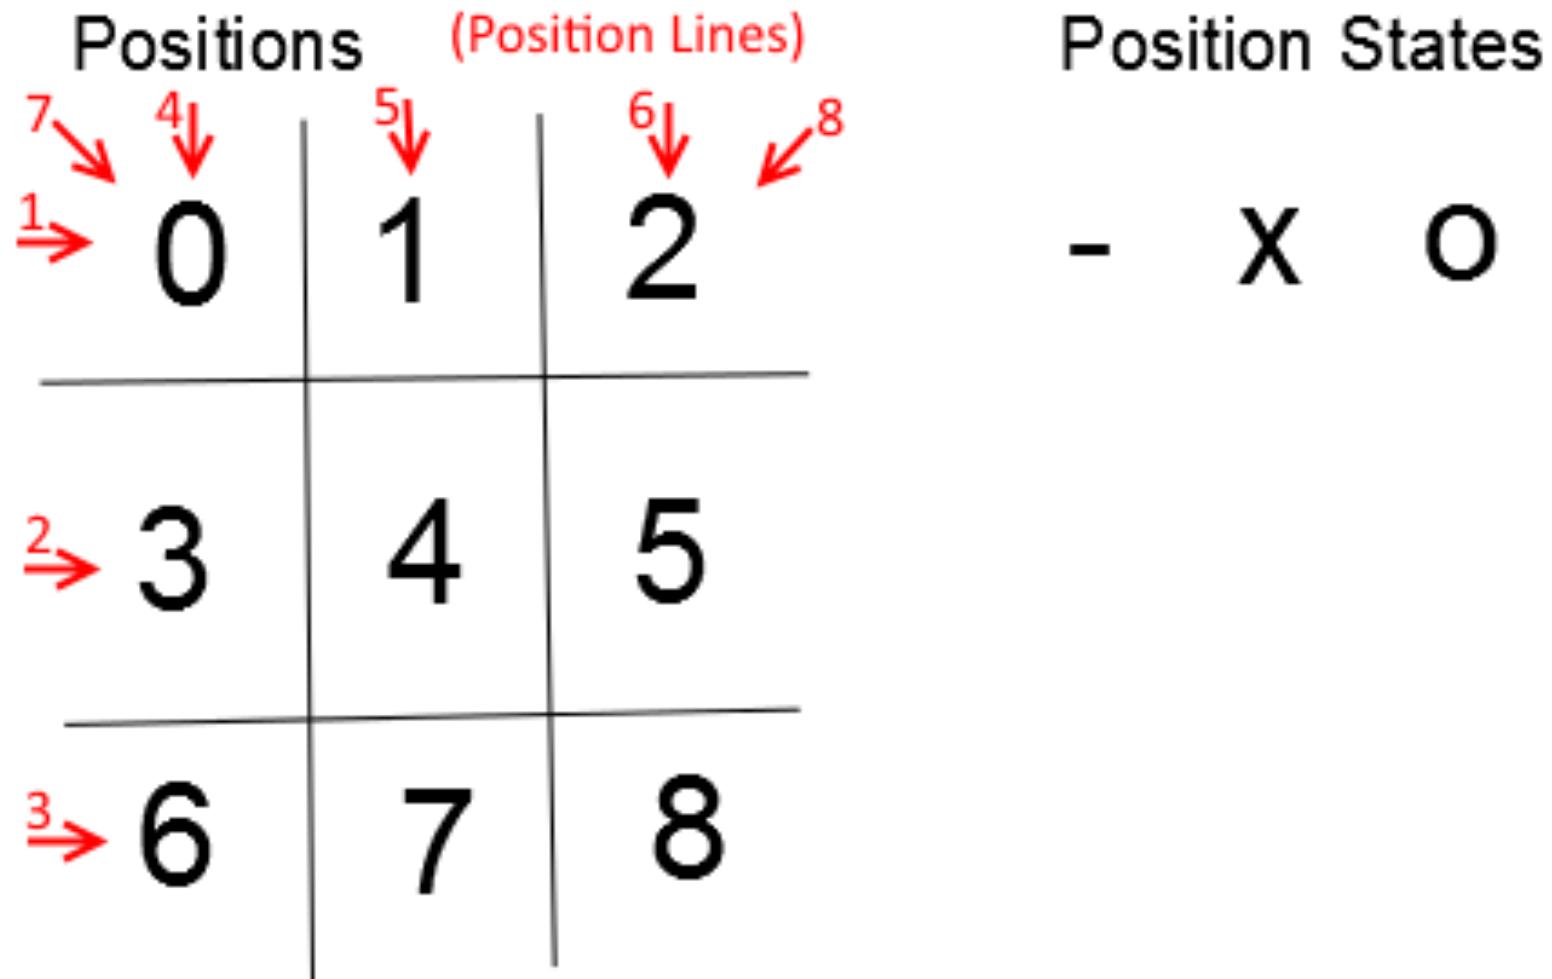
Different Games

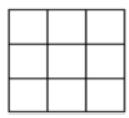


Search Problem (the search space)

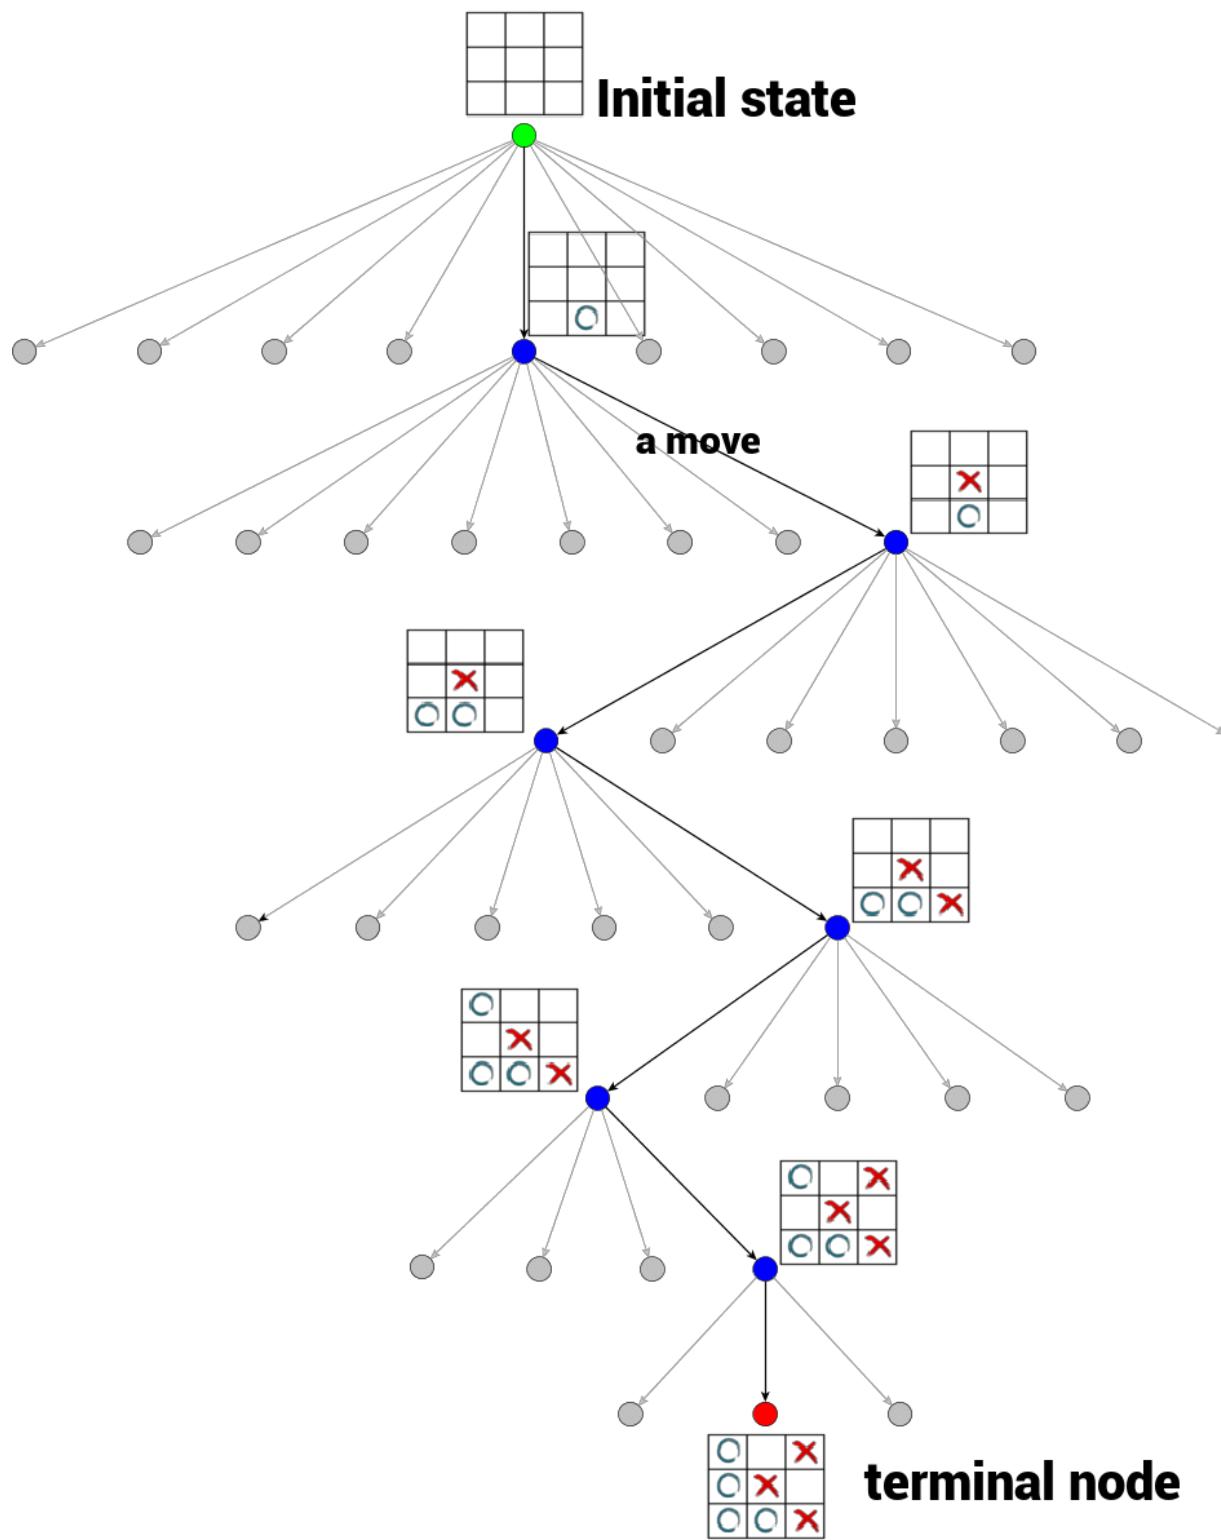


Game State (Tic Tac Toe)



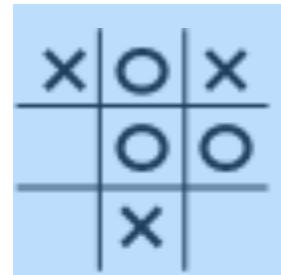


Initial state



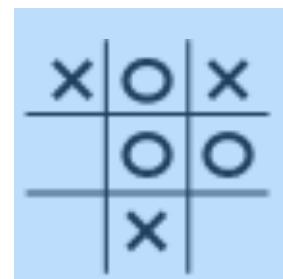
Tic Tac Toe

X's turn



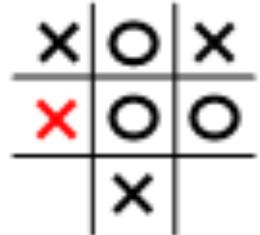
Current State

Tic Tac Toe

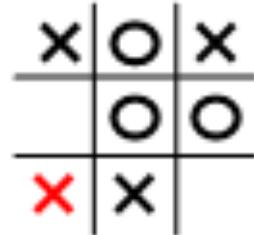


Current State/Root Node

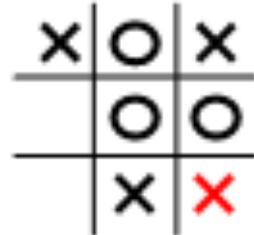
X's turn



Next State

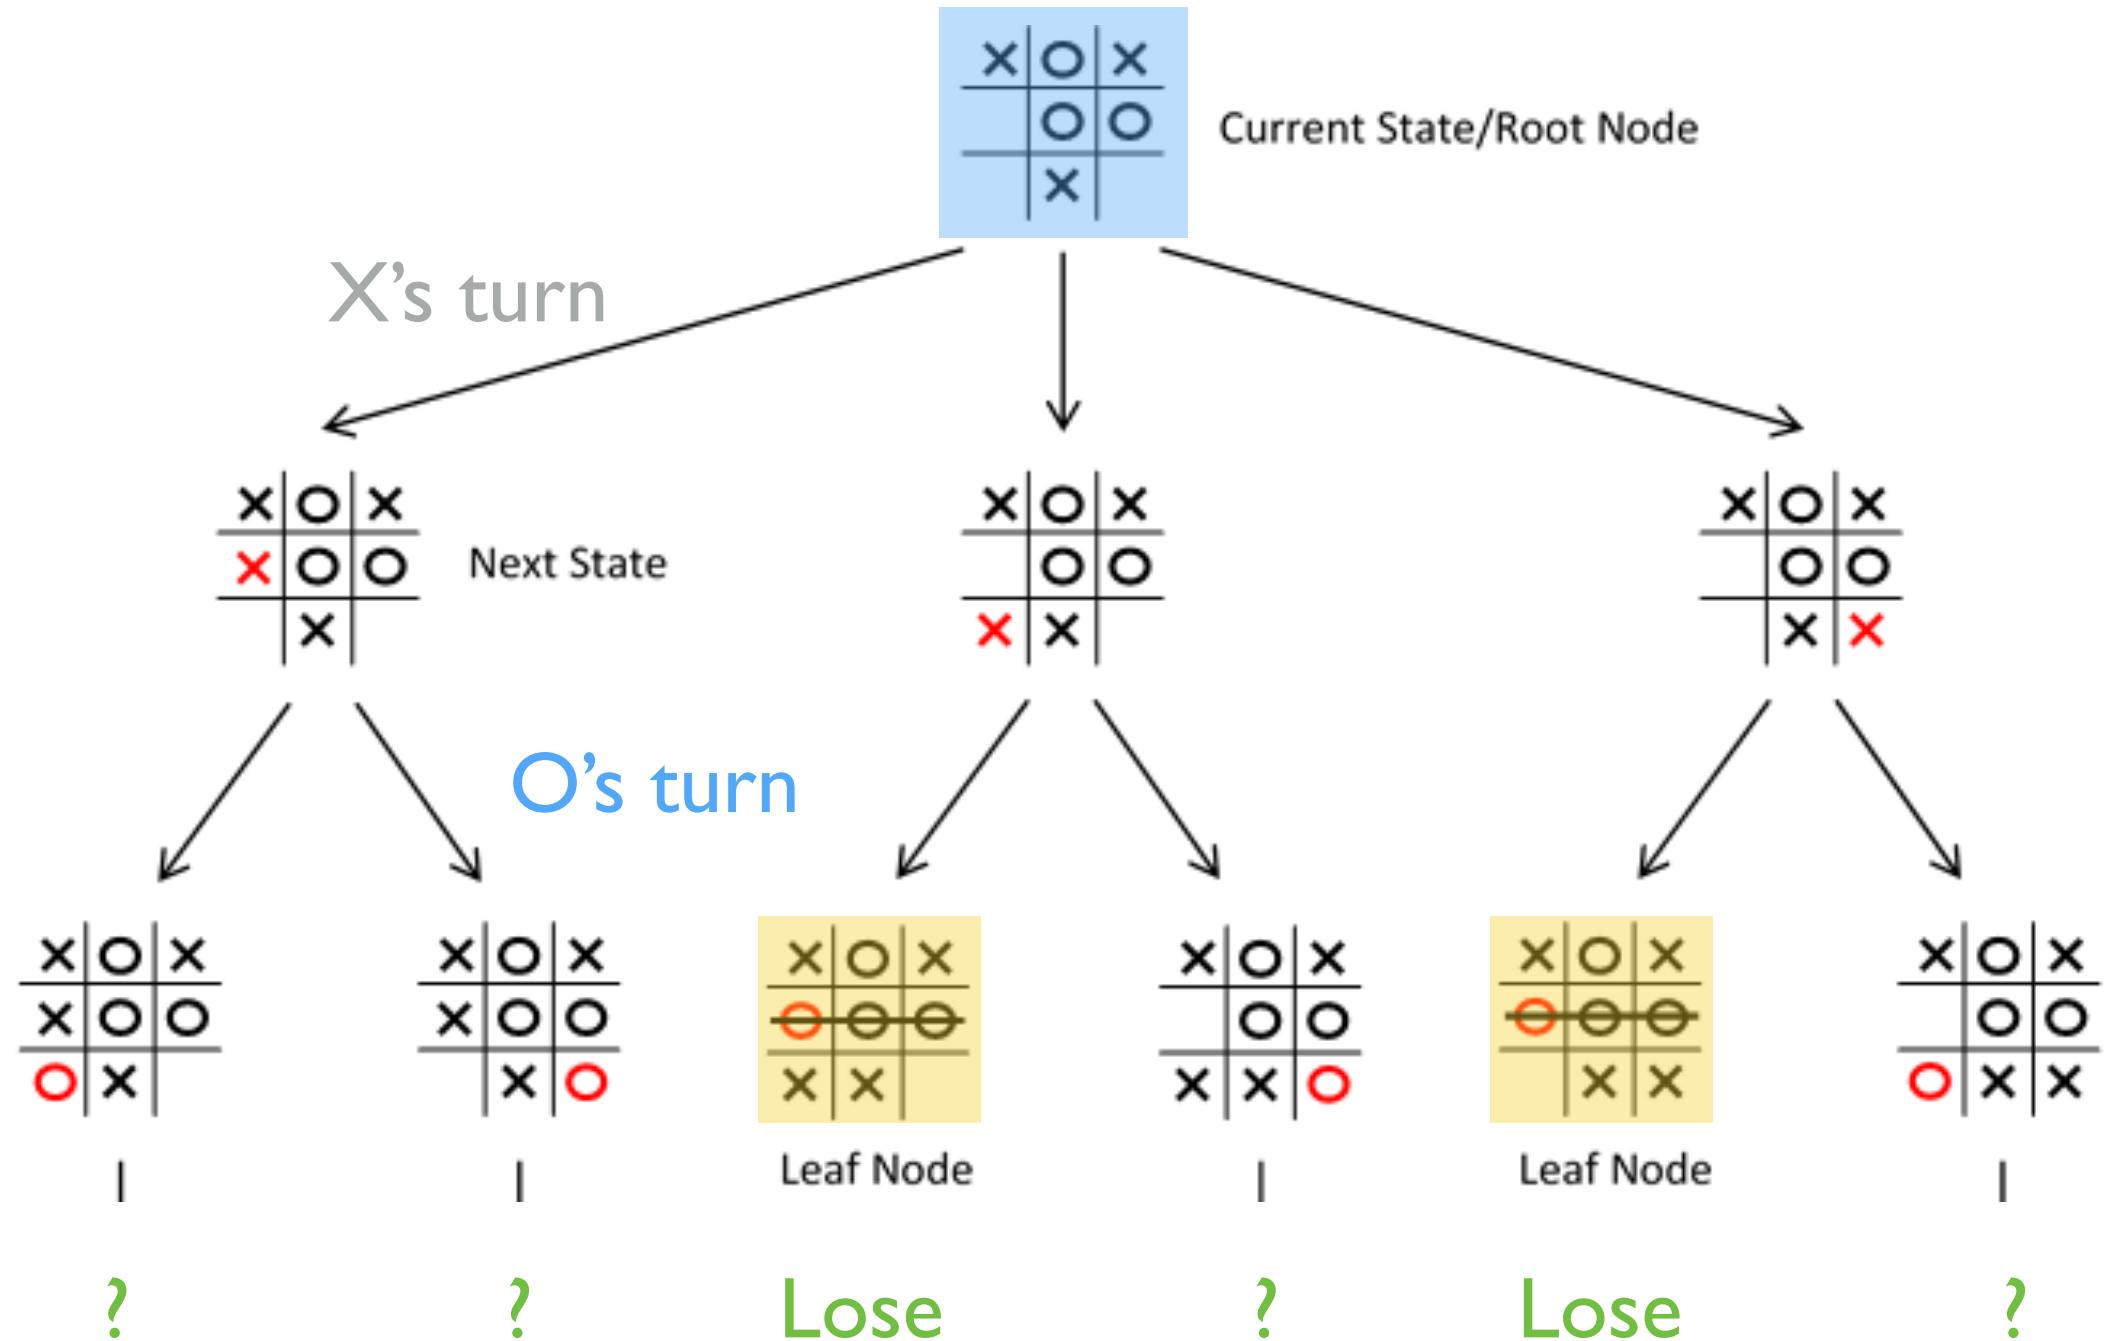


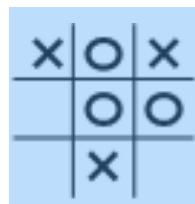
?



?

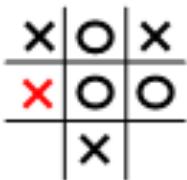
Tic Tac Toe



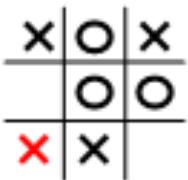


Current State/Root Node

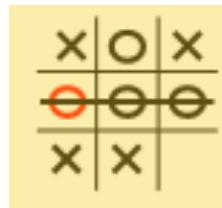
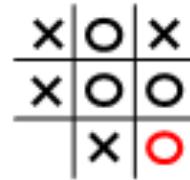
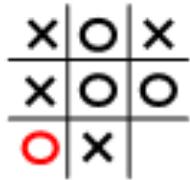
X's turn



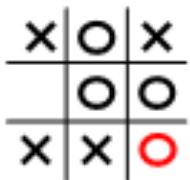
Next State



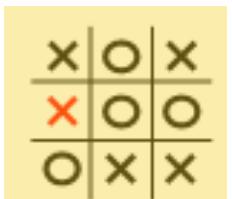
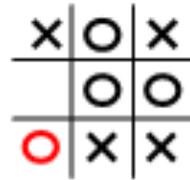
O's turn



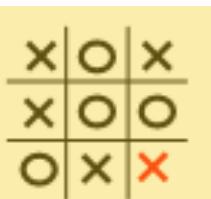
Leaf Node



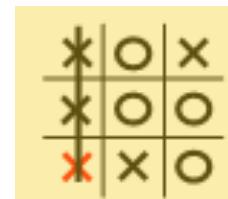
Leaf Node



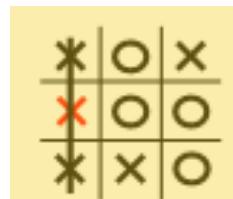
Leaf Node



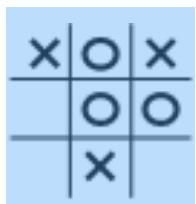
Leaf Node



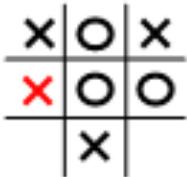
Leaf Node



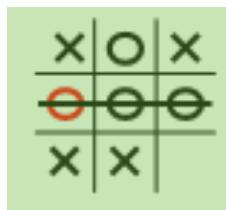
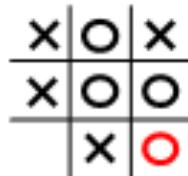
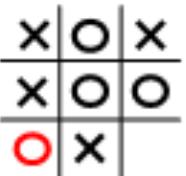
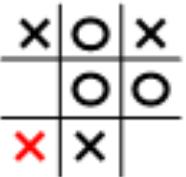
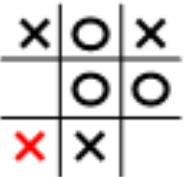
Leaf Node



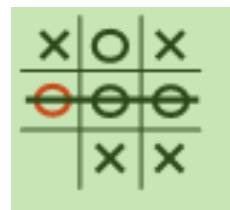
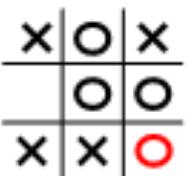
Current State/Root Node



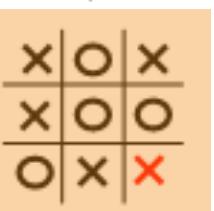
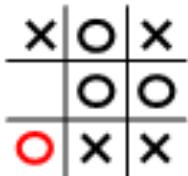
Next State



Leaf Node



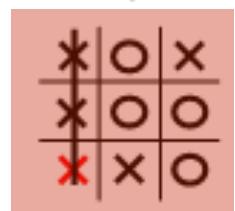
Leaf Node



Leaf Node

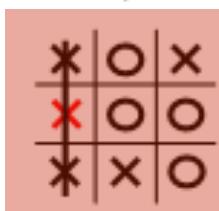
Tie

Lose



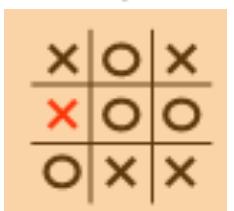
Leaf Node

Win



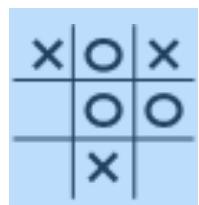
Leaf Node

Win

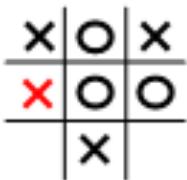


Leaf Node

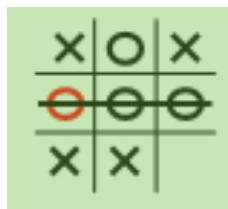
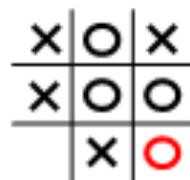
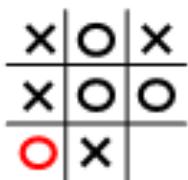
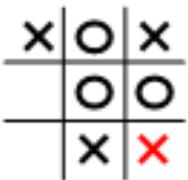
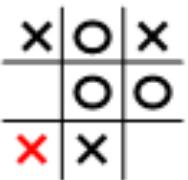
Tie



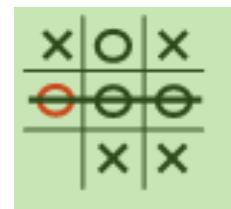
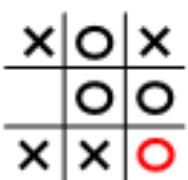
Current State/Root Node



Next State

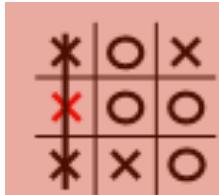


Leaf Node



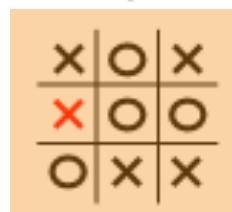
Leaf Node

-1



Leaf Node

-1



Leaf Node

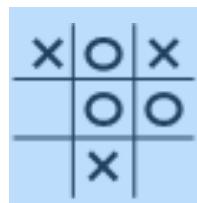
0

1

1

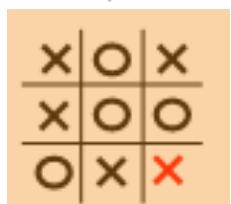
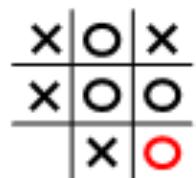
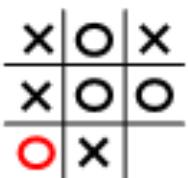
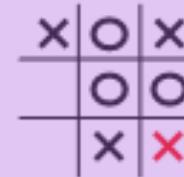
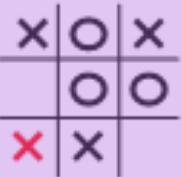
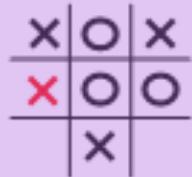
0

next move



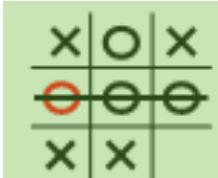
Current State/Root Node

how to choose



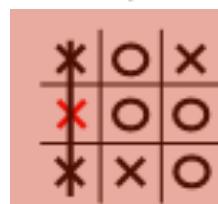
Leaf Node

0



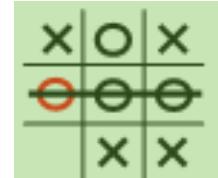
Leaf Node

-1



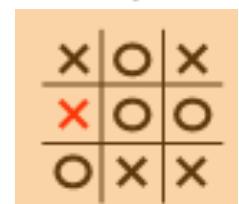
Leaf Node

1



Leaf Node

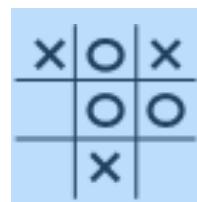
-1



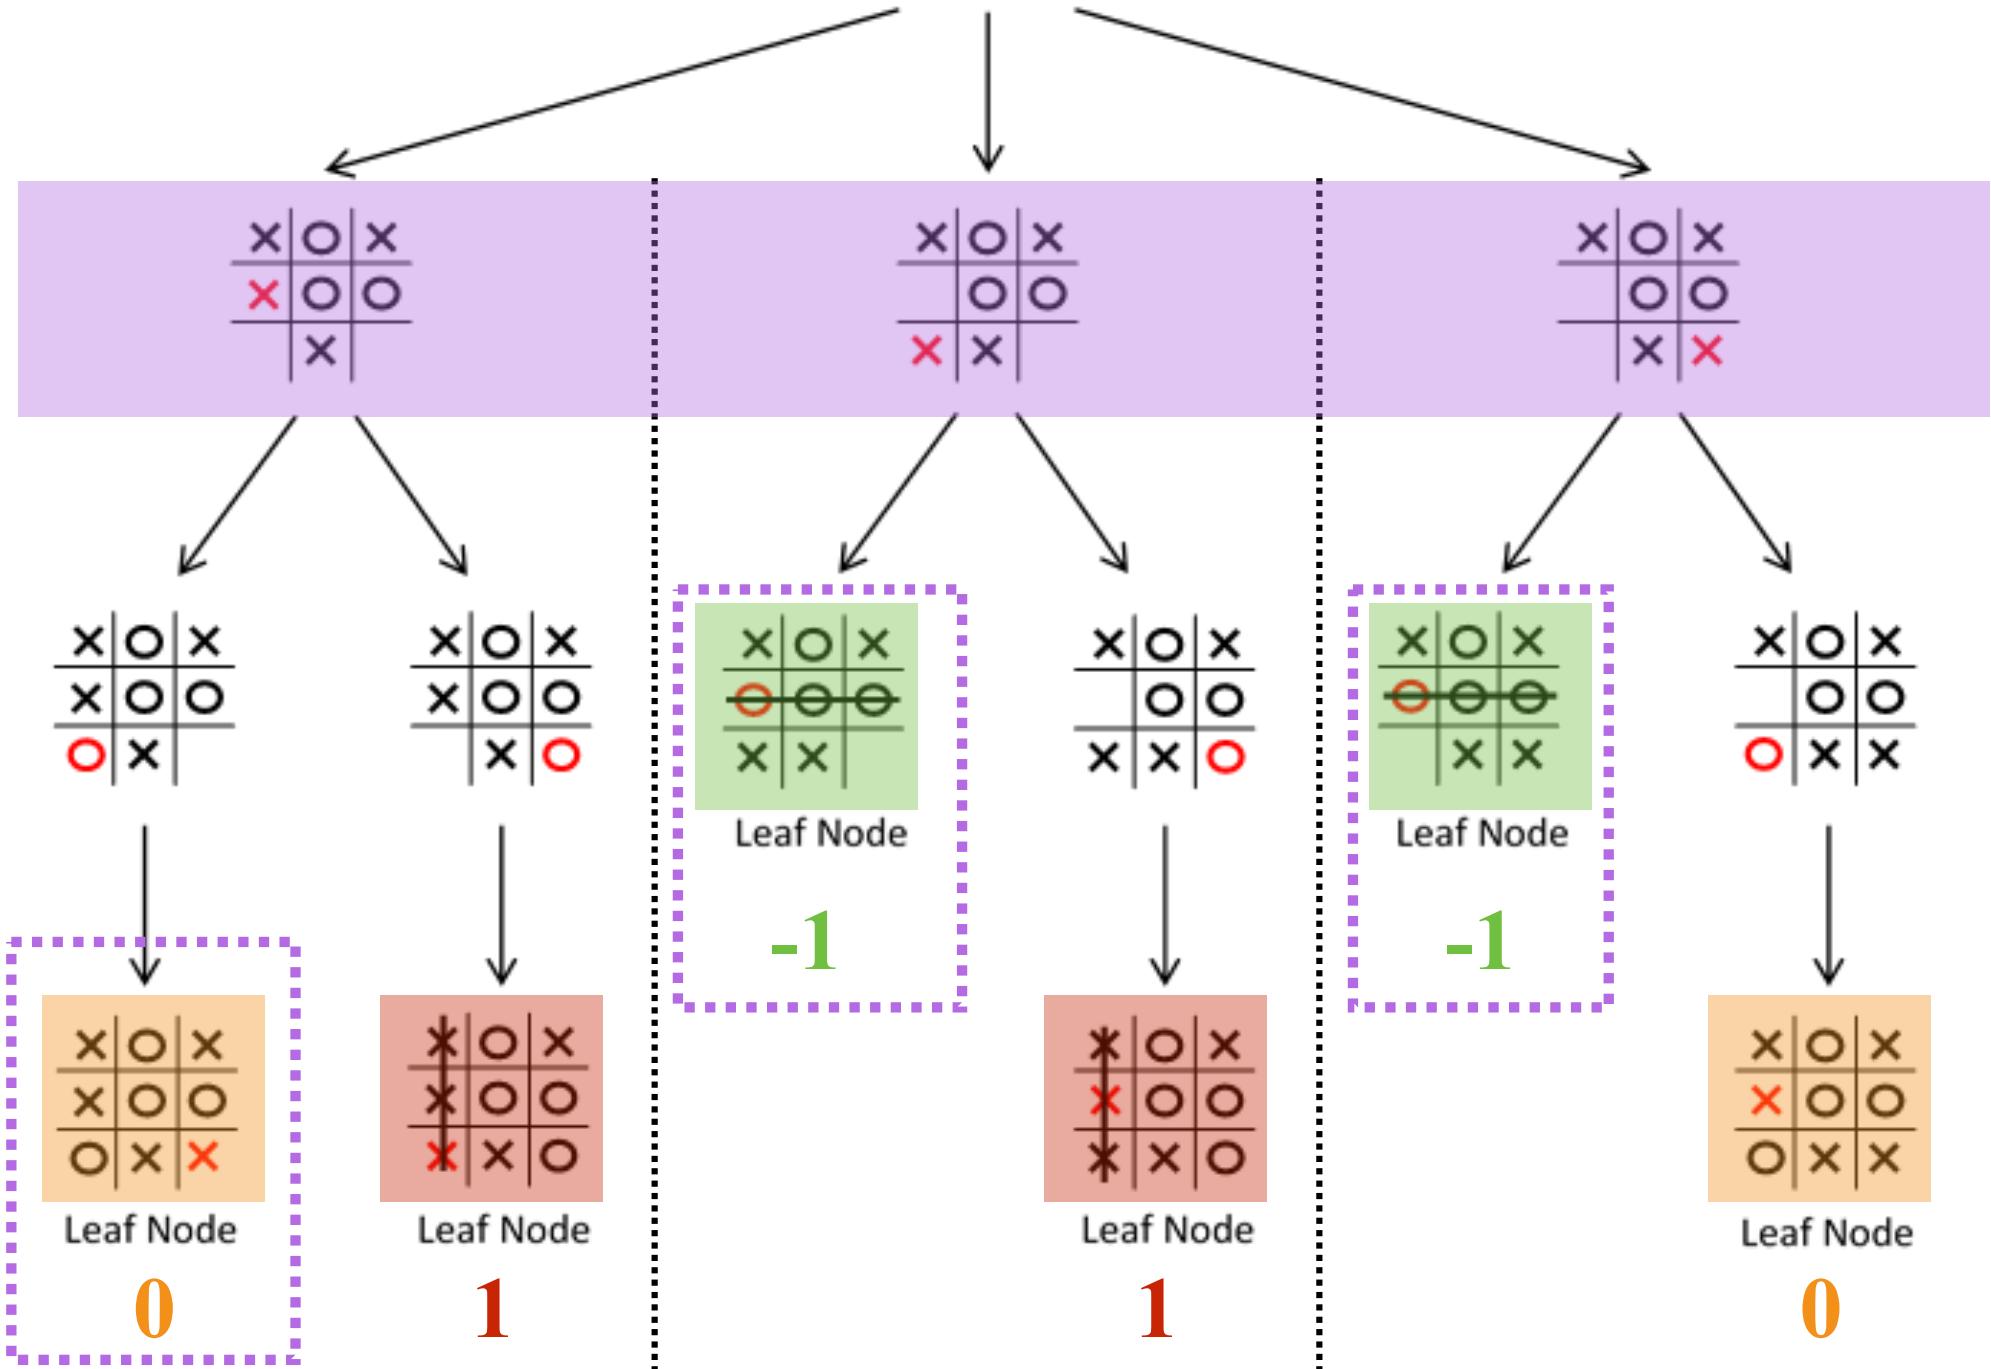
Leaf Node

0

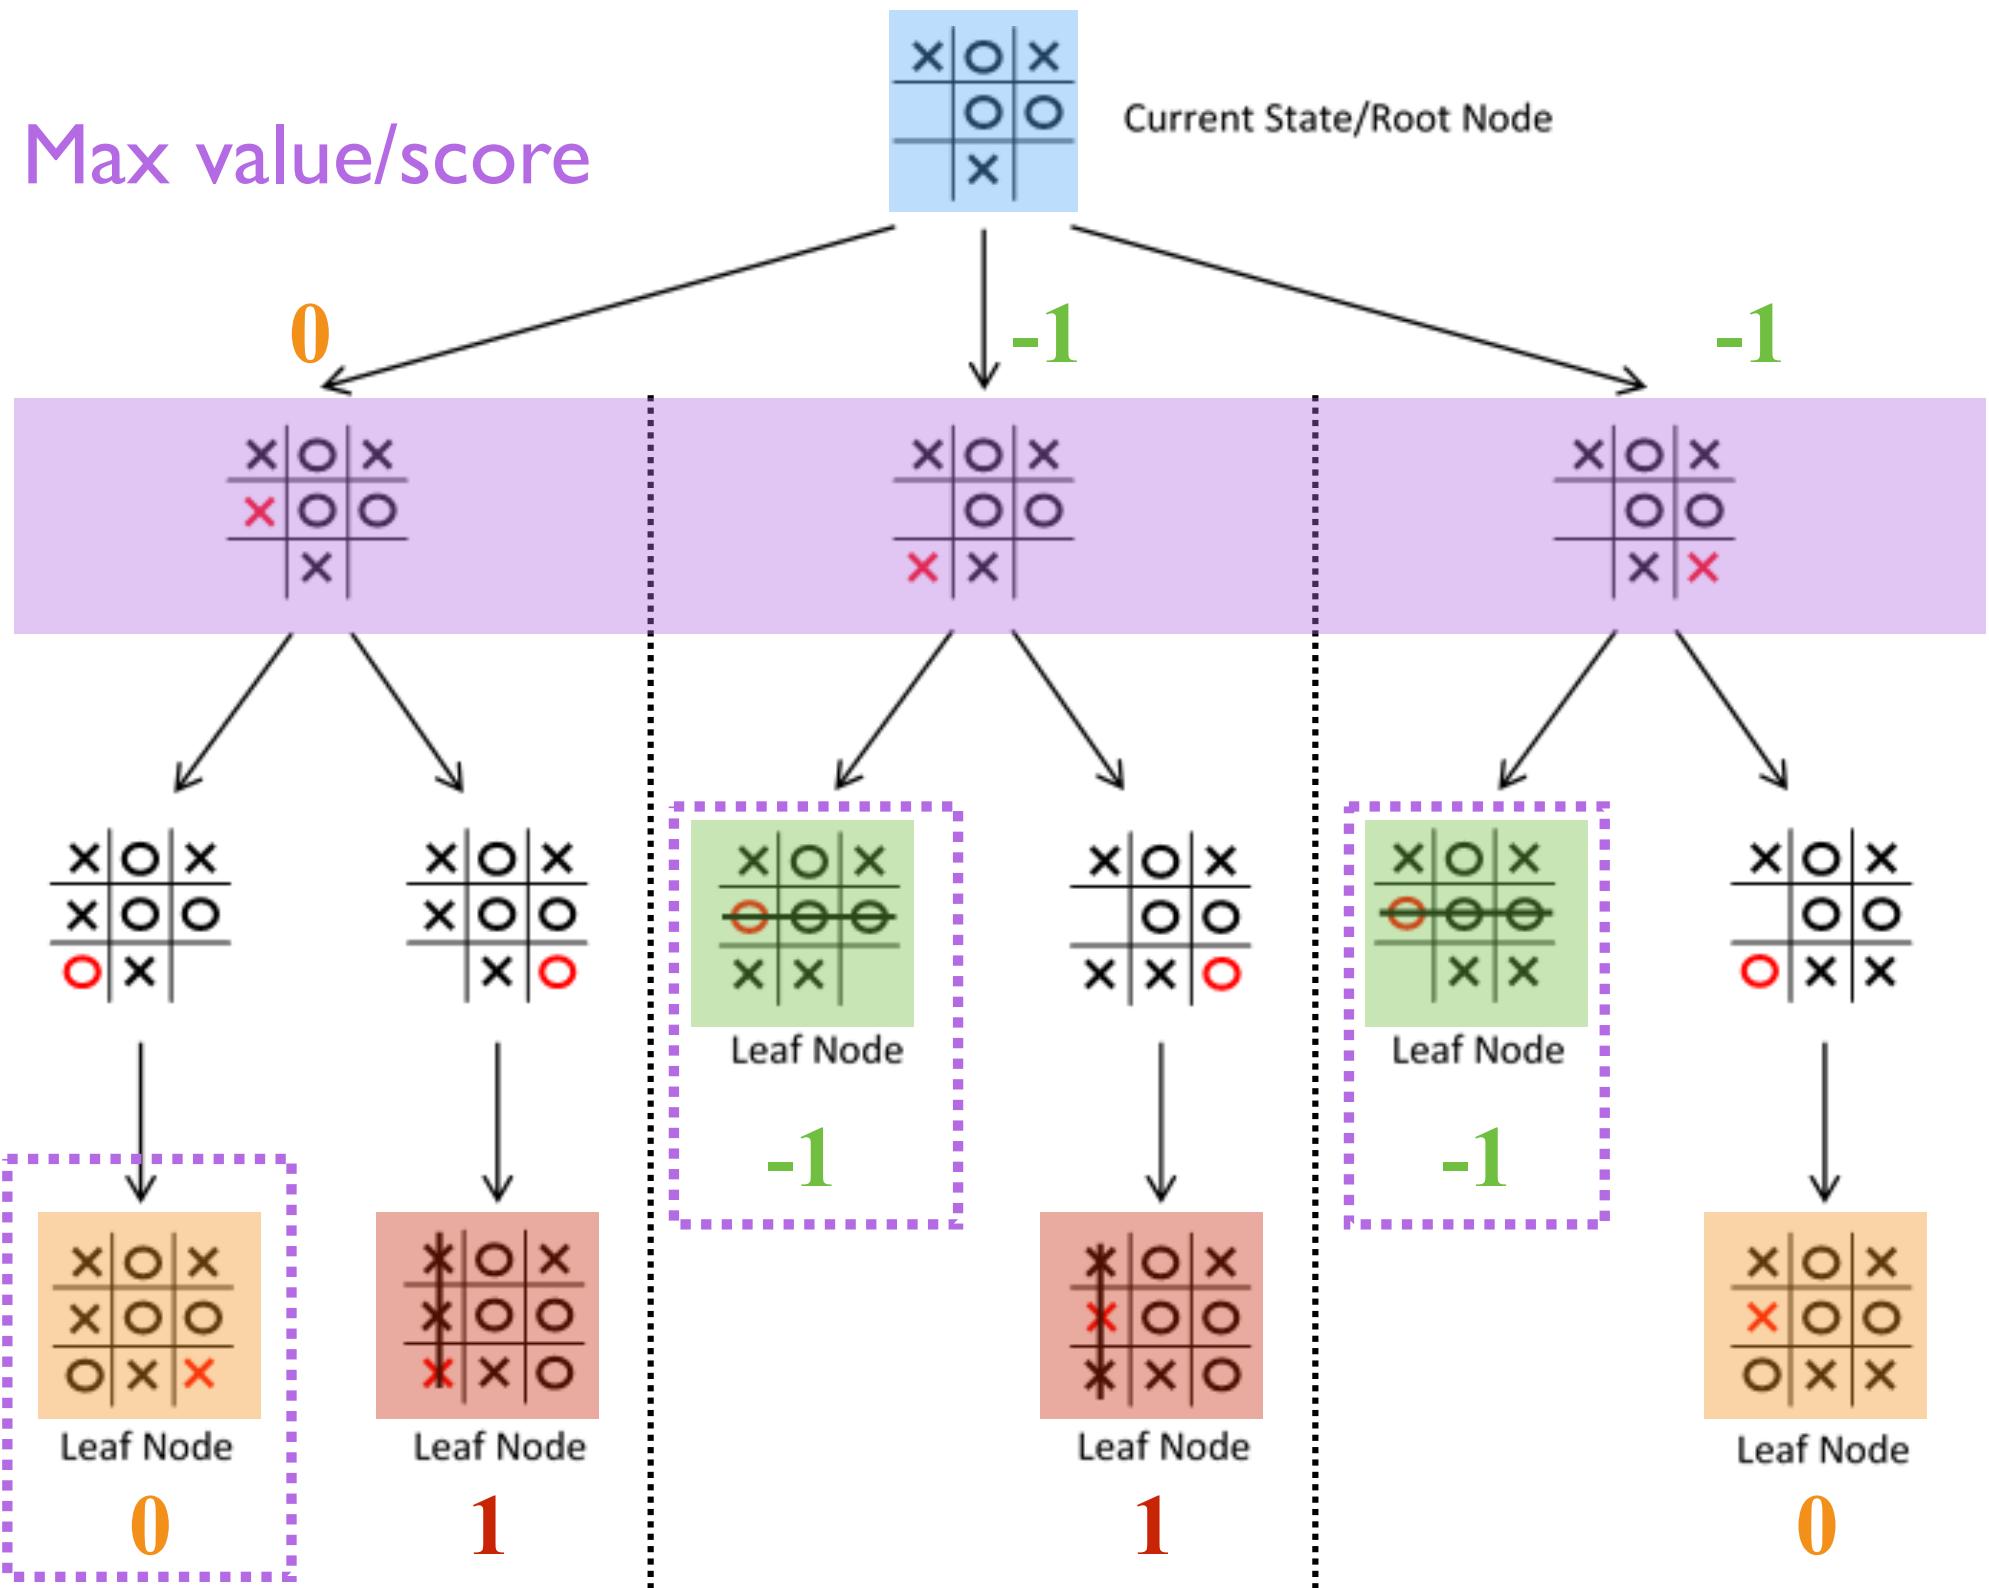
next move



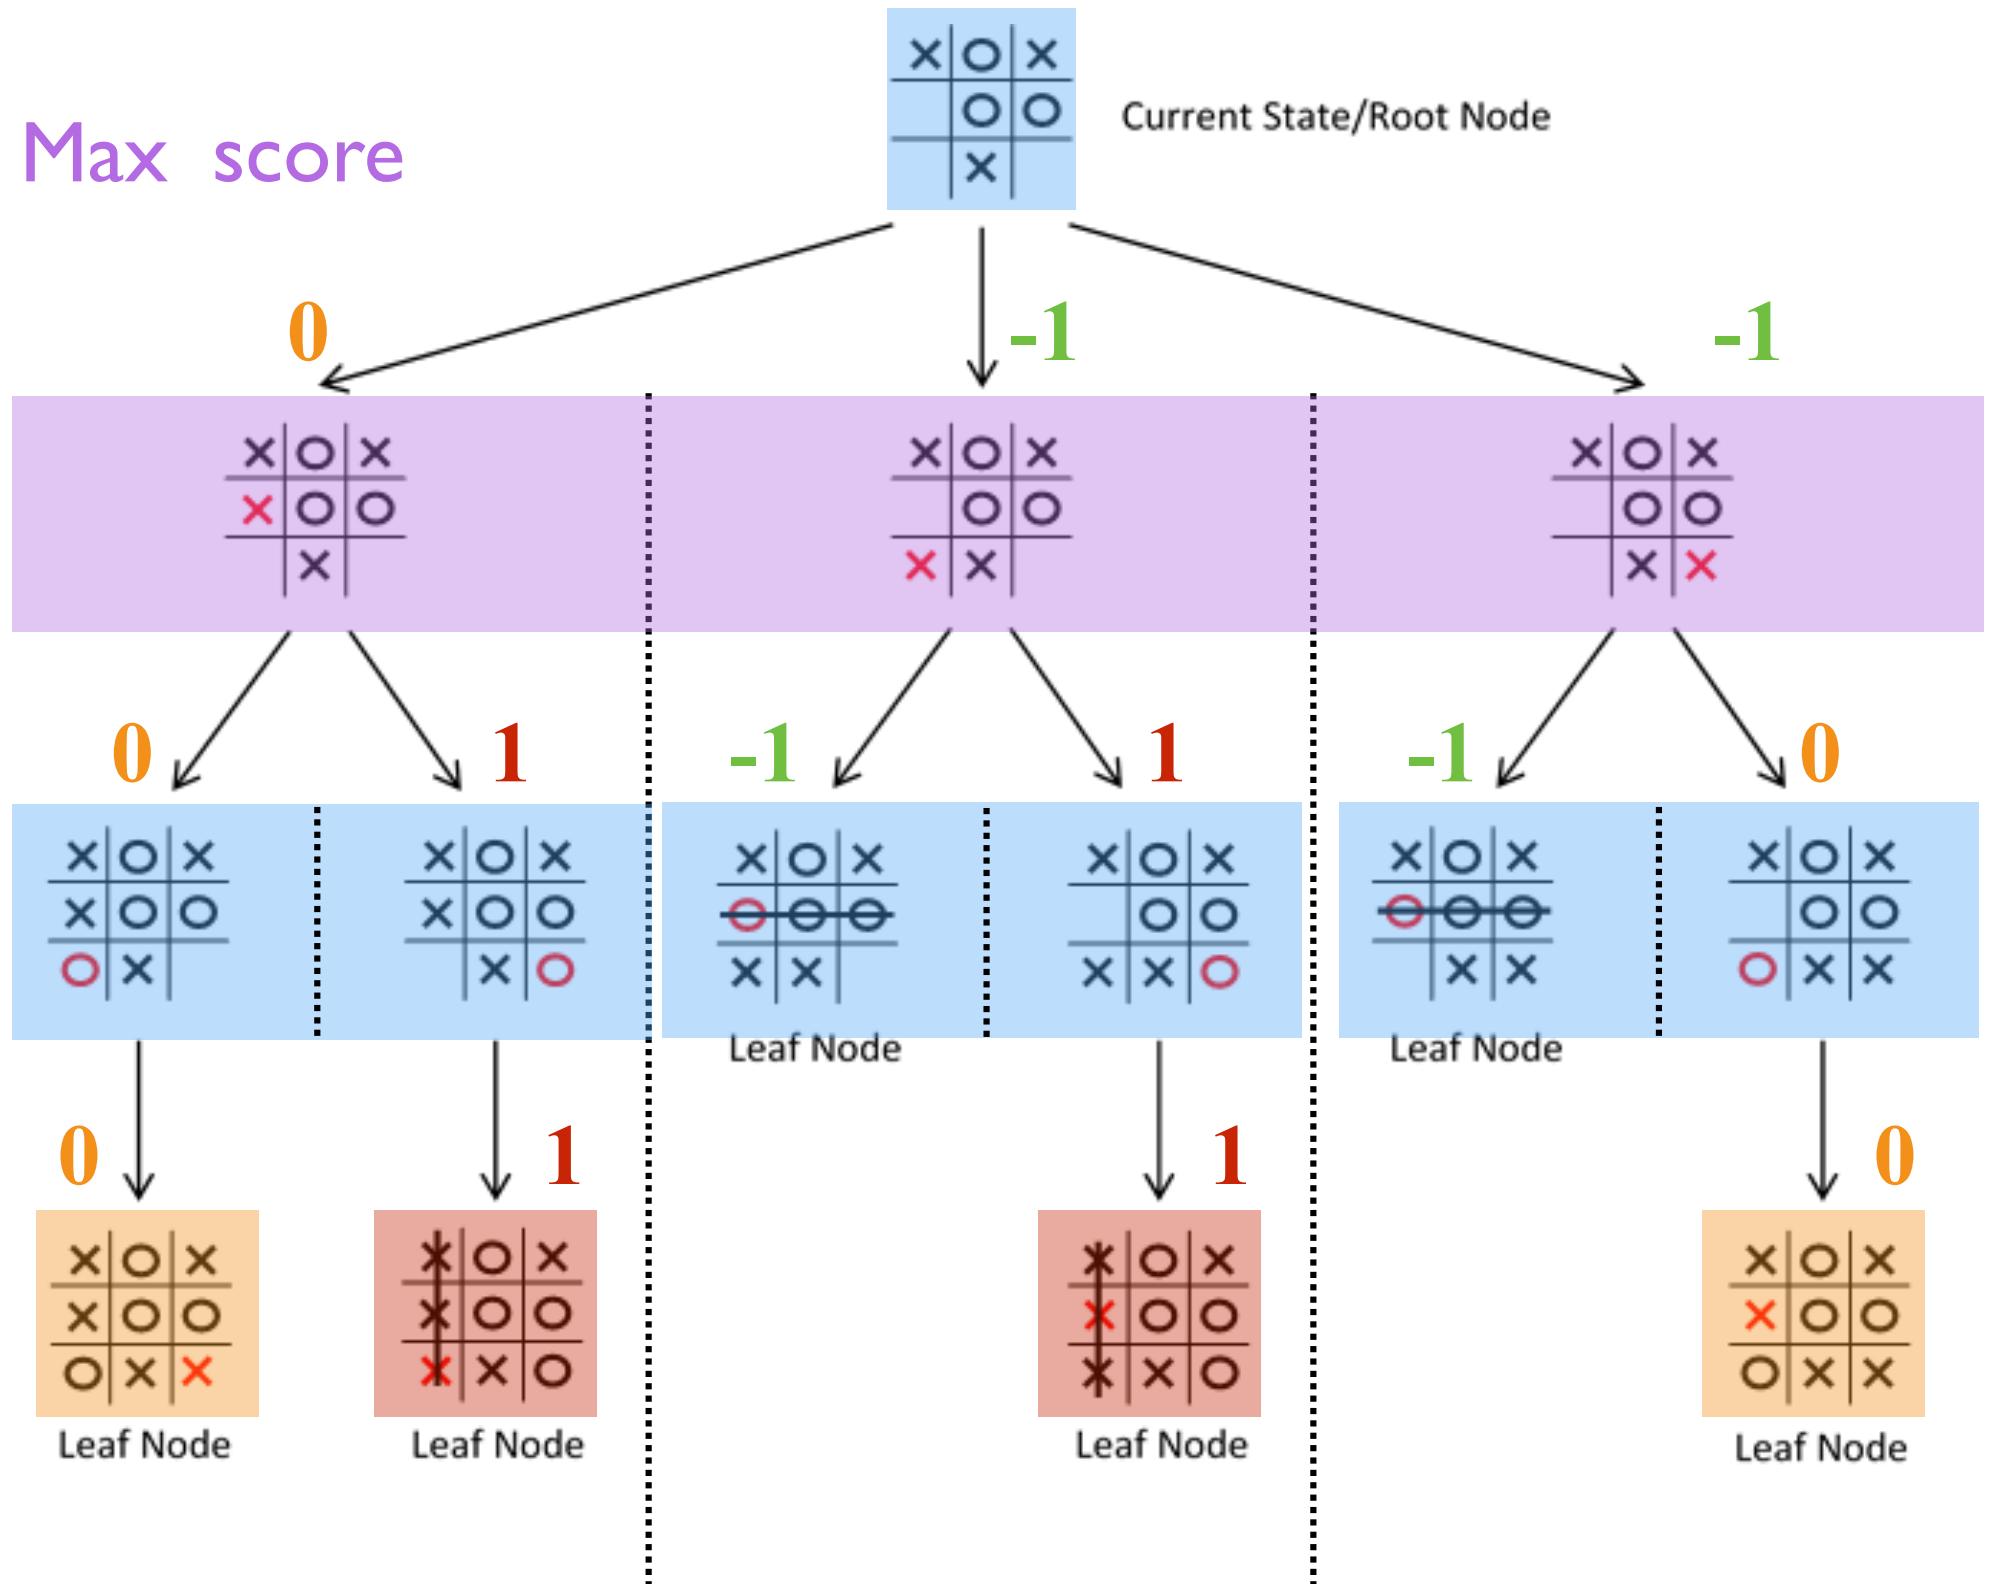
Current State/Root Node



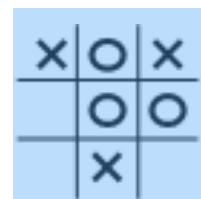
Max value/score



Max score



Max score

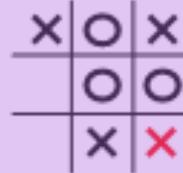
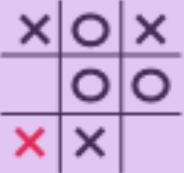
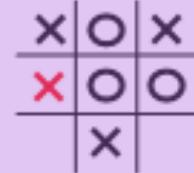


Current State/Root Node

0

-1

-1



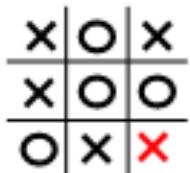
Min 0

1

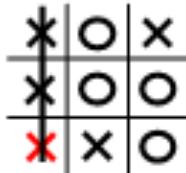


0

1



Leaf Node



Leaf Node

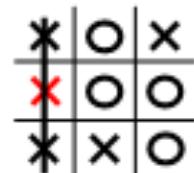
-1

1



Leaf Node

1



Leaf Node

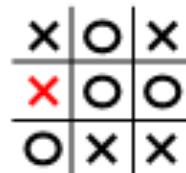
-1

0

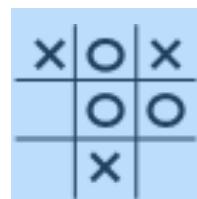


Leaf Node

0



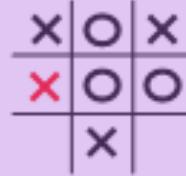
Leaf Node



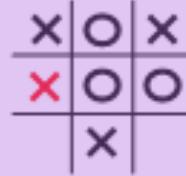
Current State/Root Node

Max score

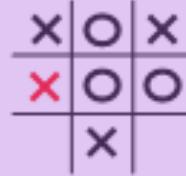
0



-1



-1



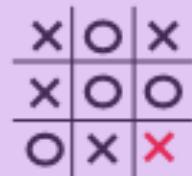
Min

0

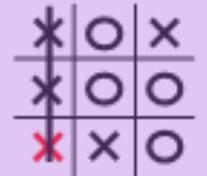
1



0



Leaf Node



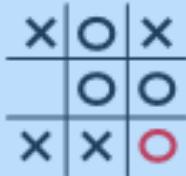
Leaf Node

-1



Leaf Node

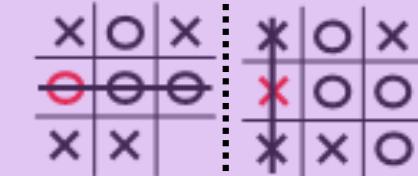
1



Max

-1

1



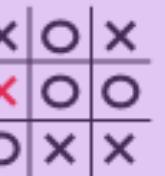
Leaf Node

-1



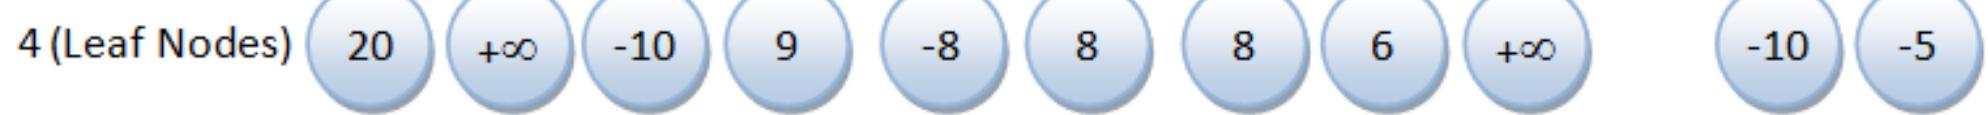
Leaf Node

0

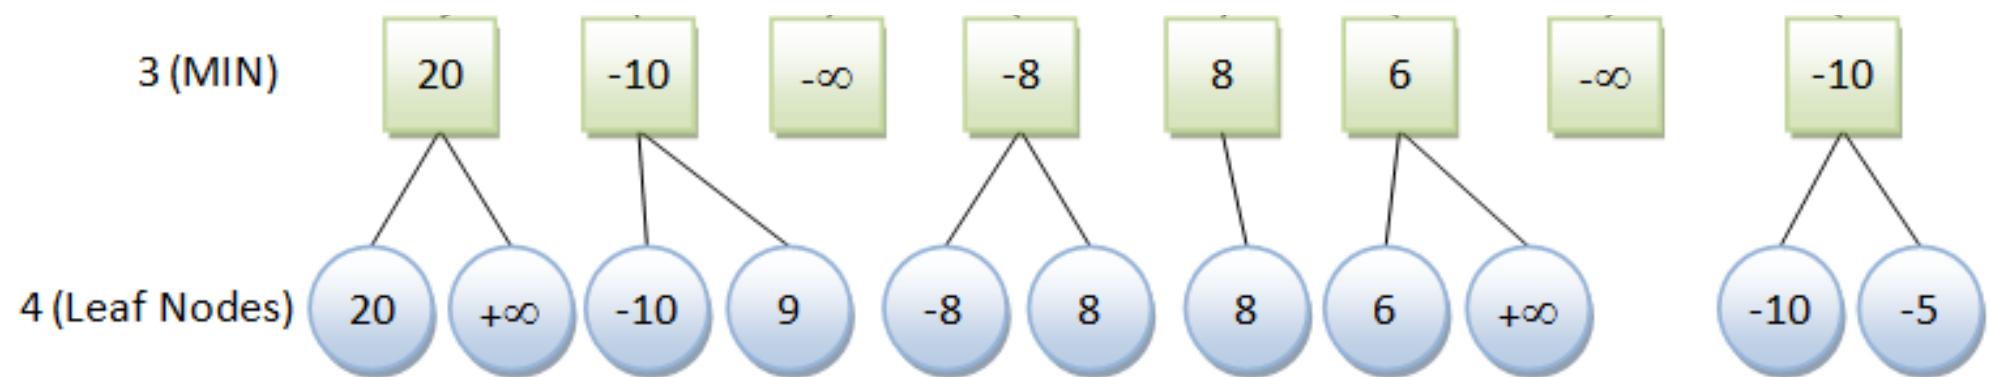


Leaf Node

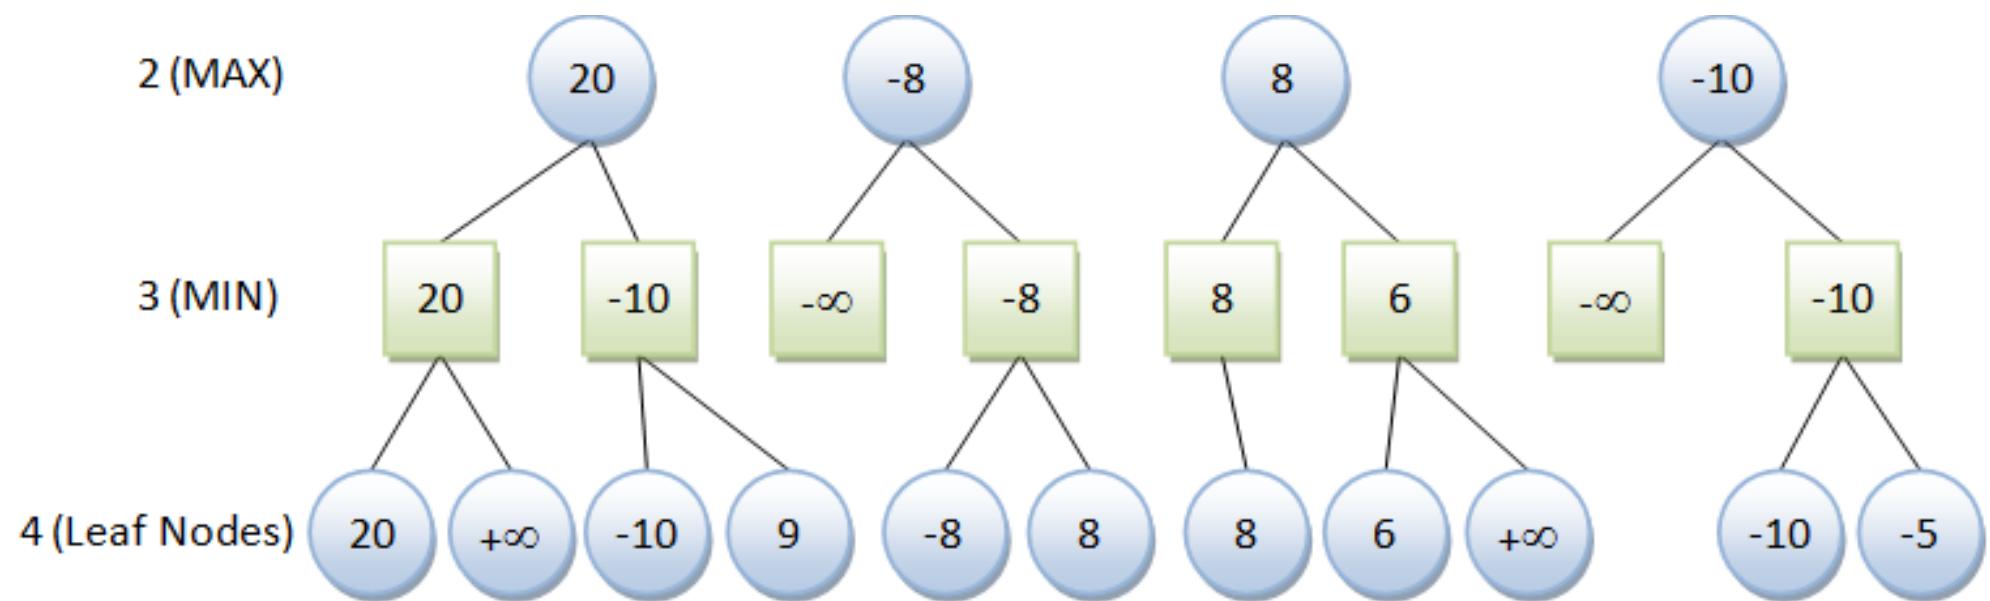
MiniMax



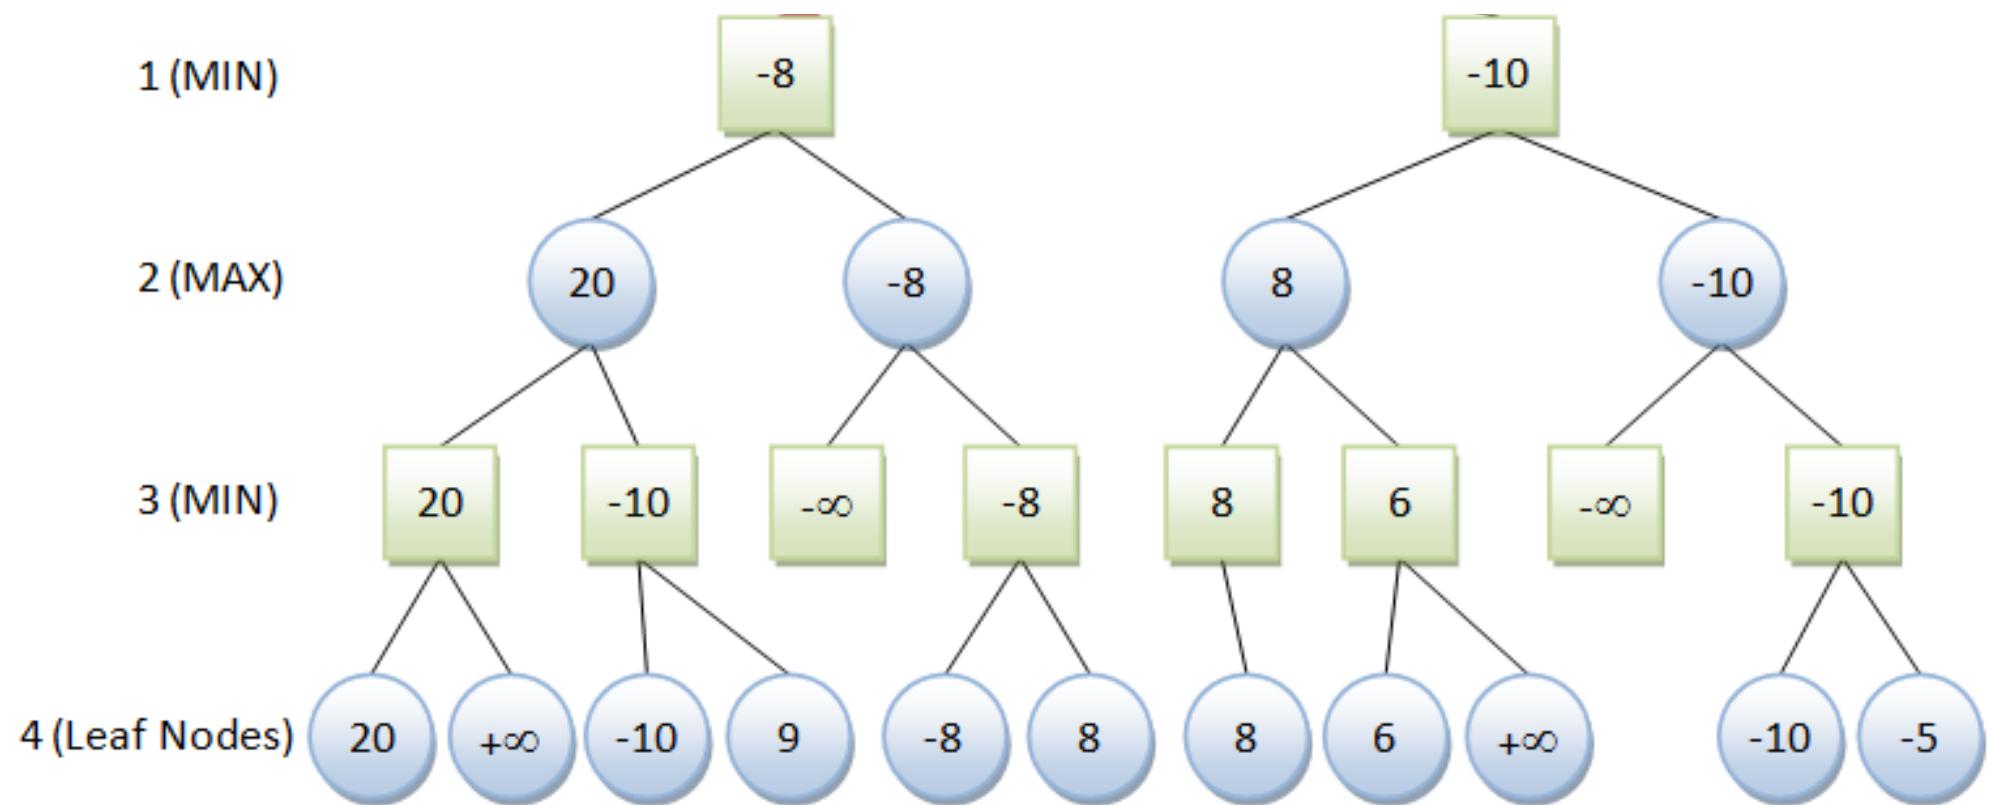
MiniMax



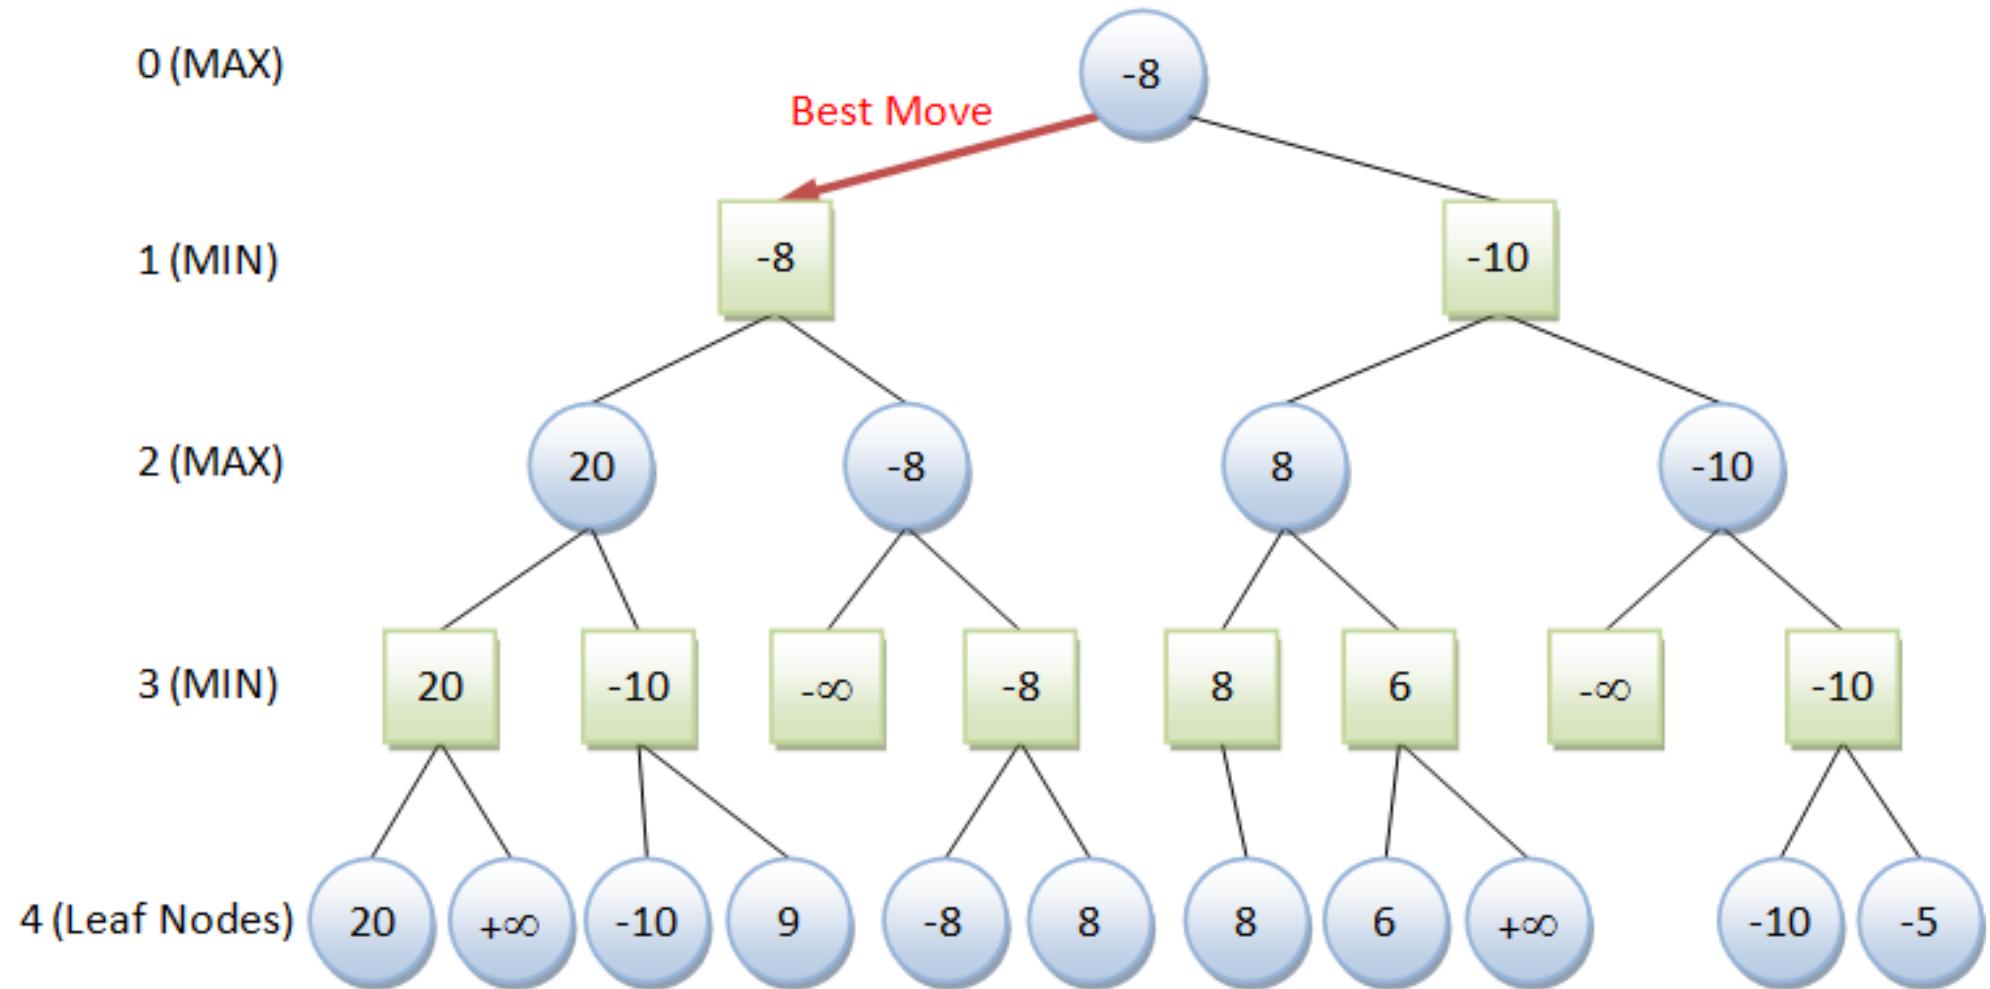
MiniMax



MiniMax



MiniMax

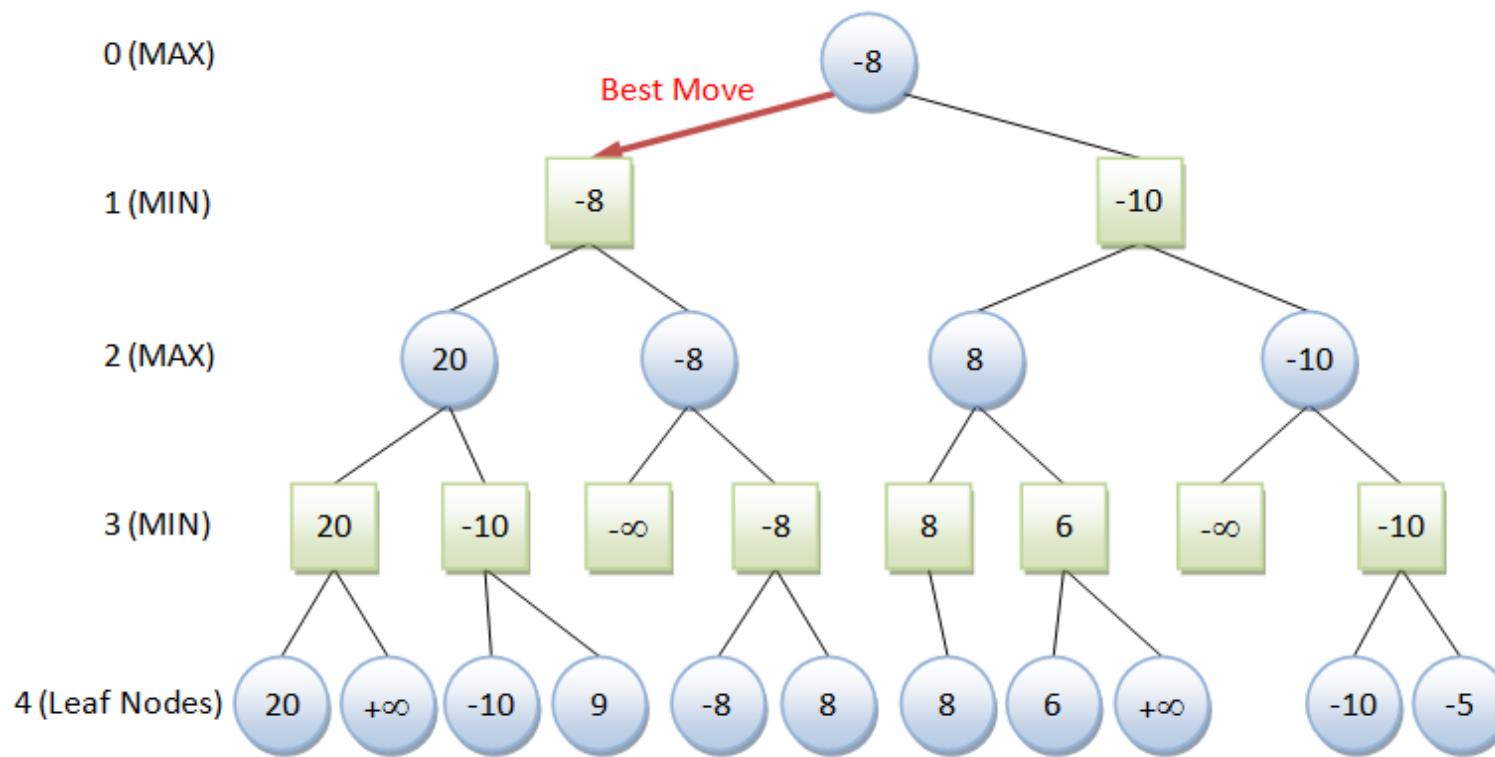


Practical Tips: MiniMax

Max Case Min Case

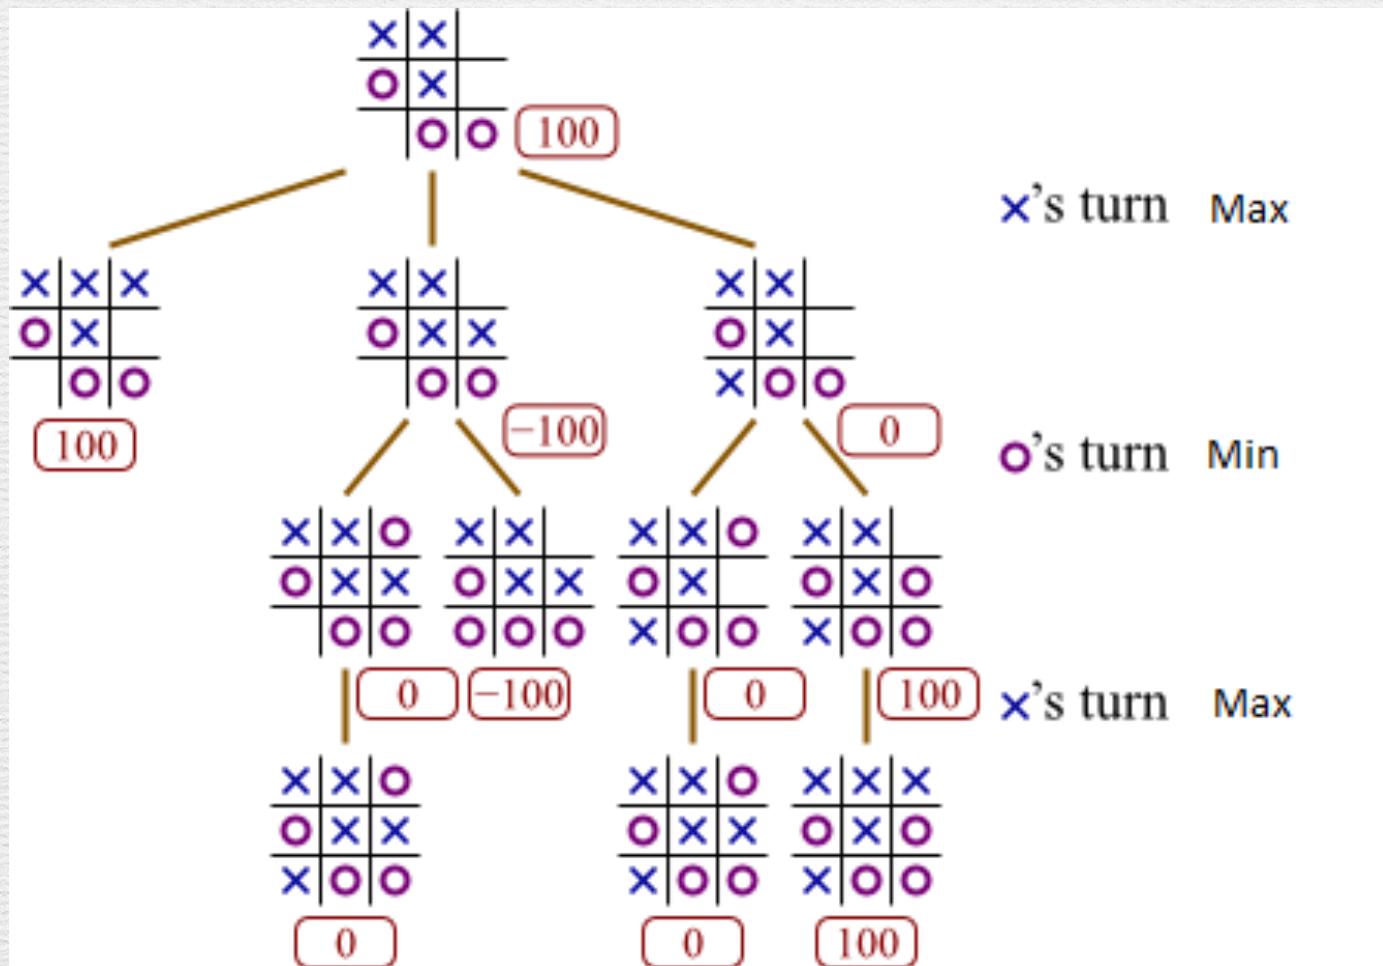
X Player

O Player



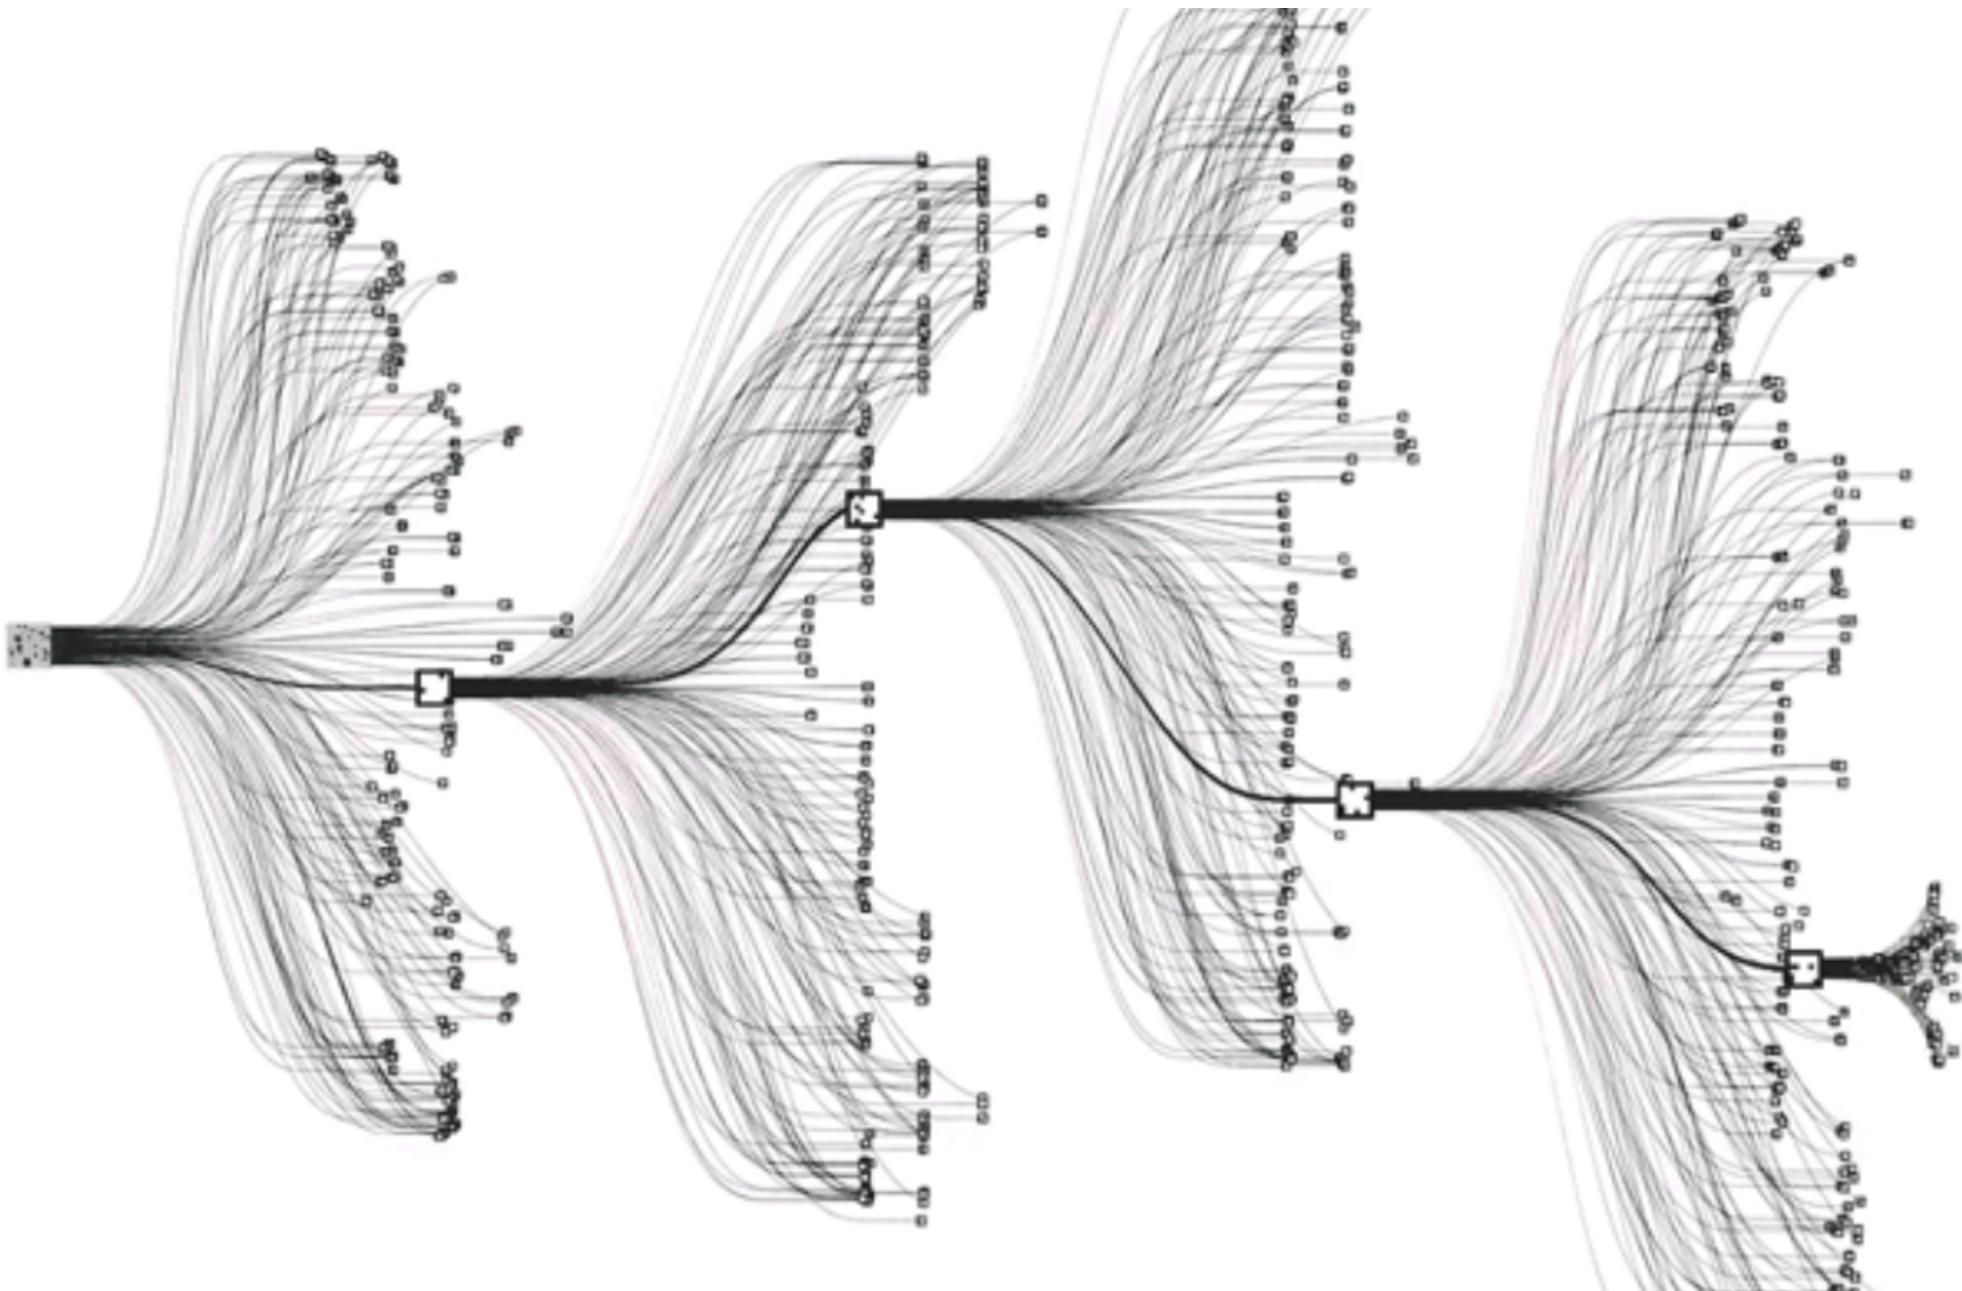
I. Generate the Search Tree

2. use MinMax Search



What is the problem?

The “Tree” of Go Game



The Size of the Tree

Tic Tac Toe:

$$b = 9, d = 9$$

Chess:

$$b = 35, d = 80$$

Go:

$$b = 250, d = 150$$

b : number of legal move per position

d : its depth (game length)

Size of the Tree in GO

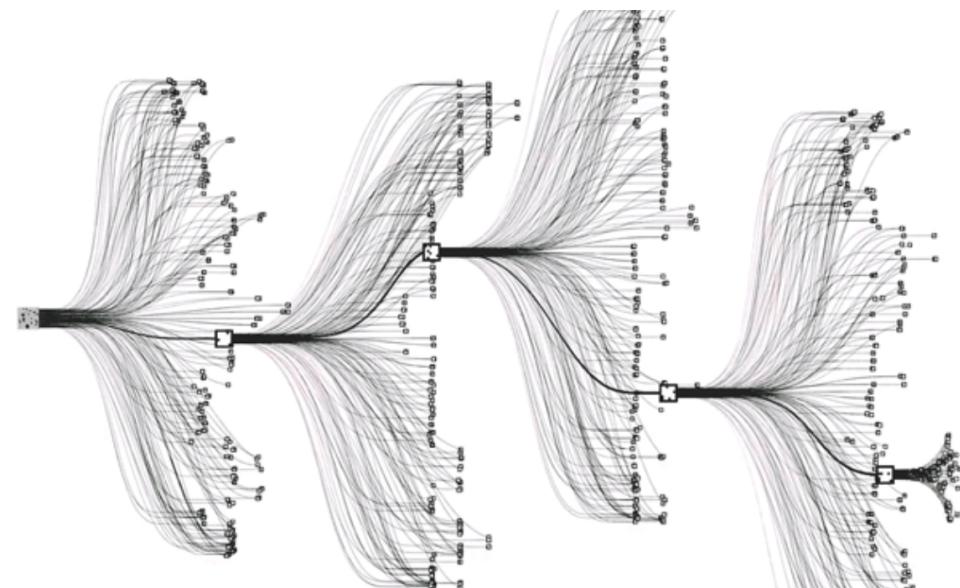
150
250

≈

300
10

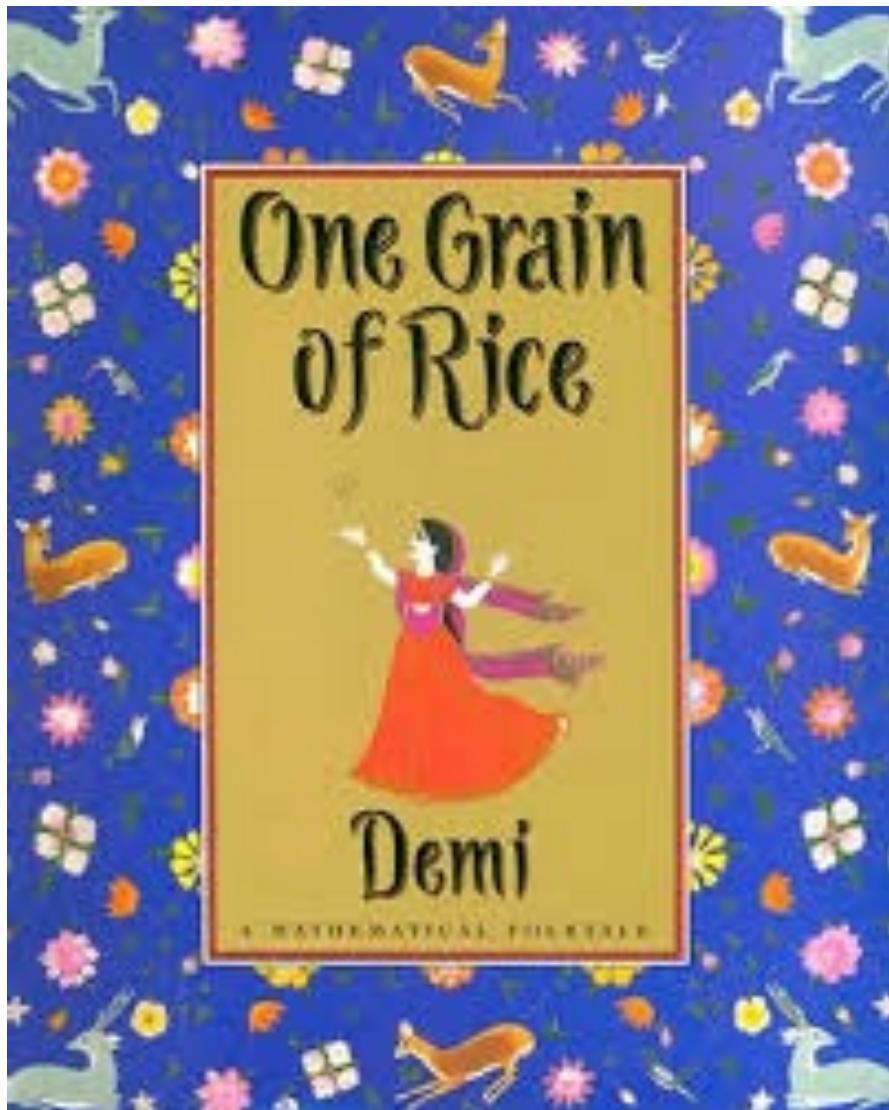
≈

1200
2



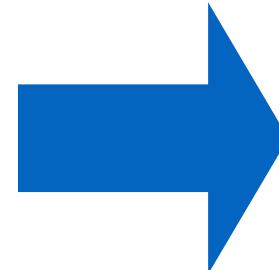
One Grain of Rice

<https://www.youtube.com/watch?v=byk3pA1GPgU>

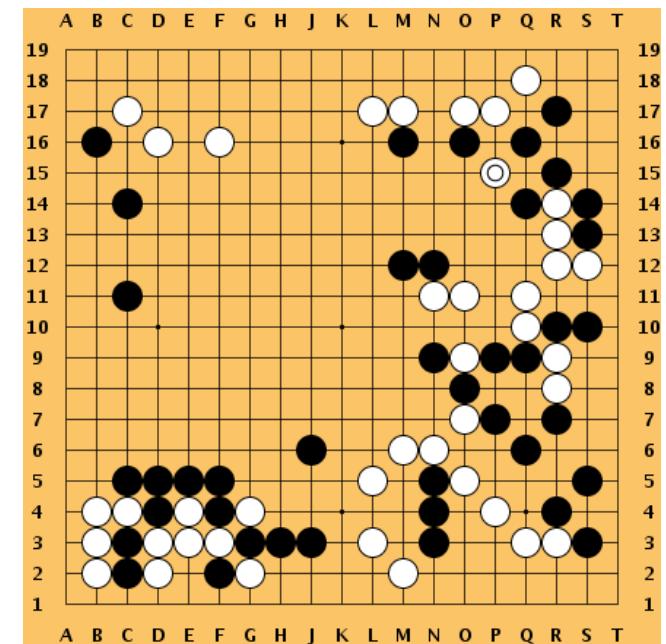


Size of the Rice in Chess

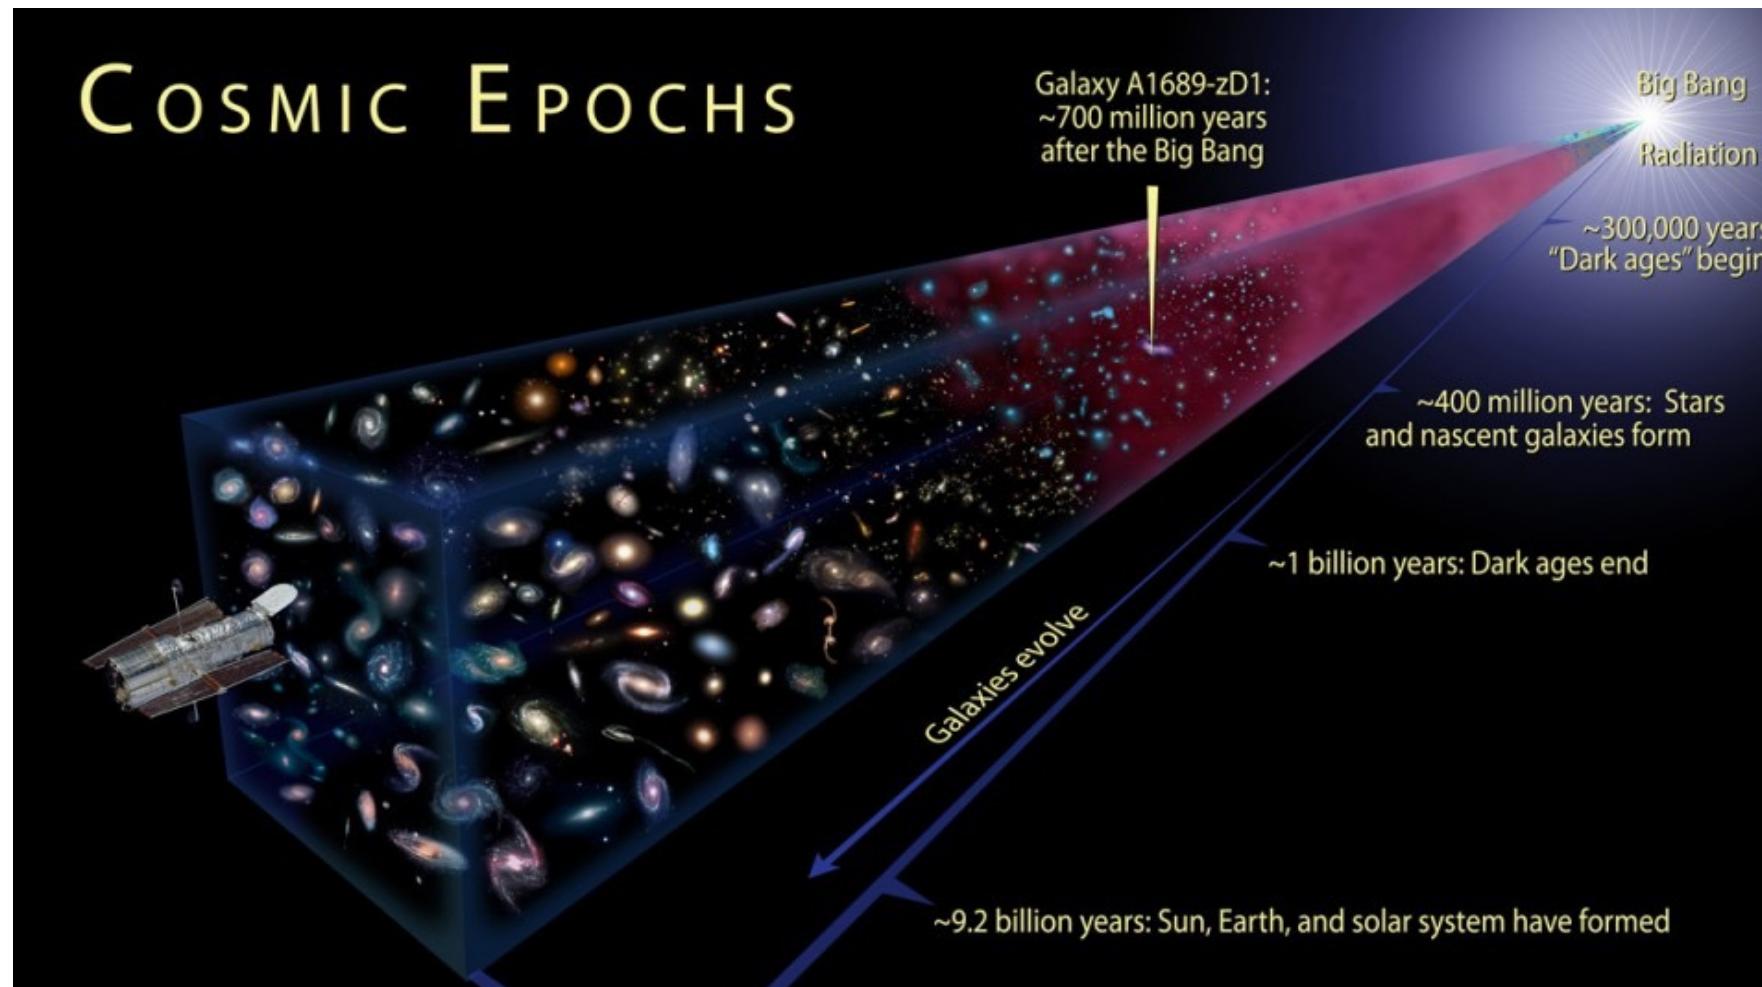
2
64



1200
2



The “Space” of GO Game



between 10^{78} to 10^{82} atoms in the known, observable universe

Numbers

1 trillion dollar

10^{18}

num Atoms in Universe

10^{78} to 10^{82}

num settings in GO game

$250^{150} = 10^{300}$
 $(10^{222} \text{ universes})$

What's the current state of Computing Power?



$$N_A = 6.02 \times 10^{23}$$

Avogadro constant

1 mol of water



Numbers

1 trillion dollar

10^{18}

num of molecules in 1 mol of Water 10^{23}
(computing power)

num Atoms in Universe

10^{78} to 10^{82}

num settings in GO game

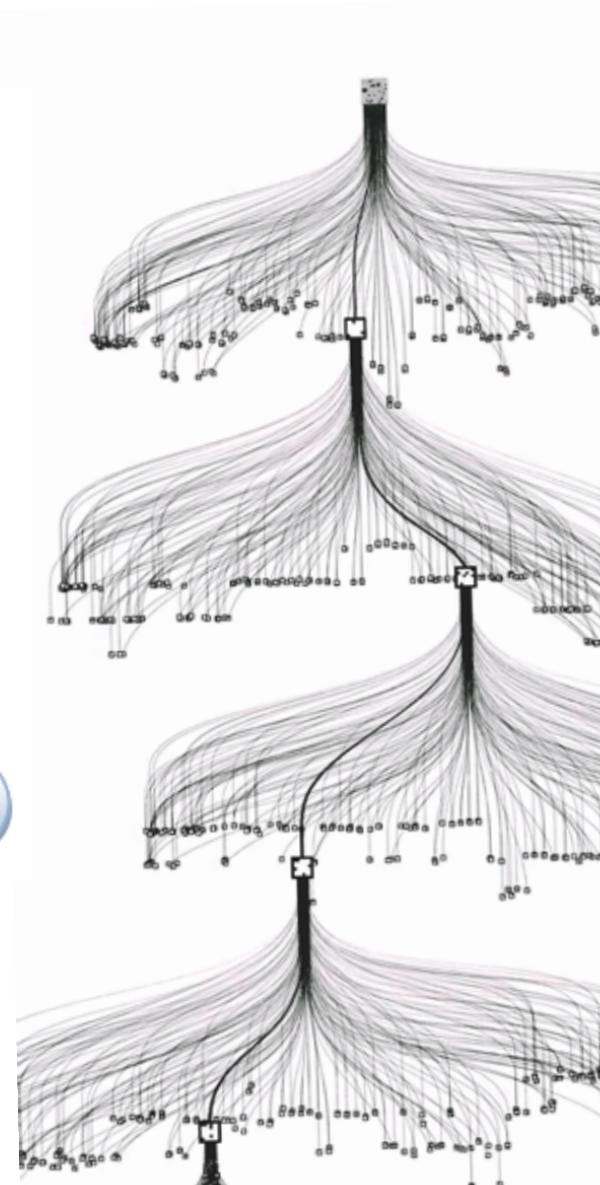
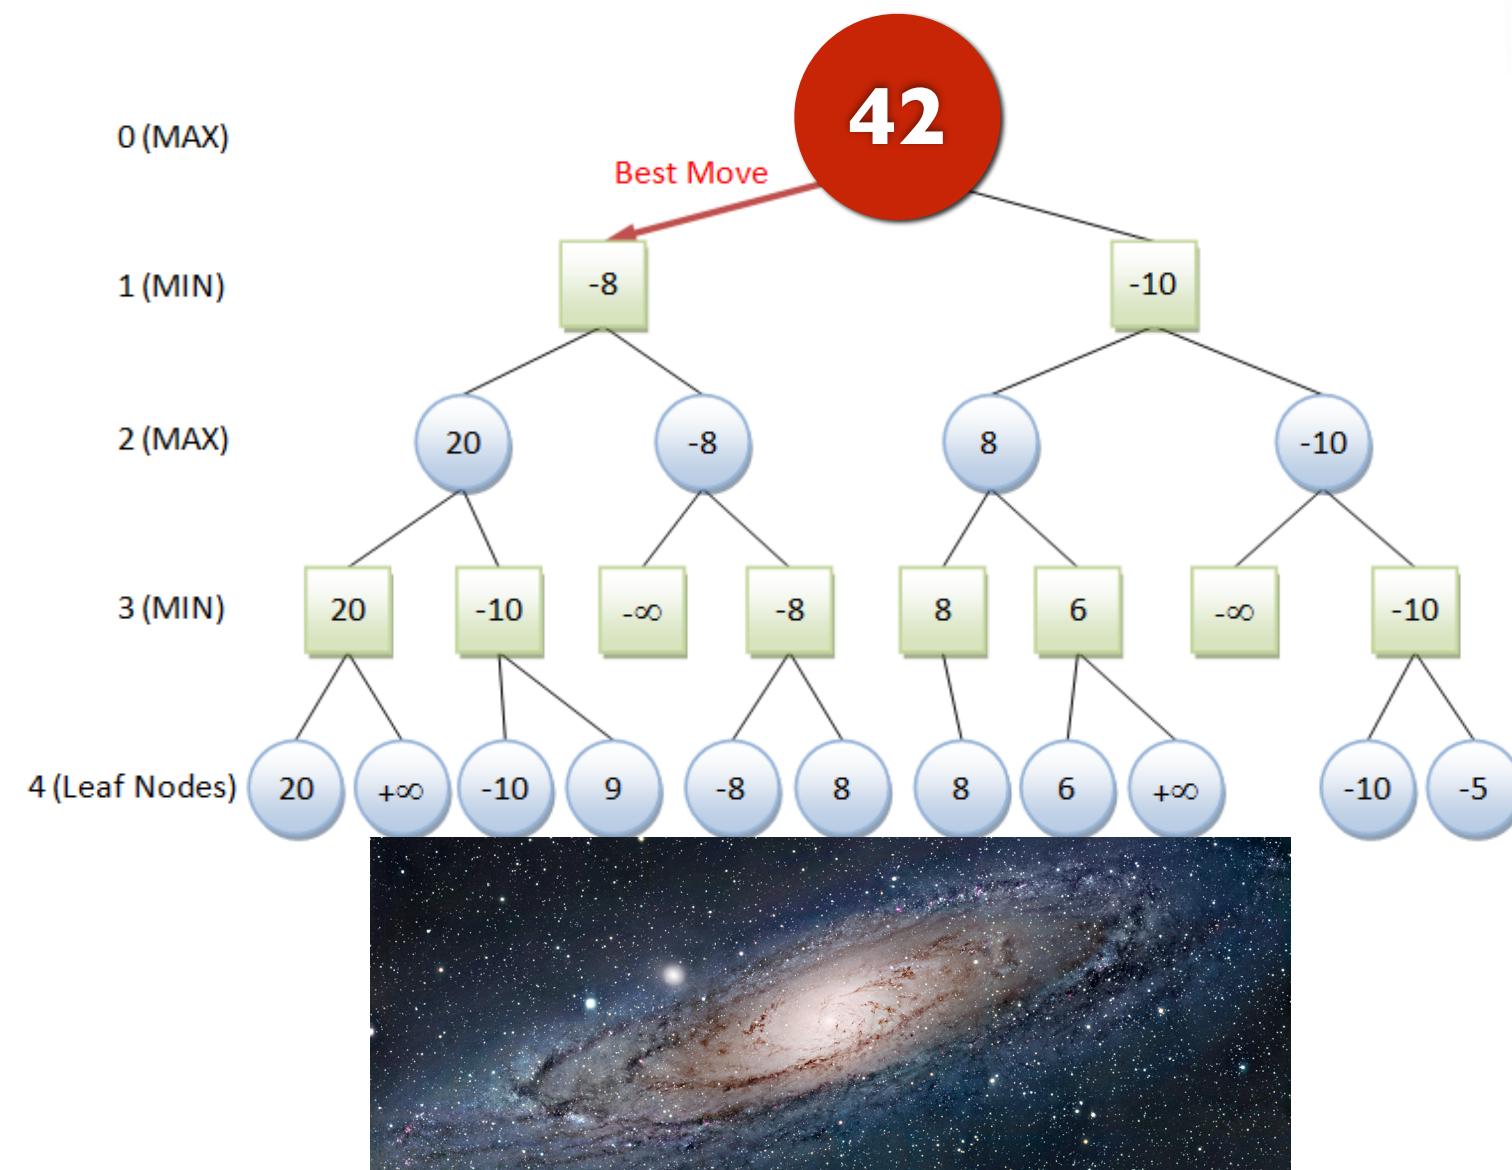
$250^{150} = 10^{300}$
(10^{222} universes)

Computers are also Hopeless to find the ultimate answer of GO



BTW: the ultimate Answer to Life is 42 <https://youtu.be/aboZctrHfK8?t=2s>

imagine that computing **this score** needs **8M** years
with a super computer that has **yet** been invented



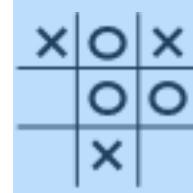
How to solve?

Good News It's all about Winning Human Players

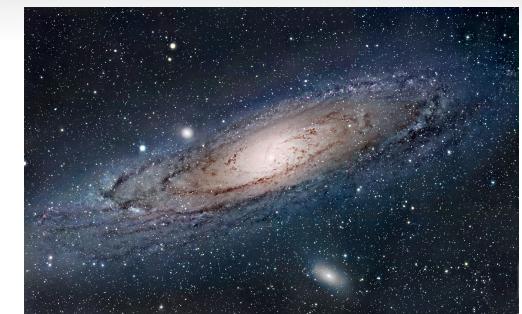
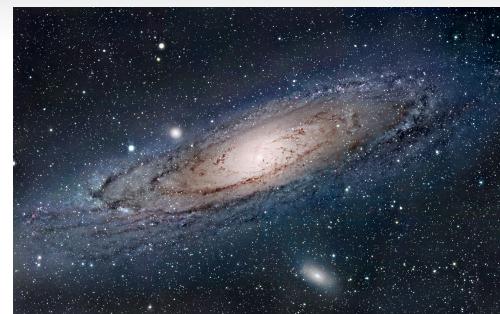
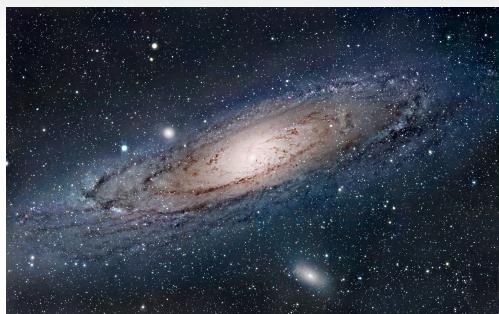
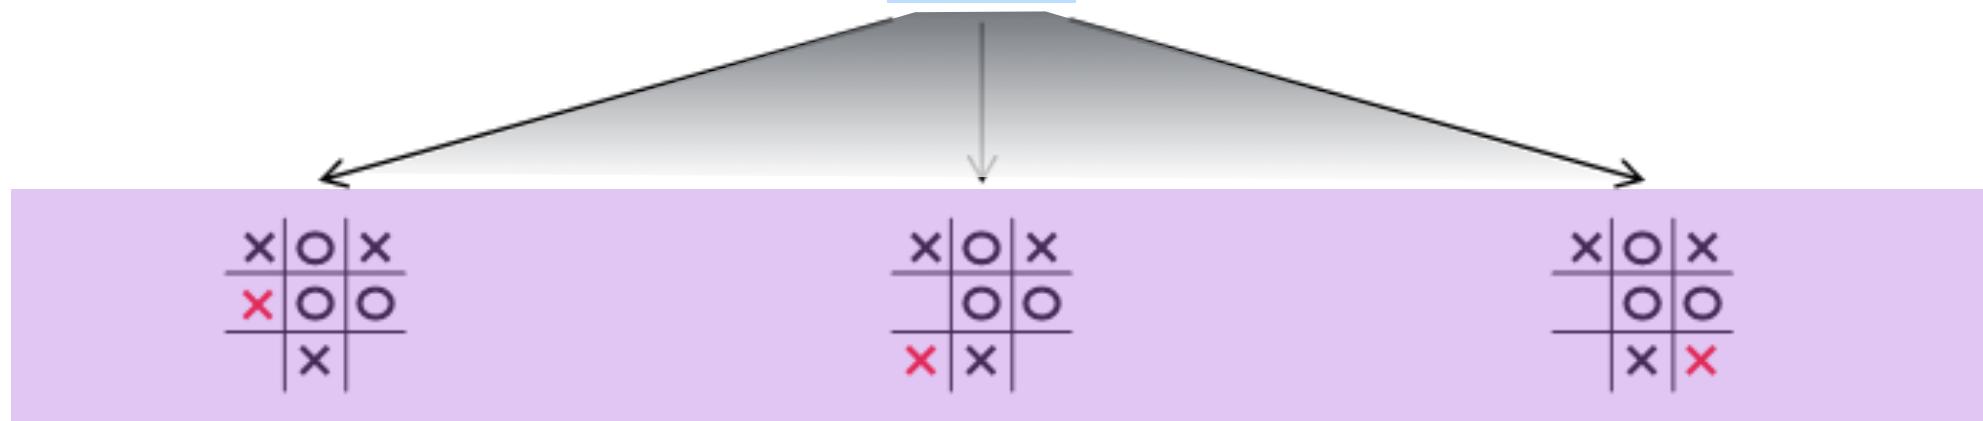


Start with the current state

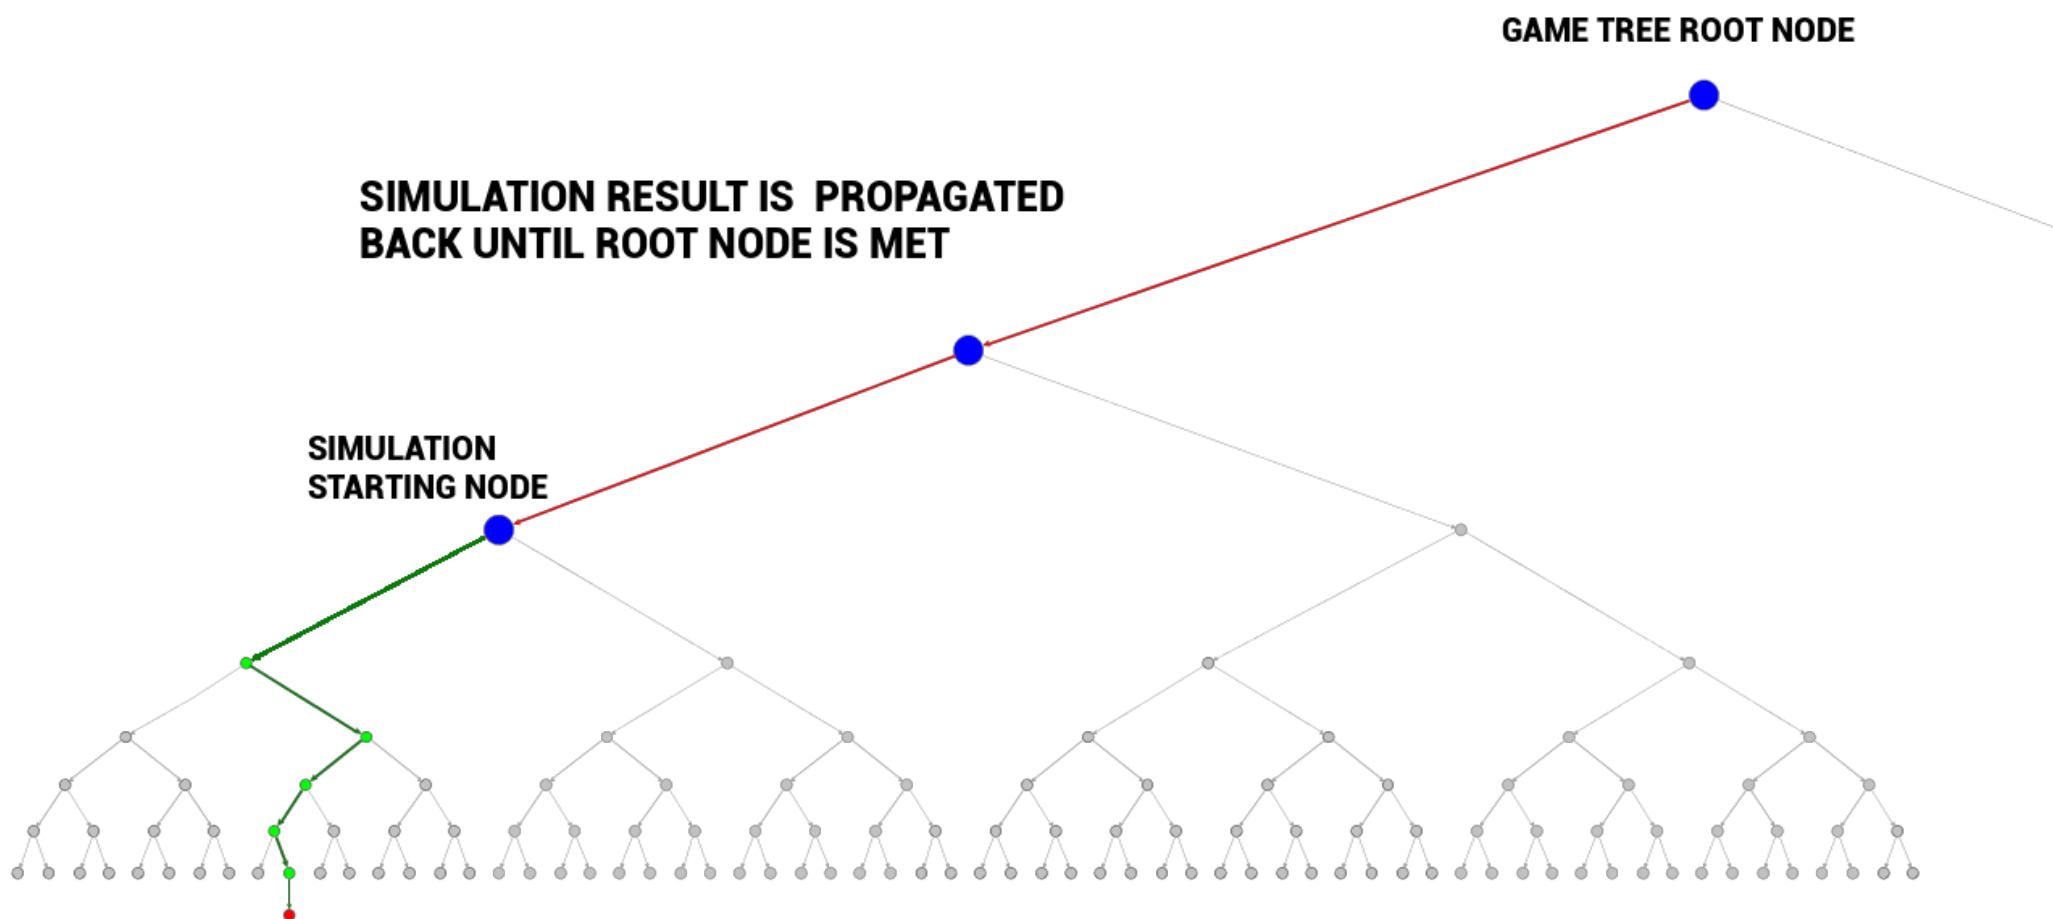
how to estimate the score of each move, without minimax search over each tree



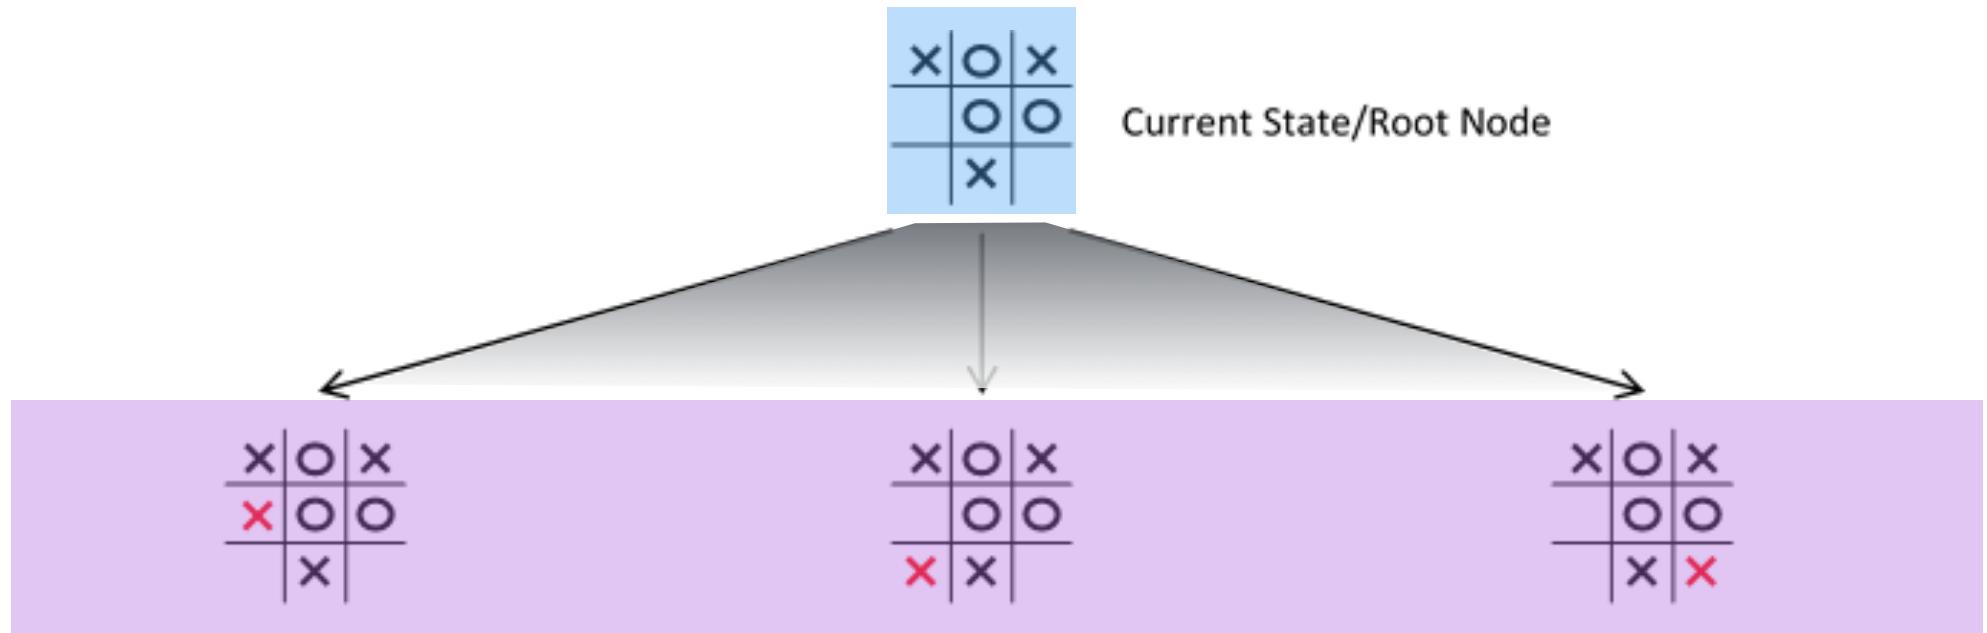
Current State/Root Node



Monte Carlo Tree Search (2005)

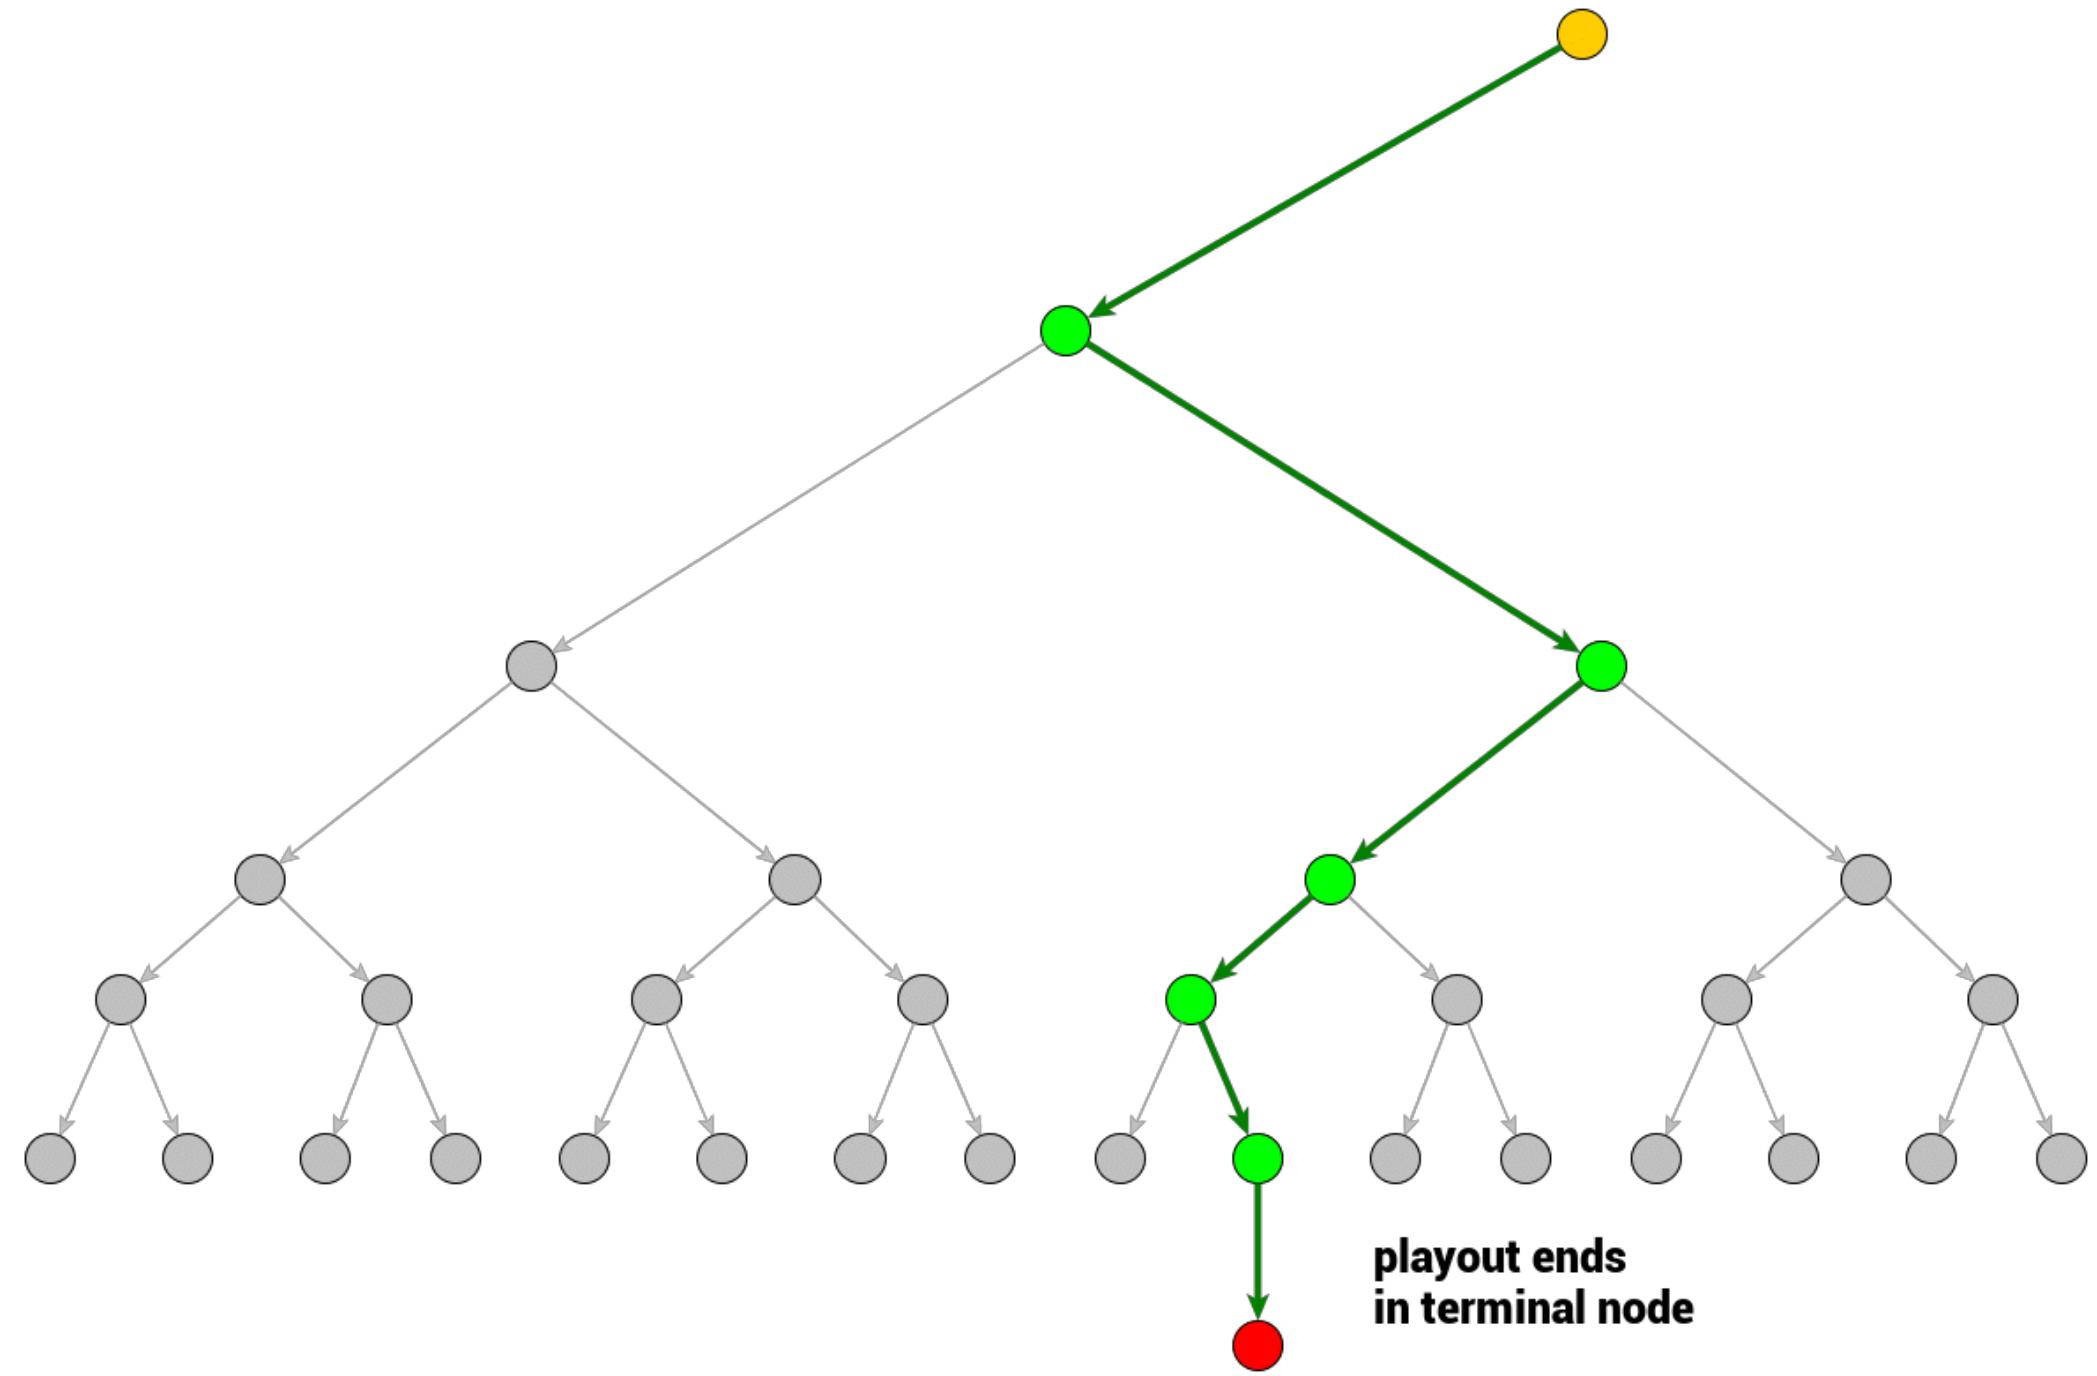


Sampling



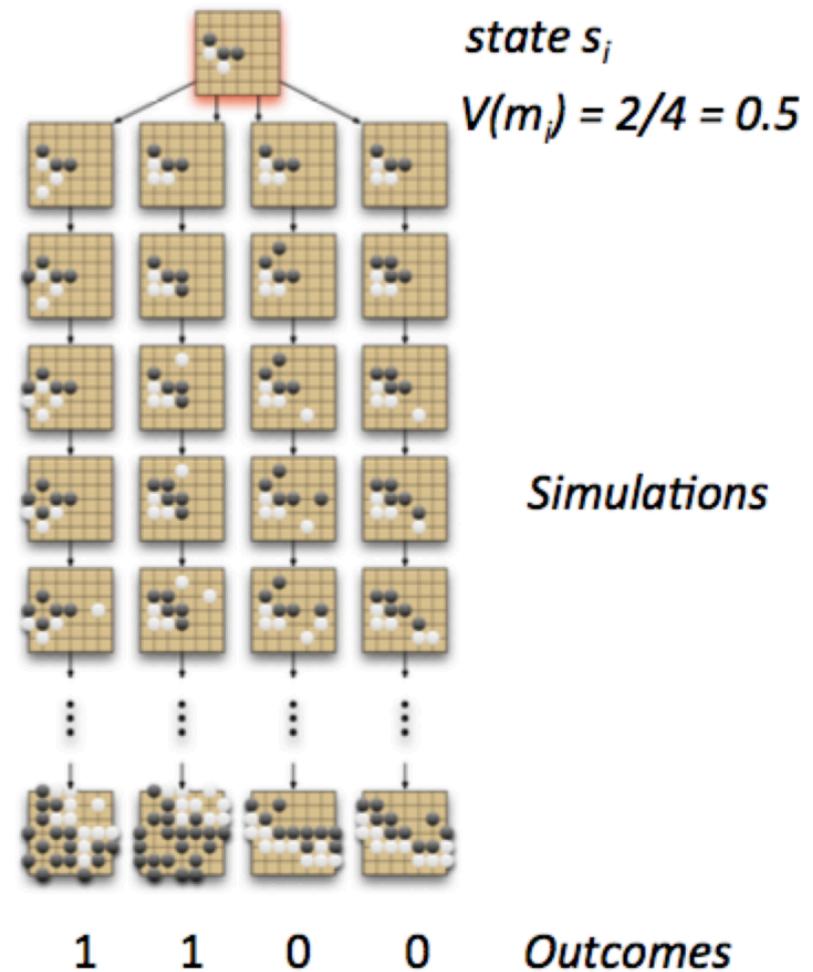
Get samples from the tree and estimate the winning probability

Monte Carlo Simulation

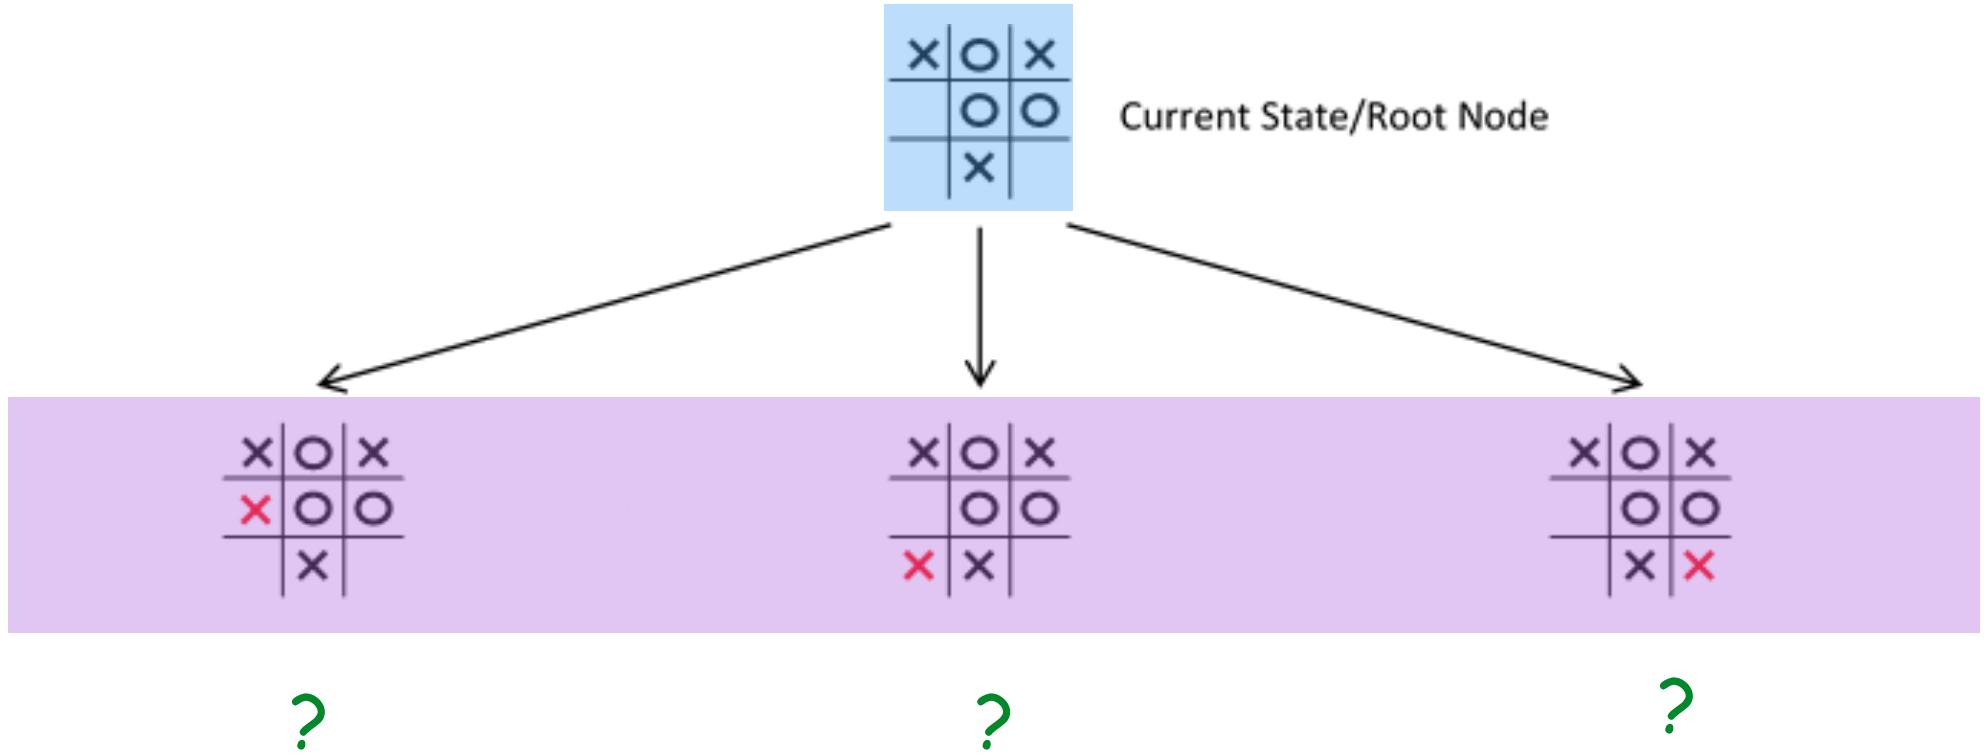


Basic Idea

- No evaluation function?
 - Simulate game using random moves
 - Score game at the end, keep winning statistics
 - Play move with best winning percentage
 - Repeat



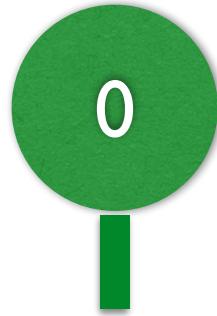
A Simpler Problem



each choice of move has a probability to win
how choose to maximize the pay-off in the long term

Slot Machine

an arm to pull



The machine has a probability to win say 0.4

When a player pull the arm, return reward 1 or 0 randomly according to the probability of winning

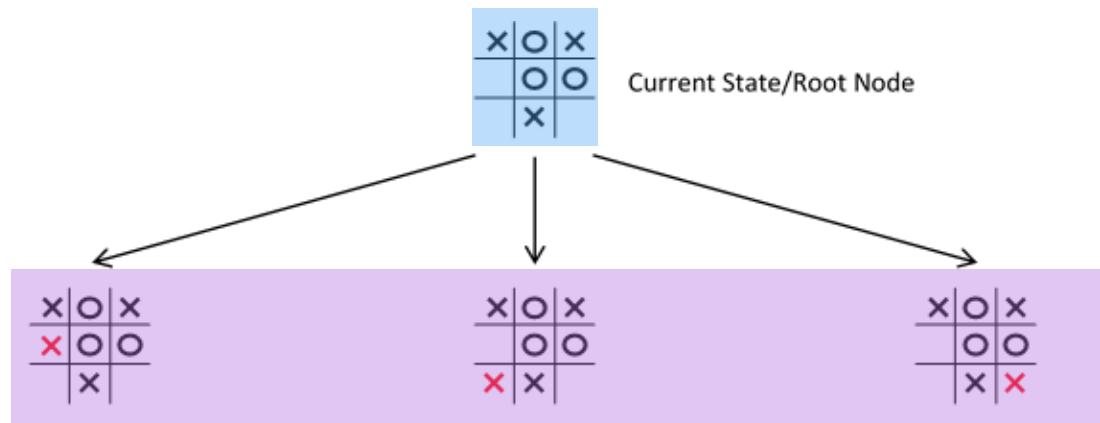


https://www.youtube.com/watch?v=9qedY_UjsIY

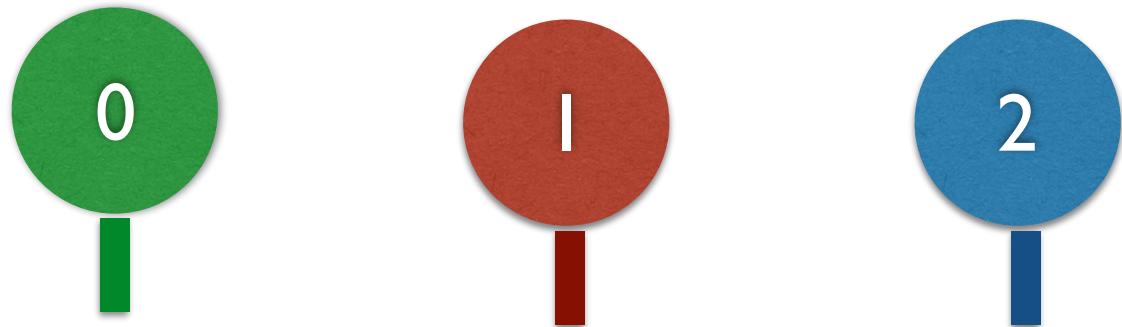
<https://www.youtube.com/watch?v=LPDI3uuRWCg>



Multi-Armed Bandit Problem



n arms to pull



each arm has a probability to win

$$p_0 = 0.8$$

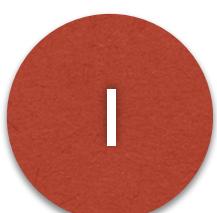
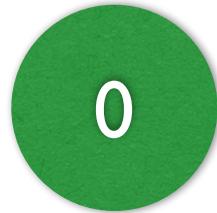
$$p_1 = 0.5$$

$$p_2 = 0.2$$

When a player pull an arm, return reward 1 or 0 randomly according to the probability of winning

Multi-Armed Bandit Problem

n arms to pull

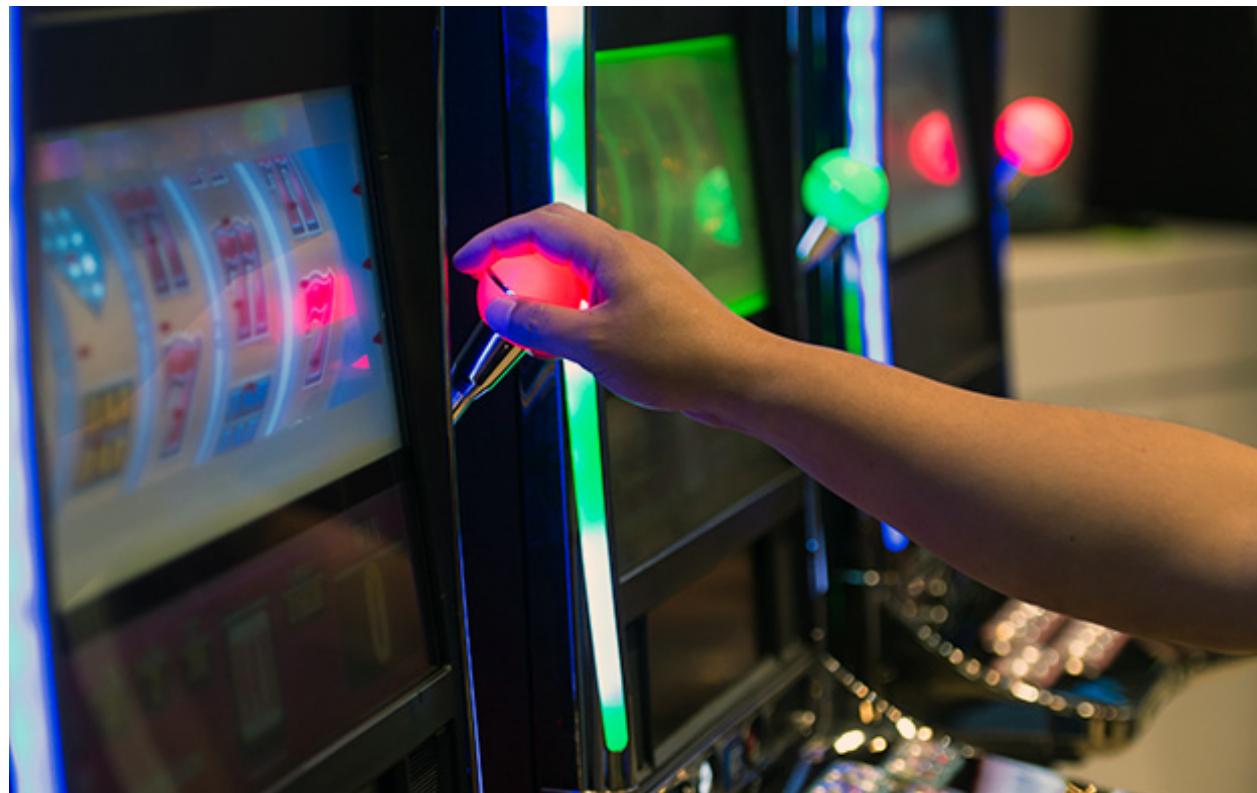


each arm has a probability to win (hidden from player)

?

?

?



Multi-Armed Bandit Problem

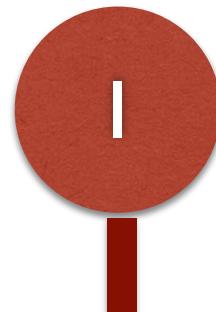
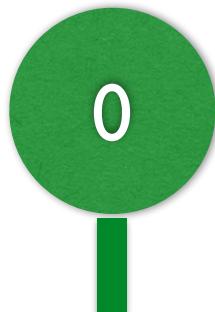


- Assumptions
 - Choice of several arms
 - Each arm pull is independent of other pulls
 - Each arm has fixed, unknown average payoff
- Which arm has the best average payoff?

How to Solve ?

Suppose p is known

n arms to pull



each arm has a probability to win

$$p_0 = 0.5$$

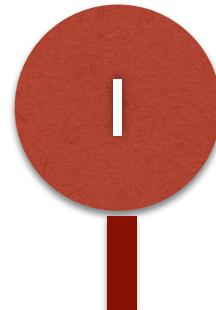
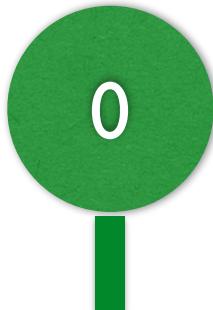
$$p_1 = 0.7$$

$$p_2 = 0.2$$

arm to pull = $\text{argmax}_i p_i$

What if p is unknown

n arms to pull



each arm has a probability to win

$$p_0 = 0.5$$

$$p_1 = 0.7$$

$$p_2 = 0.2$$

Player's estimation of winning probability:

$$Q_0 = 0.4$$

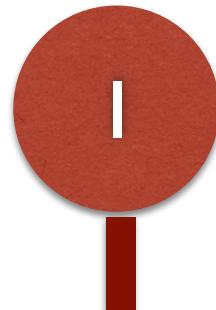
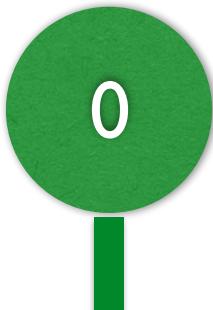
$$Q_1 = 0.6$$

$$Q_2 = 0.3$$

Each estimation is based upon the ratio
of winning after pulling each arm

Strategy 1: choose the arm with the largest estimated reward

n arms to pull



Player's estimation of winning probability:

$$Q_0 = 0.4$$

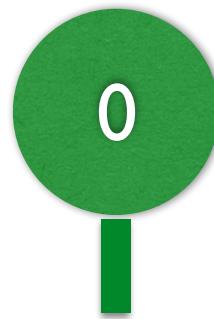
$$Q_1 = 0.6$$

$$Q_2 = 0.3$$

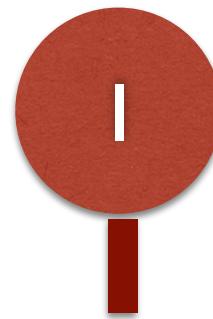
arm to pull = $\text{argmax}_i Q_i$

Problem

How to initialize Q?



$$Q_0 = 0$$



$$Q_1 = 0$$

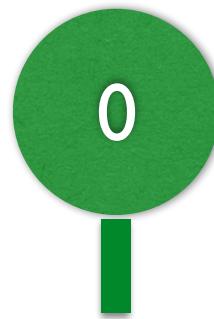


$$Q_2 = 0$$

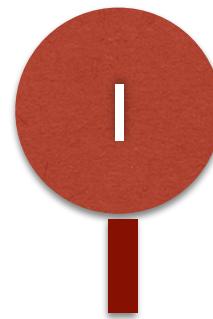
Pull arm 0

Problem

How to update Q?



$$Q_0 = 0$$



$$Q_1 = 0$$

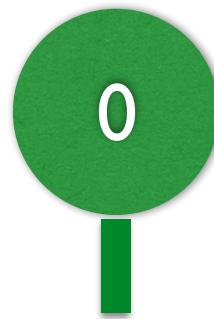


$$Q_2 = 0$$

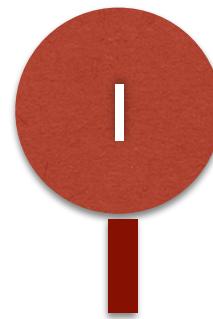
Pull arm 0
reward=1 (Win)

Problem

How to update Q?



$$Q_0 = 1.0$$



$$Q_1 = 0$$

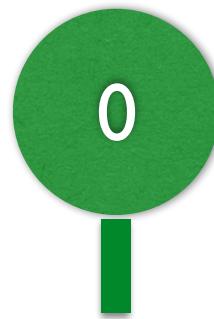


$$Q_2 = 0$$

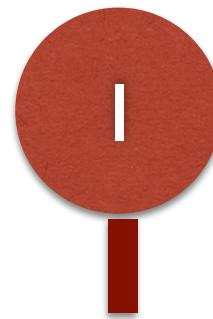
Pull arm 0
reward=1 (Win)

Problem

How to update Q?



$$Q_0 = 1.0$$



$$Q_1 = 0$$

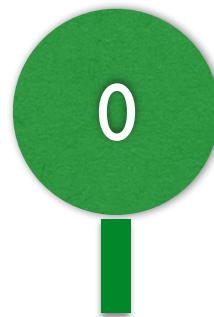


$$Q_2 = 0$$

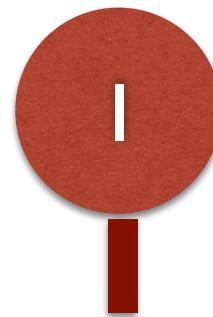
Pull arm 0
reward=1 (Win)
Pull arm 0
reward=0 (Lose)

Problem

How to update Q?



$$Q_0 = 0.5$$



$$Q_1 = 0$$



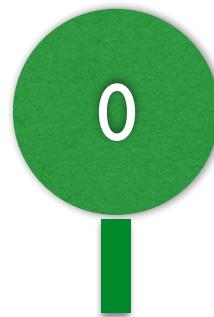
$$Q_2 = 0$$

Pull arm 0
reward=1 (Win)

Pull arm 0
reward=0 (Lose)
Pull arm 0

Problem

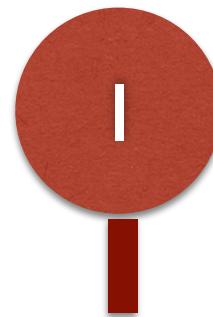
How to update Q?



$$Q_0 = 0.67$$

Pull arm 0
reward=1 (Win)

Pull arm 0
reward=0 (Lose)
Pull arm 0

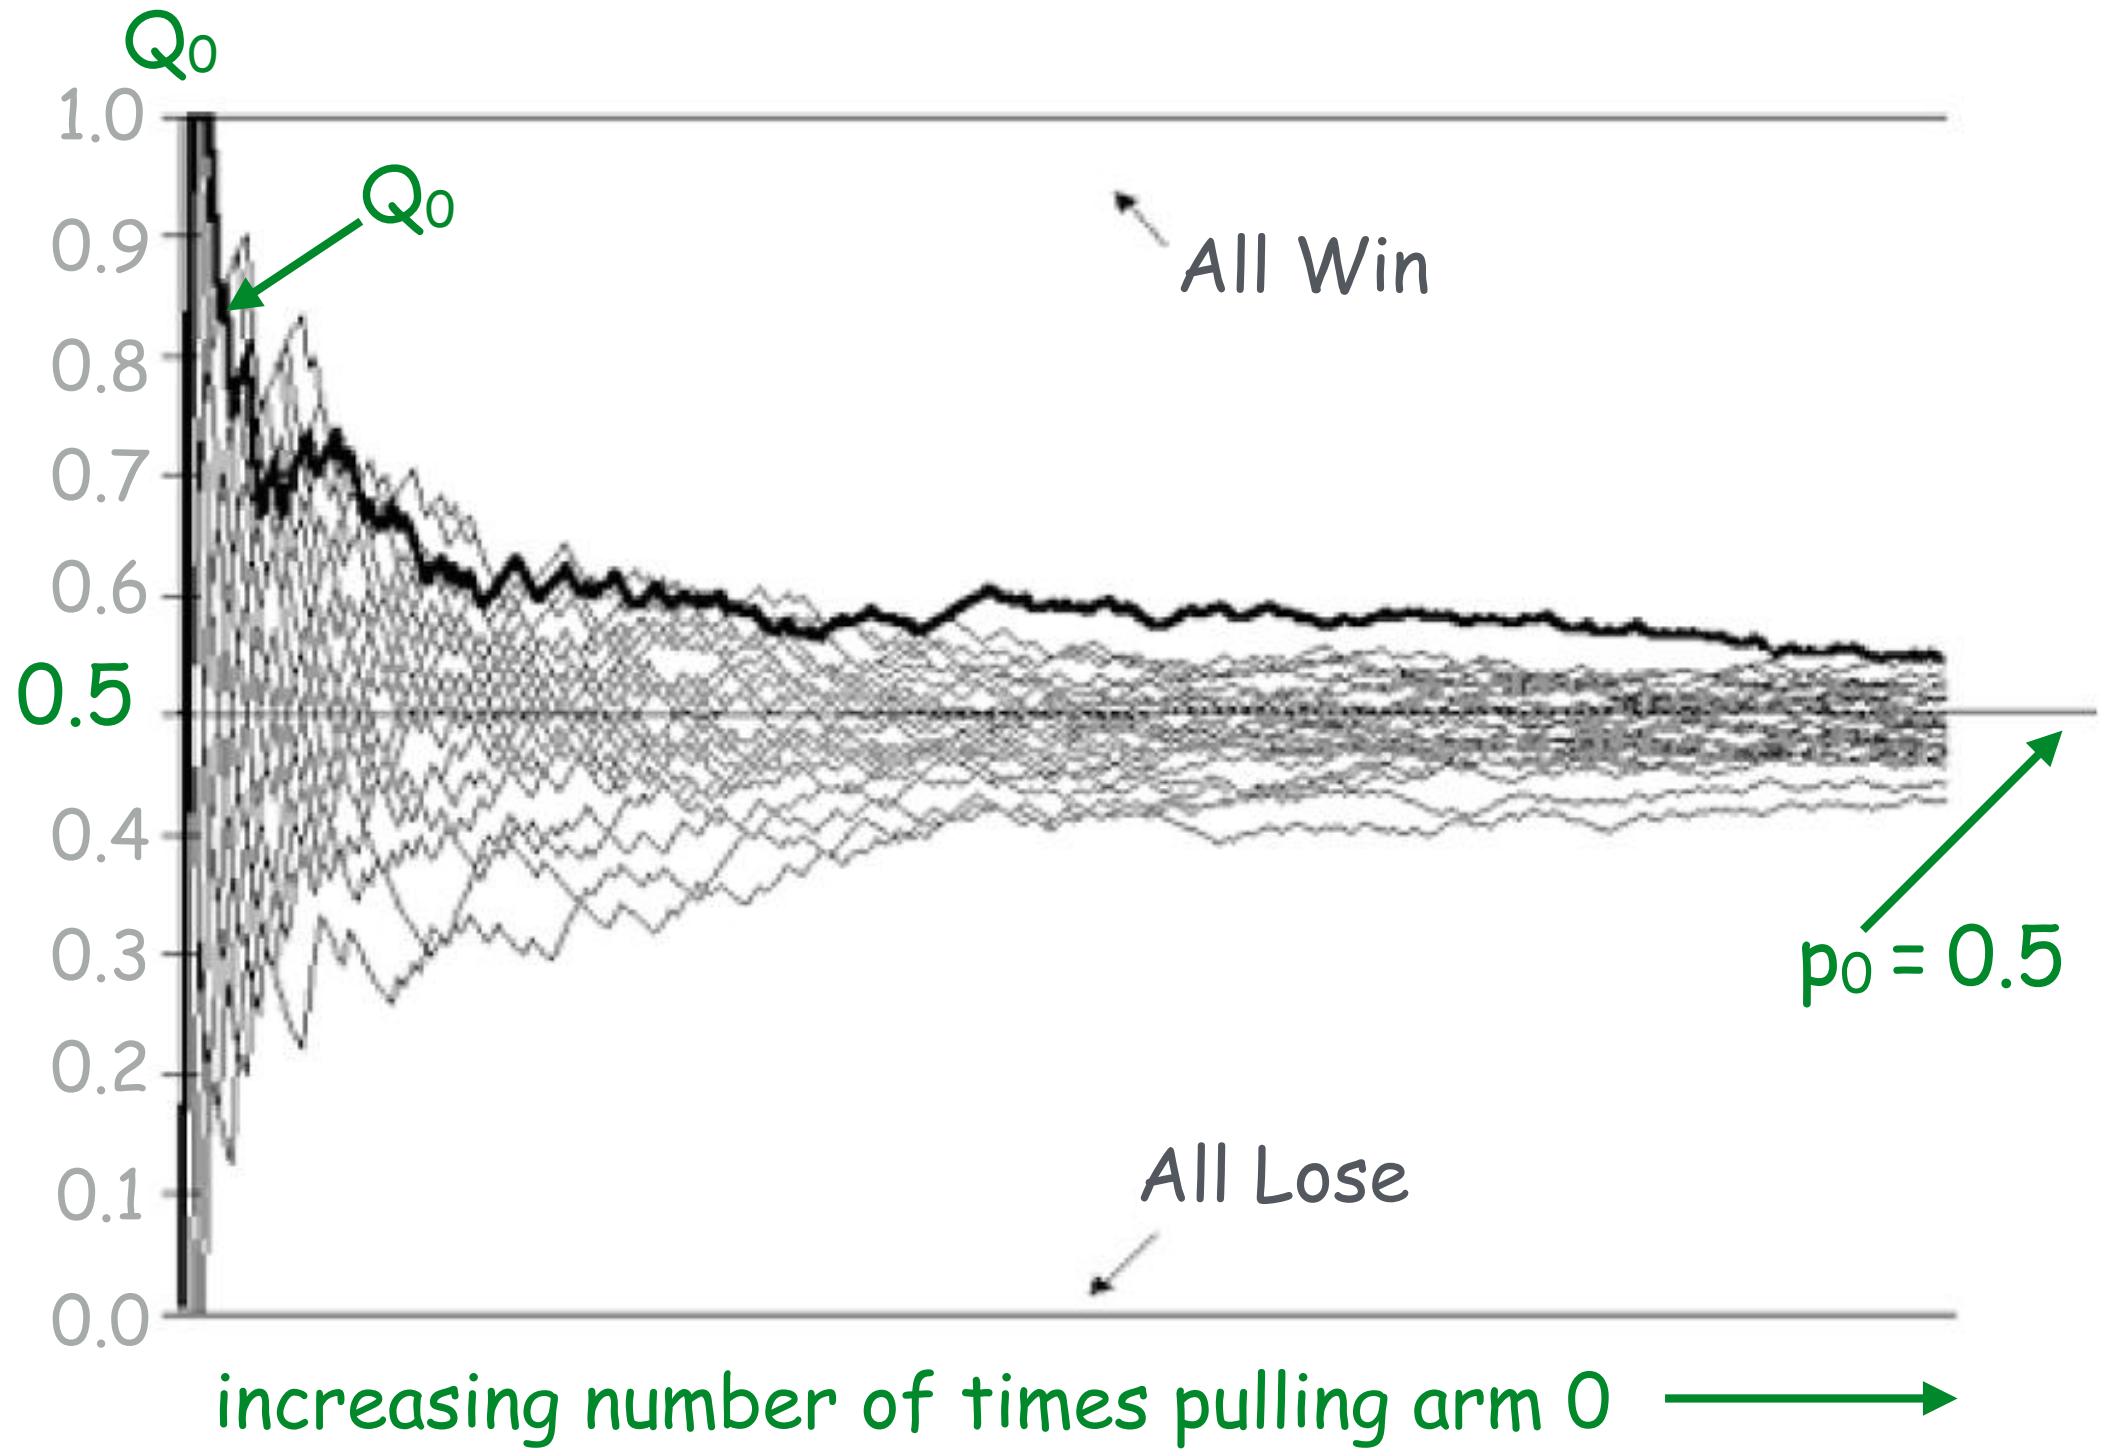


$$Q_1 = 0$$

reward=1 (Win)

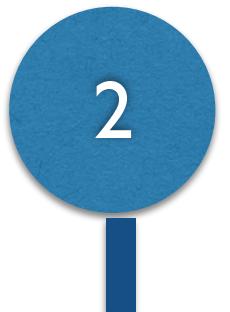
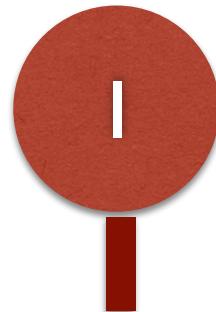
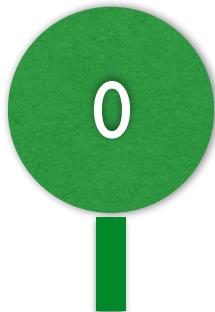


$$Q_2 = 0$$



Strategy 2: choose an arm randomly

n arms to pull



probability to pull: 1/3

1/3

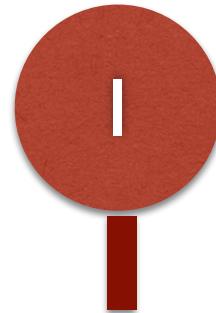
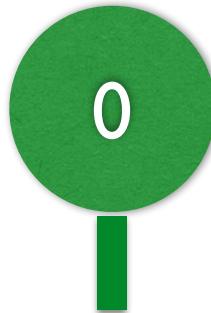
1/3

arm to pull = random

ignoring Q values

Pros and Cons

n arms to pull



each arm has a probability to win

$$p_0 = 0.5$$

$$p_1 = 0.7$$

$$p_2 = 0.2$$

Player's estimation of winning probability:

$$Q_0 = 0.5$$

$$Q_1 = 0.7$$

$$Q_2 = 0.2$$

If we pull many times, the estimated probability Q will be very close to p

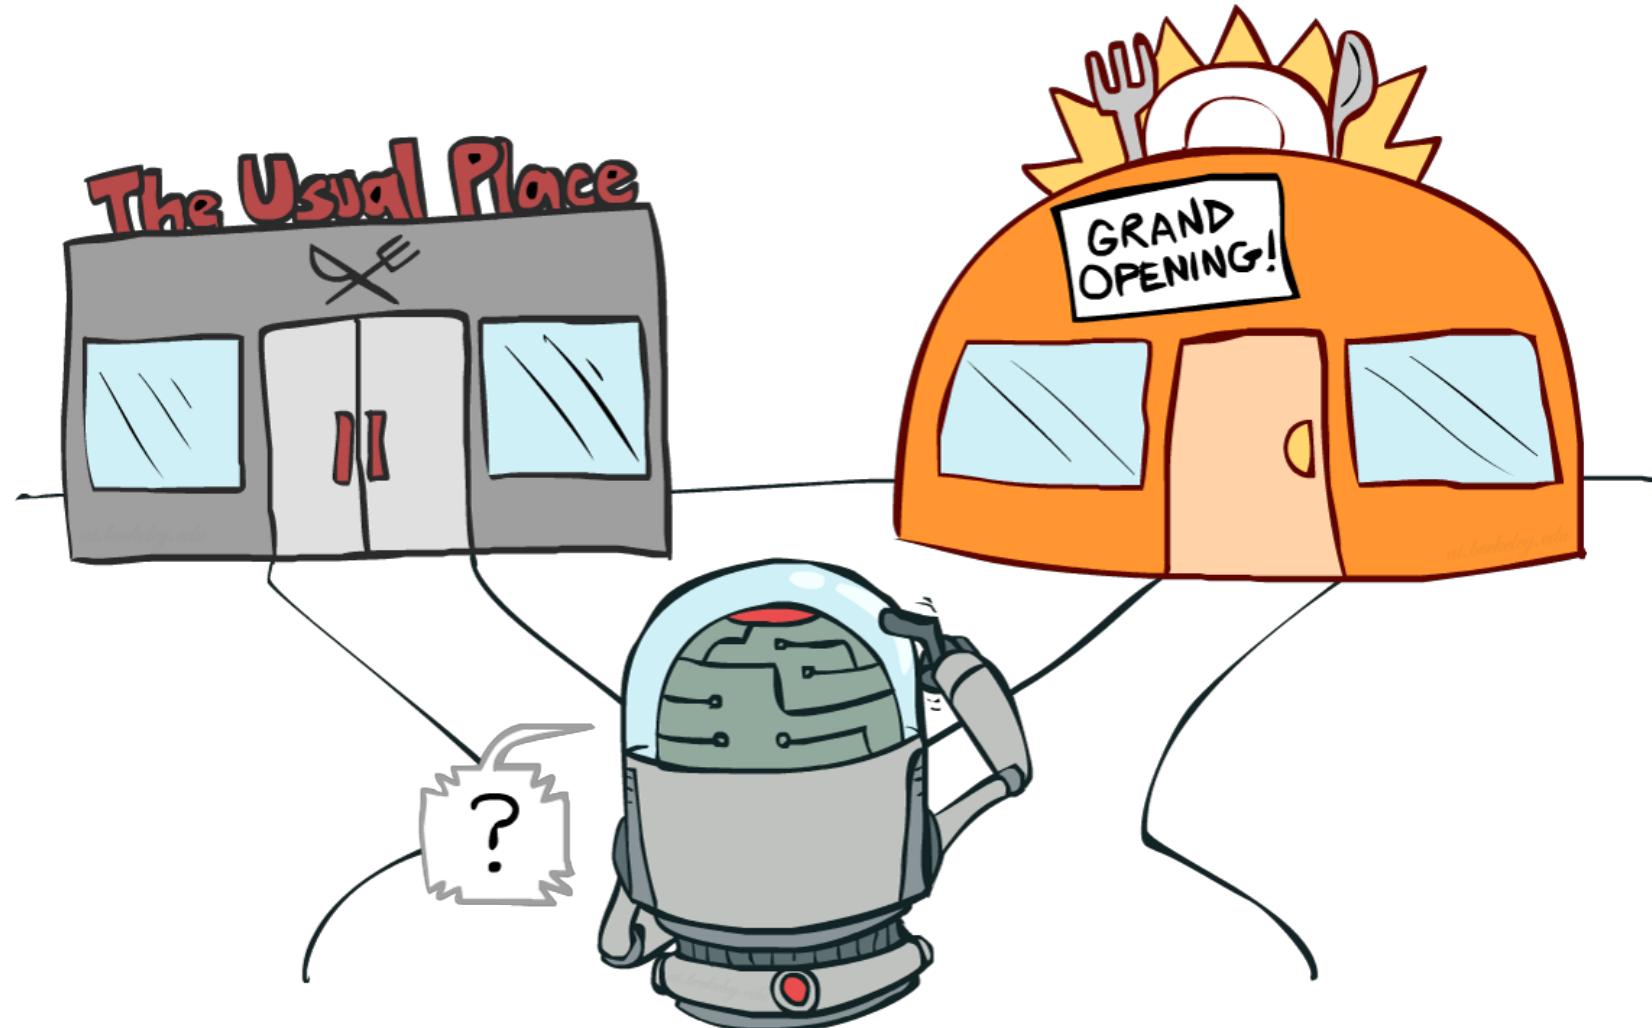
Exploration vs Exploitation

(Exploitation)

$\text{argmax}_i Q_i$

(Exploration)

random



Upper Confidence Bound

- Choose the arm that maximizes formula:

$$v_i + C \times \sqrt{\frac{\ln(N)}{n_i}}$$

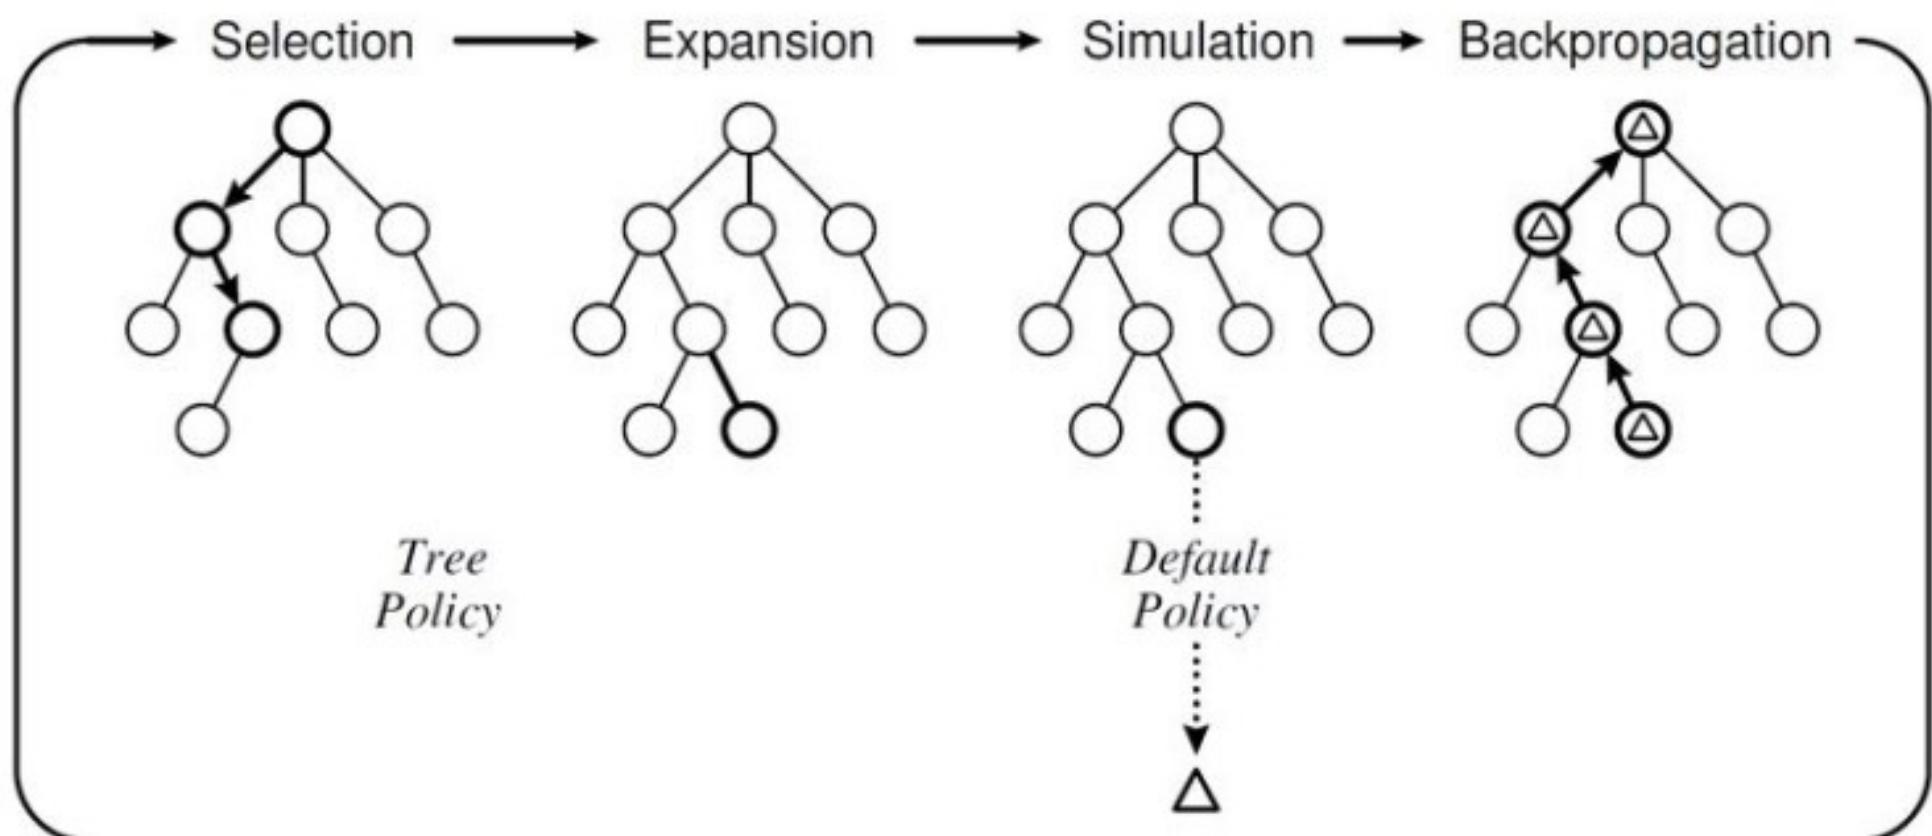
Diagram illustrating the UCB formula components:

- v_i (value estimate) - Prefers higher payoff arm
- C (tunable parameter)
- $\sqrt{\frac{\ln(N)}{n_i}}$ - Prefers less played arm

Annotations:

- v_i is associated with "value estimate".
- C is associated with "tunable parameter".
- $\ln(N)$ is associated with "total number of trials".
- n_i is associated with "num trials for arm i".

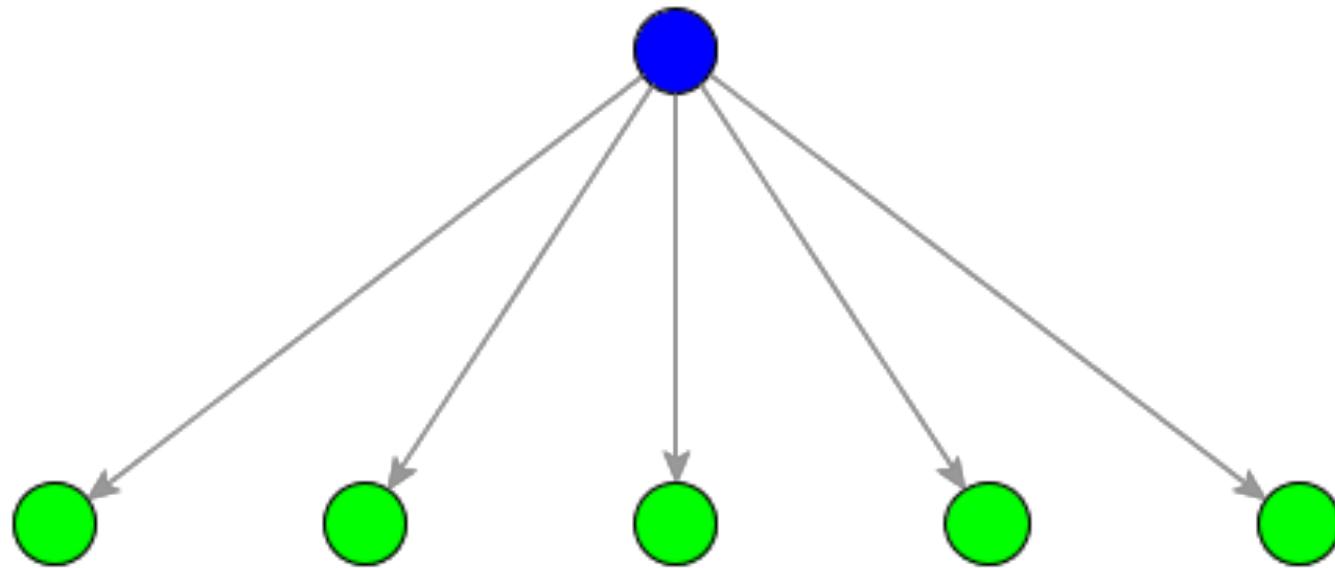
Monte-Carlo Tree-Search



Expanding a Node

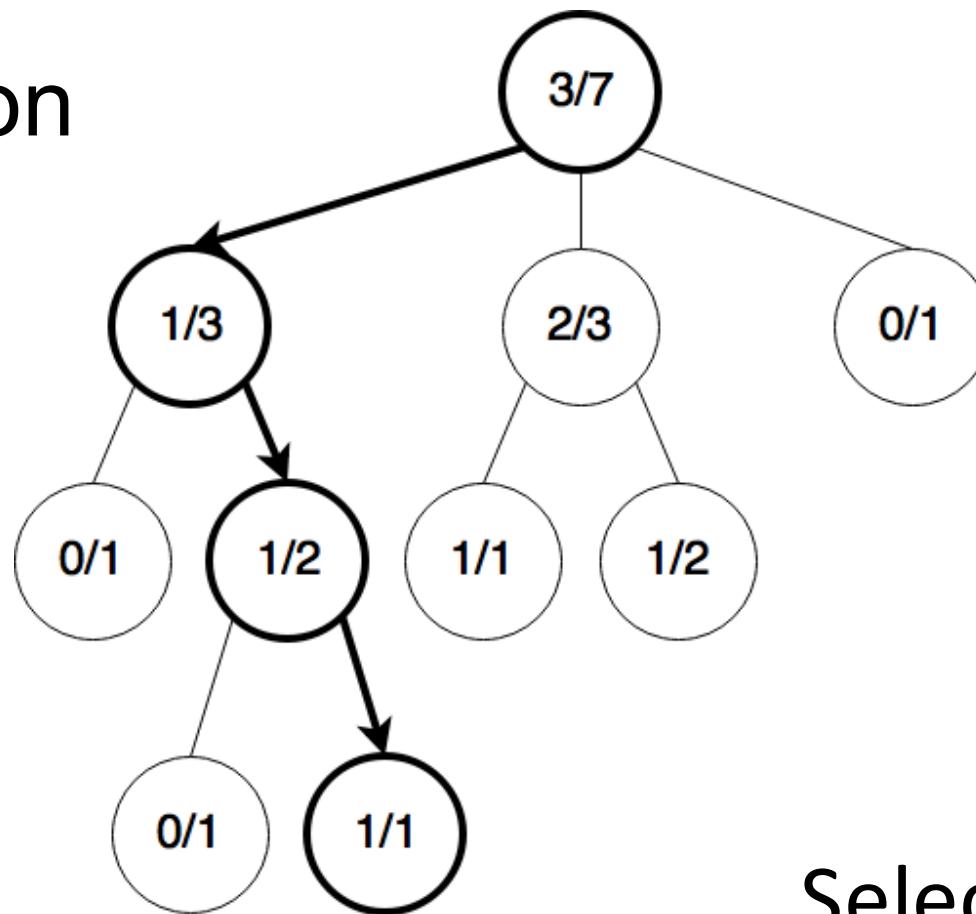


Expanding a Node



Monte Carlo Tree Search

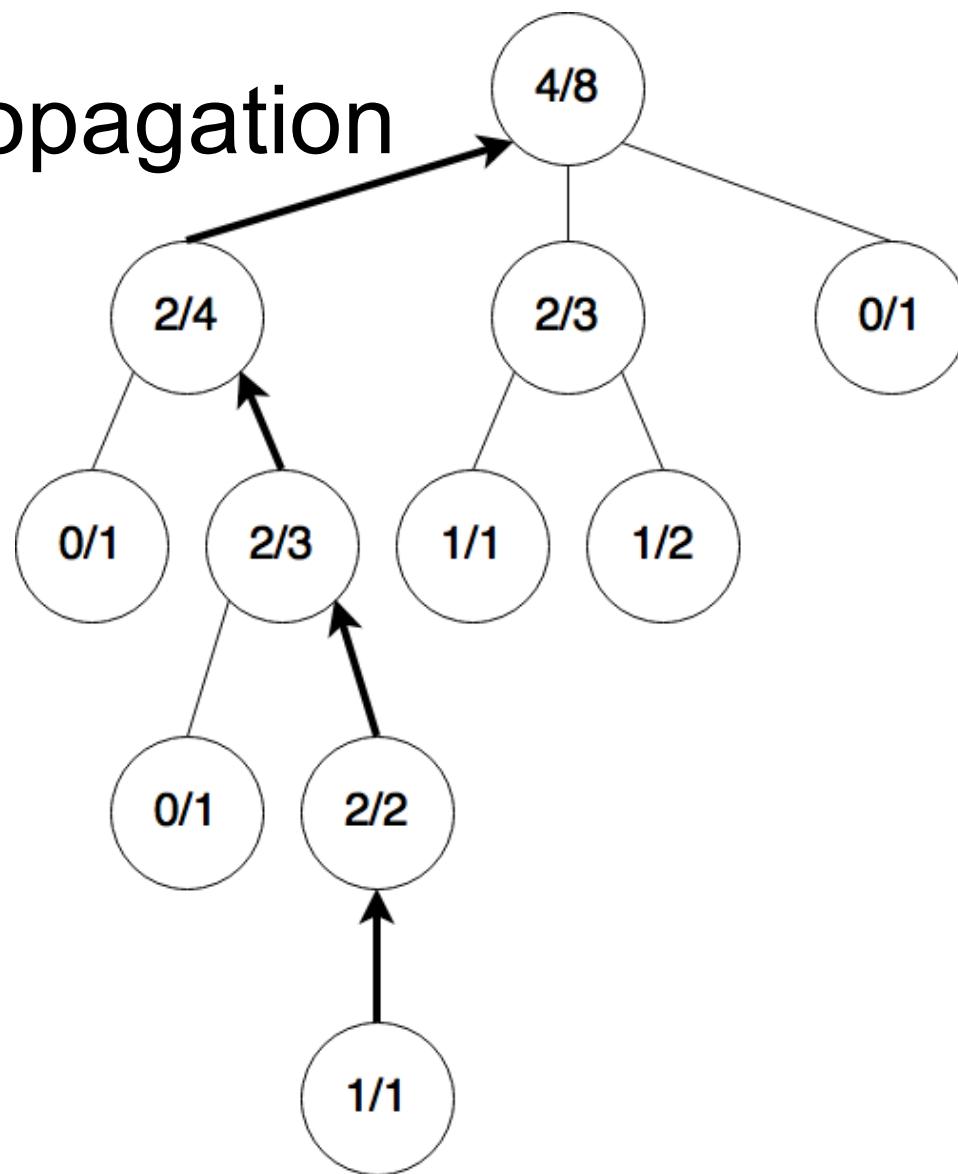
(1) Selection

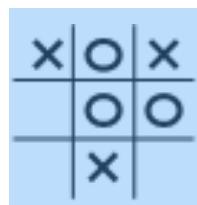


Selection policy is applied recursively until a leaf node is reached

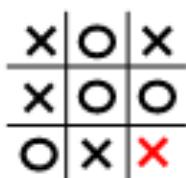
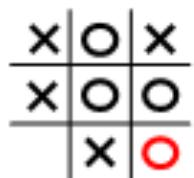
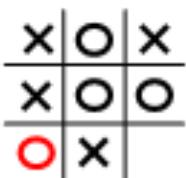
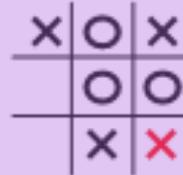
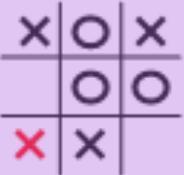
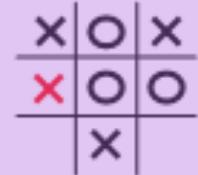
Monte Carlo Tree Search

(4) Backpropagation

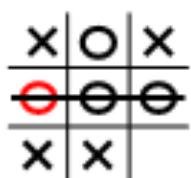




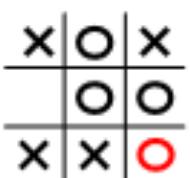
Current State/Root Node



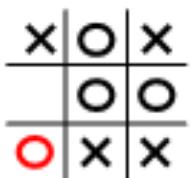
Leaf Node



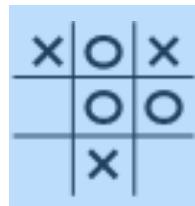
Leaf Node



Leaf Node



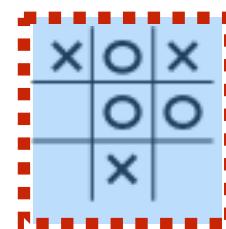
Leaf Node



Current State/Root Node



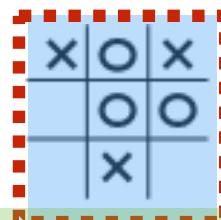
Selection (from root)



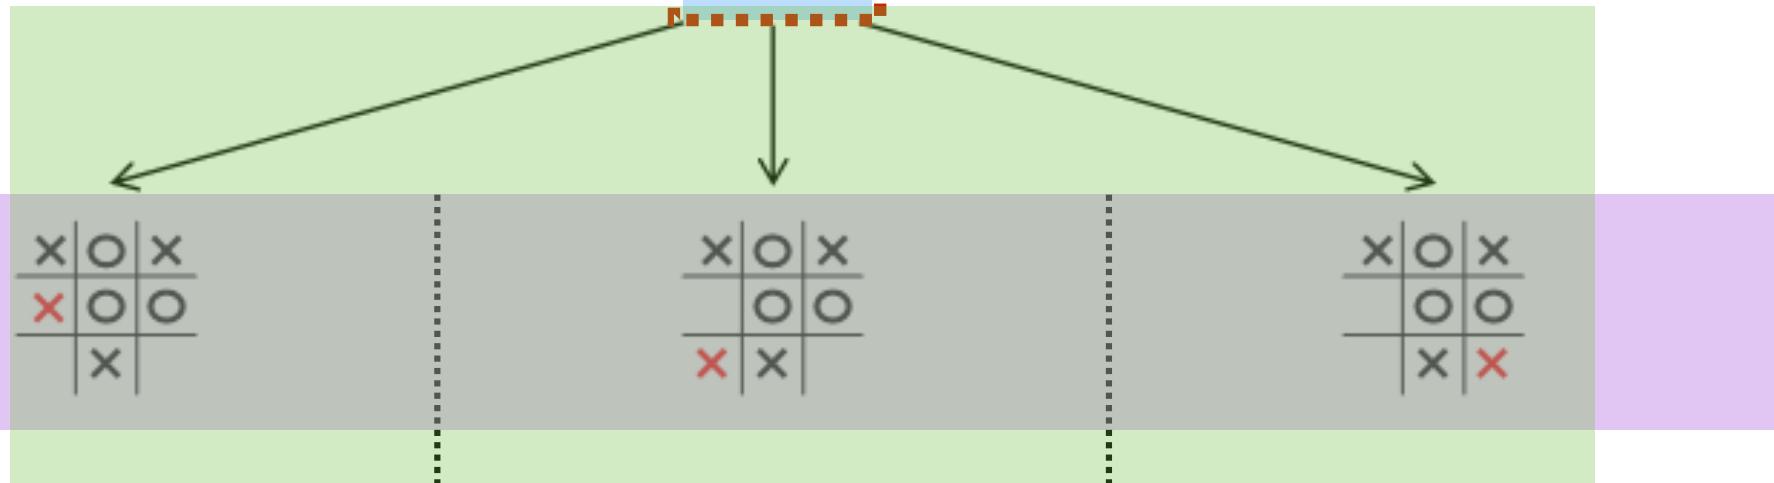
Current State/Root Node



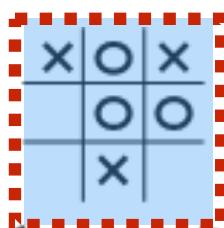
Expand



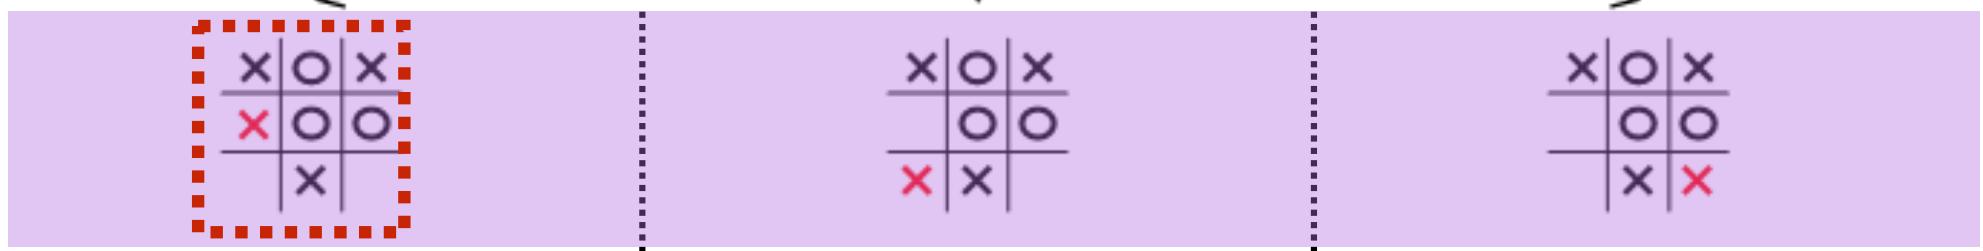
Current State/Root Node

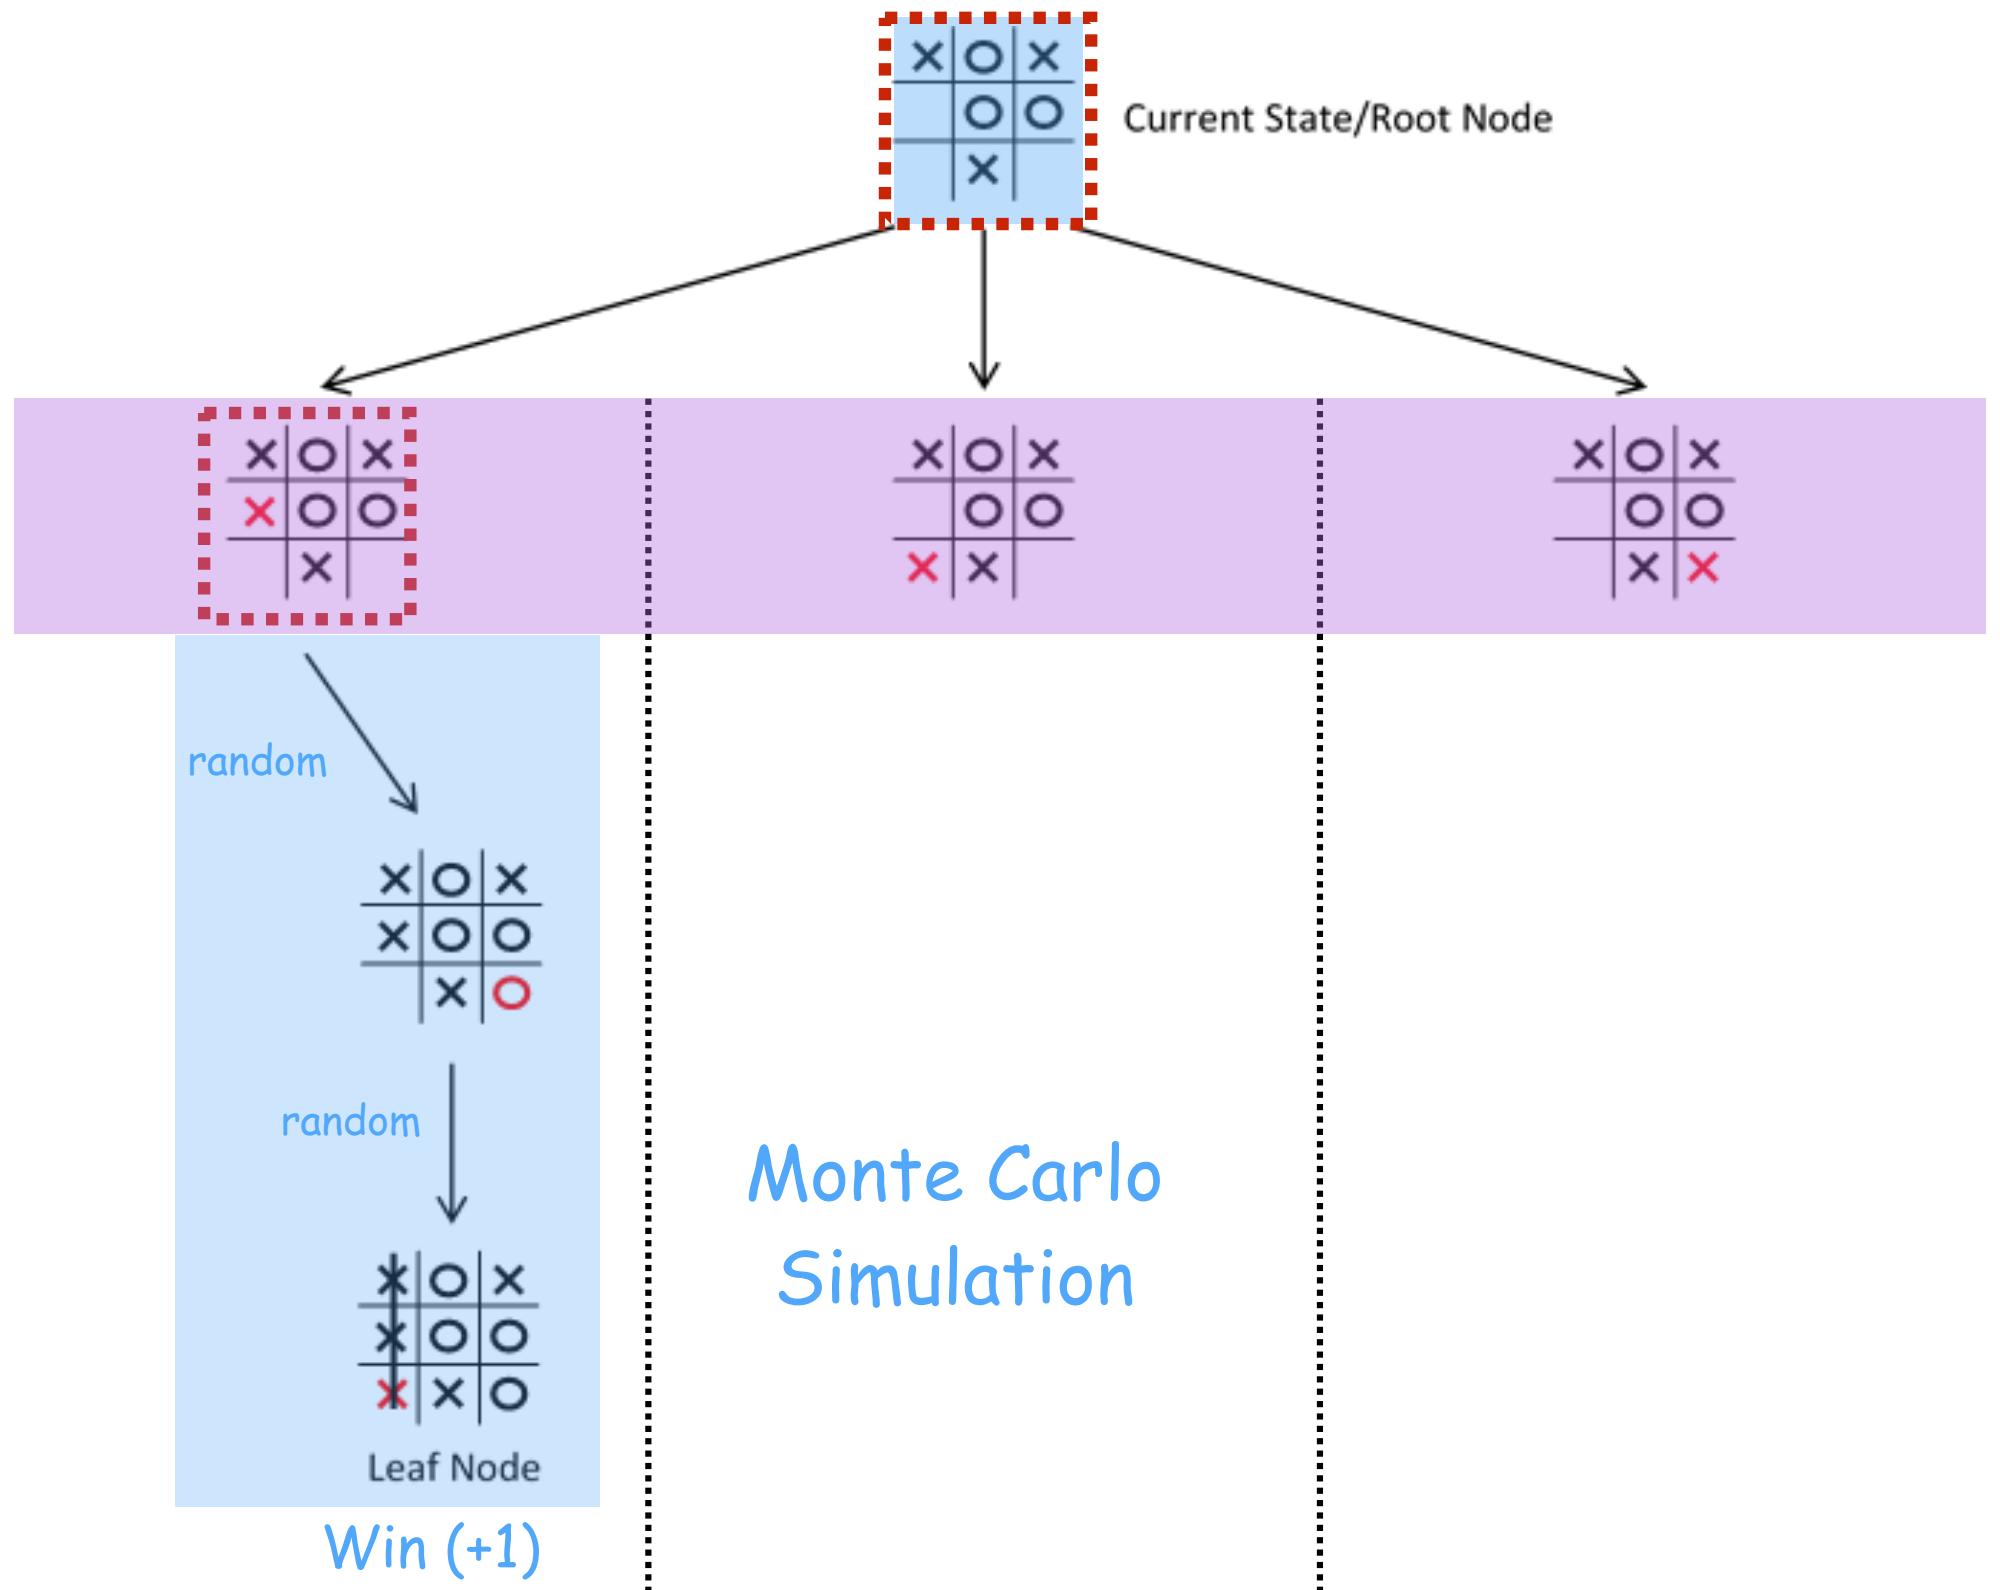


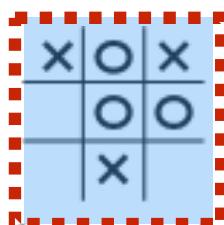
selection (expanded node)



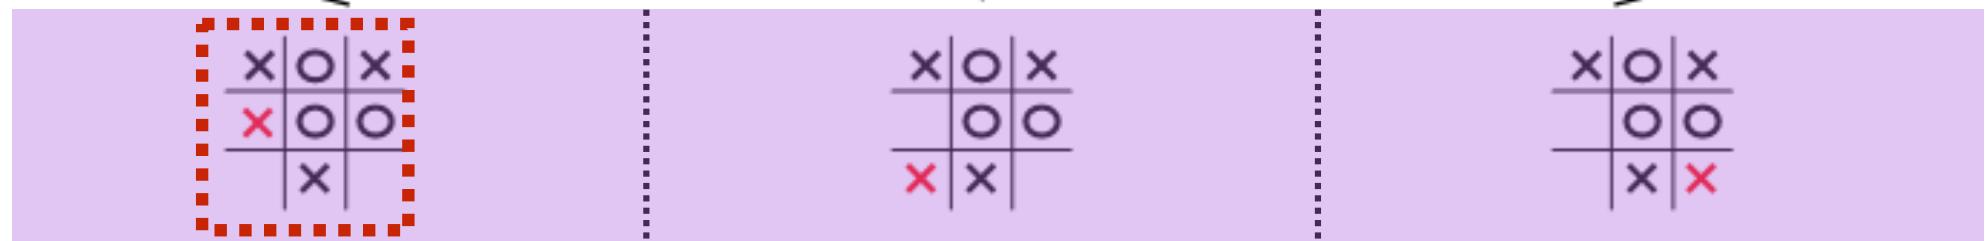
Current State/Root Node







Current State/Root Node

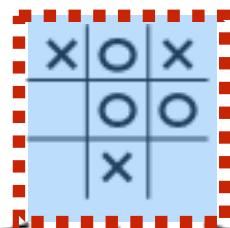


(sum of scores) $w=1$

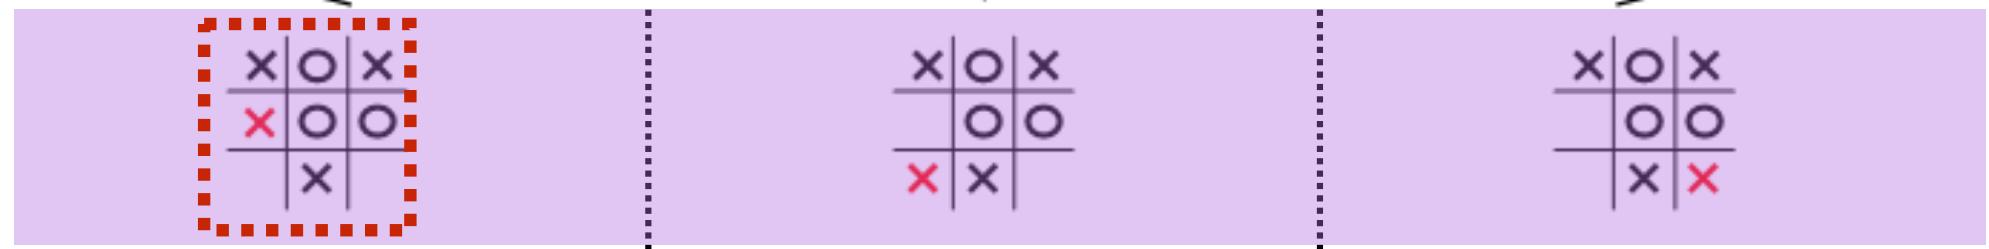
(#times selected) $n_i=1$

update node
statistics

(#times selected) $N=1$



Current State/Root Node

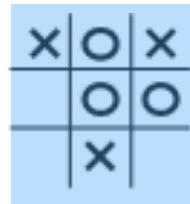


(sum of scores) $w=1$

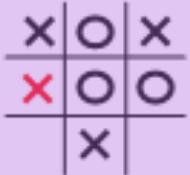
(#times selected) $n_i=1$

Back
Propagate

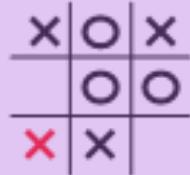
$N=1$



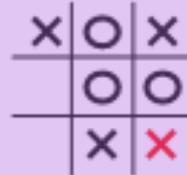
Current State/Root Node



$w=1$
 $n_i=1$



$w=0$
 $n_i=0$

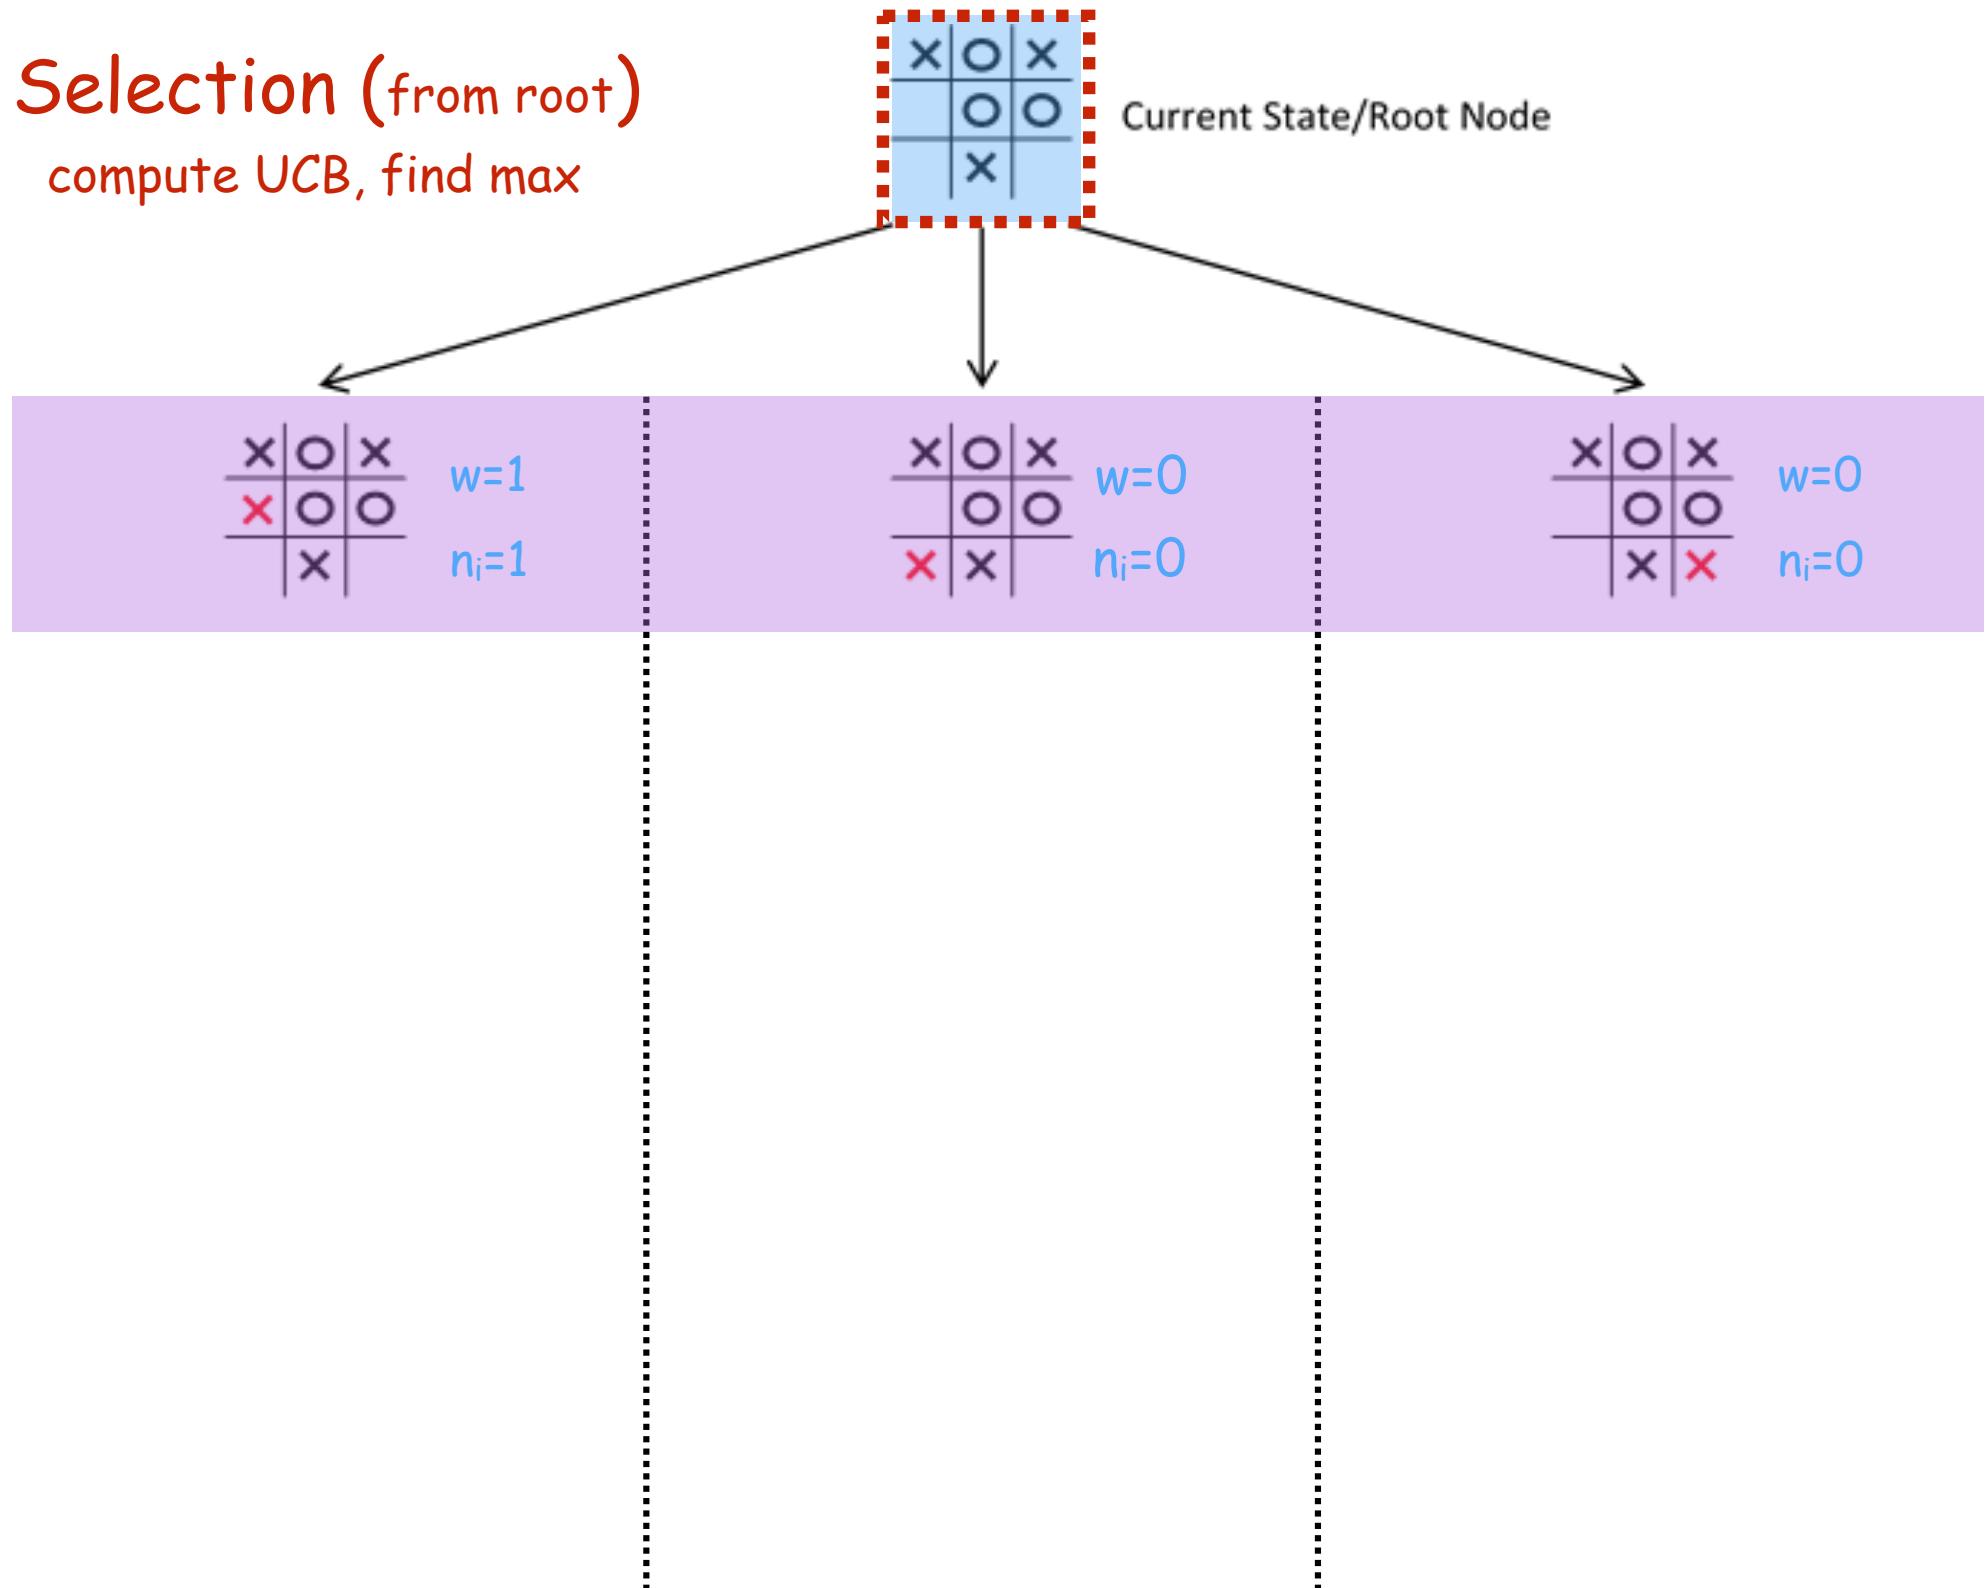


$w=0$
 $n_i=0$

After 1 iteration

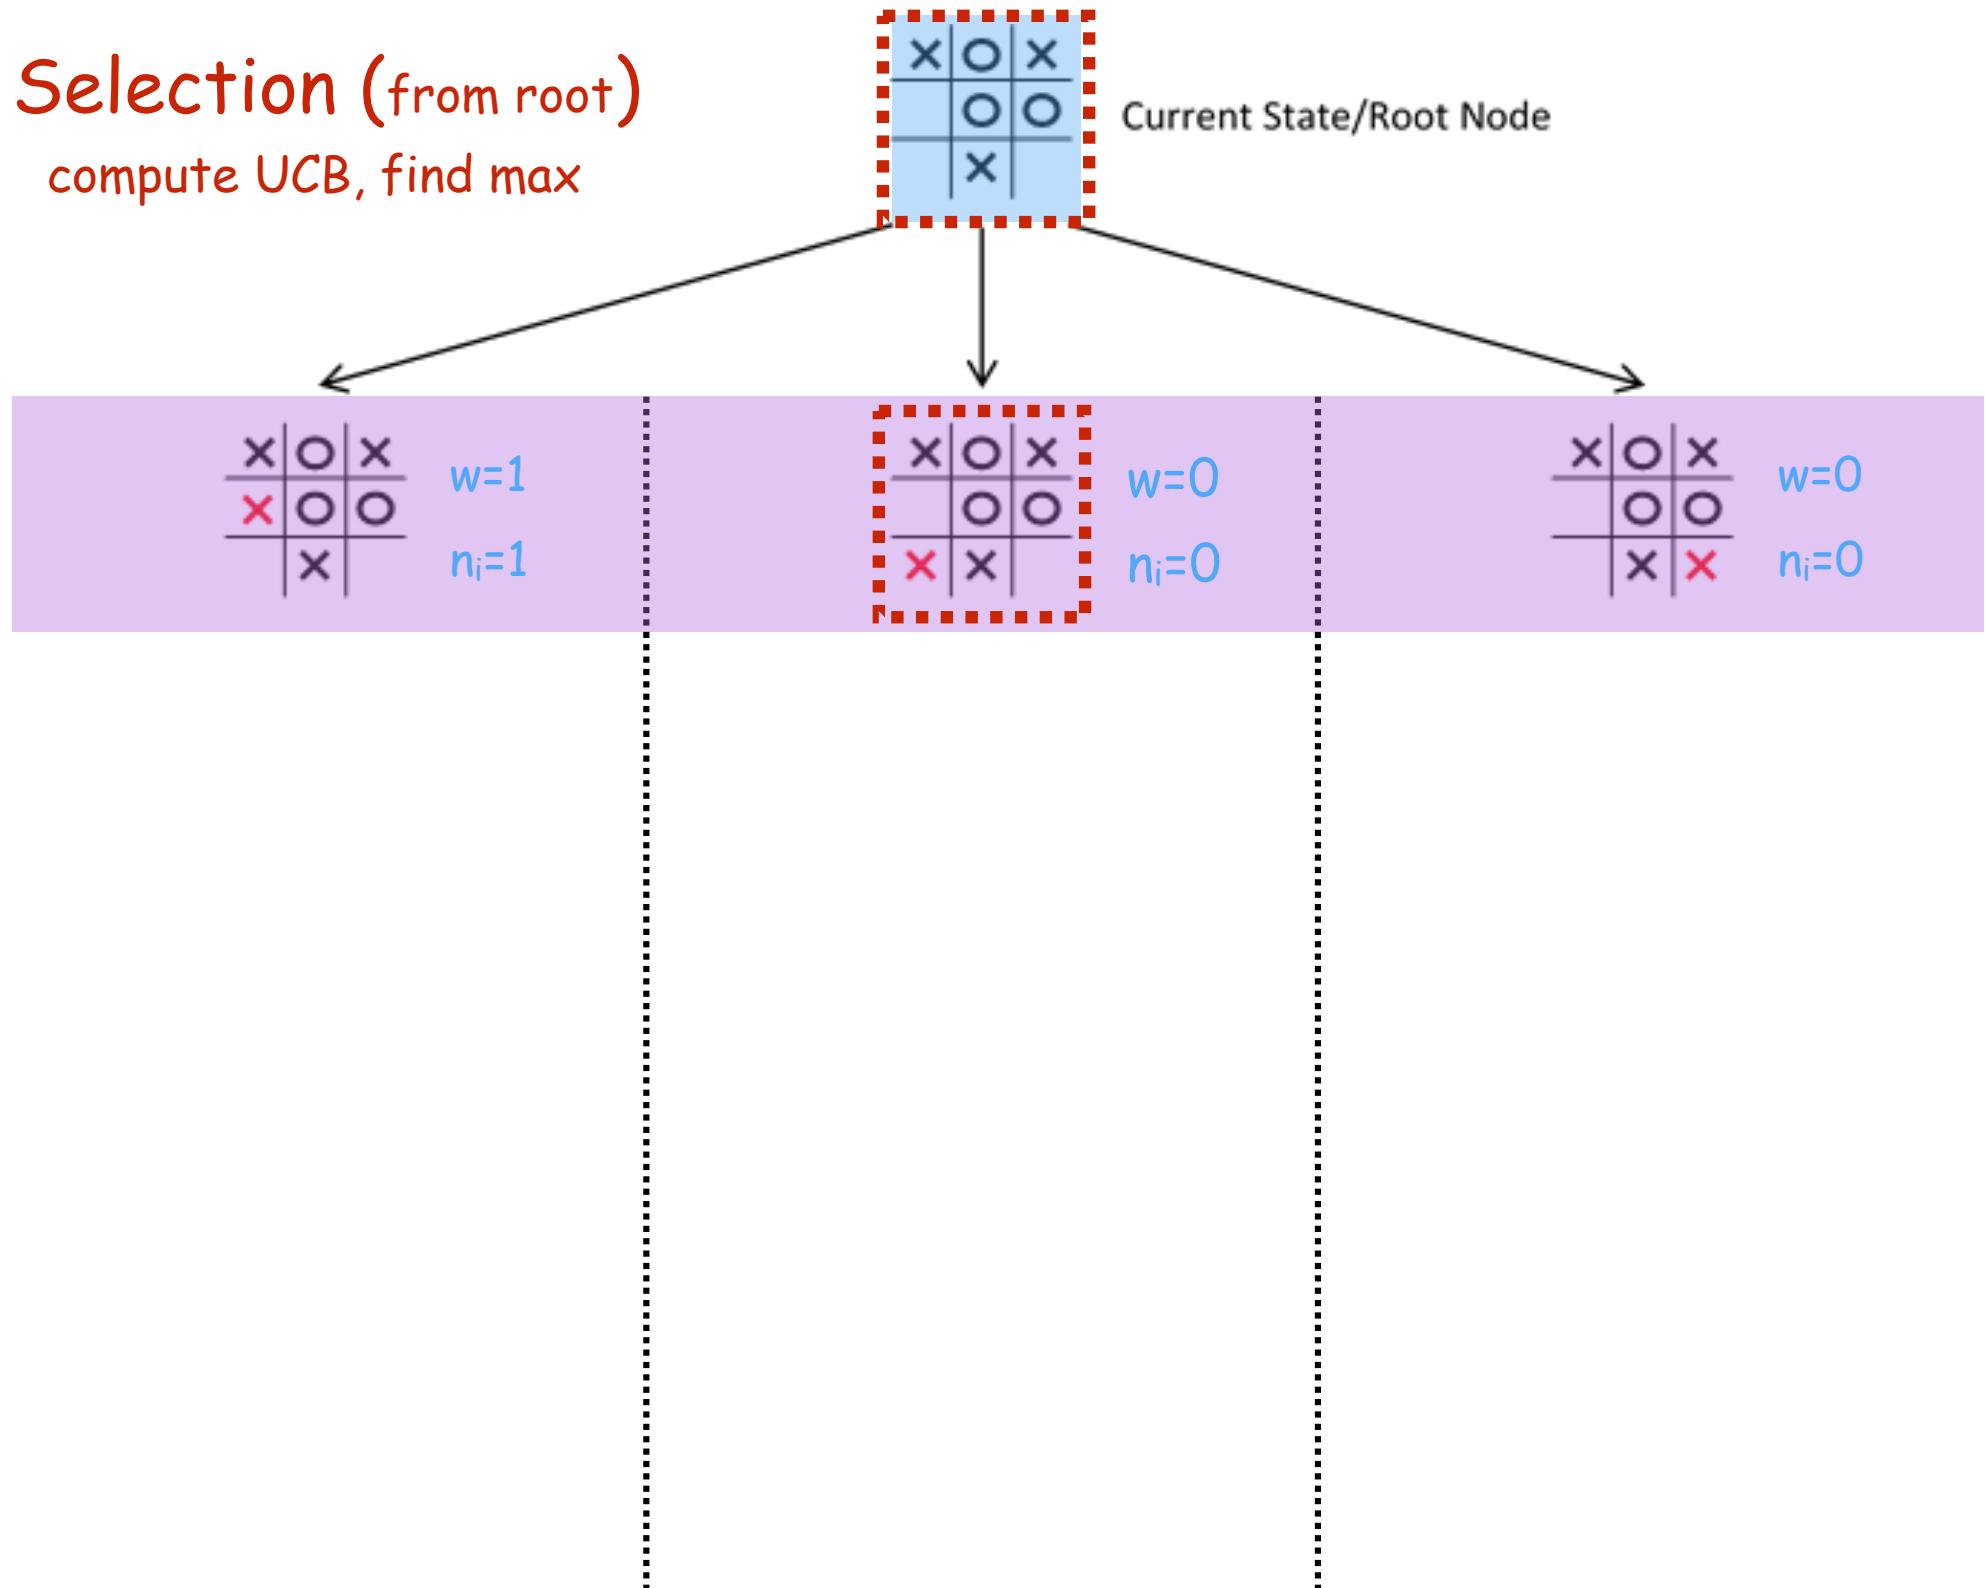
Selection (from root)

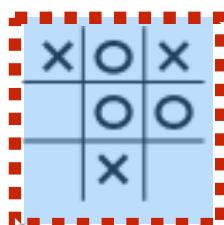
compute UCB, find max



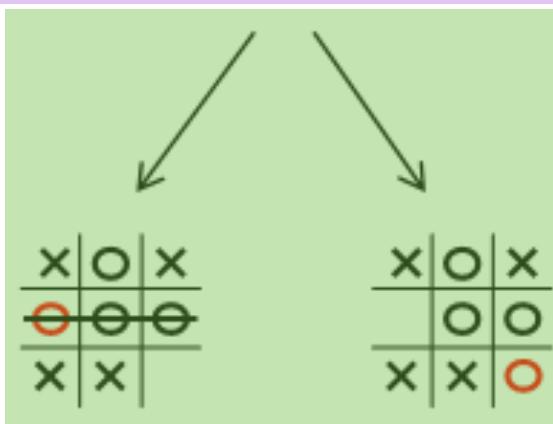
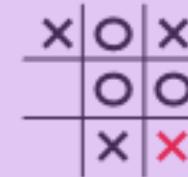
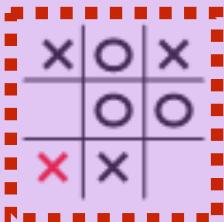
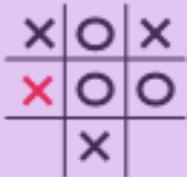
Selection (from root)

compute UCB, find max

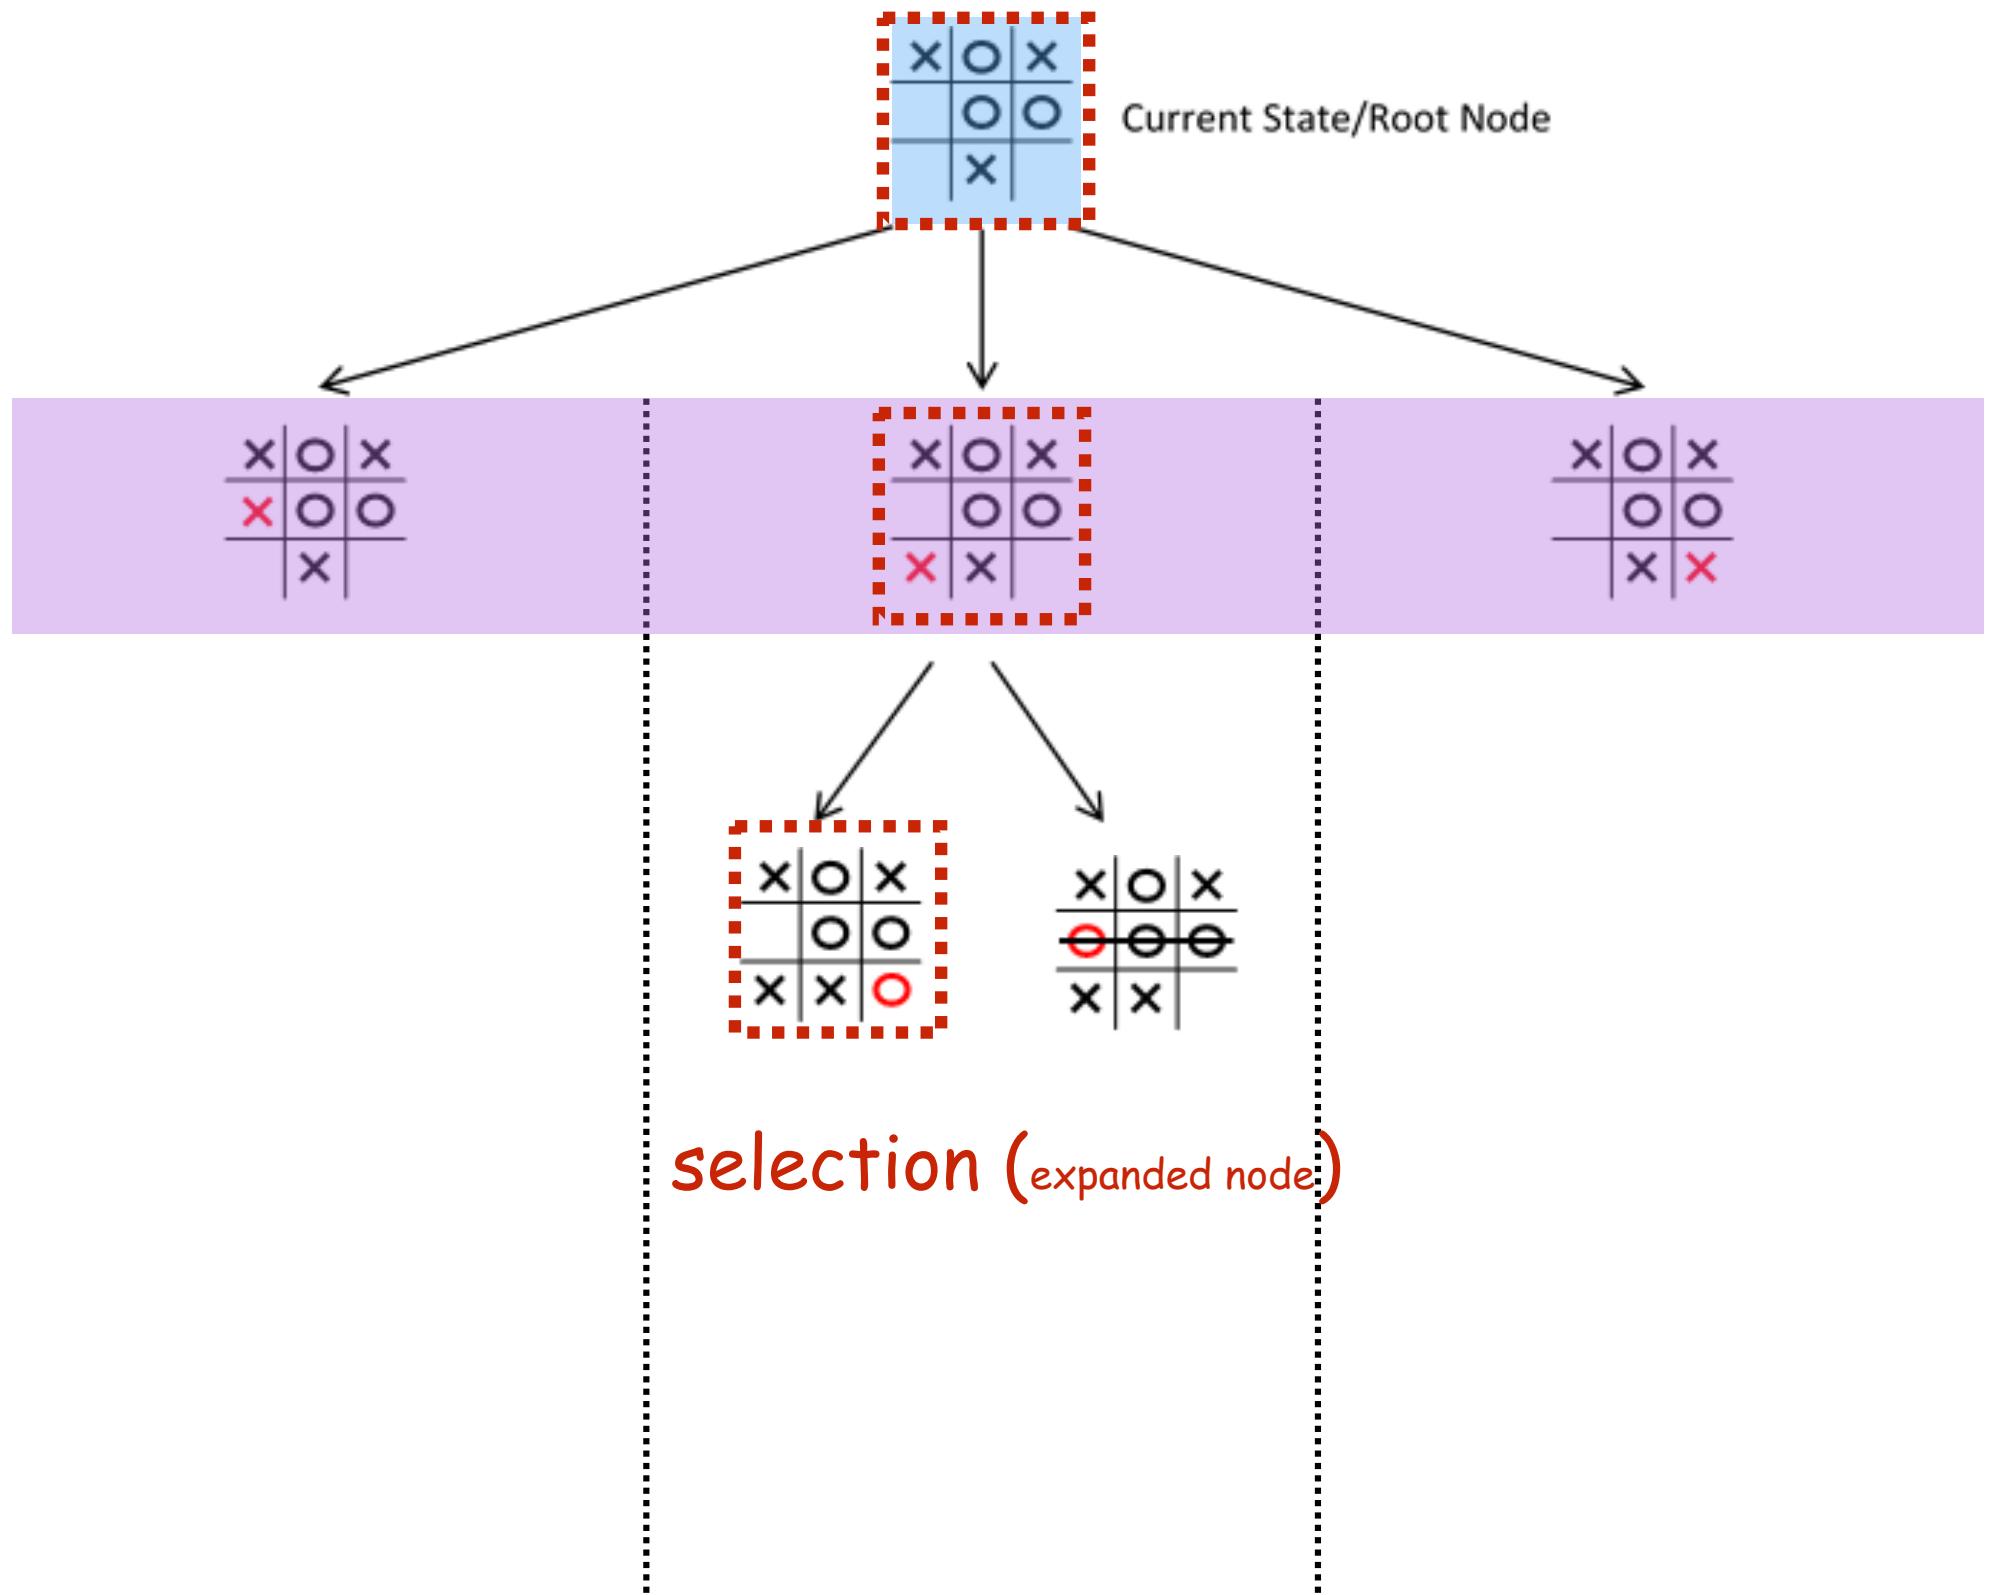


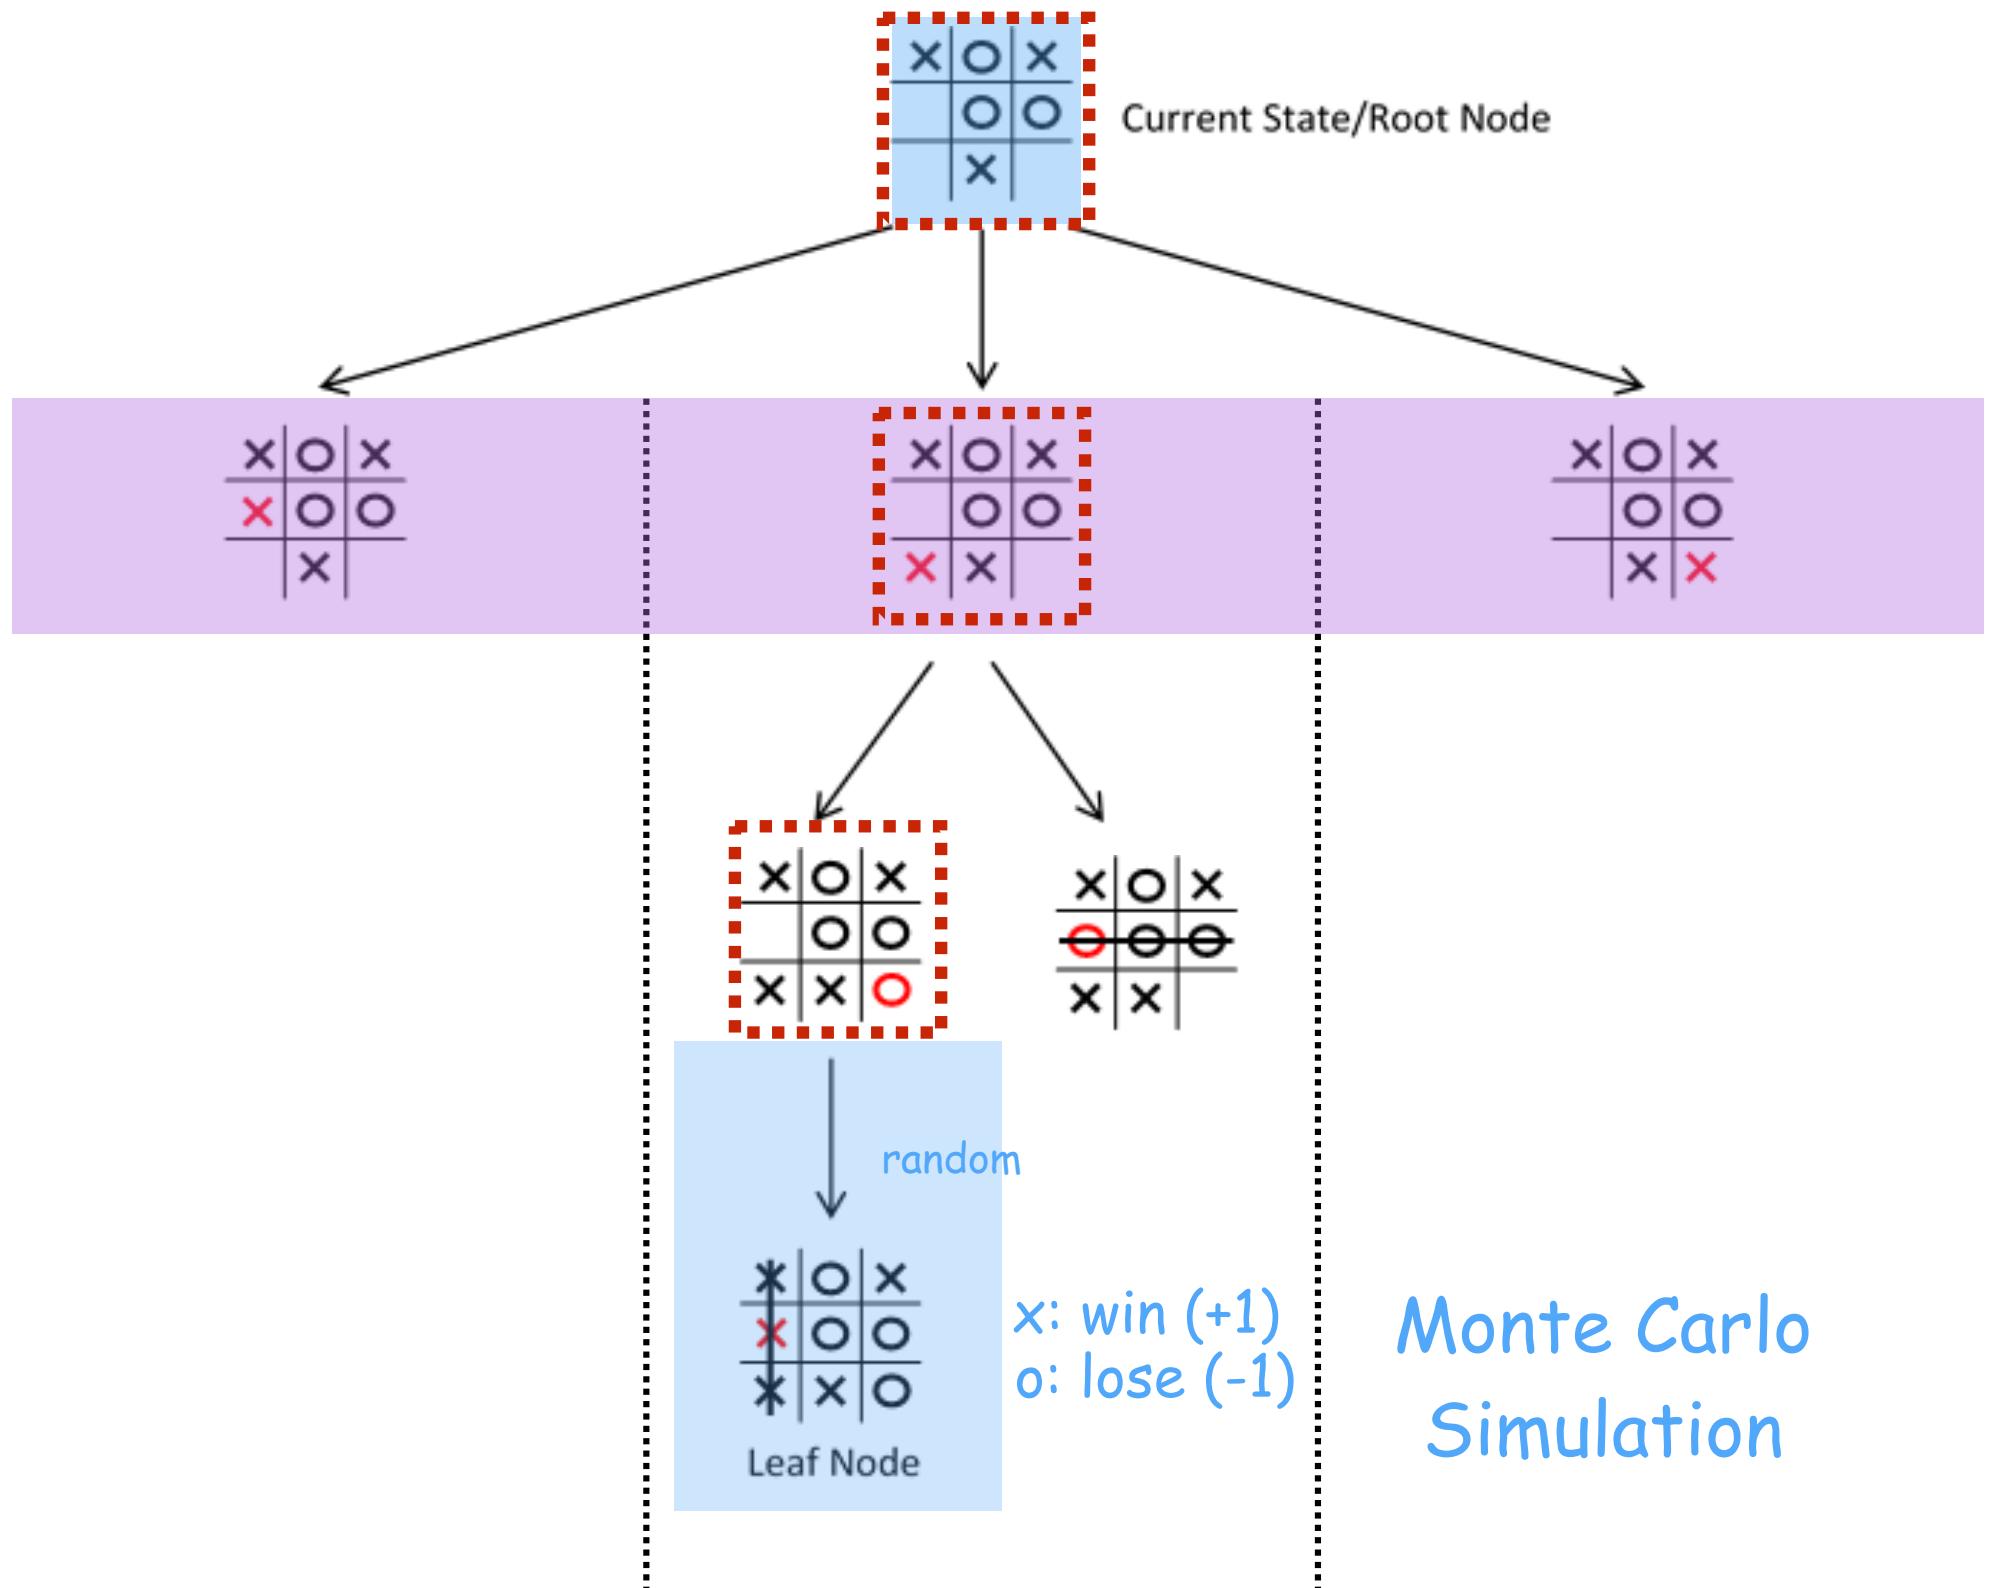


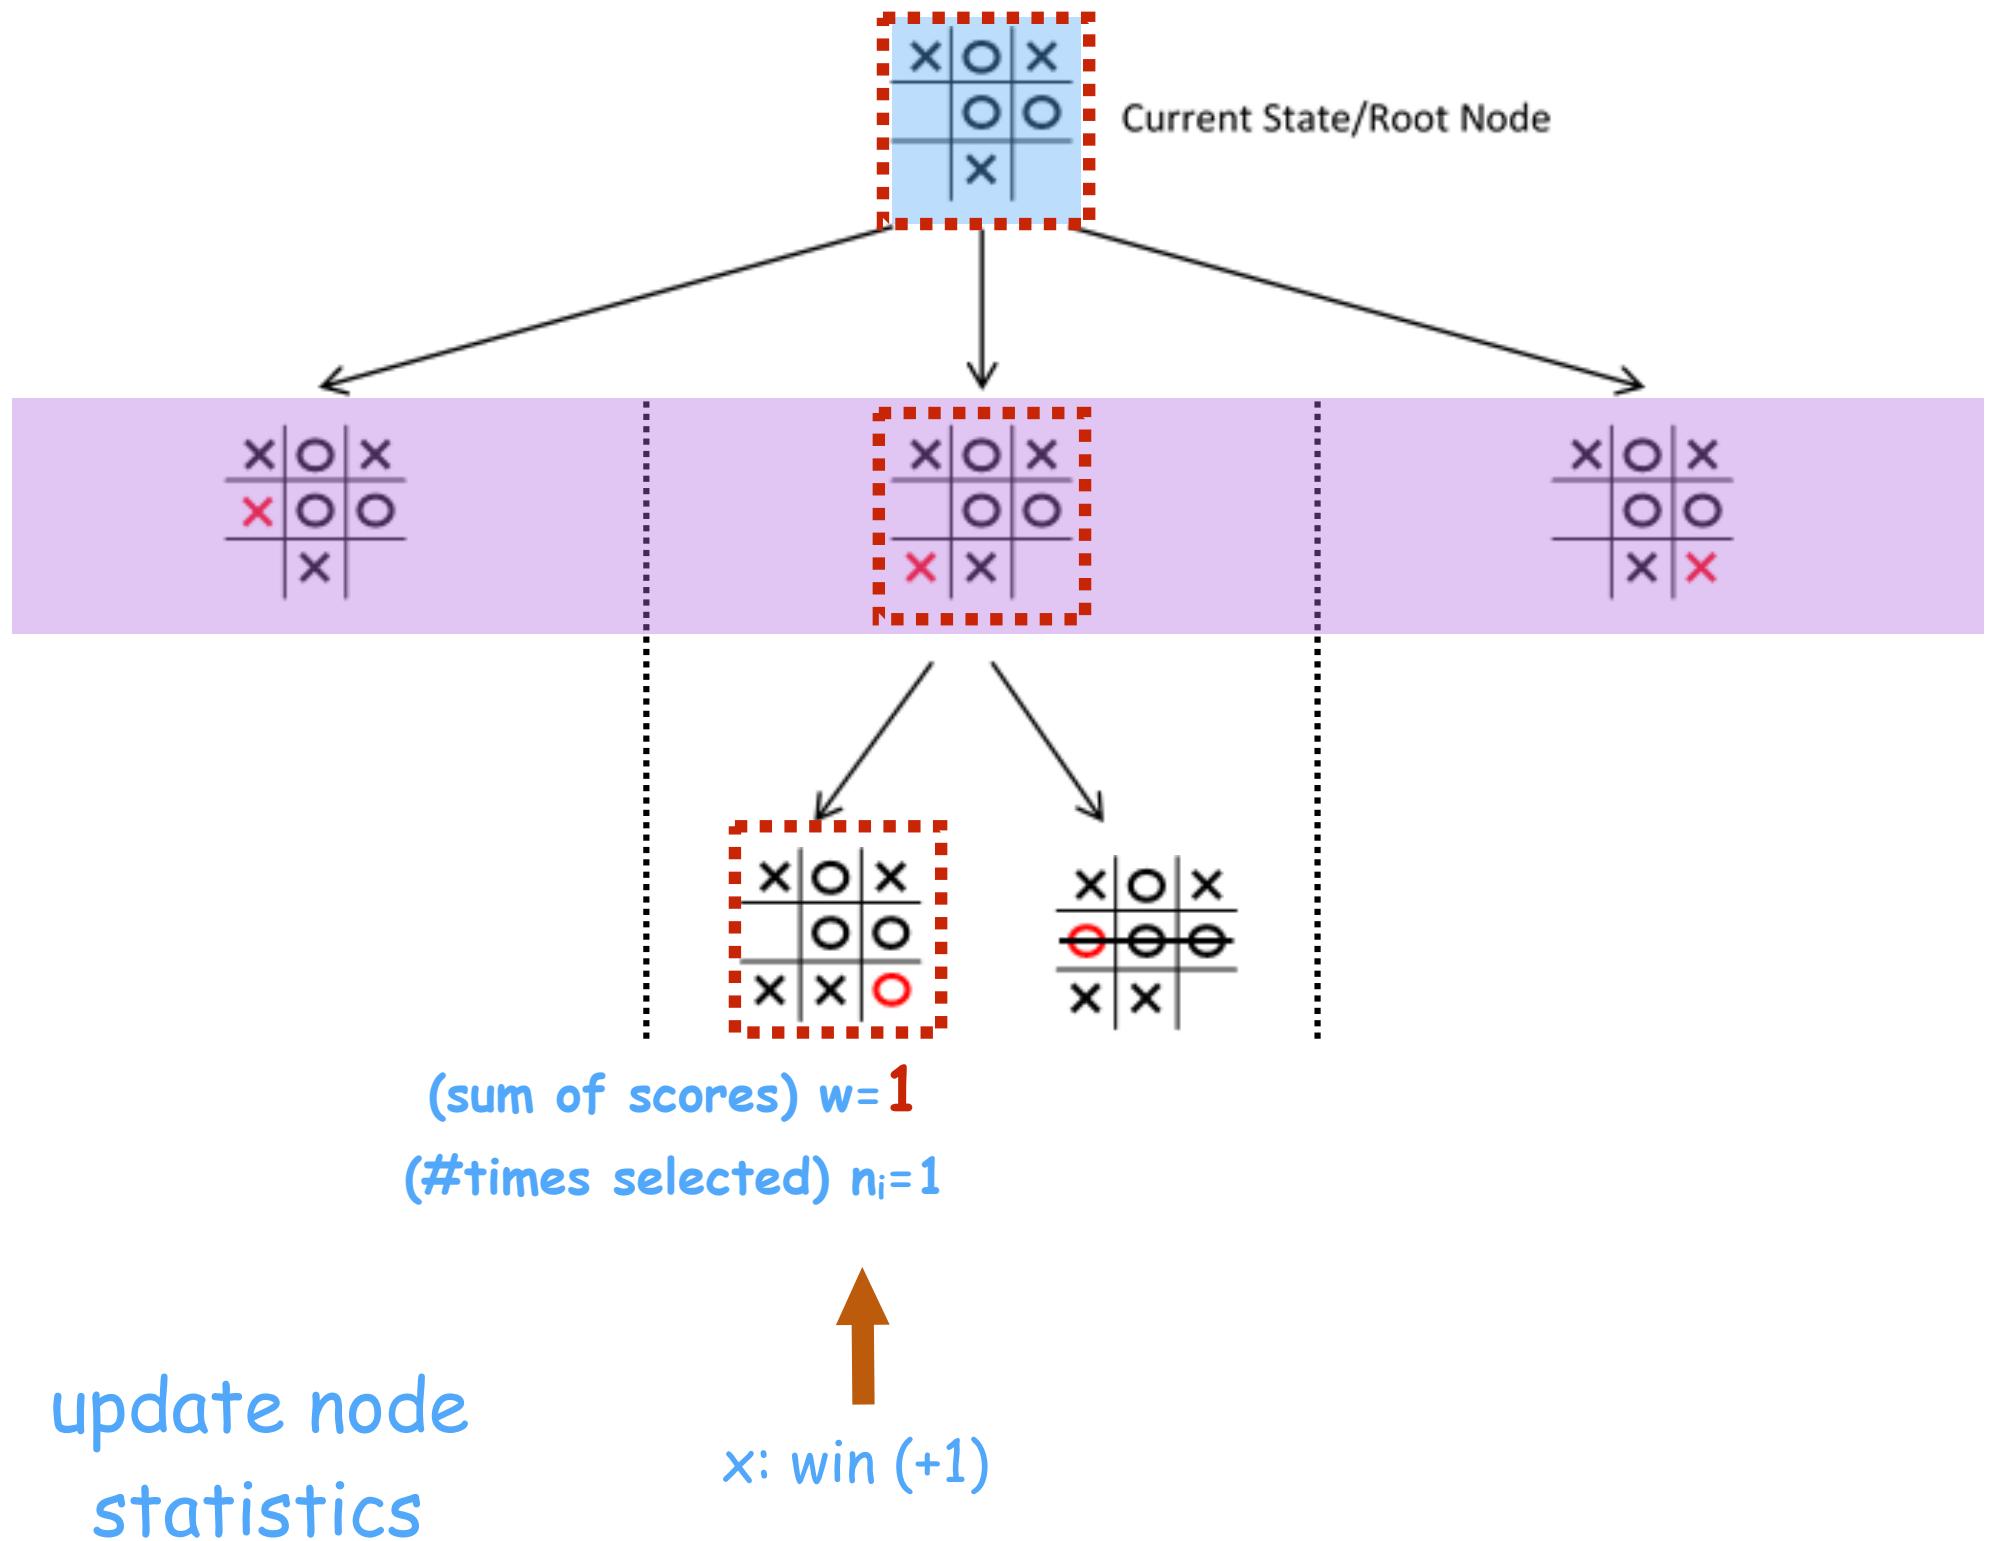
Current State/Root Node

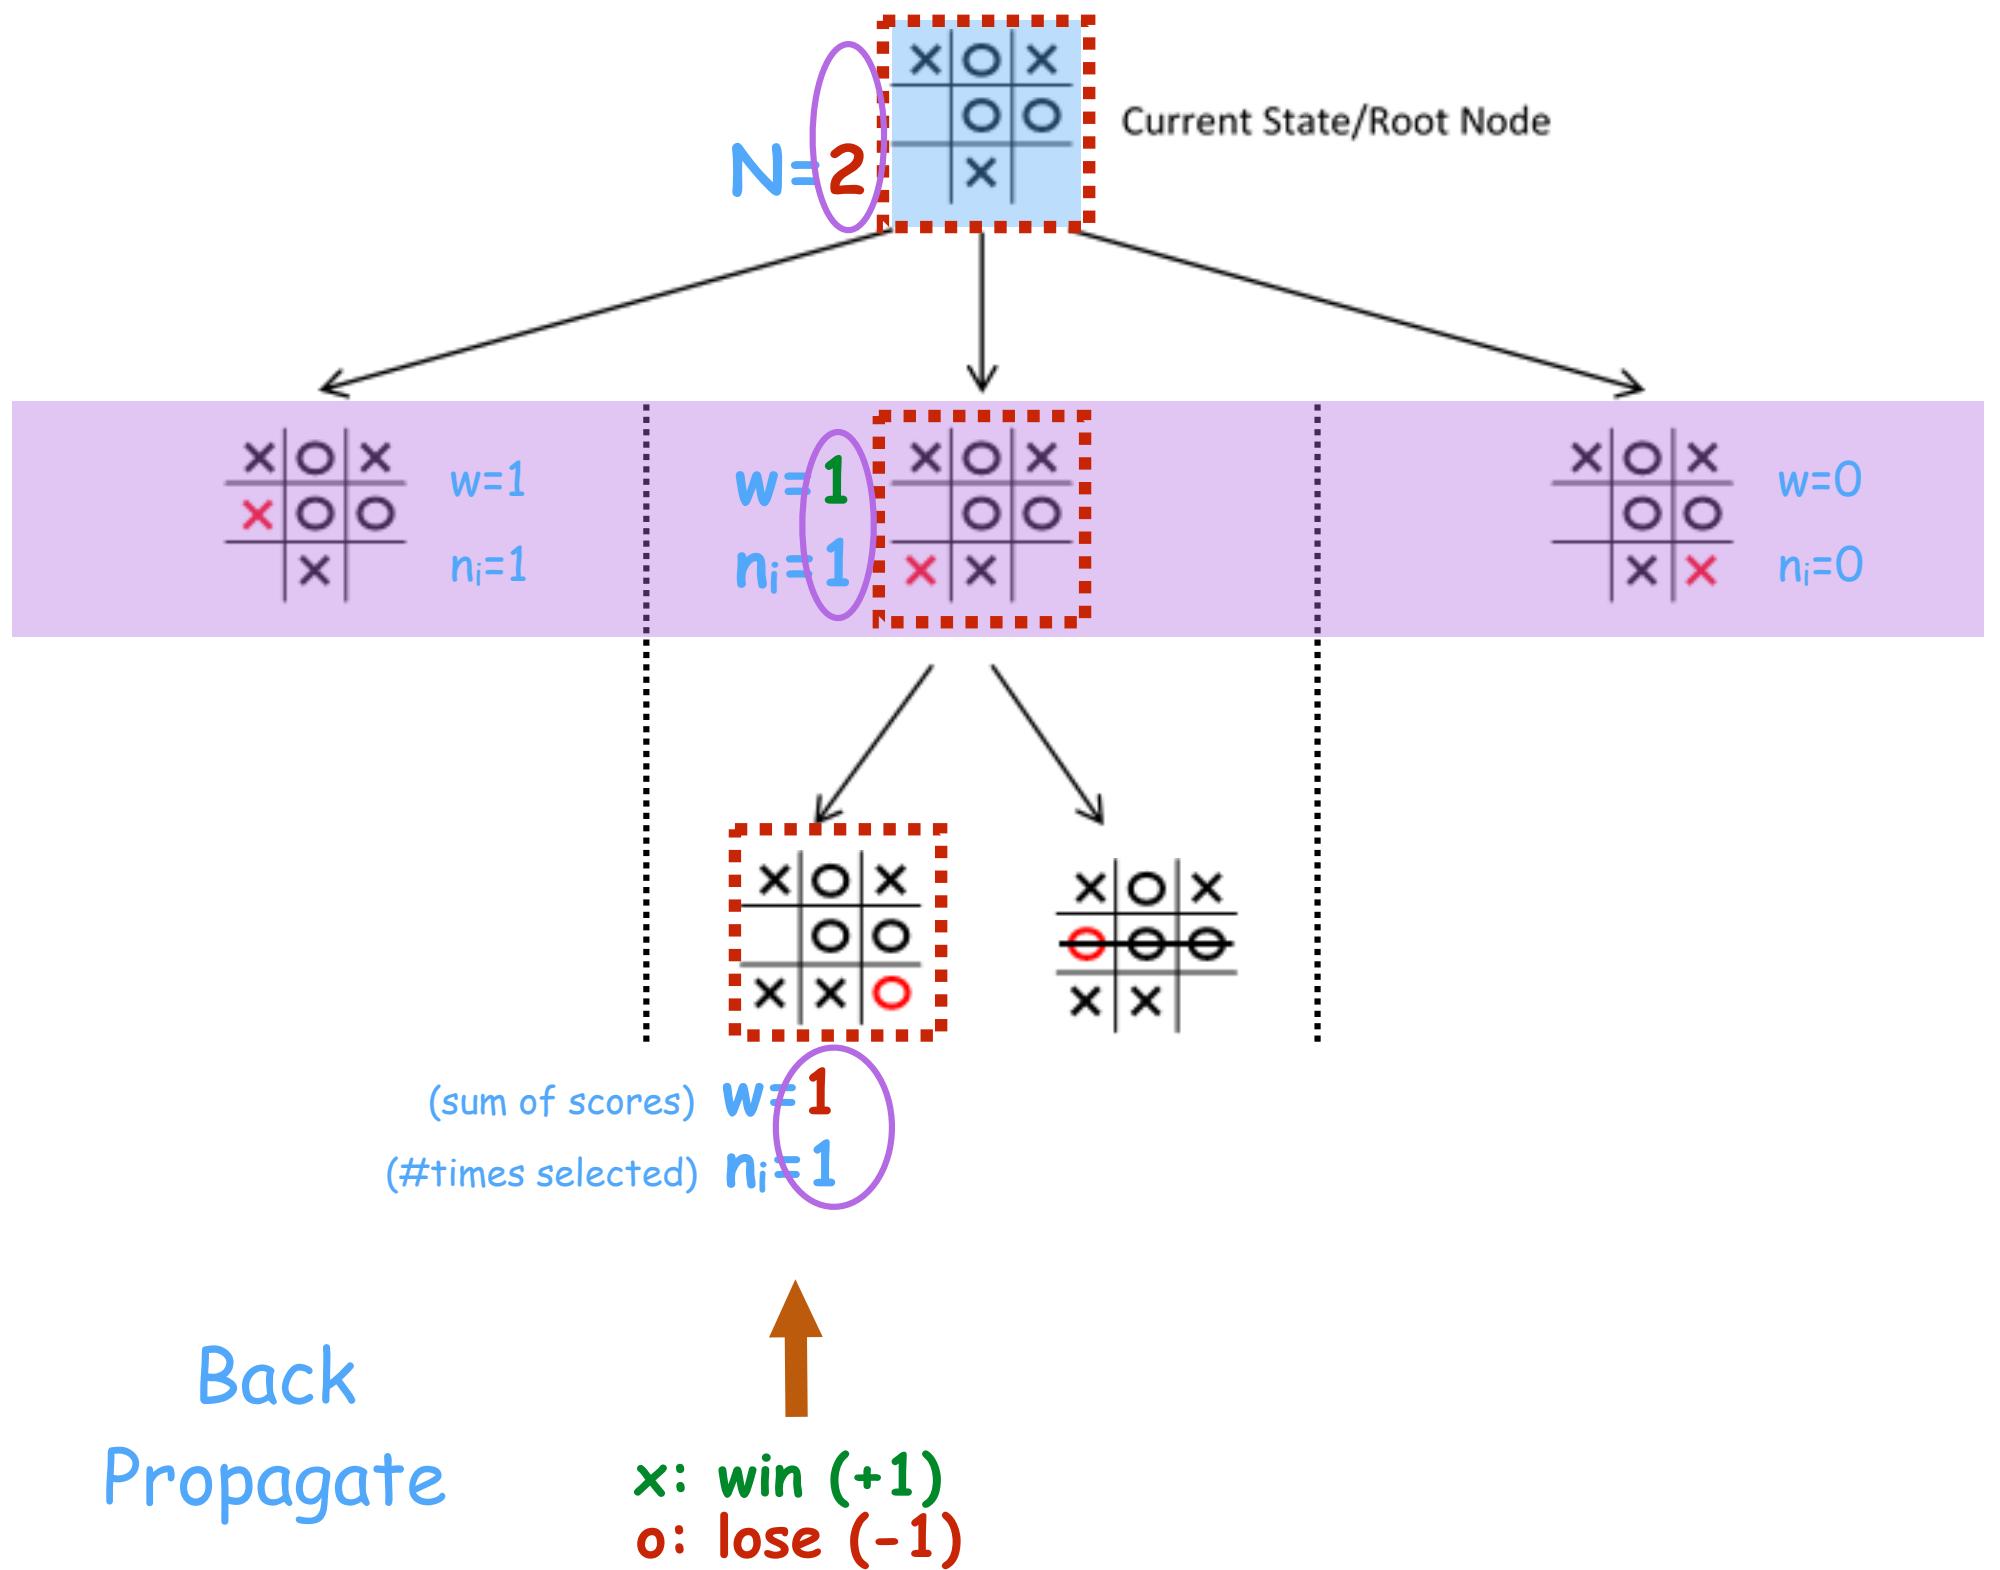


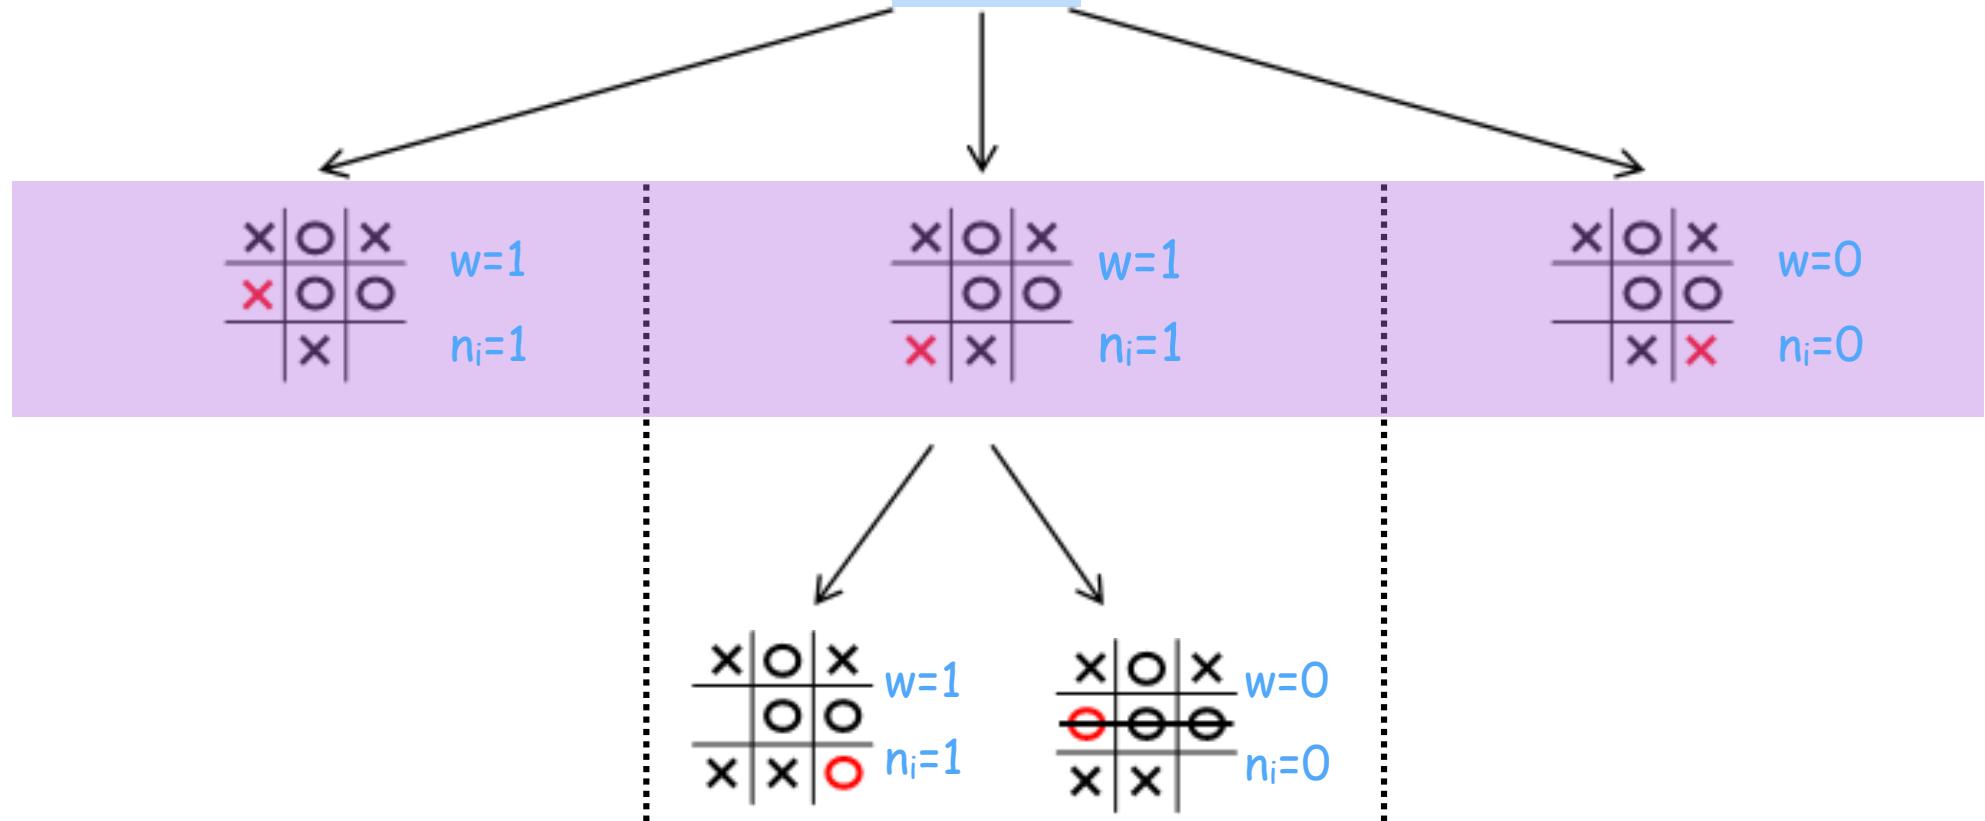
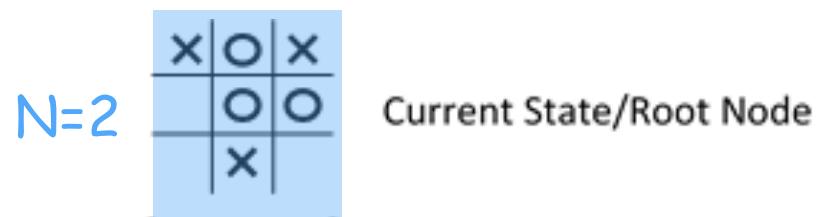
Expand





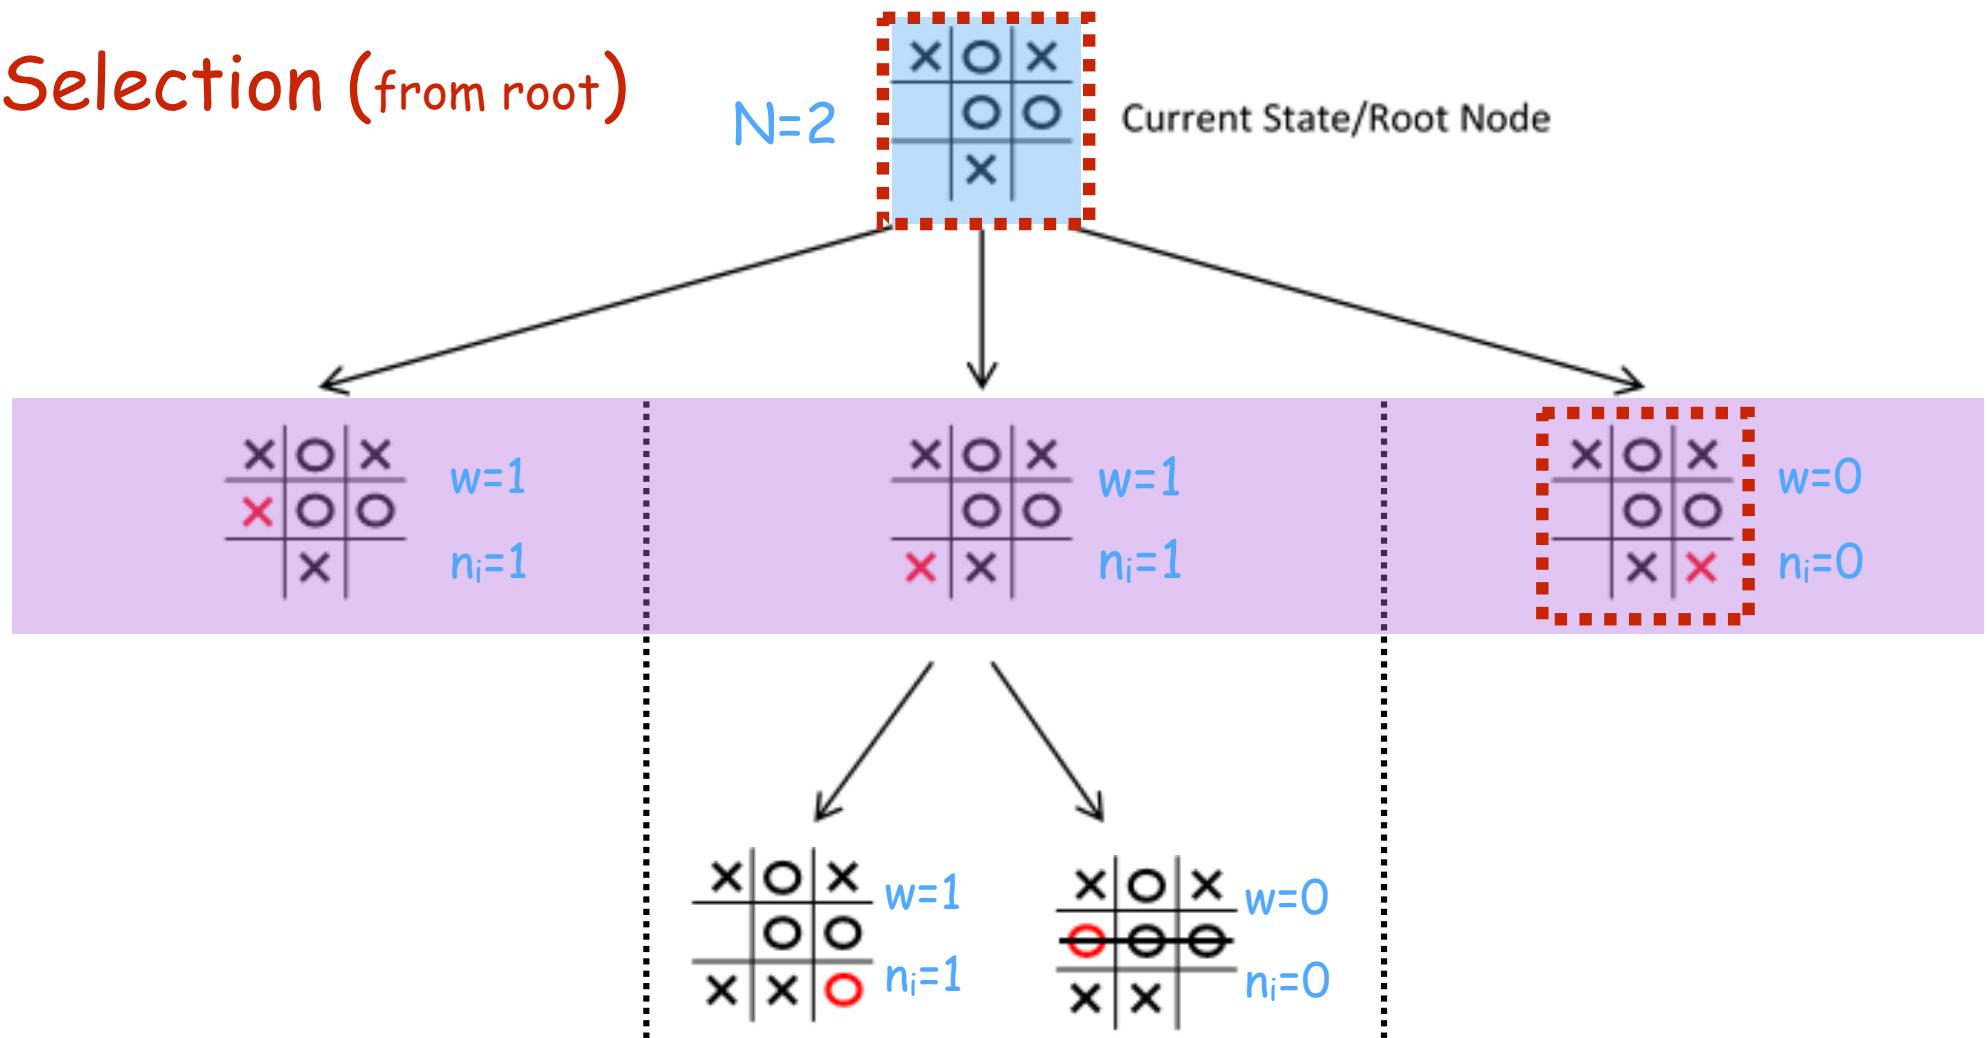


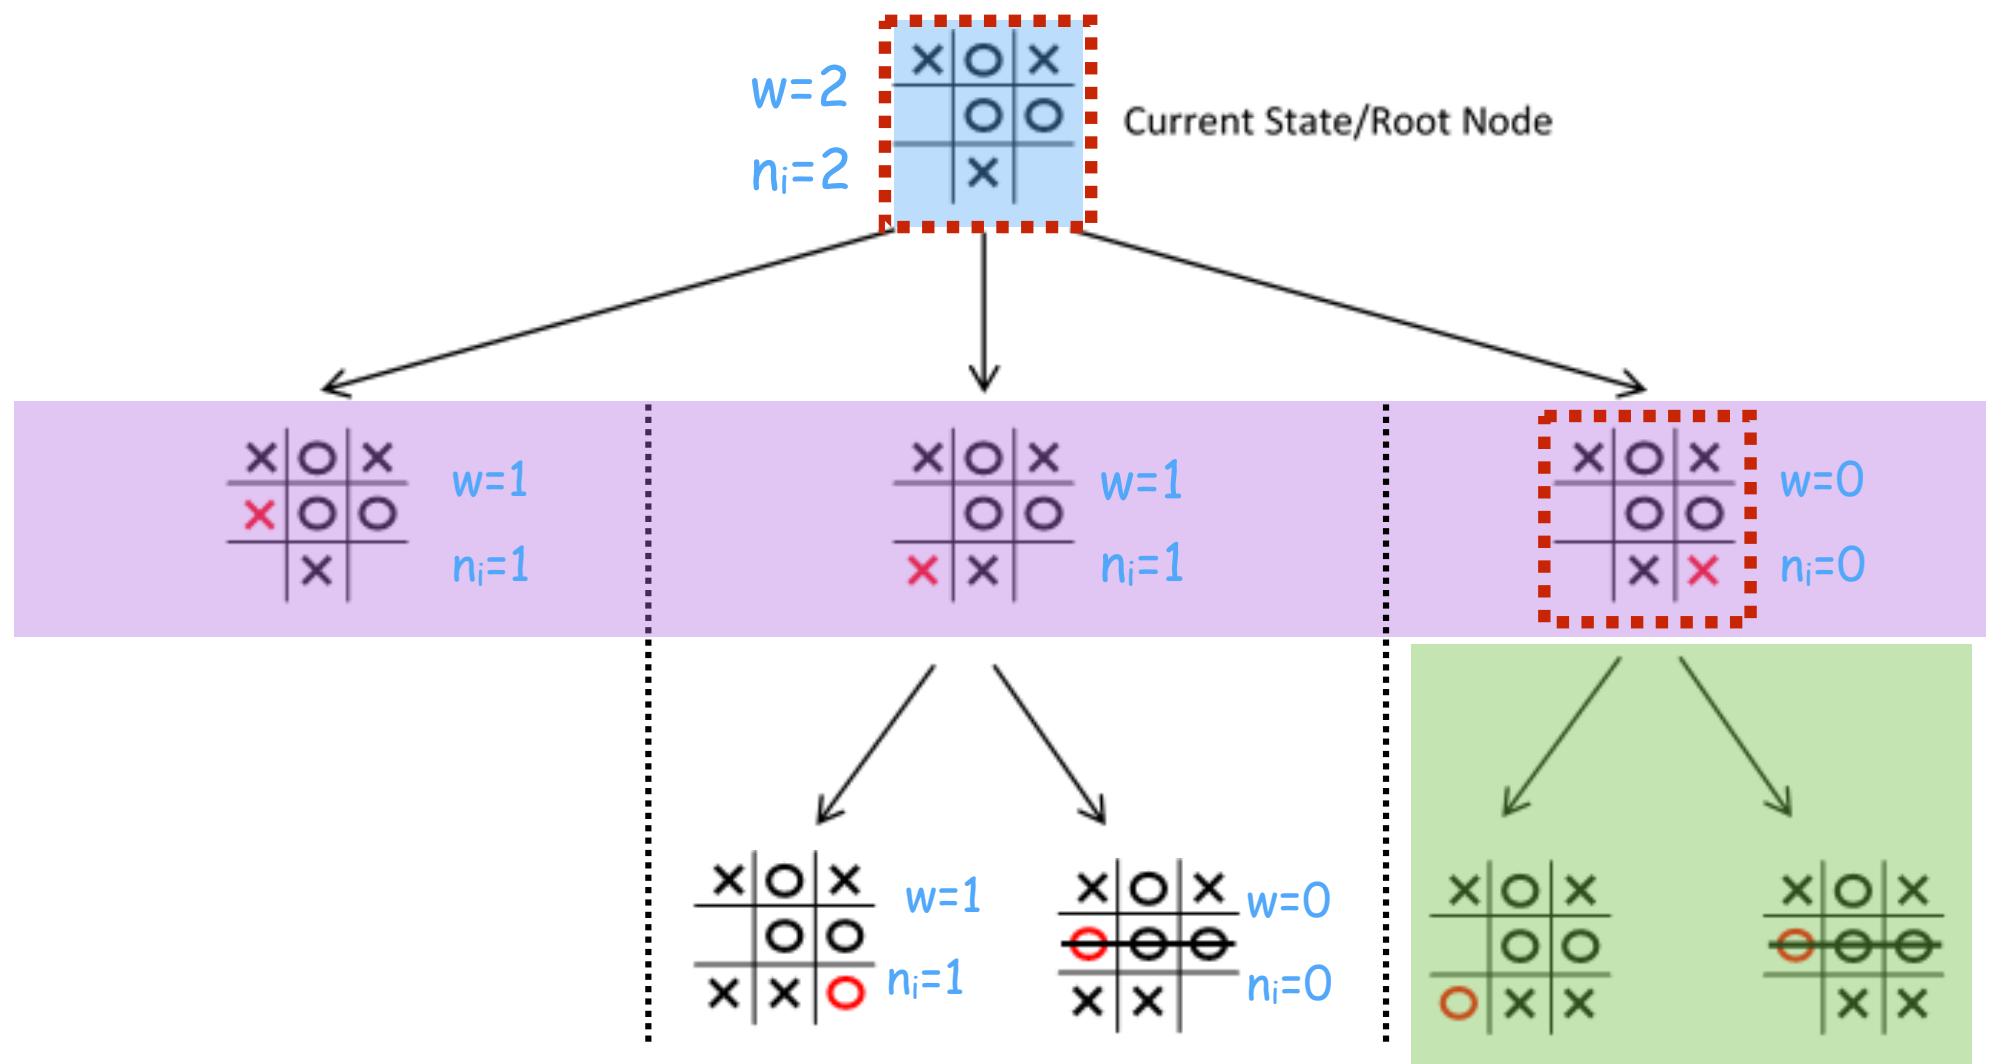




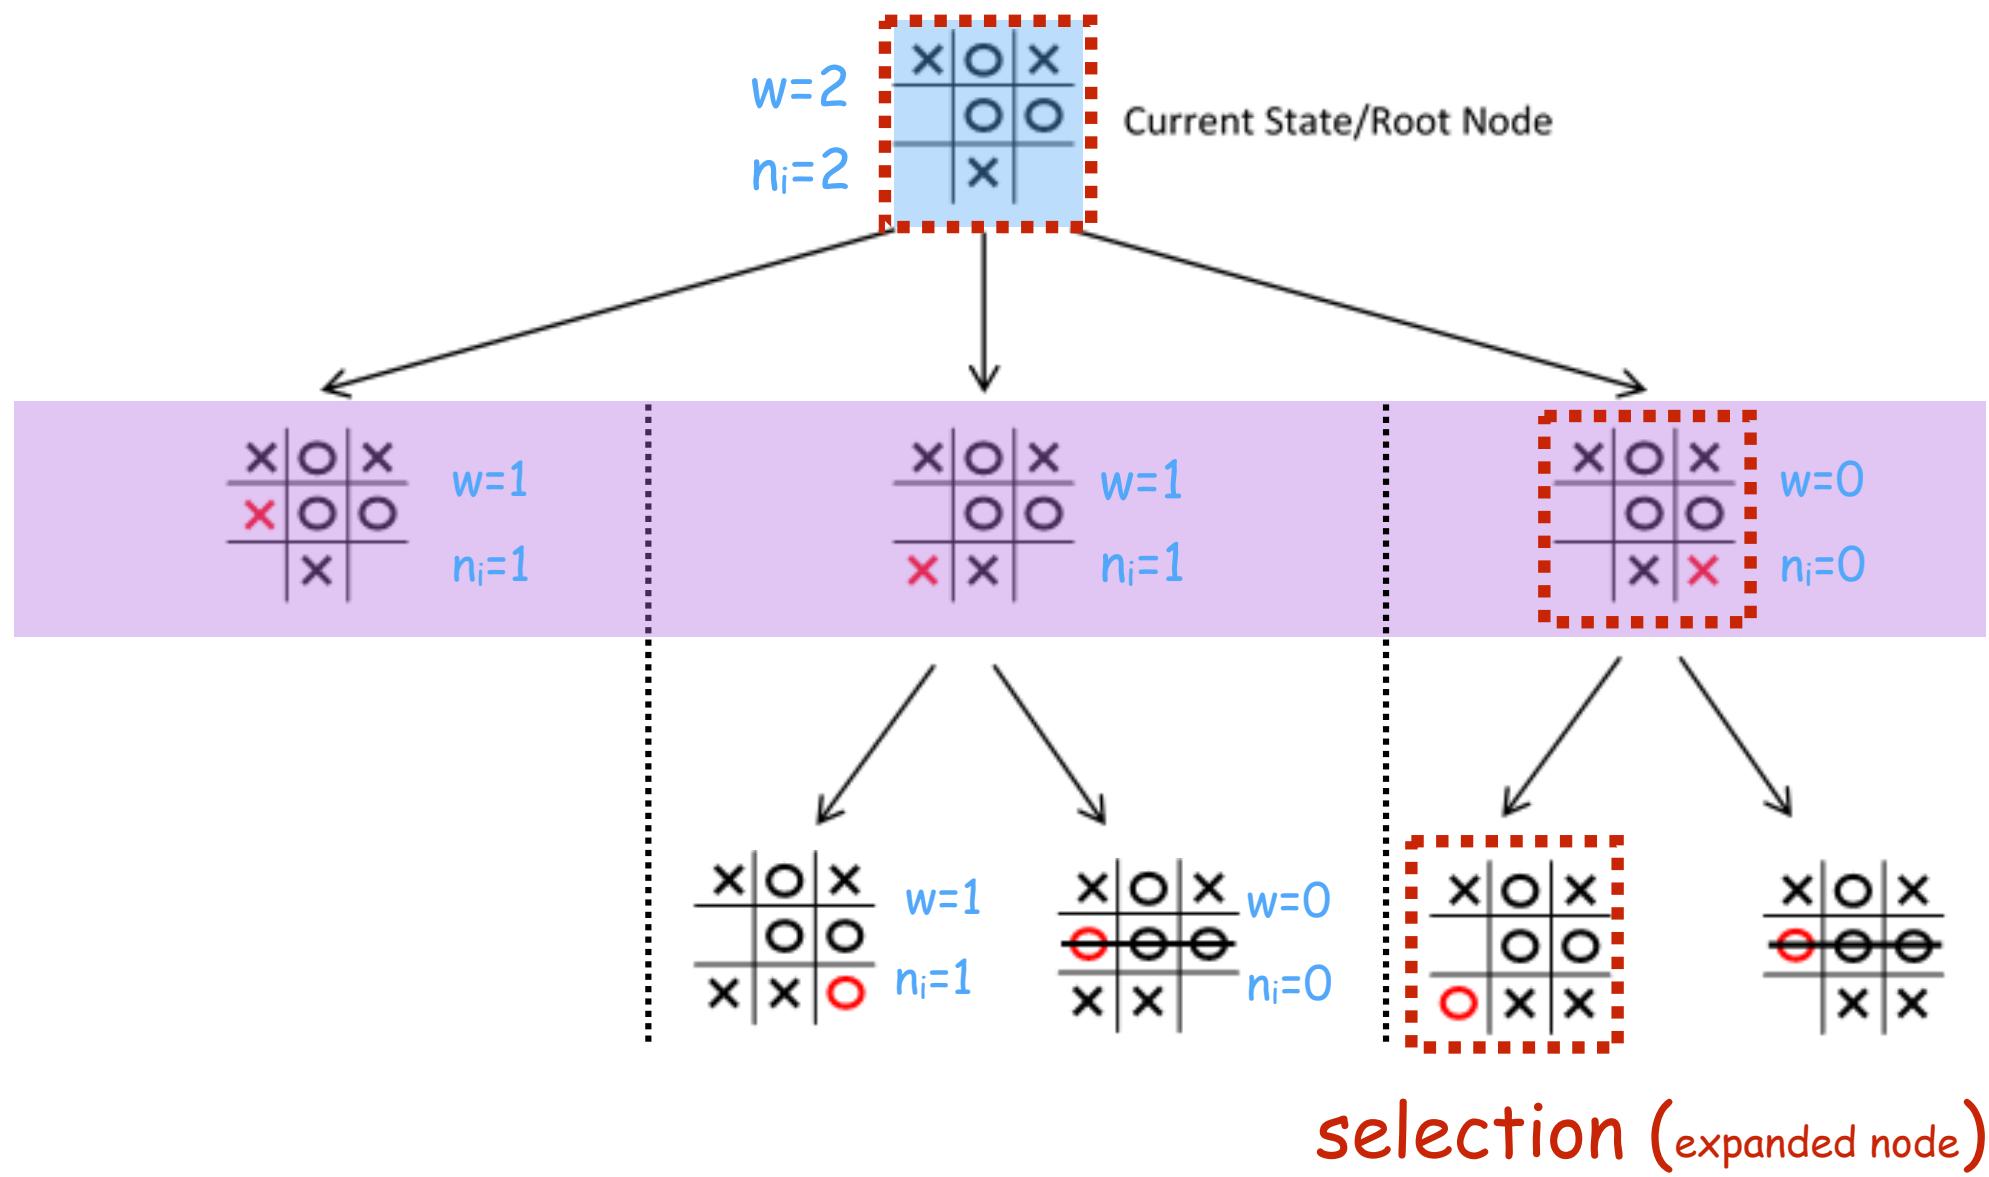
After 2 iterations

Selection (from root)

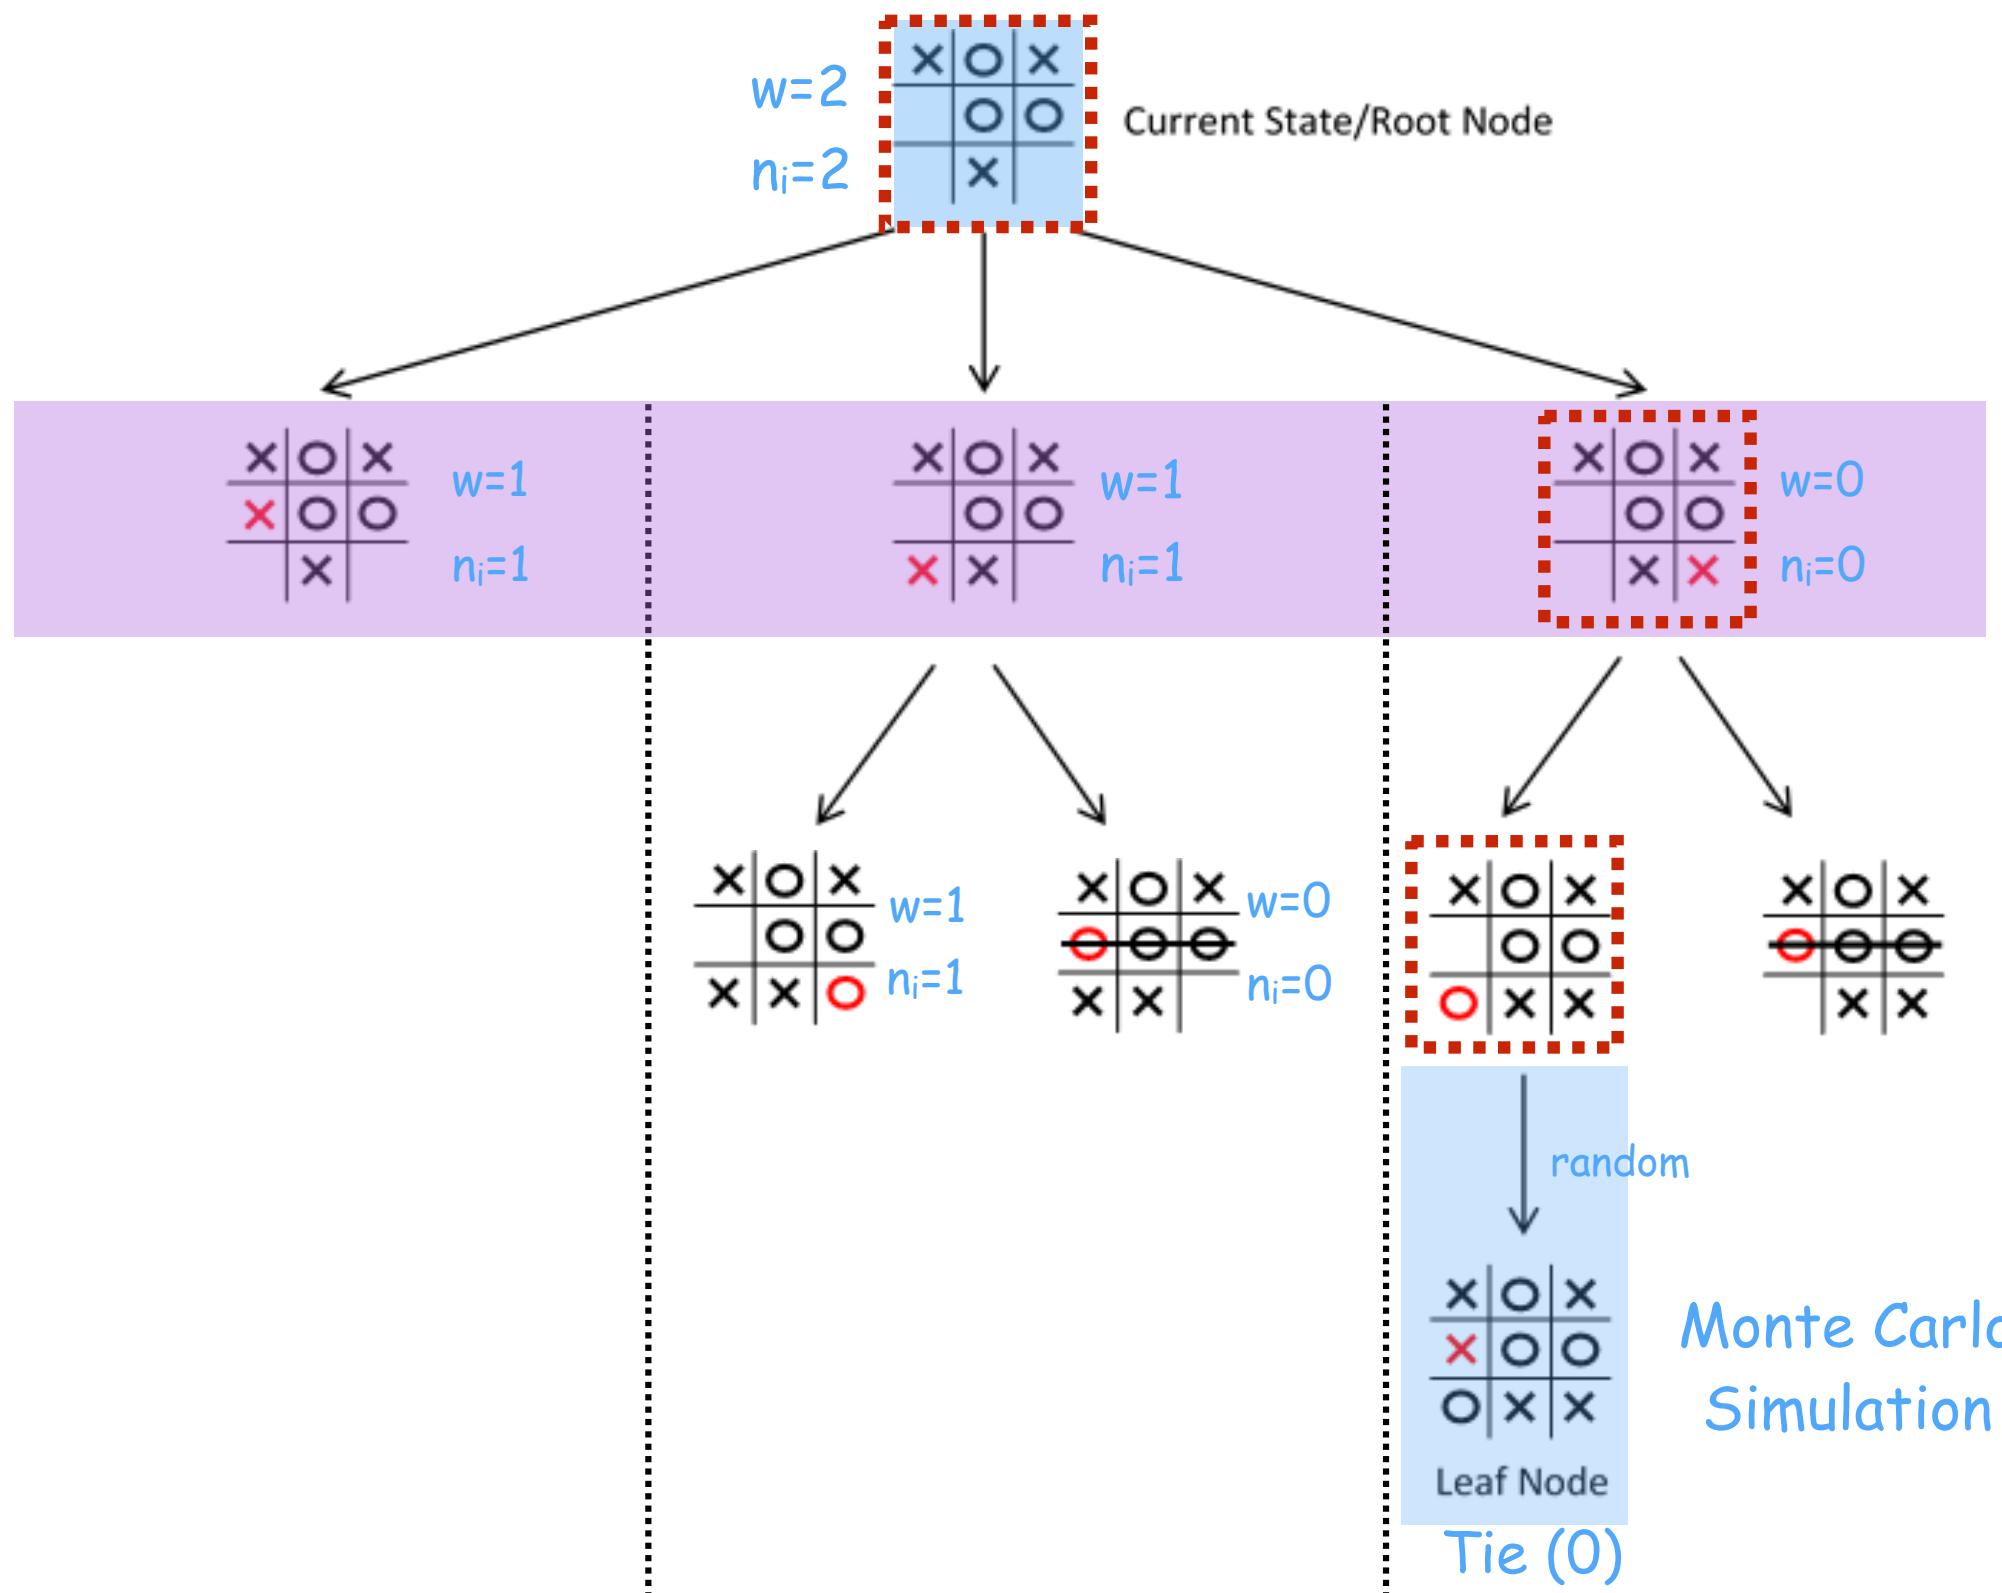


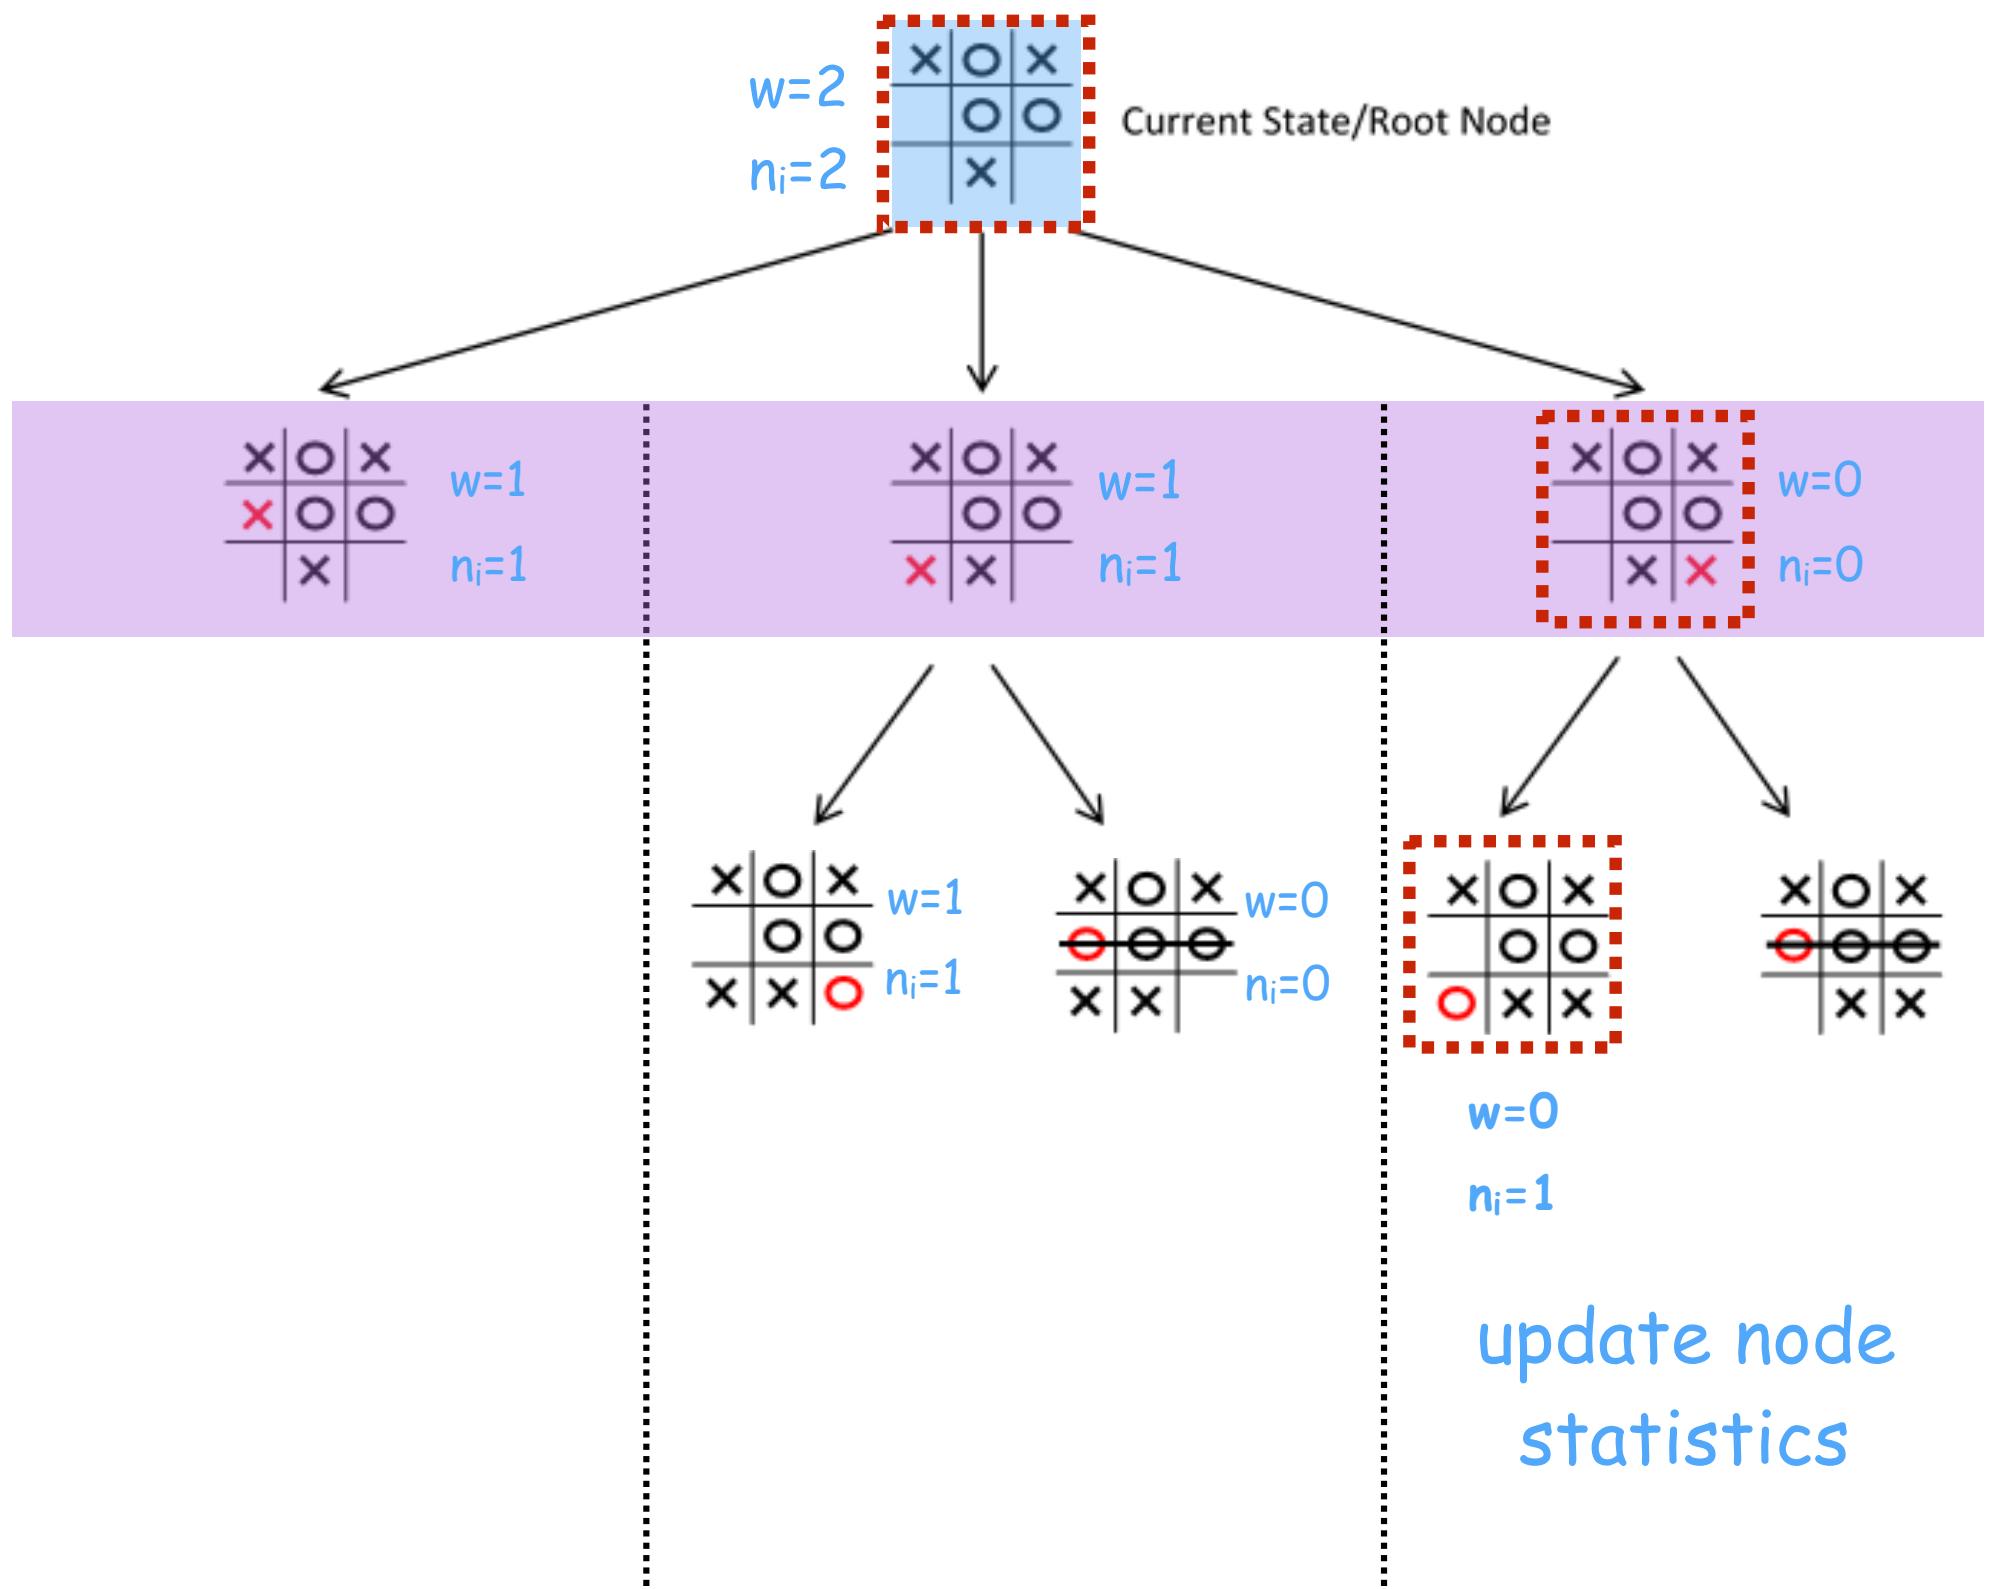


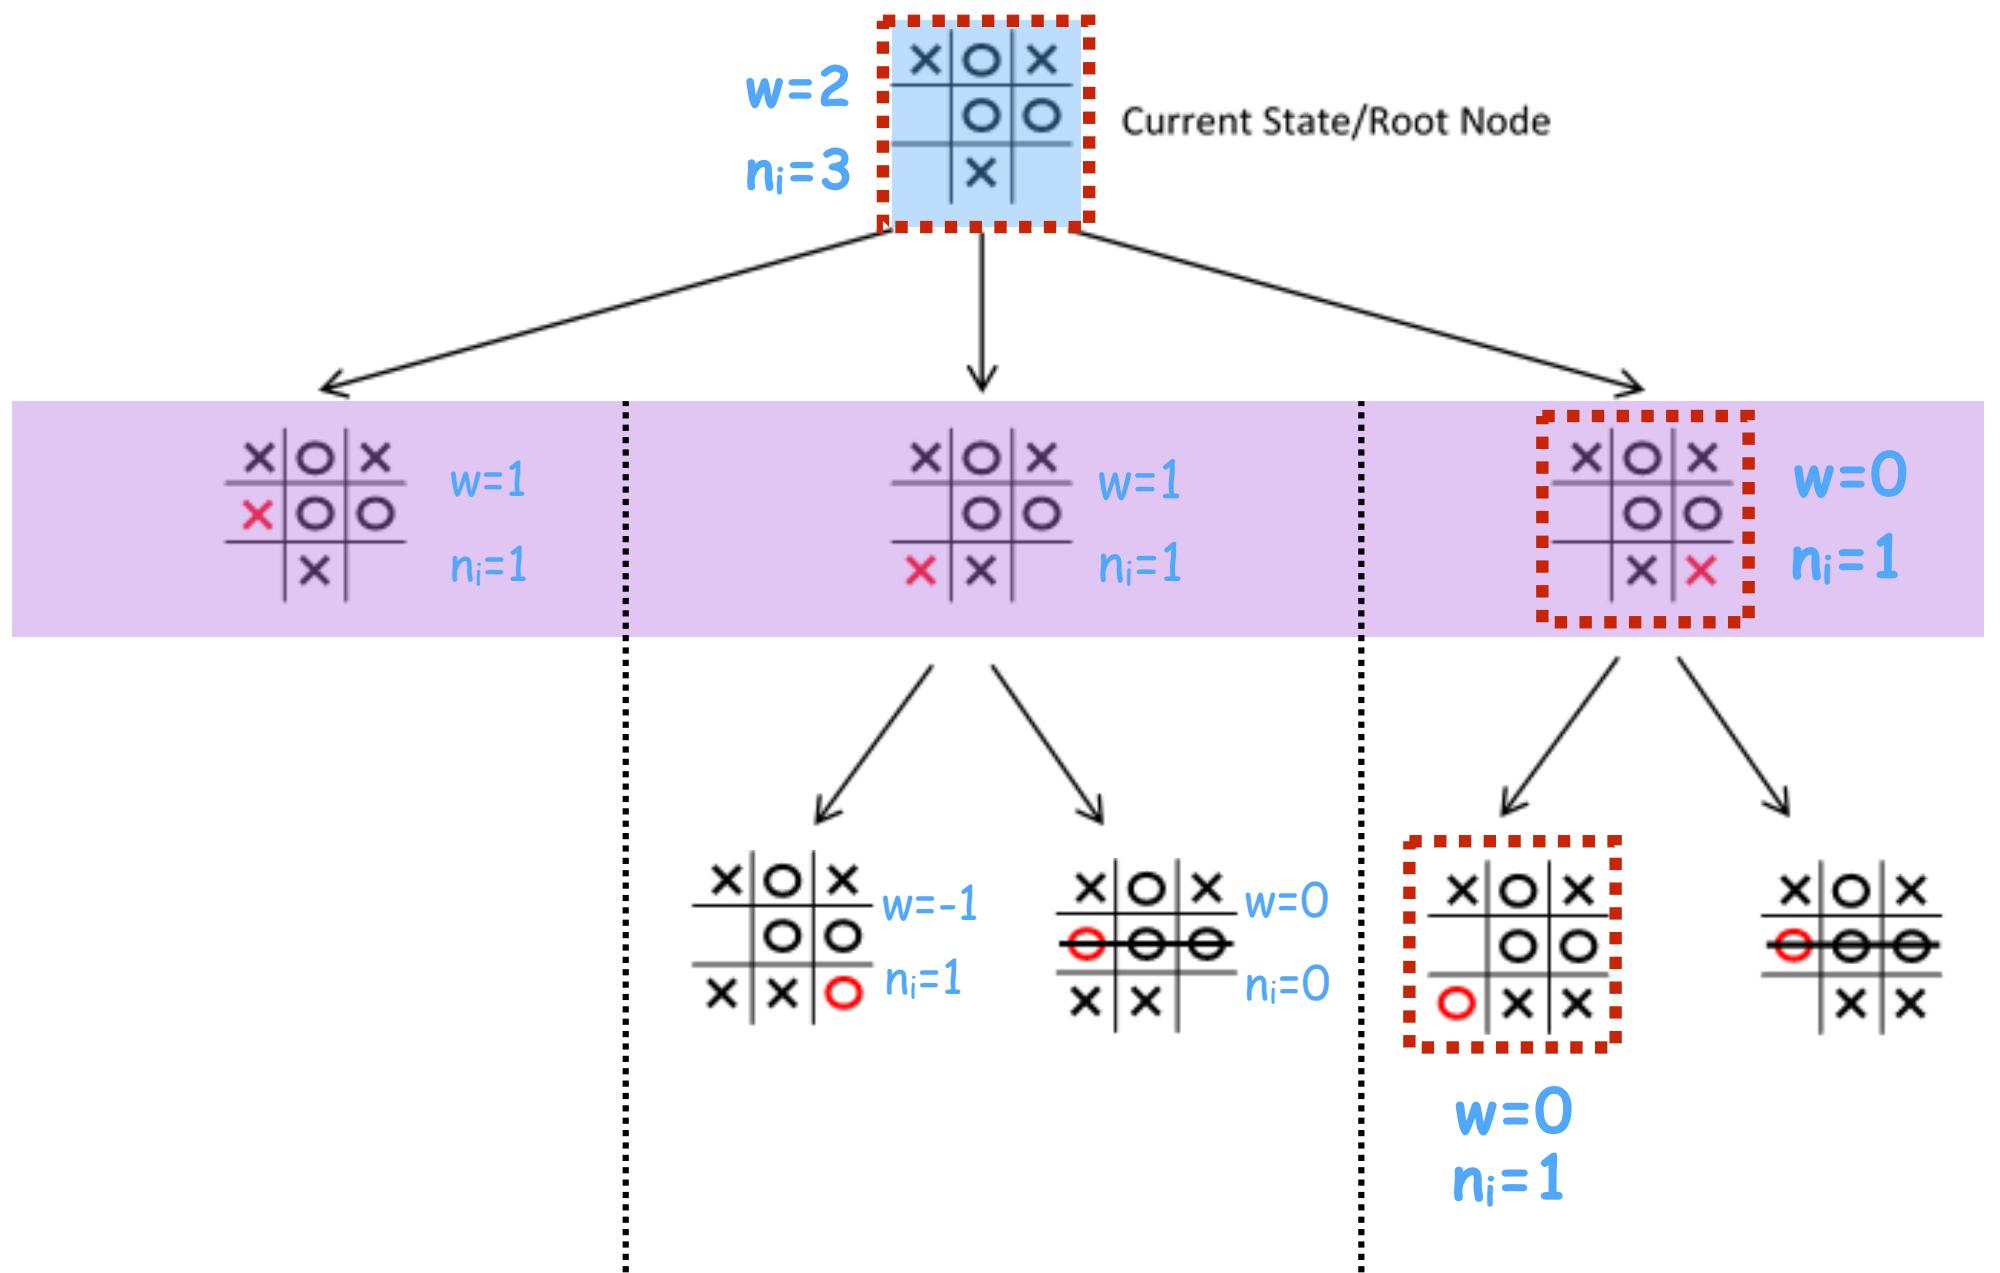
Expand



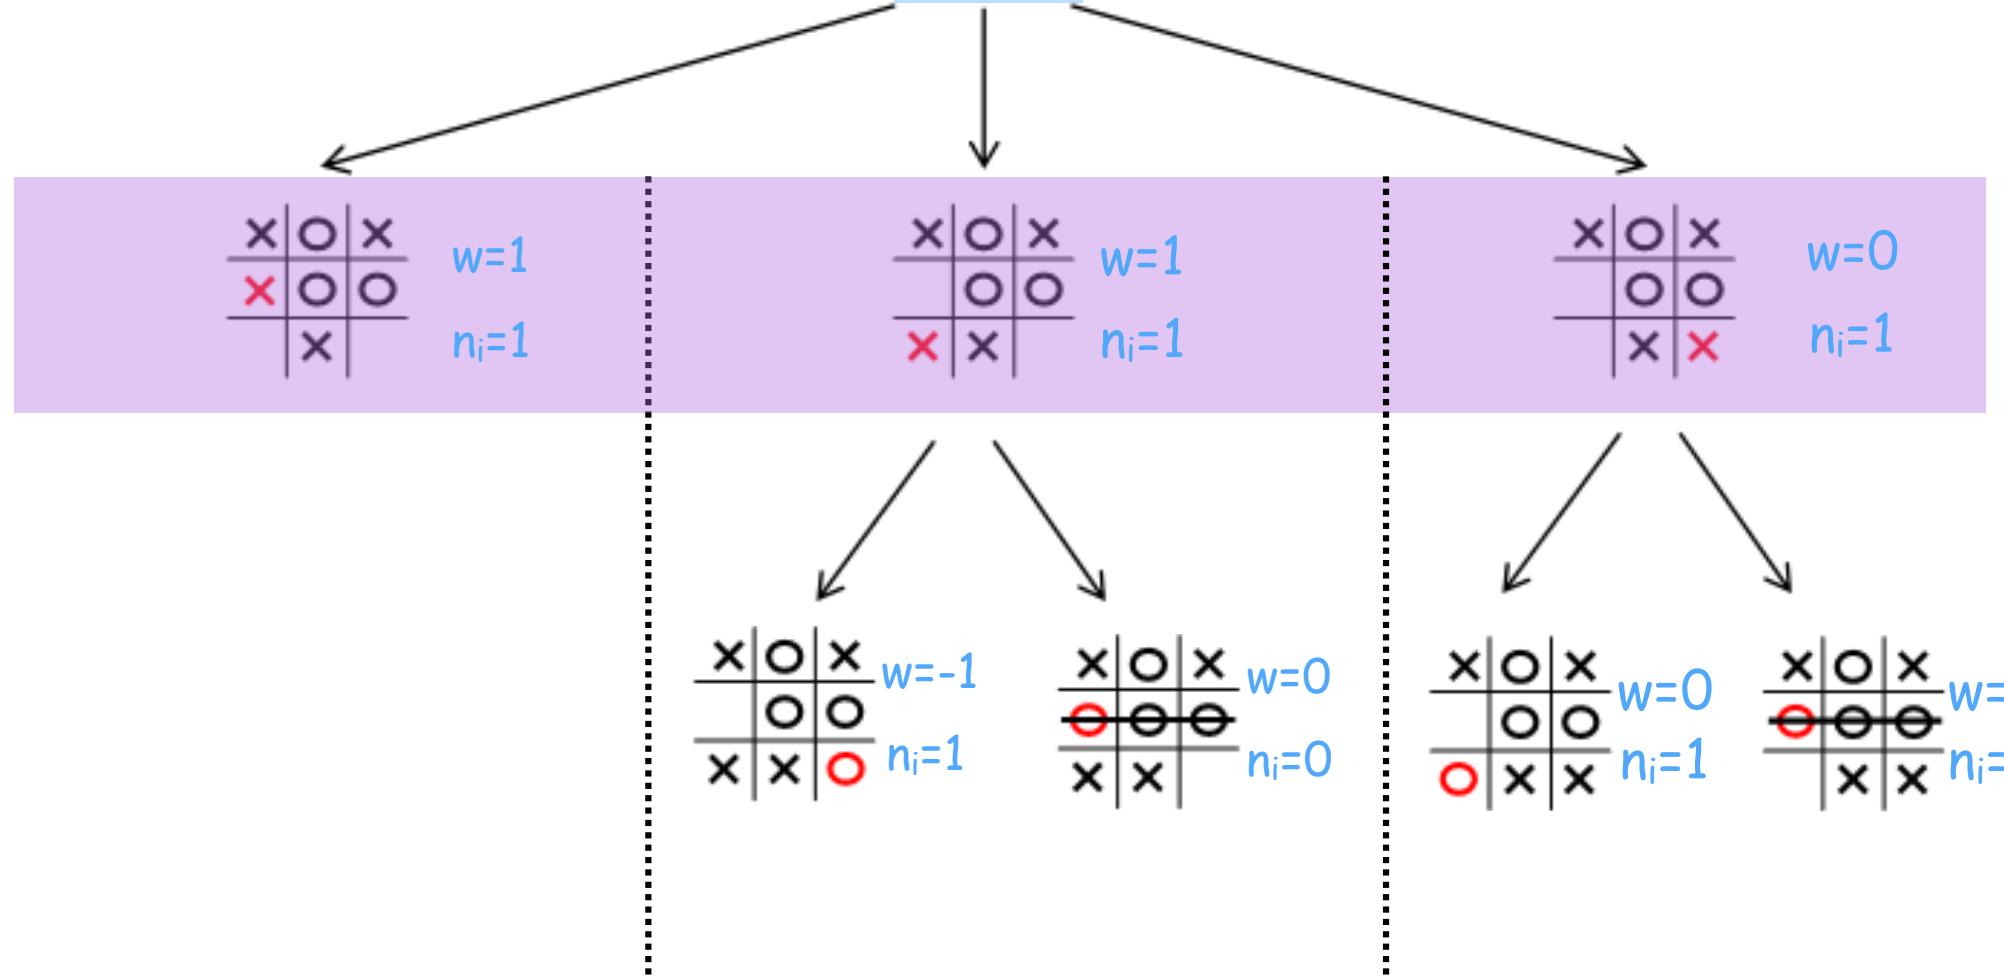
selection (expanded node)



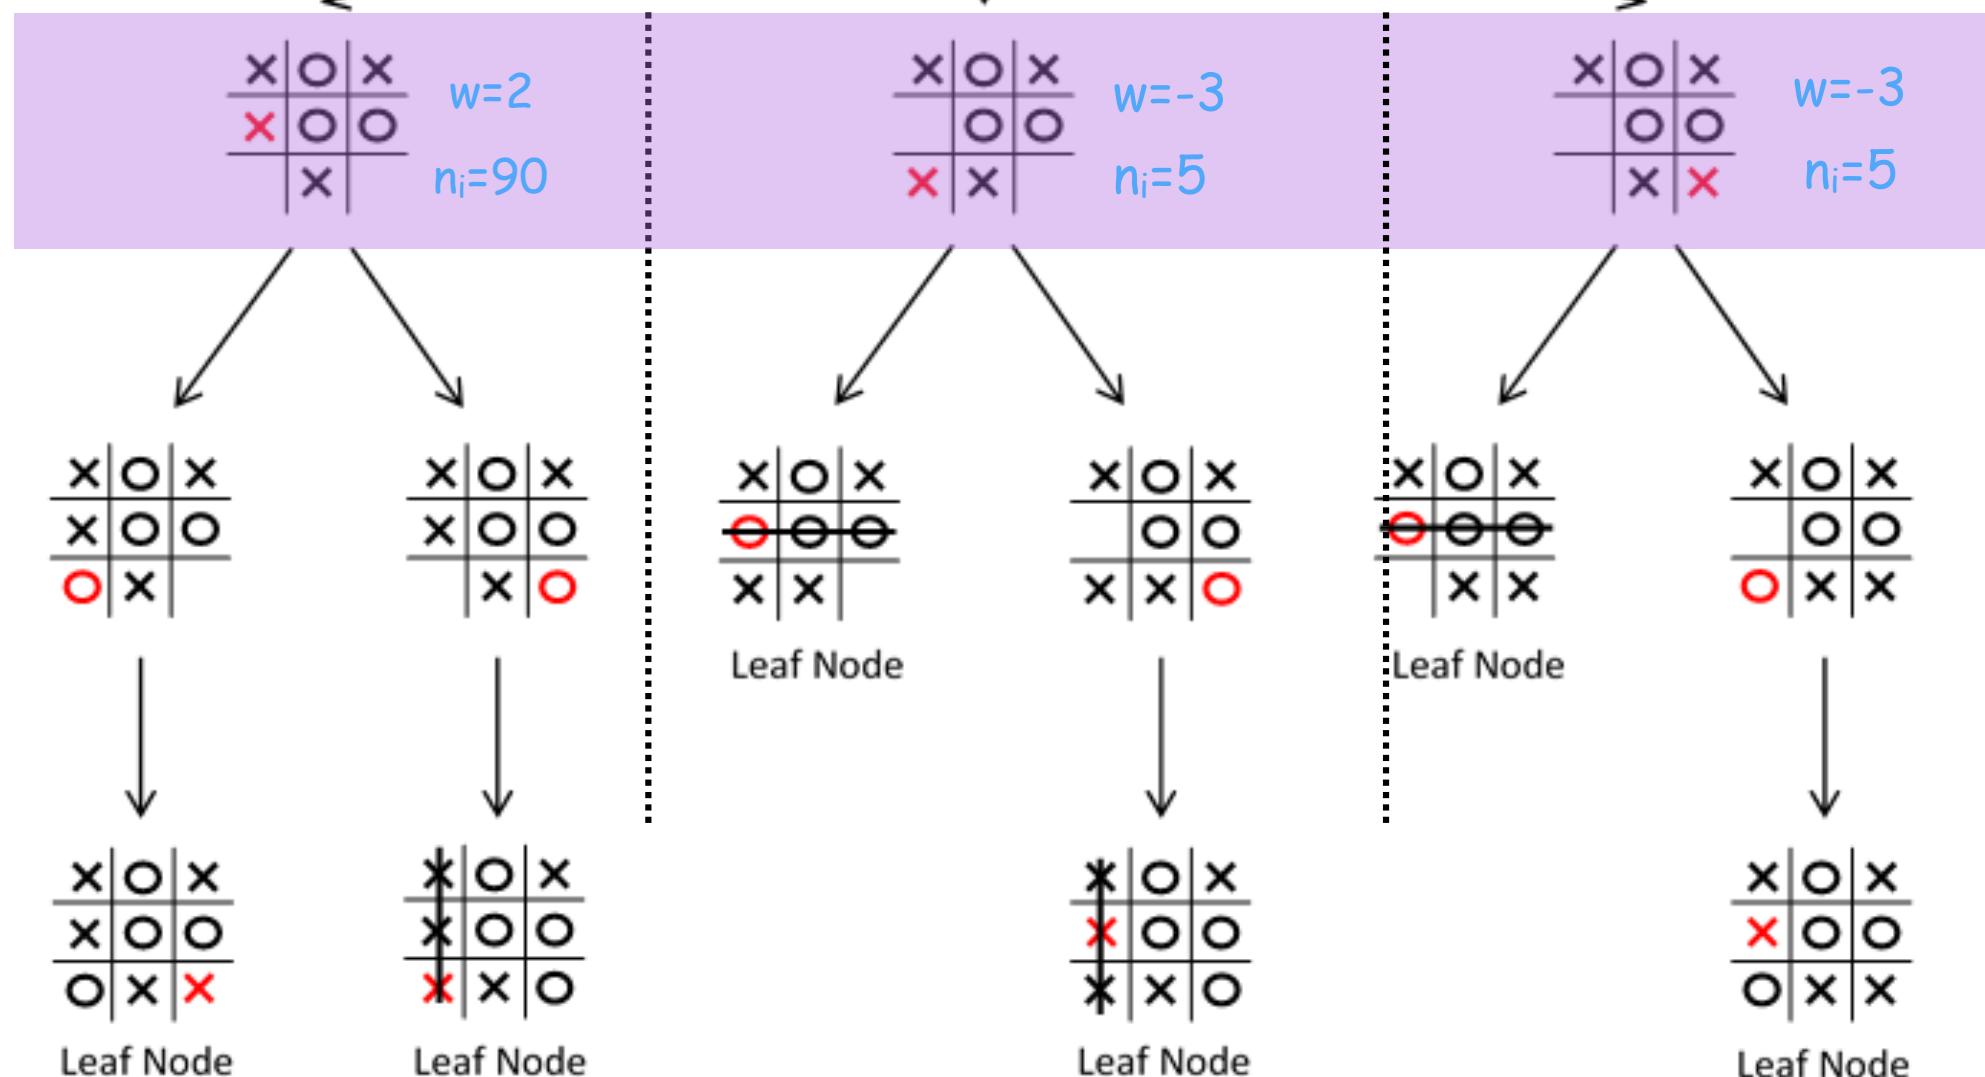
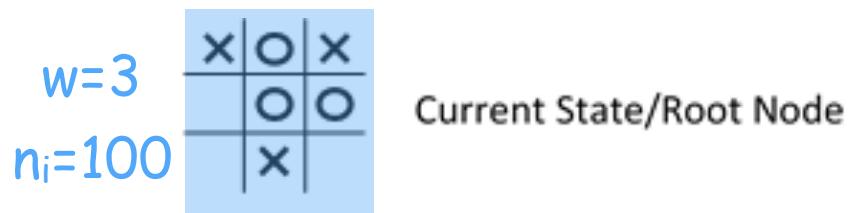




Back
Propagate



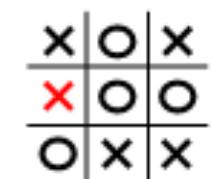
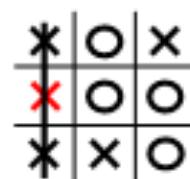
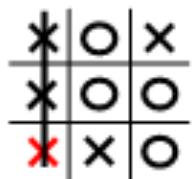
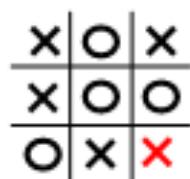
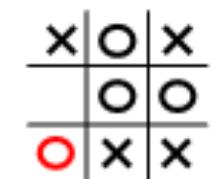
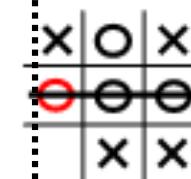
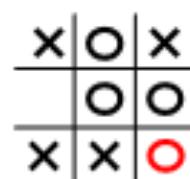
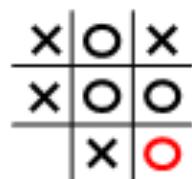
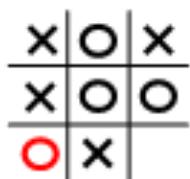
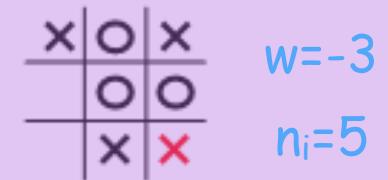
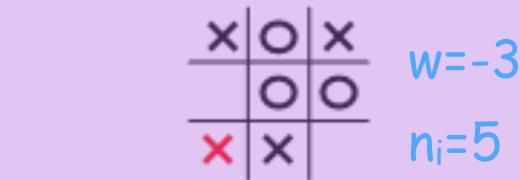
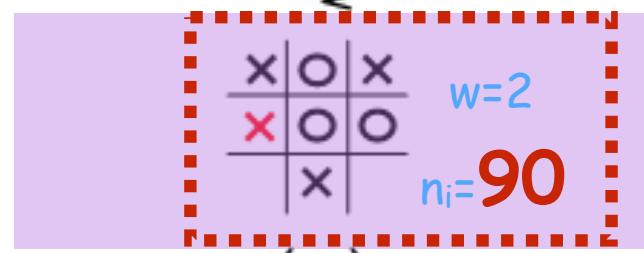
After 3 iterations



After 100 iterations

Best Move

(most visited node)



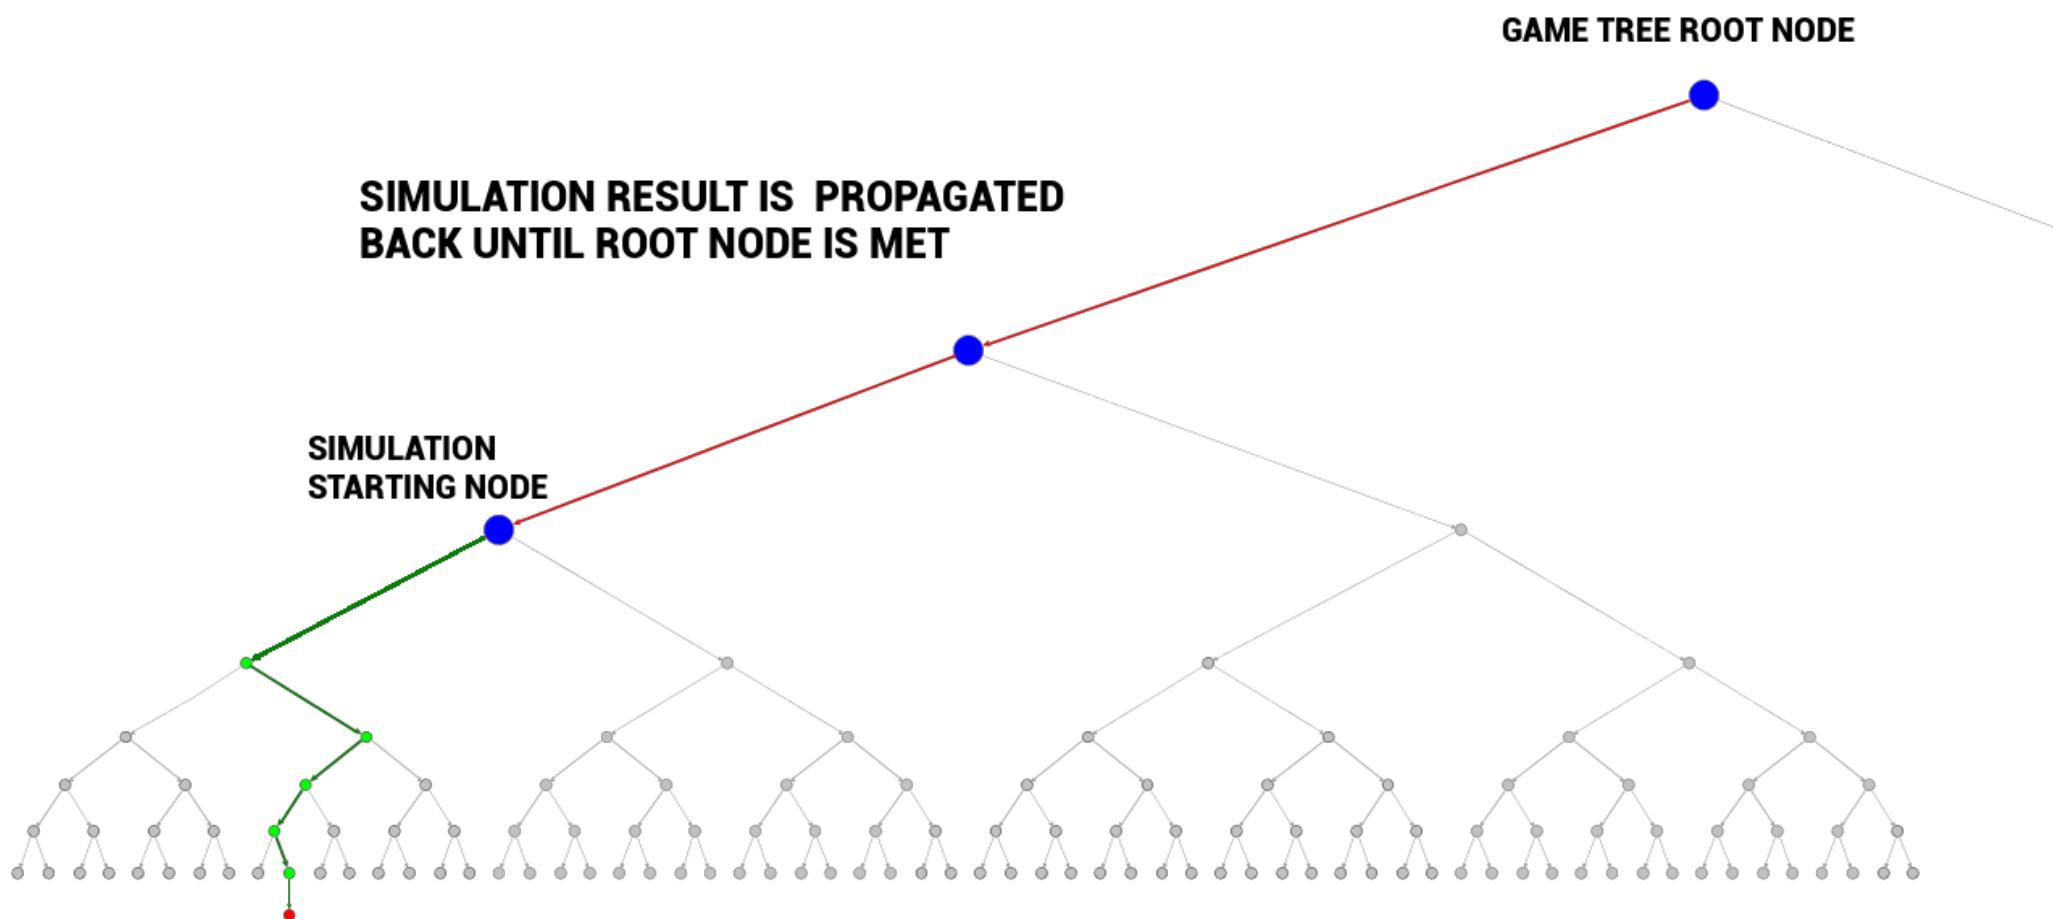
Leaf Node

Leaf Node

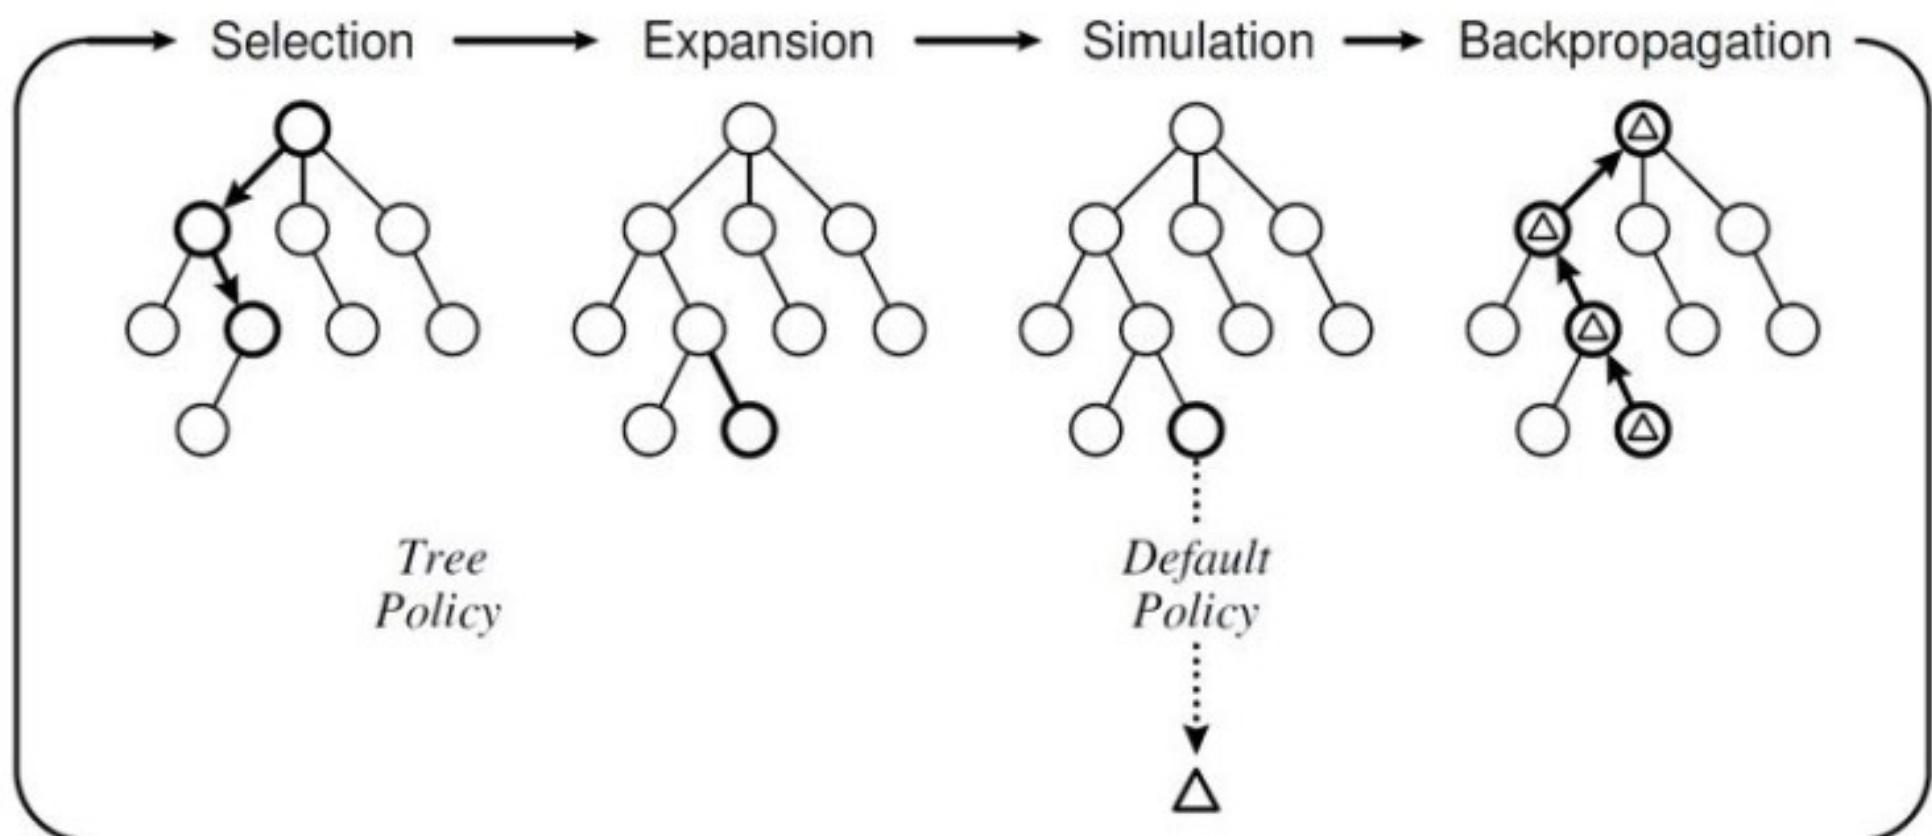
Leaf Node

Leaf Node

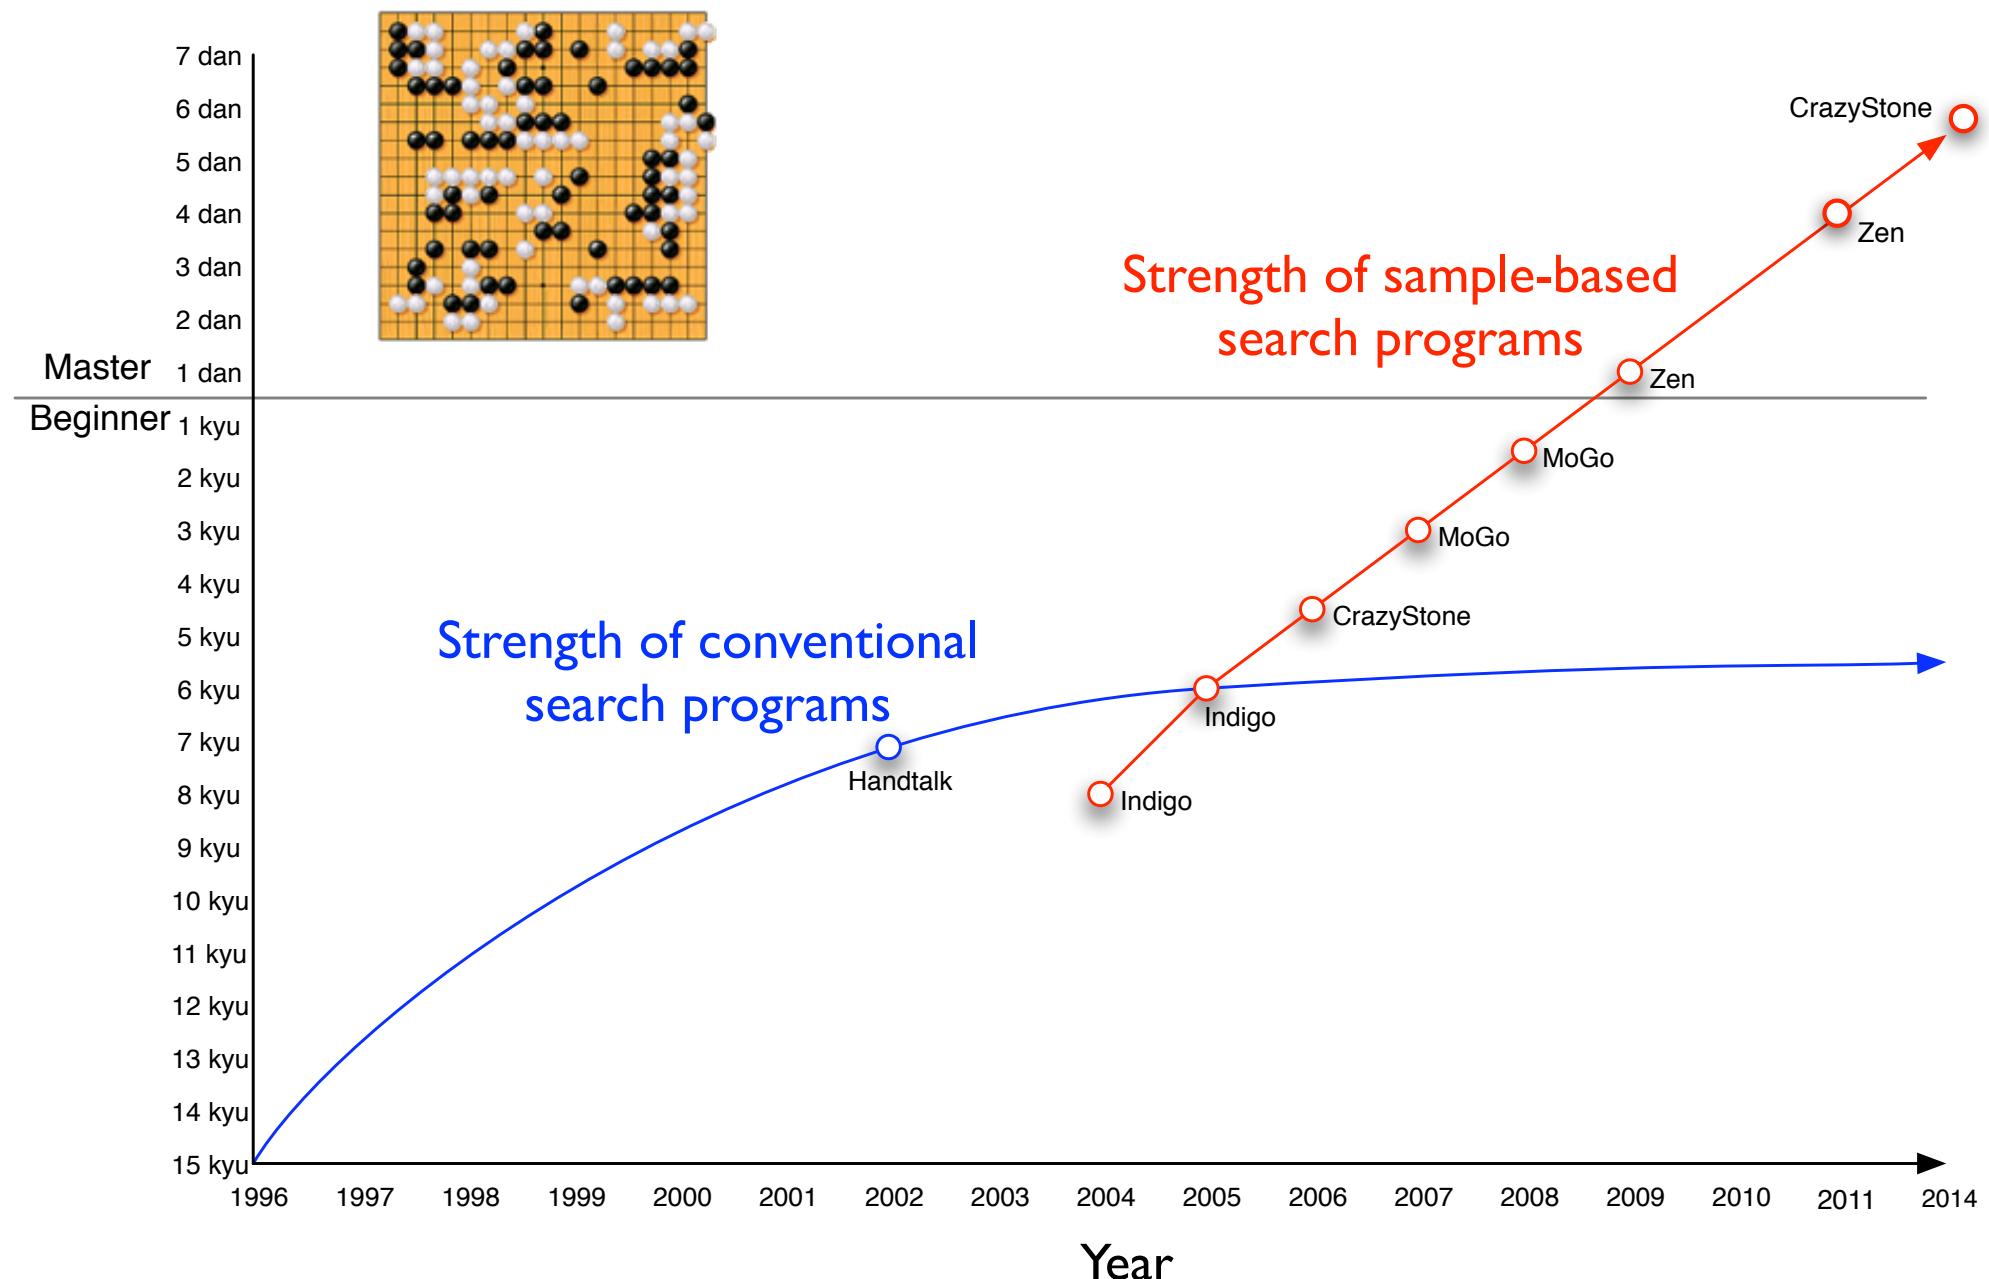
Monte Carlo Tree Search (2005)



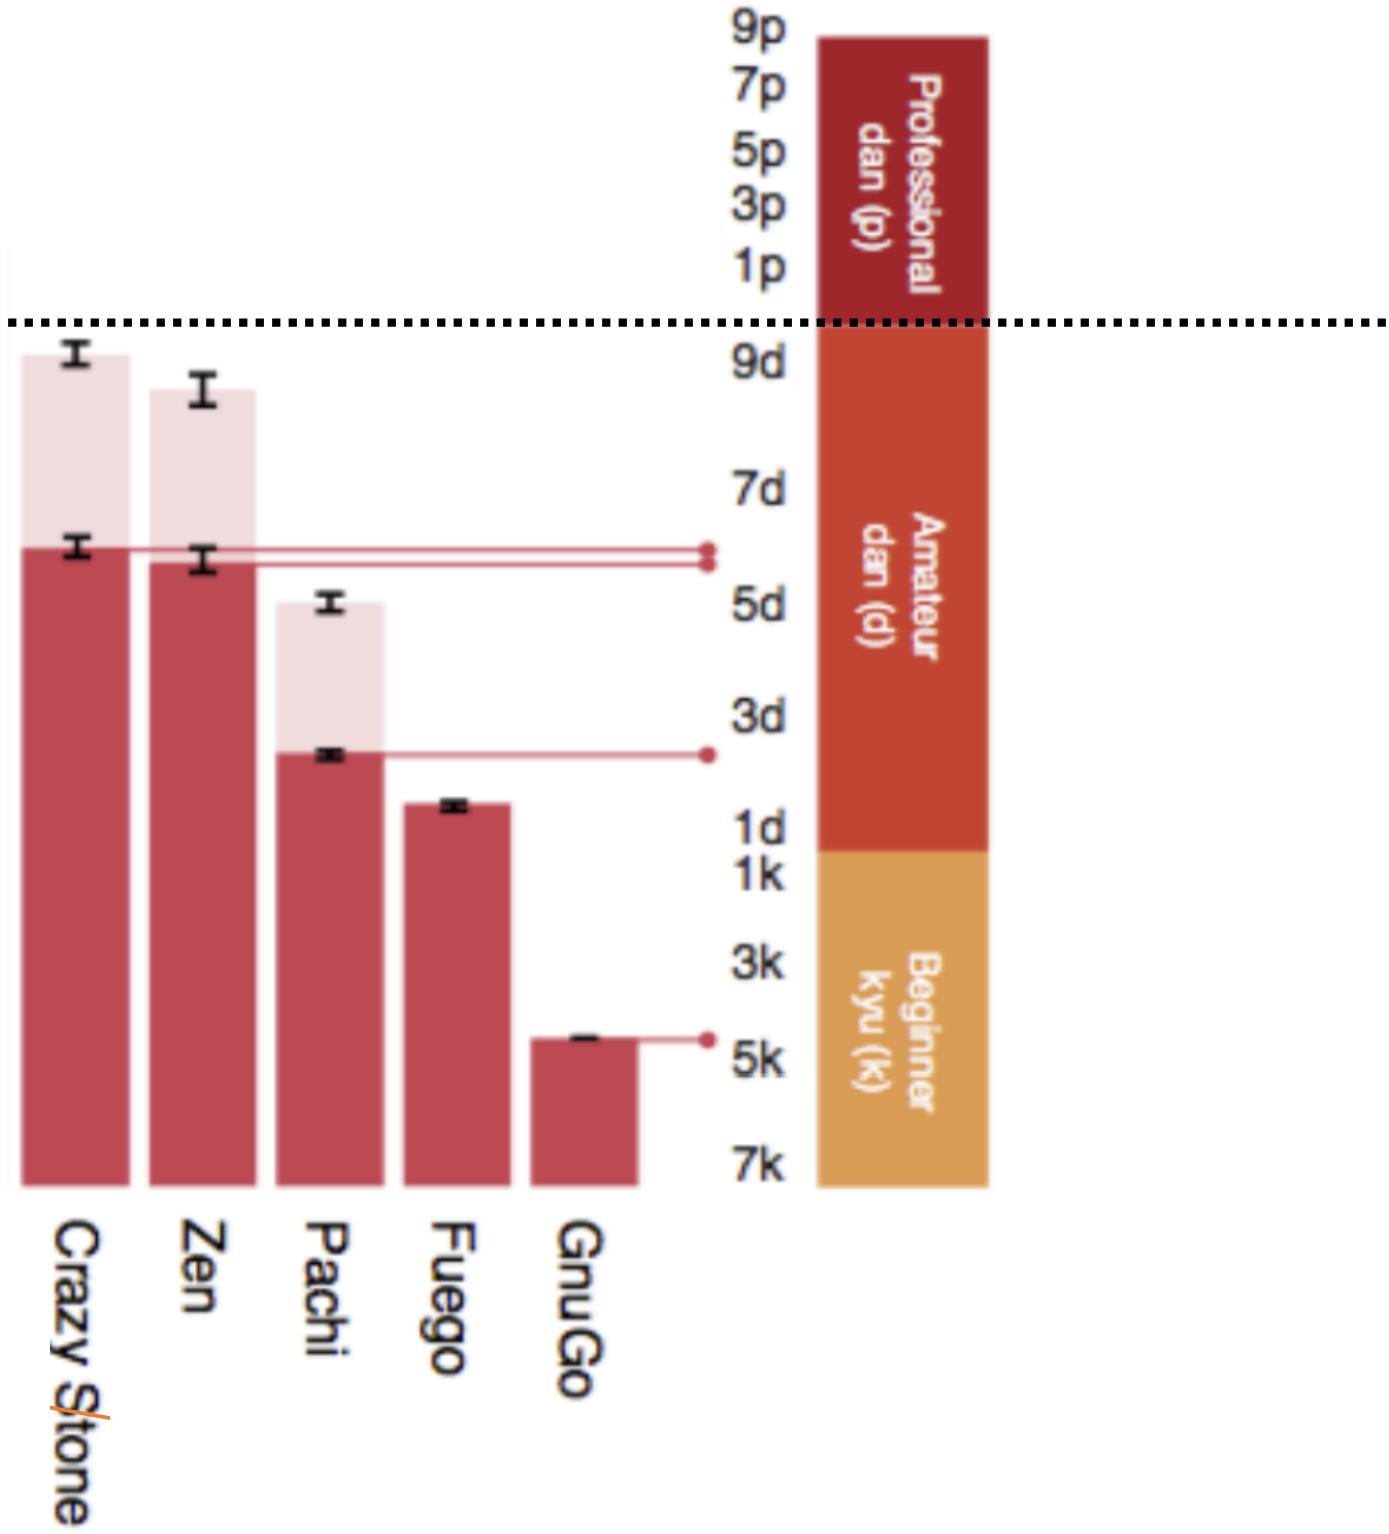
Monte-Carlo Tree-Search



Steady, exponential improvement (since MCTS, 2005) in the strength of the best computer Go programs



if 4 handicap moves
are given

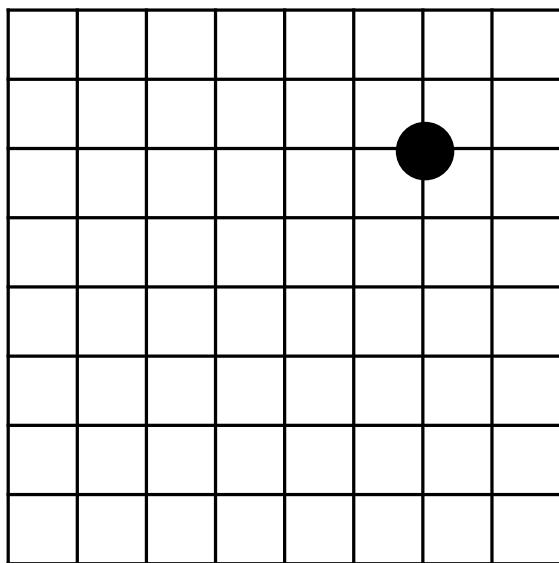


Professional Player vs MCTS

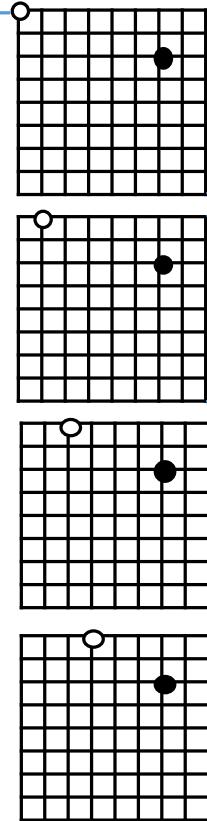


The Huge Tree

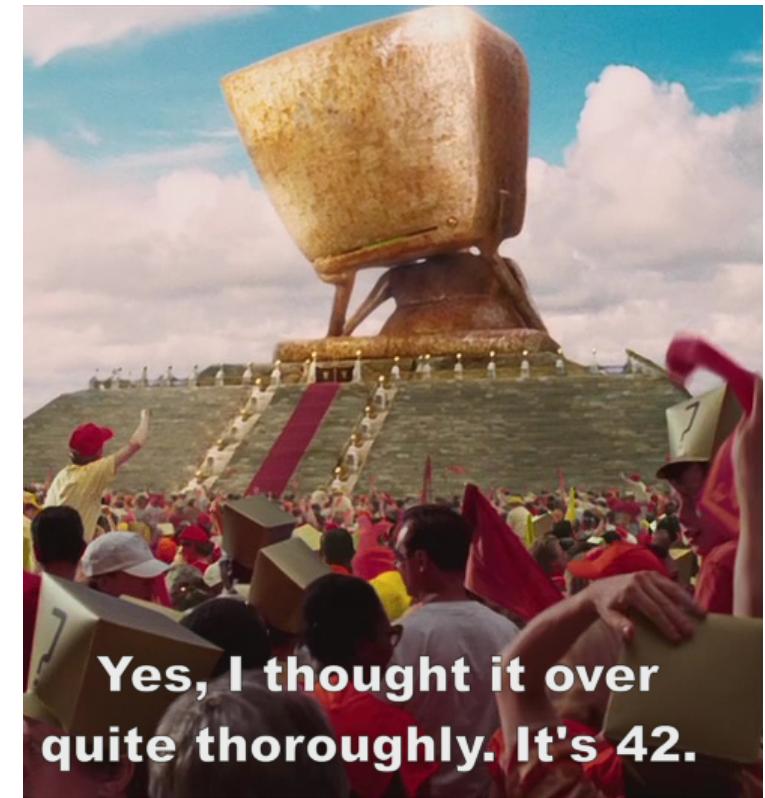
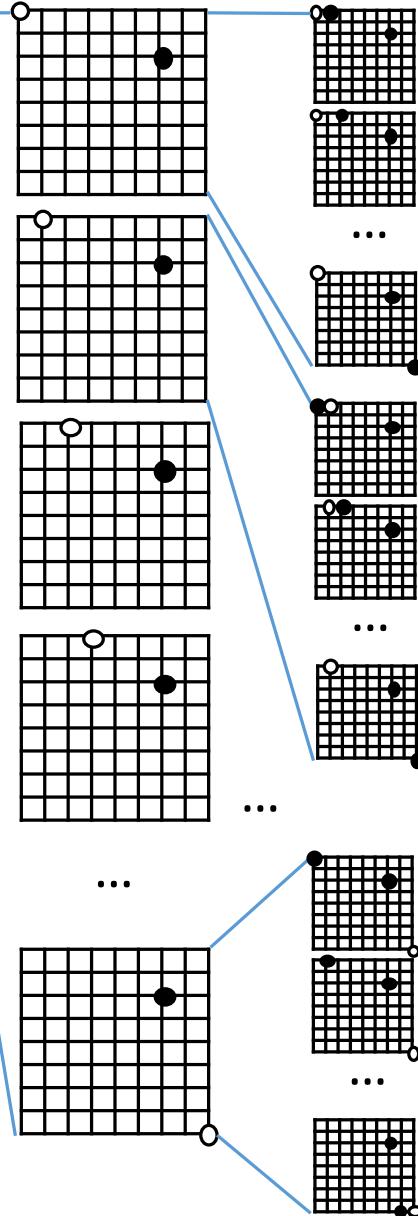
$d = 1$

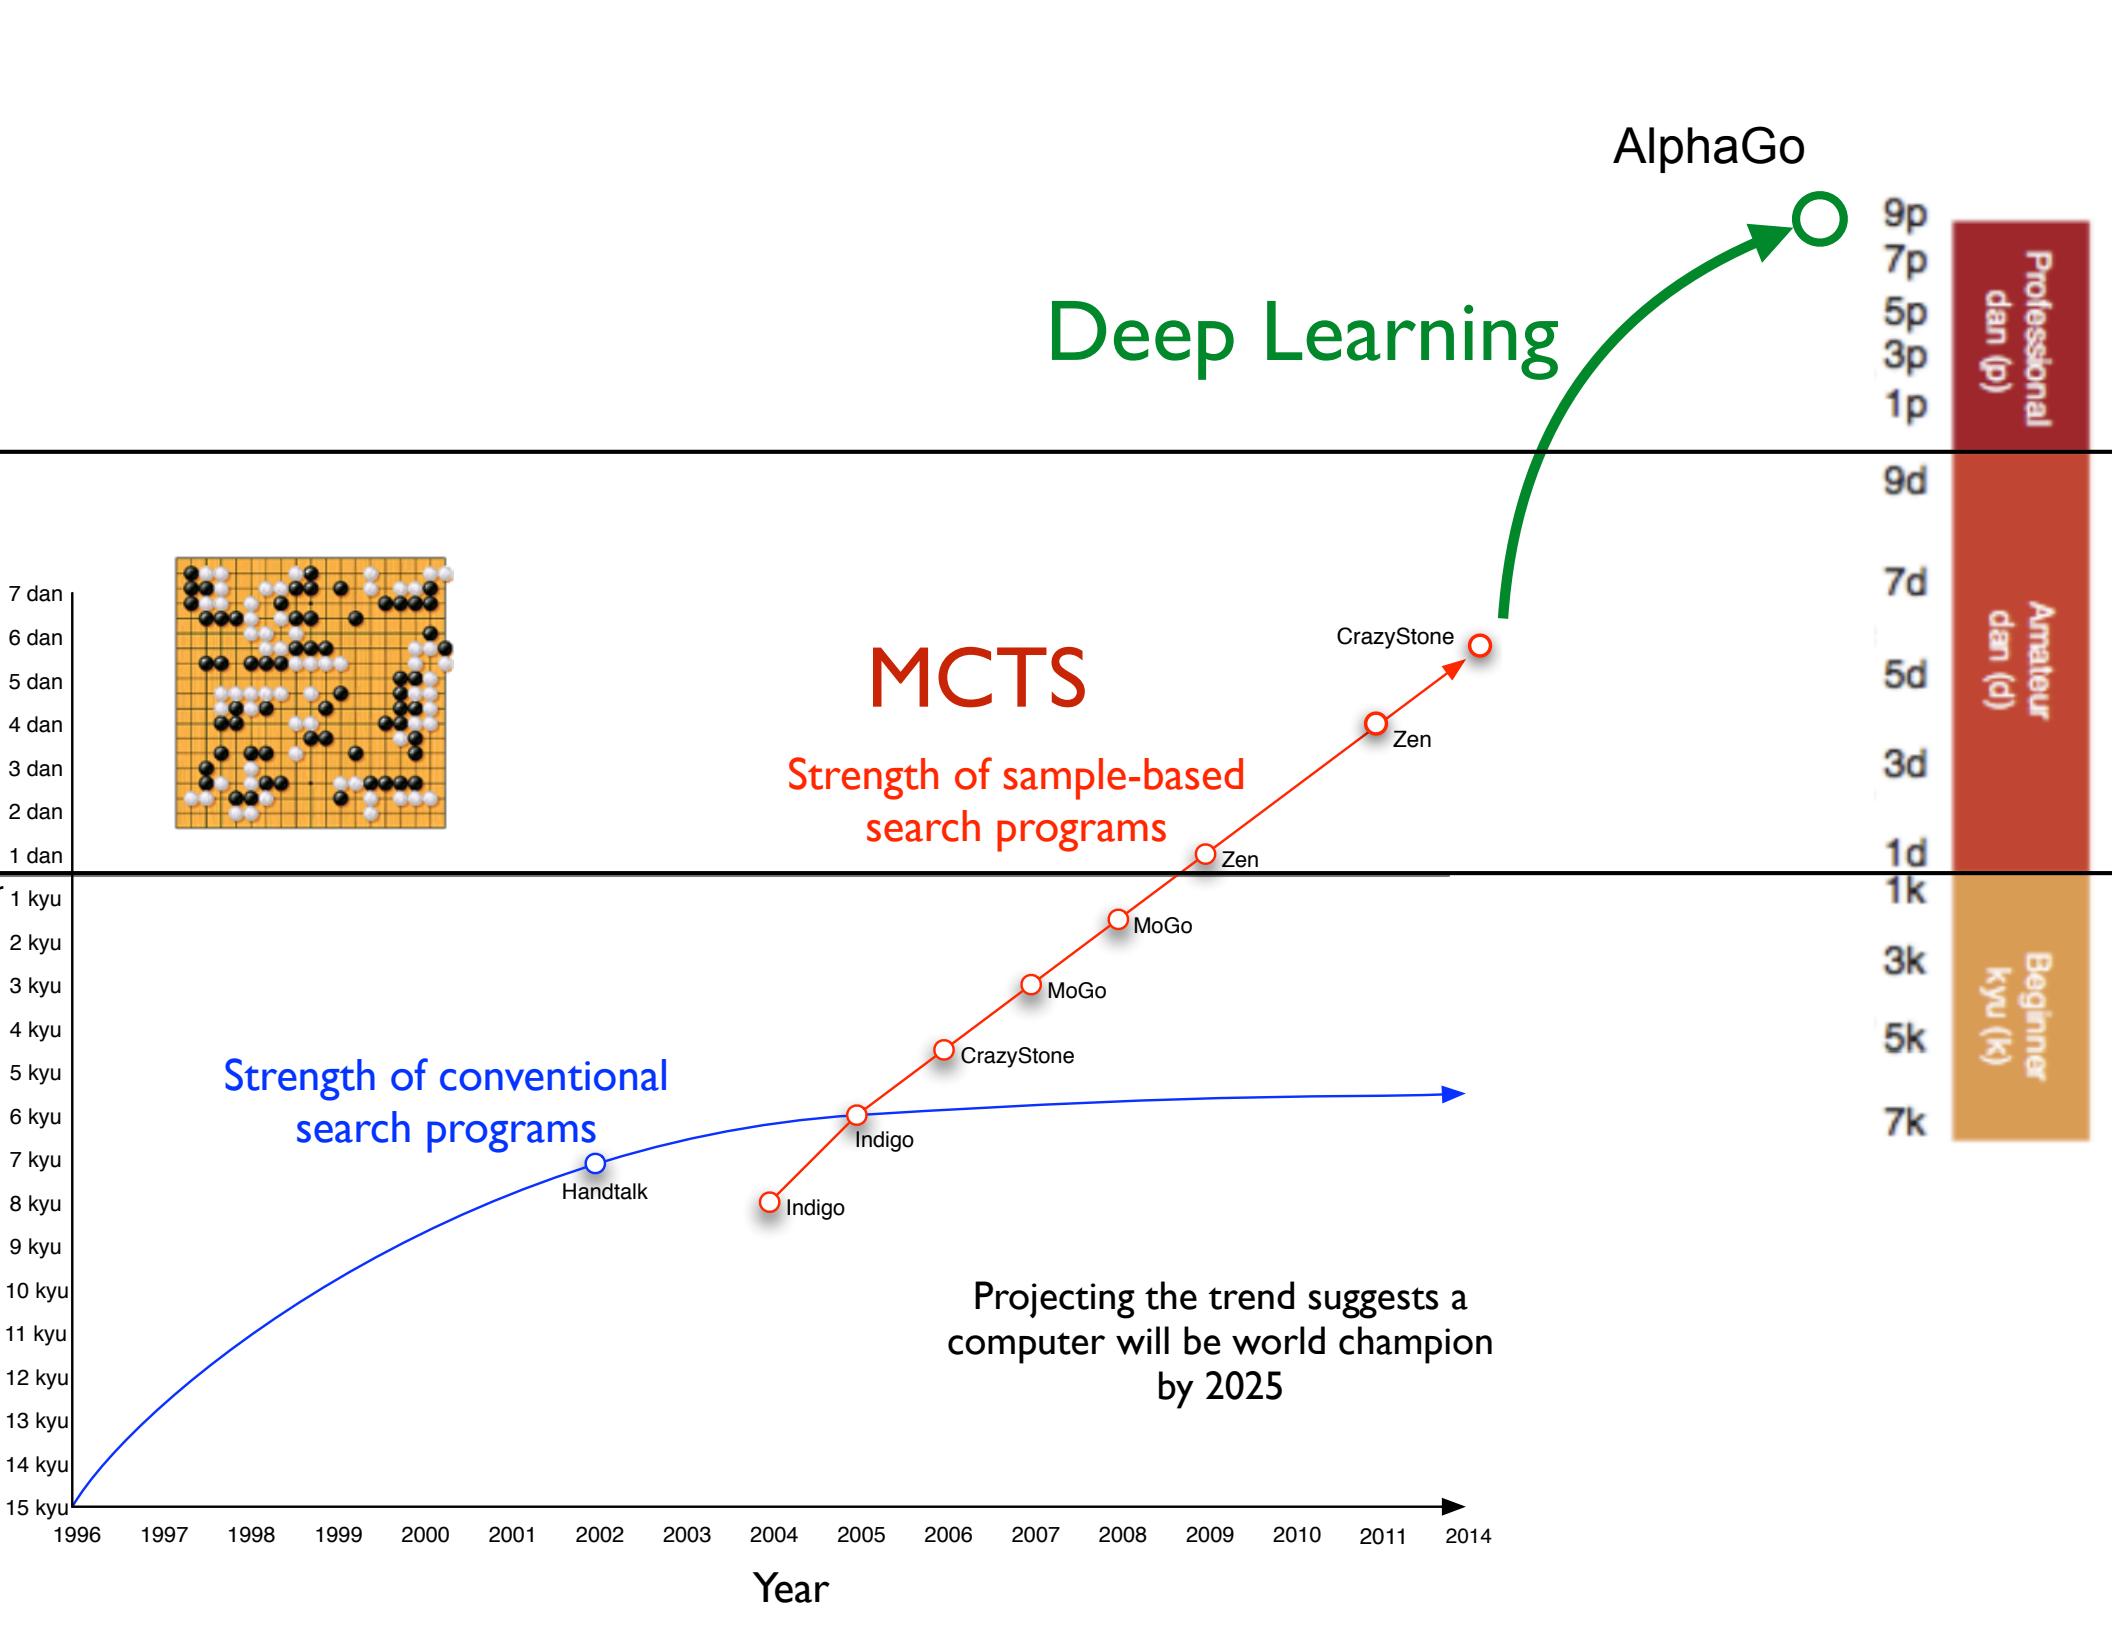


$d = 2$

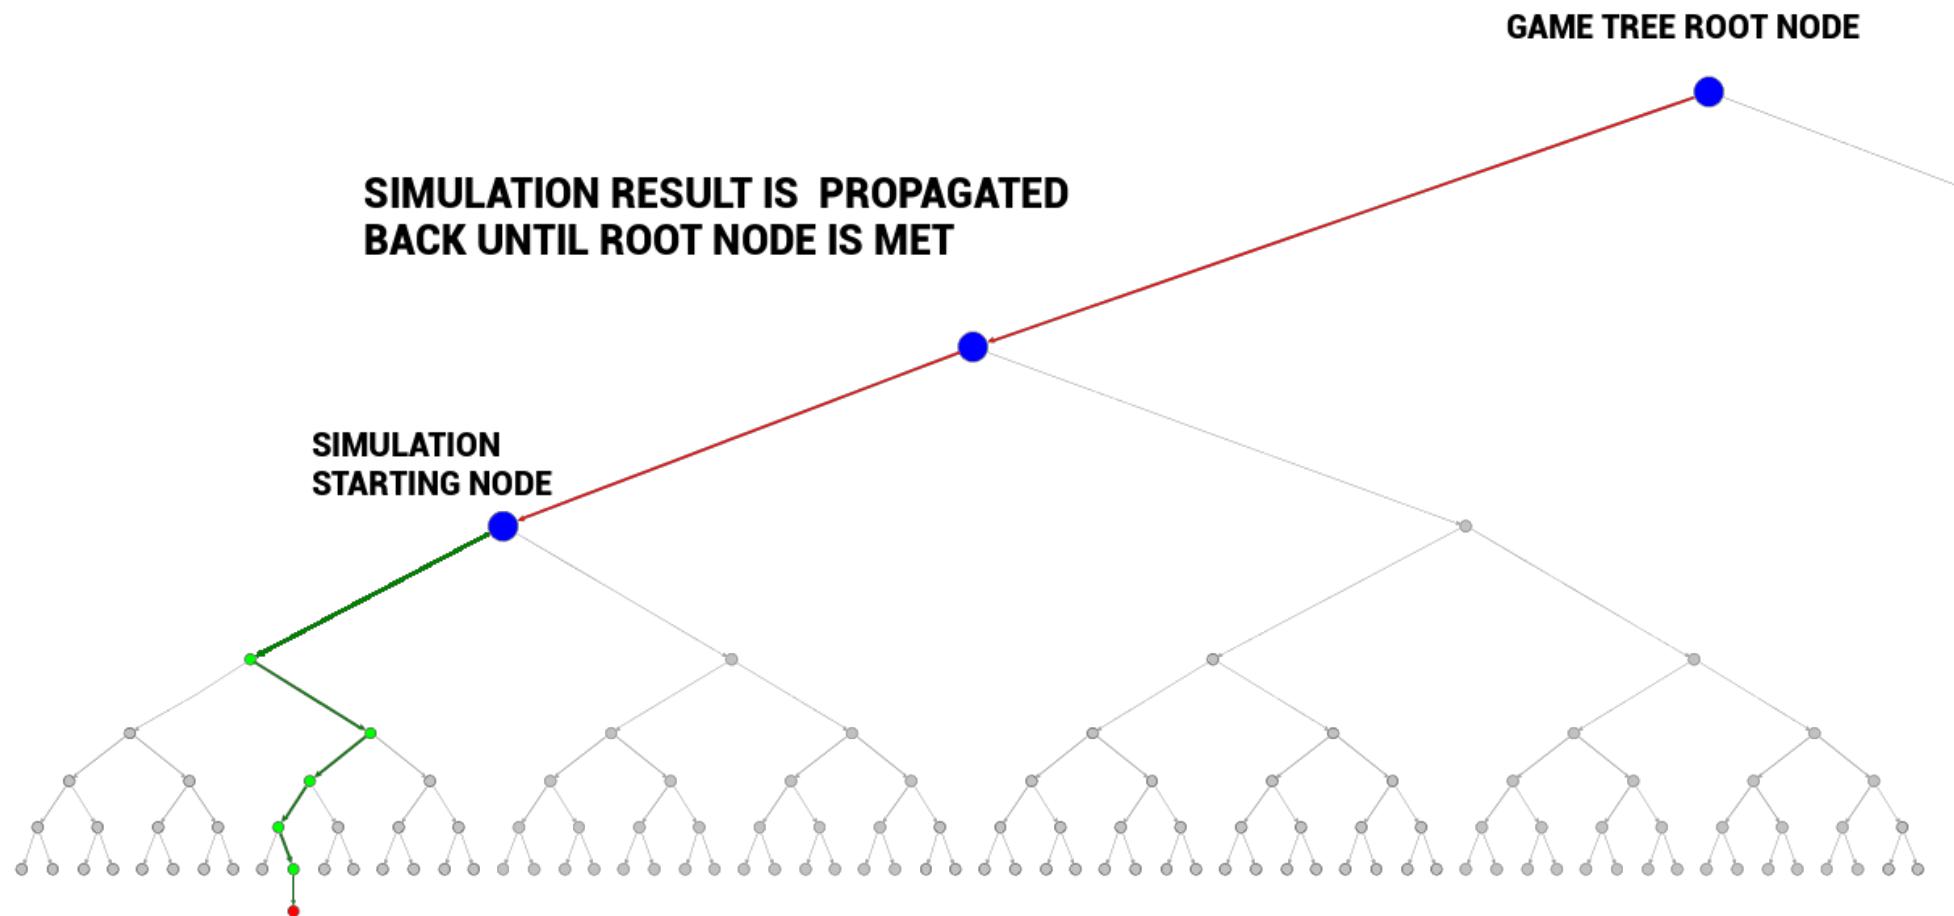


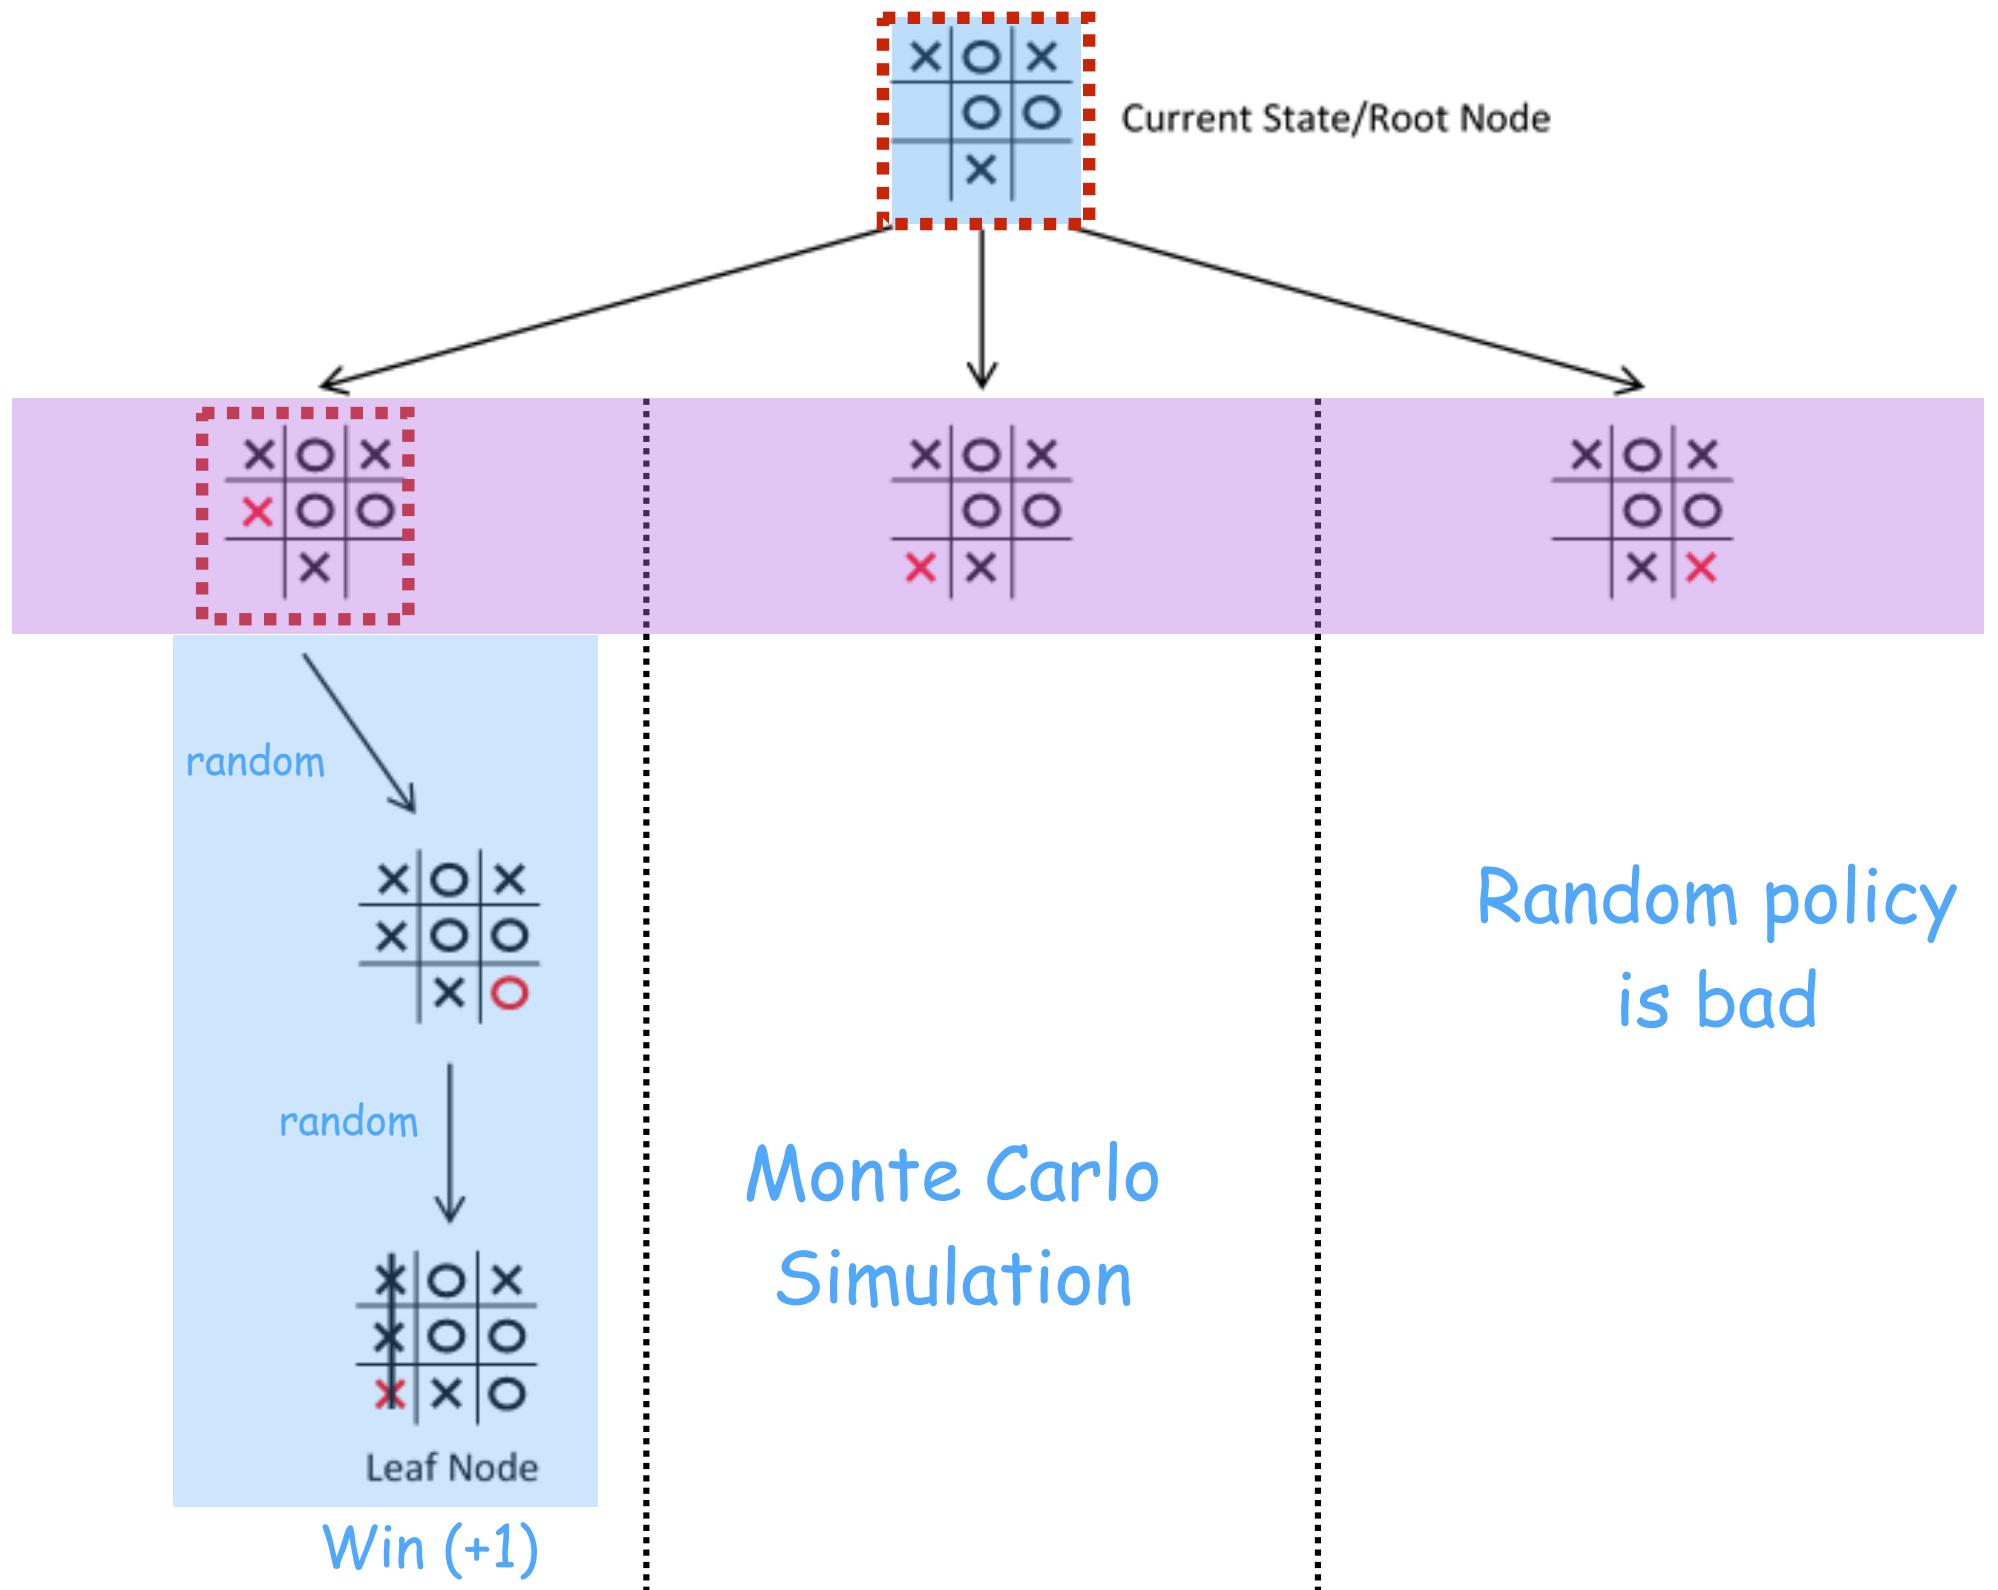
$d = 3$



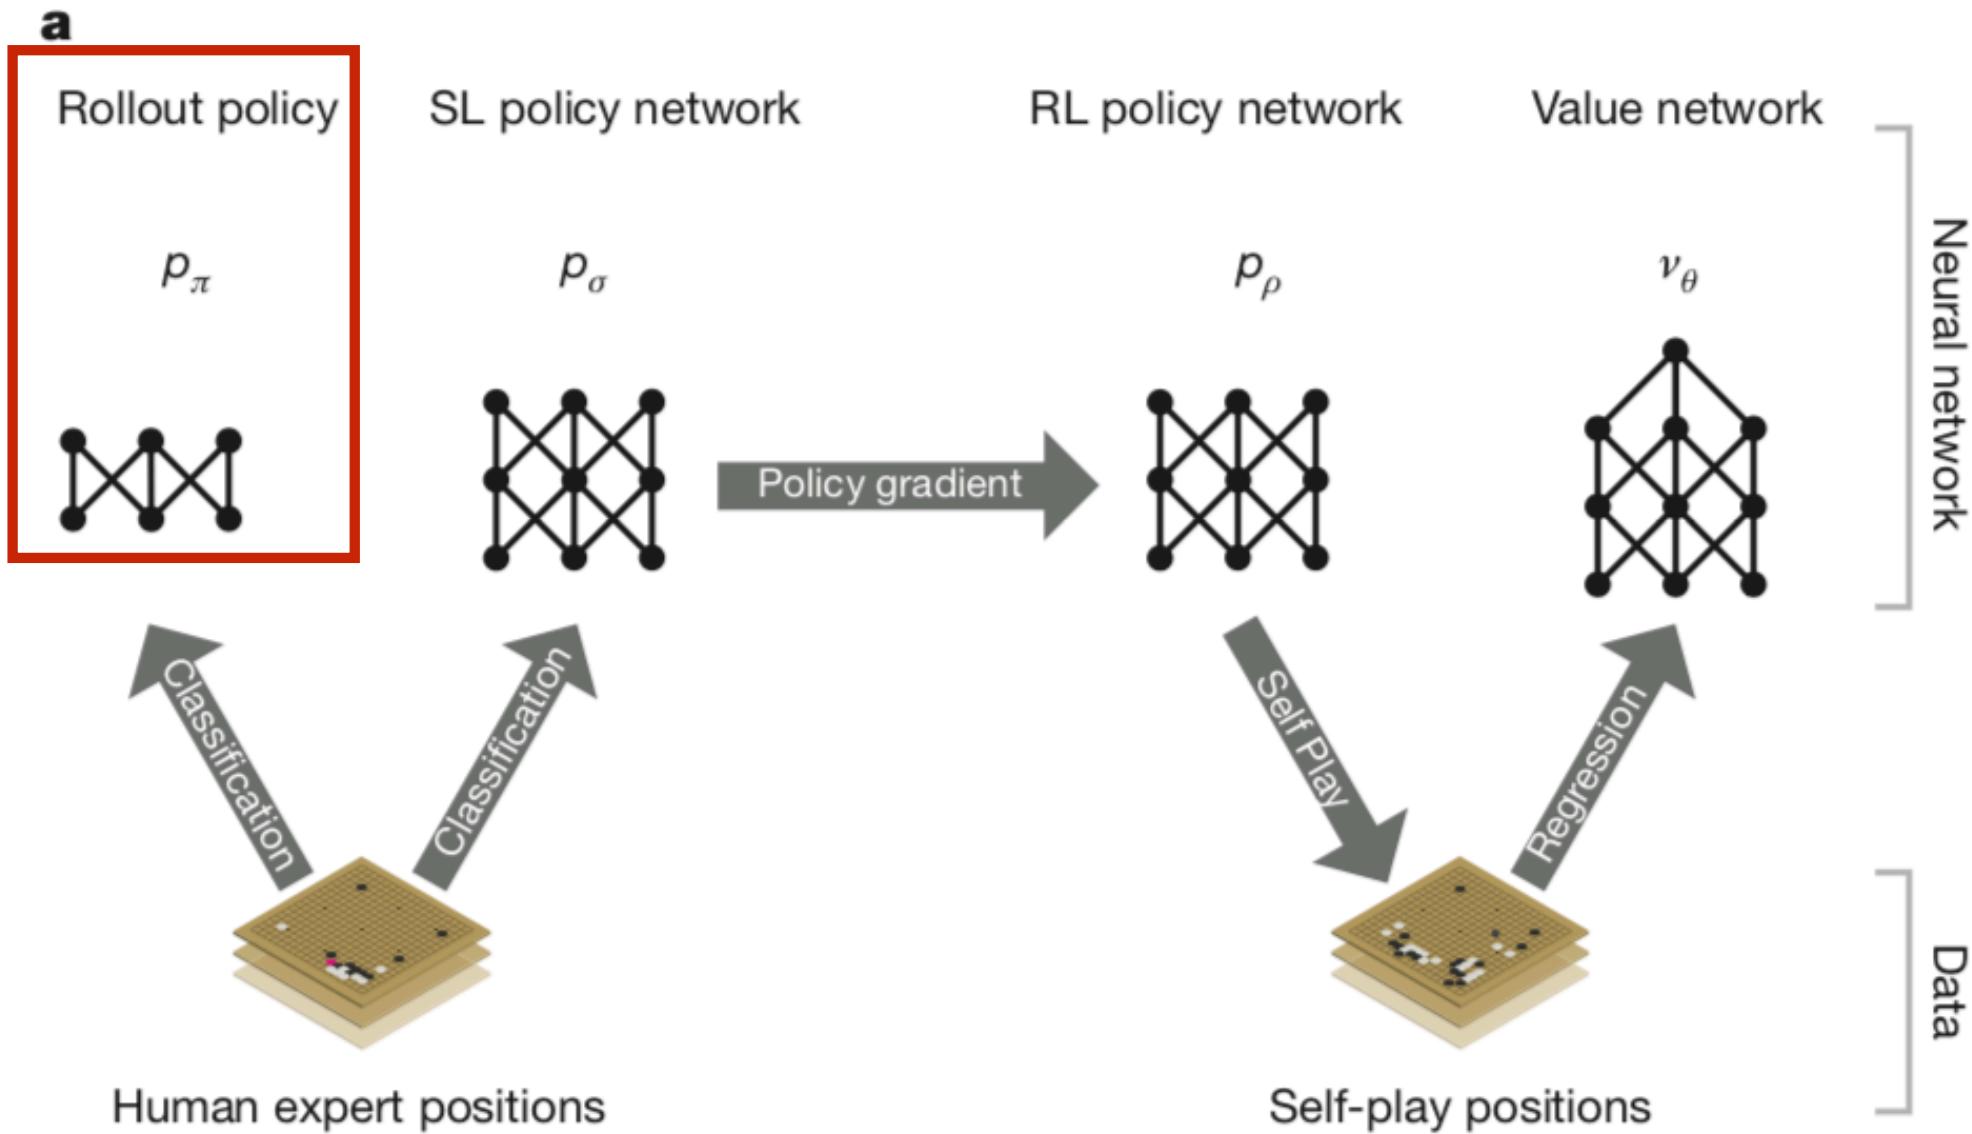


What to improve in MCTS ?

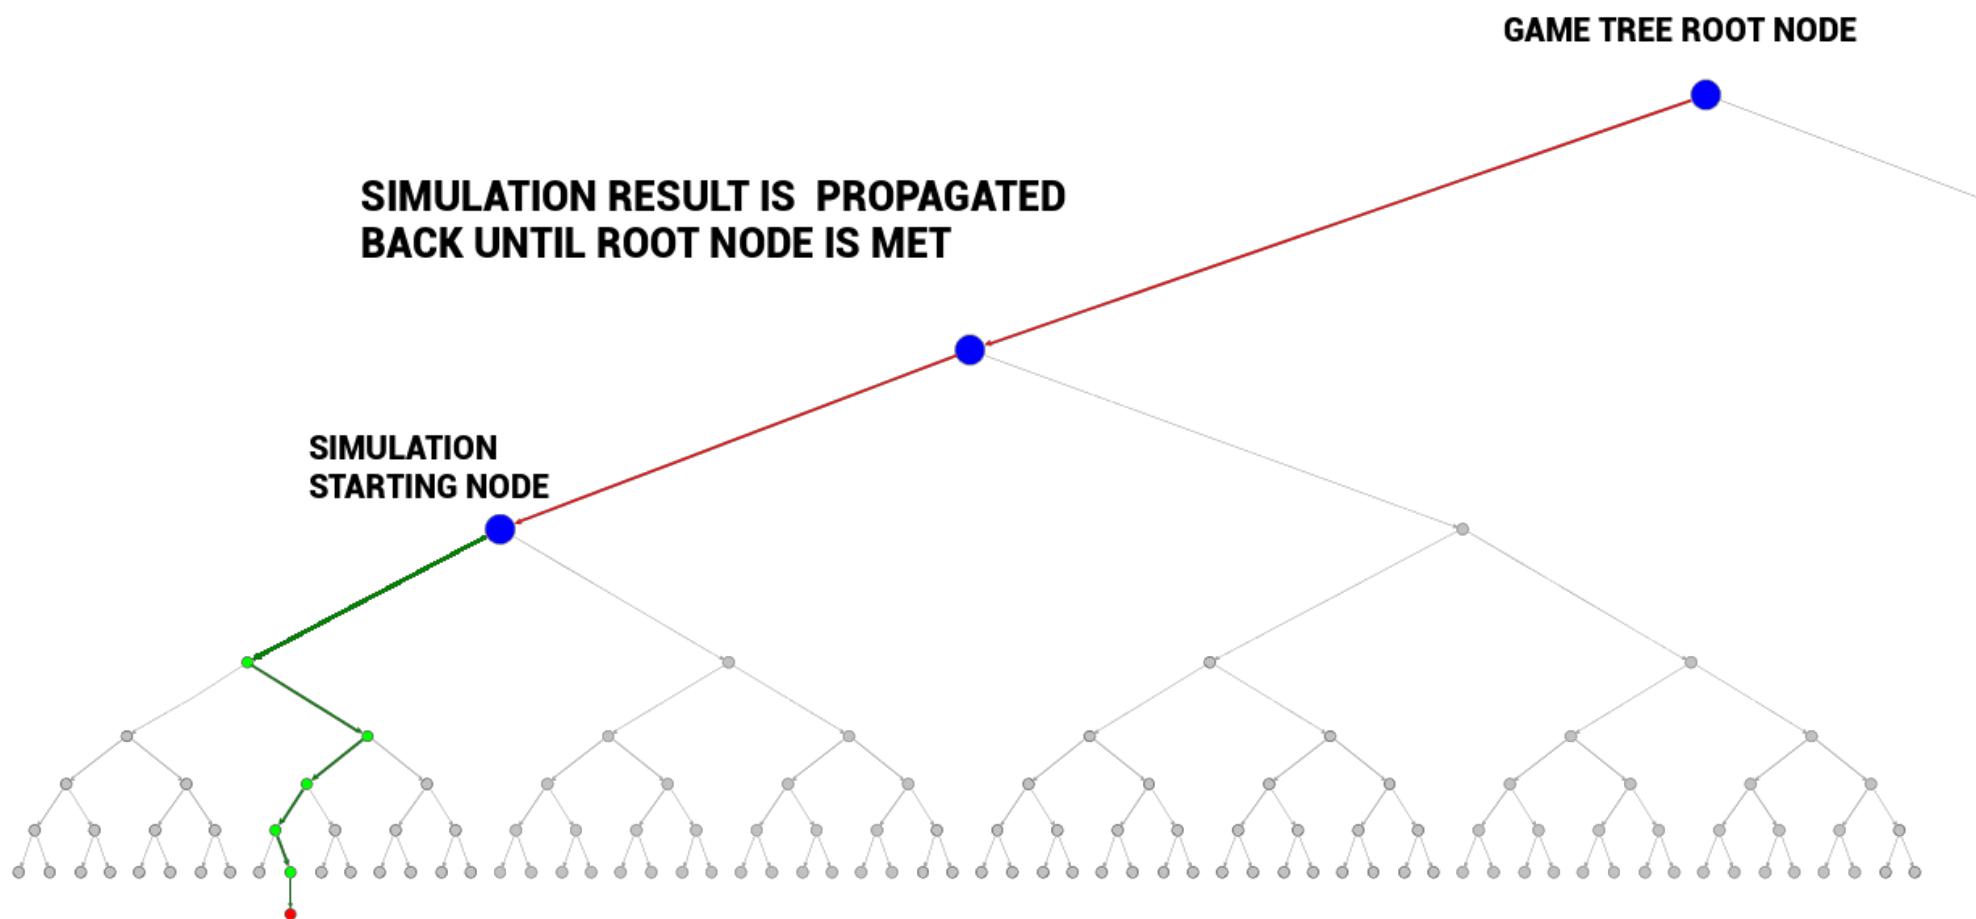




AlphaGo



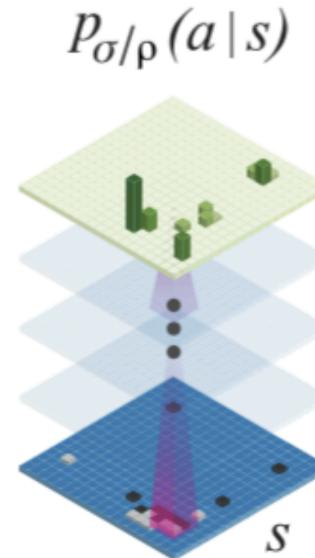
Selection



Reducing Search Space

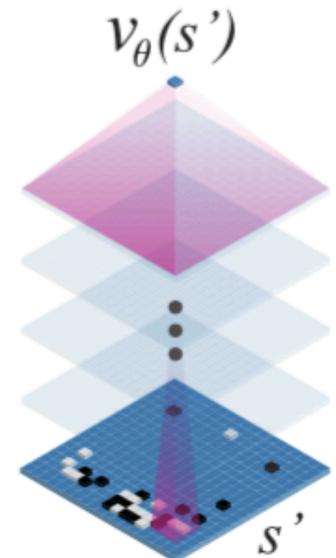
1. Reducing “action candidates”
(Breadth Reduction)

Policy Network



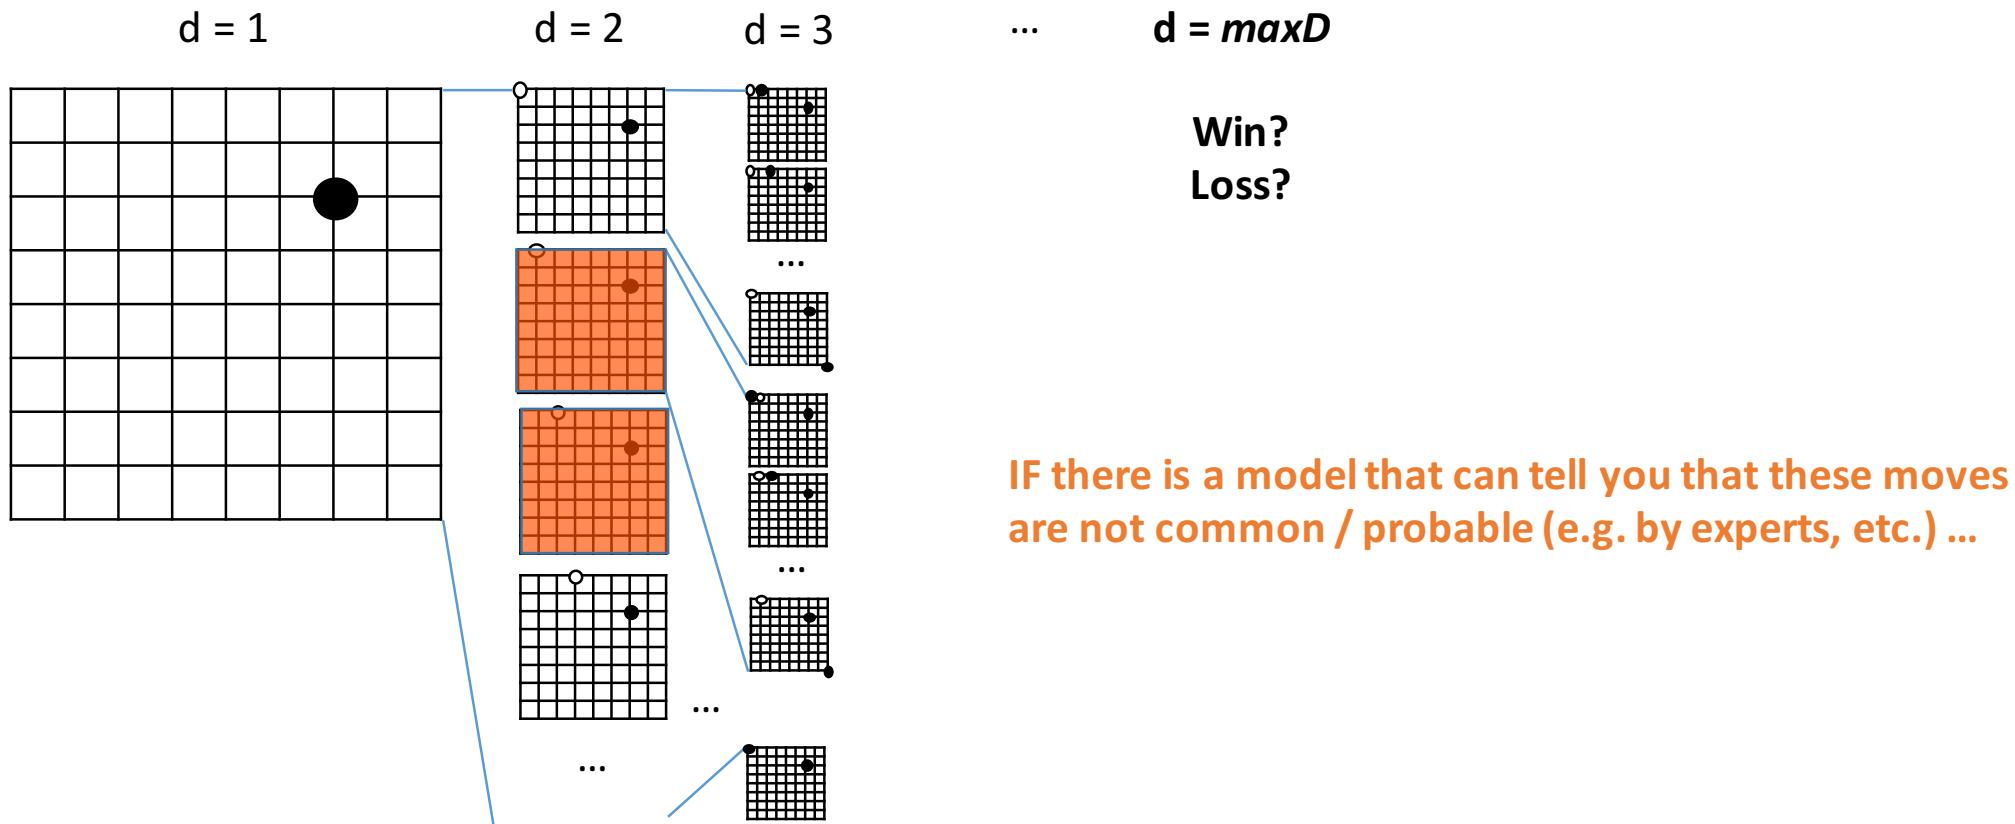
2. Board Evaluation (Depth Reduction)

Value Network



Reducing num of Choices in each decision

1. Reducing “action candidates” (Breadth Reduction)



1. Reducing “action candidates”

(1) Imitating expert moves (supervised learning)

Current Board

00 000 0000
00 000 **1**000
0**-1**00 **1**-**1**00
0**1** 00**1**-**1**000
00 00-**1**0000
00 000 0000
0**-1**000 0000
00 000 0000

Deep Learning
(13 Layer CNN)

Next Action

0000000 000
0000000 000
0000000 000
000000.20.100
00000 **0.4**0.200
000000.1 000
000000 000
000000 000

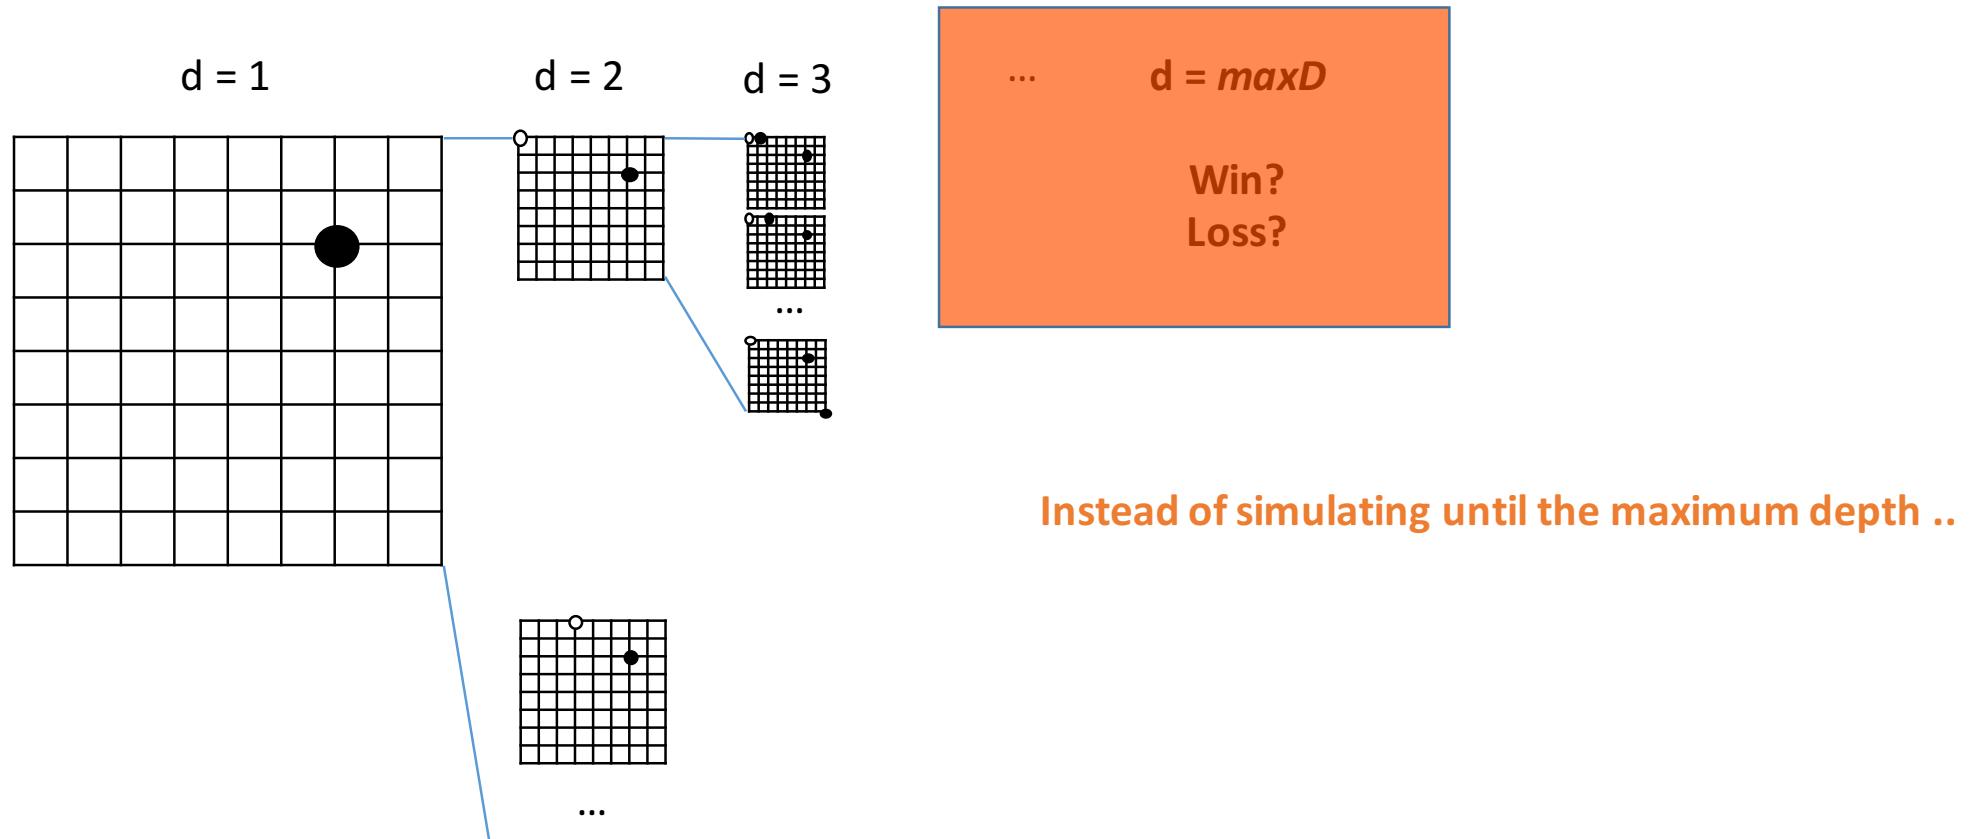
s

$g: s \rightarrow p(a|s)$

$p(a|s)$ argmax a

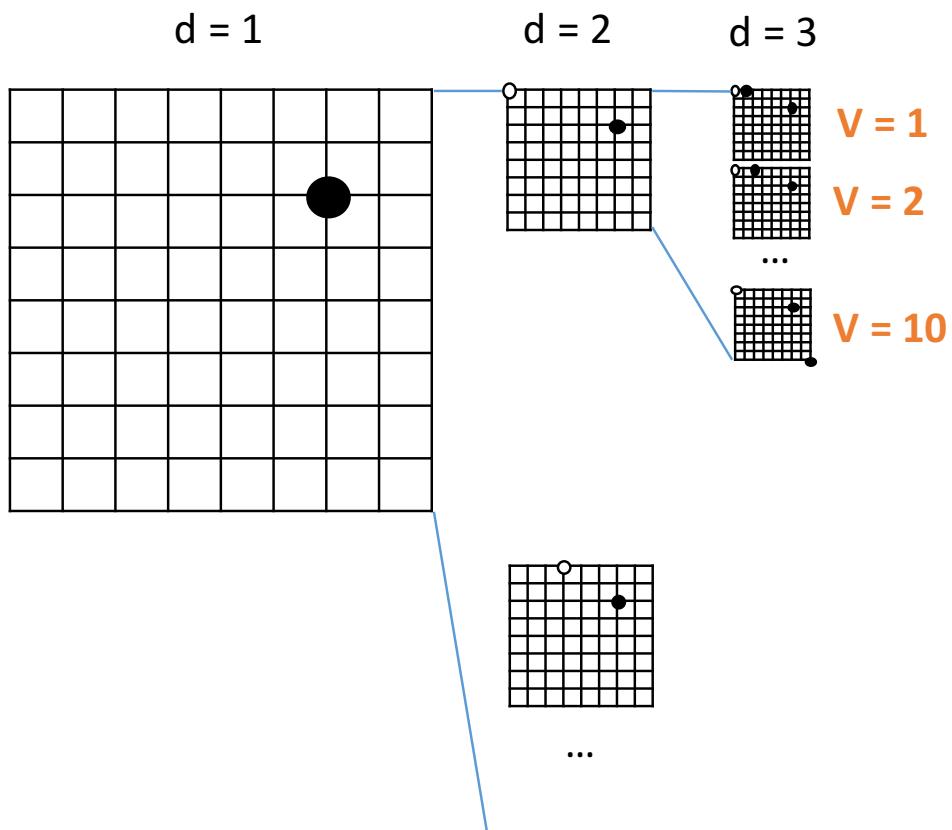
Reducing Depth of Thought

2. Position evaluation ahead of time (Depth Reduction)



Reducing Depth of Thought

2. Position evaluation ahead of time (Depth Reduction)



IF there is a function that can measure:
 $V(s)$: “board evaluation of state s ”



Game records

For other sources of game records, see the [list](#) in the links section.

On this page, you can download game records of top amateur games played on the [K Go Server](#) (KGS, formerly known as the Kiseido Go Server). I am grateful to Bill Shubert, who created KGS, for the permission to use these files, and for making them available to me in an easy way.

The games in the archives below are those where either one of the players (or both) is 7d or stronger, or both are 6d. All comments are stripped from these games, and all games with variations are omitted. They are suitable for use with [Komiblio](#).

Need still more games? Have a look at the [KGS games played by 4d+ players](#).

There are several versions of each archive, compressed in different ways: a .tar.gz, .tar.bz2, and .zip version; please choose the one which is most suitable for you. The content of the uncompressed archives is completely identical.

Source	Time period	Number of games	Archive format	File size	Link
KGS	2017_01	733	.zip	0.7 MB	Download
			.tar.gz	0.3 MB	Download
			.tar.bz2	0.2 MB	Download
2016_12	1208		.zip	1.0 MB	Download
			.tar.gz	0.5 MB	Download
			.tar.bz2	0.3 MB	Download
2016_11	980		.zip	0.9 MB	Download
			.tar.gz	0.4 MB	Download



<https://u-go.net/gamerecords/>

<https://github.com/hughperkins/kgsgo-dataset-preprocessor>

Winning Human Players

1. Reducing “action candidates”

(2) Improving through self-plays (reinforcement learning)

Improving by playing against itself

**Expert Moves
Imitator Model
(w/ CNN)**

vs

**Expert Moves
Imitator Model
(w/ CNN)**



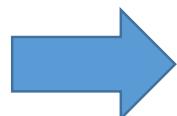
1. Reducing “action candidates”

(2) Improving through self-plays (reinforcement learning)

**Updated Model
ver 1.3**

VS

**Updated Model
ver 1.7**



Return: board positions, win/lose info

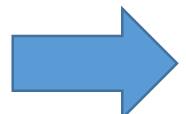
1. Reducing “action candidates”

(2) Improving through self-plays (reinforcement learning)

**Updated Model
ver 1.5**

VS

**Updated Model
ver 2.0**



Return: board positions, win/lose info

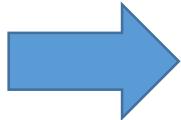
1. Reducing “action candidates”

(2) Improving through self-plays (reinforcement learning)

**Updated Model
ver 3204.1**

VS

**Updated Model
ver 46235.2**



Return: board positions, win/lose info

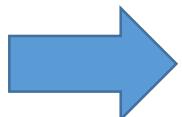
1. Reducing “action candidates”

(2) Improving through self-plays (reinforcement learning)

**Expert Moves
Imitator Model**

VS

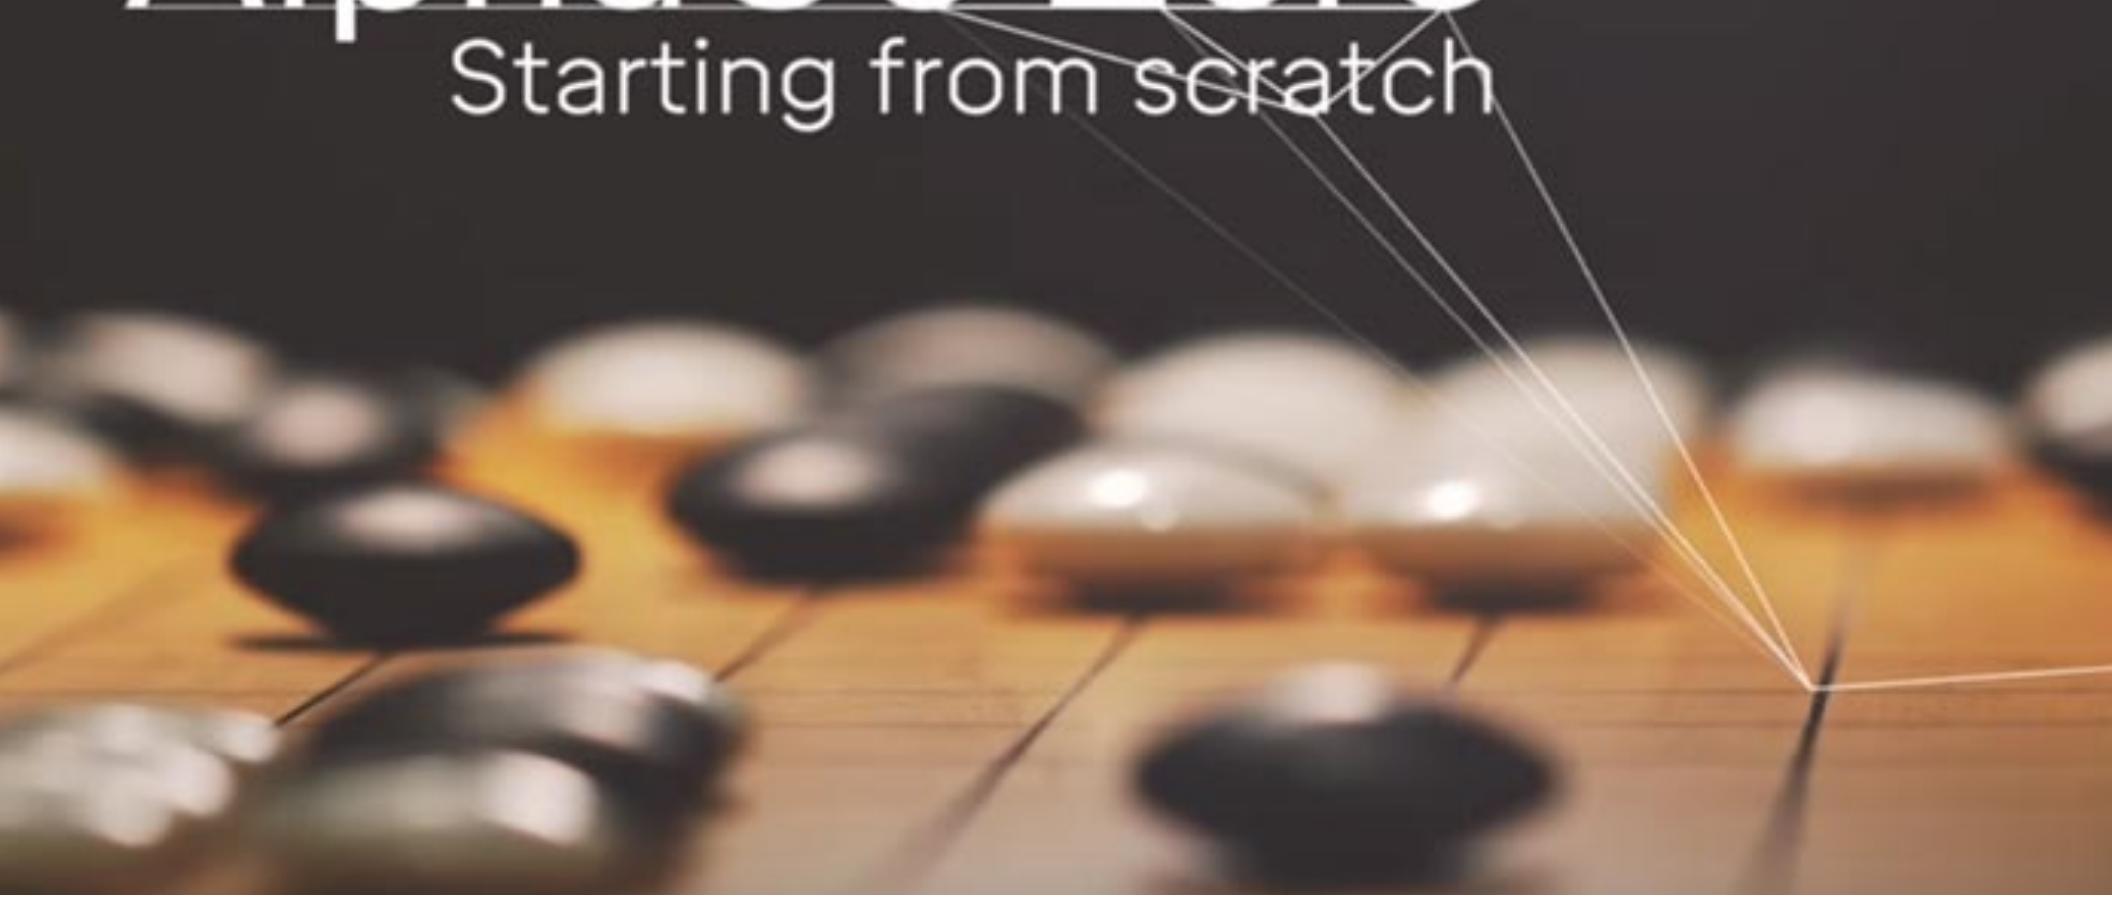
**Updated Model
ver 1,000,000**

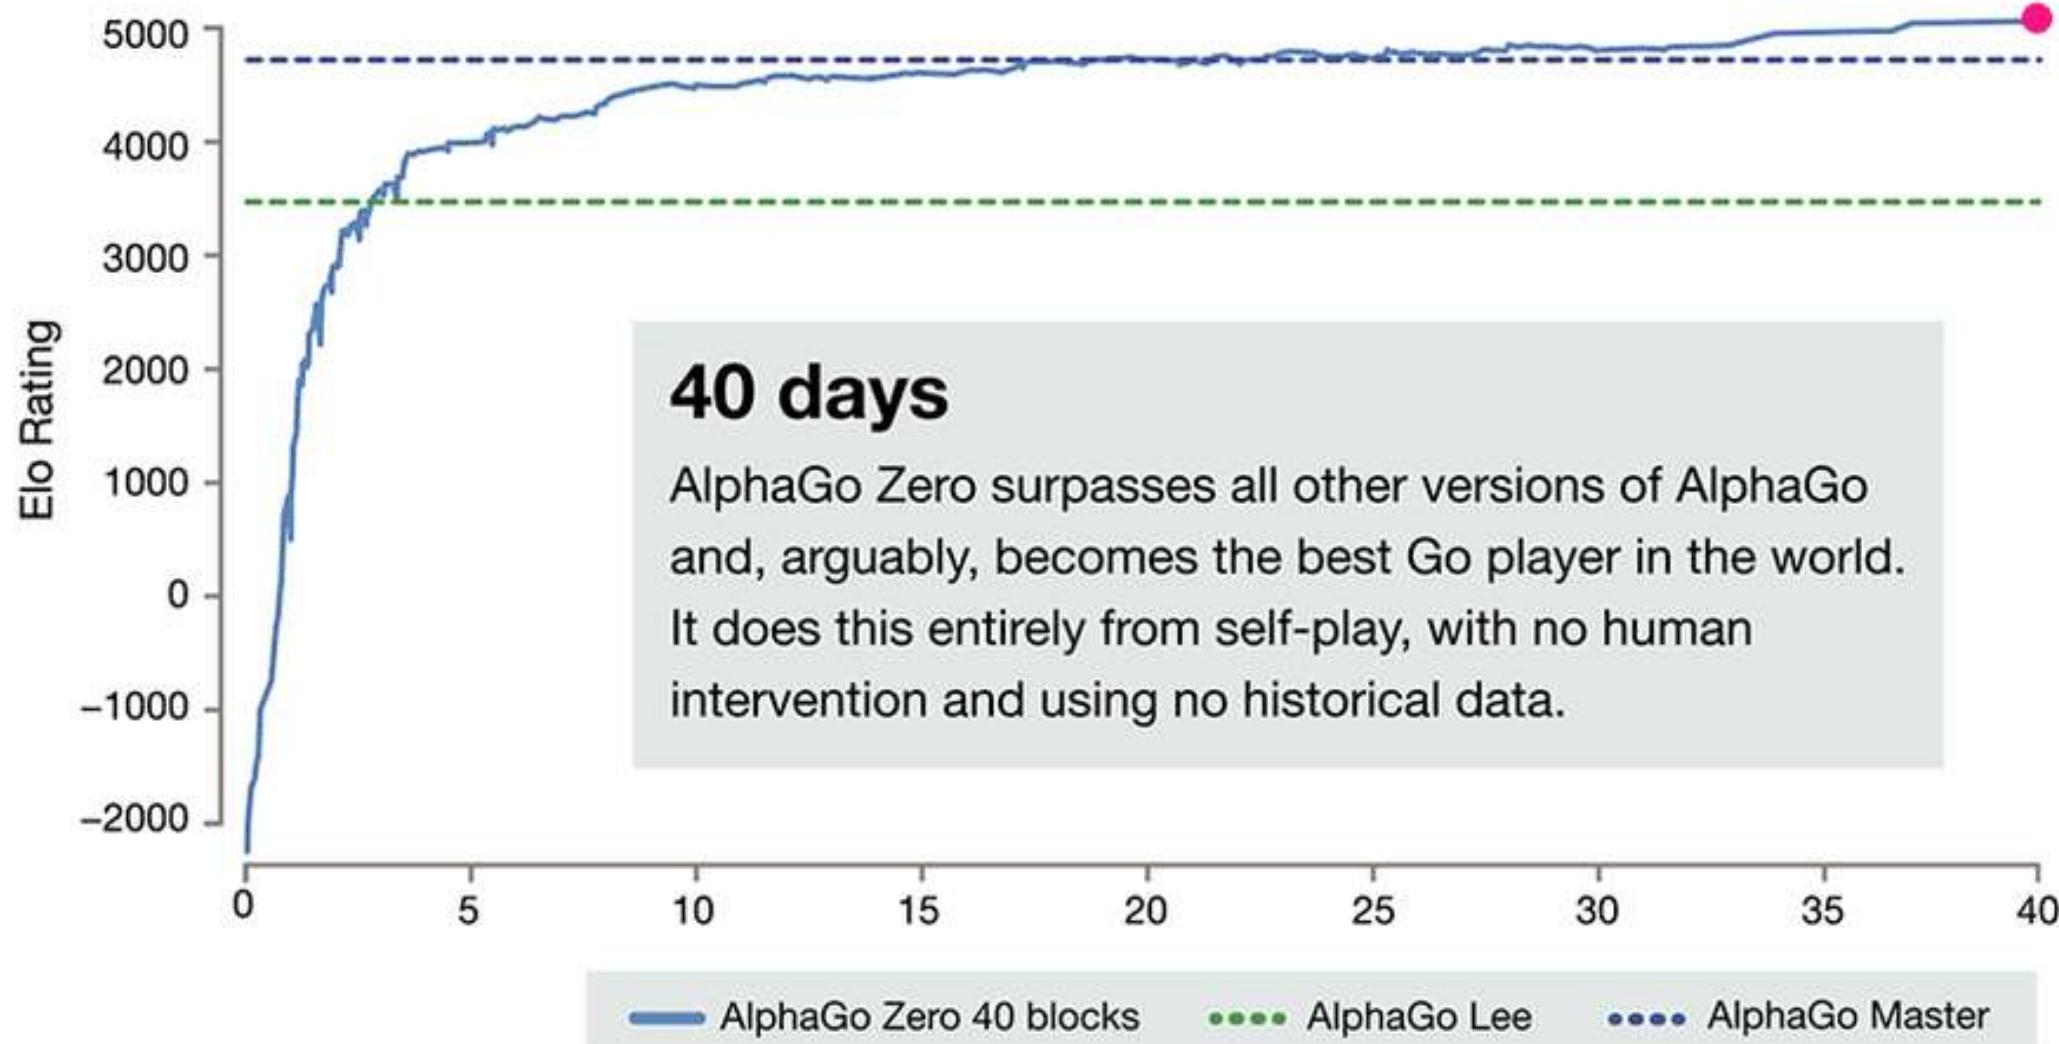


**The final model wins 80% of the time
when playing against the first model**

AlphaGo Zero

Starting from scratch





Lee Sedol 9-dan vs AlphaGo

Energy Consumption

Lee Sedol	AlphaGo
<ul style="list-style-type: none">- Recommended calories for a man per day : ~2,500 kCal- Assumption: Lee consumes the entire amount of per-day calories in this one game $2,500 \text{ kCal} * 4,184 \text{ J/kCal}$ <p>$\approx 10M \text{ [J]}$</p>	<ul style="list-style-type: none">- Assumption: CPU: ~100 W, GPU: ~300 W- 1,202 CPUs, 176 GPUs $170,000 \text{ J/sec} * 5 \text{ hr} * 3,600 \text{ sec/hr}$ <p>$\approx 3,000M \text{ [J]}$</p>

A very, very rough calculation ;)

Seoul AlphaGo

Deep Reinforcement Learning (as Nature AlphaGo)

- Improved value network
- Improved policy network
- Improved search
- Improved hardware (TPU vs GPU)



